00. Index: The AI-Native Transition research pack

Produced overnight, 6 November 2026. Markdown files are the source of truth; PDFs are in *pdf/*, with the combined *pdf/AI-native-transition-full.pdf* in the order 99, 00, 01 to 15.

Basis. The cloud environment blocked page fetches for almost every host, so most findings rest on search-engine extracts that cite primary sources. Pages on anthropic.com, microsoft.com and raw.githubusercontent.com were read directly. Every track file says at the top what is missing. See `verification.md` and `15-open-questions.md` before citing anything publicly.

Start here: the 10 most important findings

1. **Workers are ahead of employers, not resisting.** 78% of AI users bring their own tools and 52% hide use (Microsoft/LinkedIn 2024, n=31,000). Twice as many skilled users are blocked by their organisation as ready organisations are held back by people (WTI 2026, n=20,000). Files: 05, 02, 01.
2. **Agent reliability is a technology limit, not only a readiness problem.** CRMArena-Pro success falls from about 58% to 35% when tasks go multi-turn. tau-bench success halves when every one of several tries must succeed. Files: 06, 11, 04.
3. **Measured gains are small so far.** Denmark: about 3% time savings, zero effect on earnings and hours (Humlum and Vestergaard). US TFP growth slowed to 0.8% in 2025. Files: 11, 05, 08.
4. **Breadth is up, depth and profit are flat.** About 3% of US small employers have AI "fully integrated" (Fed SBCS). The UK's depth proxy barely moved since 2023 (ONS). McKinsey's EBIT-impact share is flat at about 37%. Files: 03, 04.
5. **"Share of pilots reaching production" ranges from 5% to 46%** depending on who asks what. "95% of pilots fail" is a misreading of a preliminary, self-selected report. Files: 04, 12.
6. **Human oversight is mostly deference.** Analysts followed an agent's coin-flipped lean 86% of the time (Anthropic RCT, n=50). Users accept 93% of permission prompts. Files: 07.
7. **Context helps only when people write it.** LLM-generated context files reduced coding-agent success; developer-written ones helped slightly (Gloaguen et al. 2026). No framework defines "AI-native" by shared context. Files: 07, 02, 01.
8. **Harm is hitting entry-level hiring before transformation arrives.** GenAI adopters hired about 80% fewer juniors per quarter; junior employment fell about 9% (Hosseini and Lichtinger). No aggregate unemployment effect is detectable. Files: 08.
9. **"Freedom to be more human" is not supported yet.** AI time savings turned into task expansion (Berkeley Haas ethnography); AI collaboration lowered intrinsic motivation (Scientific Reports 2025). Files: 09, 11.
10. **Public opinion is turning against AI, led by the young.** 52% of US adults are more concerned than excited (Pew, June 2026), and under-30s moved most. Unions are winning contract terms, not stopping deployment. Regulation is delayed (EU high-risk duties to Dec 2027 and Aug 2028). Files: 10.

Files

File	What it is
99-morning-brief.md	Under 800 words: five findings that should change your writing, the strongest evidence against the thesis, three next questions.
00-INDEX.md	This file.
01-thesis-stress-test.md	The thesis split into six claims, each judged: what holds, what breaks, what's contested, plus a suggested rewrite.
02-ai-native-definitions.md	Track 1. Where the term comes from (telecom 2020 onward, VC from 2023), how maturity models differ, named AI-first companies (Klarna, Duolingo, Shopify, IBM, Salesforce), and AI-washing enforcement.
03-smb.md	Track 2. Official vs vendor adoption figures and why they differ by 14% to 76%, depth of use, barriers, costs, who implements, micro-firm evidence.
04-enterprise.md	Track 3. Pilot-to-production figures traced, vendor disclosures, the shared-context problem technically (RAG, text-to-SQL, MCP security), bank deployments, controlled ROI evidence.
05-inertia-and-education.md	Track 4. The "70% of change fails" myth, hidden AI use, manager support, training experiments, gender and age gaps, middle-management cuts.
06-what-goes-wrong.md	Track 5. Named failures with sources, the RPA and AI failure-rate myths, process mining and tacit knowledge, agent reliability and benchmark fragility.
07-working-with-agents.md	Track 6. Teaming experiments, oversight and deference, division of labour, roles cut vs reshaped, context files and skills, liability and audit (Moffatt, EU AI Act Arts 14 and 26, Product Liability Directive).
08-displacement.md	Track 7. Measured employment effects (US, UK, euro area), the junior-hiring effect, layoff announcements, exposure measures vs effects.
09-what-will-people-do.md	Track 8. Historical displacement, UBI and four-day-week trials, Marienthal and the meaning of work, freed-time evidence, solo operators and care work.
10-social-political.md	Track 9. Public opinion by country and age, union cases, EU/UK/US regulation enforced vs announced, surplus capture, the anti-AI case, the AI-advice grifter layer.
11-contrarian.md	Track 10. Slower and less profitable, agents further from production, task change not job loss, and how "more human" can backfire.
12-facts-and-figures.md	About 330 figures with source, URL, date and reliability note; a verified headline table; DO NOT CITE section.
13-quotes.md	41 verbatim quotes, all re-fetched and matched; plus candidate quotes not yet verifiable. Vendor-skewed by necessity.
14-reading-list.md	35 sources ranked, one line each, marked read or not opened.
15-open-questions.md	What was thin, why, and what to research next.
verification.md	Method and results of the verification pass: 29 numbers, 42 quotes.
progress.log	Timestamped log of every step and judgement call.
pdf/	PDF of every Markdown file plus the combined PDF.