13. Quotes

Every quote below was copied from a page fetched in this session and re-fetched and matched word for word in the verification pass (see verification.md). Format: quote (speaker, source, date) URL. Context: one line.

Read this before using these quotes. Only three hosts could be fetched in this session: anthropic.com, microsoft.com and raw.githubusercontent.com. So every verified quote comes from Anthropic, Microsoft or a benchmark repository. Anthropic and Microsoft sell AI and have an interest in adoption. This is a sampling artefact of the session's network limits, not a judgement that these are the best or most balanced quotes. Before publishing, balance them with the candidate quotes at the end once their sources can be opened.

AI-native, shared context and the Frontier Firm (from track file 02)

- "The upshot is that deploying AI for complex tasks might be constrained more by access to information than on underlying model capabilities." (Anthropic Economic Index report, Anthropic, 15 Sep 2025) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report. Context: the authors' reading of input-length data from 1M API transcripts. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "This could represent a barrier to broader enterprise deployment for some important tasks that rely on dispersed context that is not already centralized or digitized. Correcting for this bottleneck may require firms to restructure their organization, invest in new data infrastructure, and centralize information for effective model deployment." (Anthropic Economic Index, 15 Sep 2025) same URL. Context: the closest primary statement to the user's shared-context definition. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "For software development with centralized code repositories, access to this information is in principle straightforward." (Anthropic Economic Index, 15 Sep 2025) same URL. Context: why coding leads agent adoption. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "I work way more with Claude than with any of my colleagues." (Anonymous Anthropic engineer, "How AI is transforming work at Anthropic", 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: reduced collaboration. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "More junior people don't come to me with questions as often." (Anonymous senior Anthropic engineer, same source, 2 Dec 2025) same URL. Context: mentorship loss. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Structured around on-demand intelligence and powered by "hybrid" teams of humans + agents, these companies scale rapidly, operate with agility, and generate value faster." (Microsoft, 2025 Work Trend Index, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: the Frontier Firm definition. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "In many cases, people are ready. The systems around them are not." (Microsoft, 2026 Work Trend Index, 5 May 2026) https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization. Context: framing for the

"Transformation Paradox". Verified: re-fetched and matched in the verification pass, 2026-11-06.

- "They capture these signals and encode them into shared routines, improving future work while preserving accountability and control." (Microsoft, 2026 Work Trend Index, 5 May 2026) same URL. Context: what Frontier Firms do with agent outputs; the nearest Microsoft gets to "shared context". Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Organizations that can answer these questions are building Owned Intelligence" (Microsoft, 2026 Work Trend Index, 5 May 2026) same URL. Context: the page then defines Owned Intelligence as institutional know-how unique to the firm. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Adoption, hidden use and stigma (from track file 05)

- "Without guidance or clearance from the top, employees are taking things into their own hands and keeping AI use under wraps: 78% of AI users are bringing their own AI tools to work (BYOAI)" (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: framing of shadow AI. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "52% of people who use AI at work are reluctant to admit to using it for their most important tasks." (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: hidden use. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Only 39% of people globally who use AI at work have gotten AI training from their company." (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: training deficit. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Last year, employees led the AI wave. This year, it's flipped." (Microsoft, Work Trend Index 2025, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: leaders now ahead of employees on agent familiarity. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "69% of professionals mentioned the social stigma that can come with using AI tools at work" (Anthropic, "Introducing Anthropic Interviewer", 4 Dec 2025) https://www.anthropic.com/research/anthropic-interviewer. Context: 1,250 AI-conducted interviews; vendor source. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "A colleague recently said they hate AI and I just said nothing. I don't tell anyone my process because I know how a lot of people feel about AI." (Fact-checker interviewed by Anthropic Interviewer, 4 Dec 2025) https://www.anthropic.com/research/anthropic-interviewer. Context: hidden AI use because of peer stigma. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "When producing output is so easy and fast, it gets harder and harder to actually take the time to learn something." (Anthropic employee, quoted in "How AI is transforming work at Anthropic", 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: skill-atrophy concern in an AI-native firm. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Agent reliability and benchmarks (from track file 06)

- "**Key finding**: AI agent time horizons have been doubling approximately every 7 months." (METR, eval-analysis-public repository README, accessed 2026-11-06) https://raw.githubusercontent.com/METR/eval-analysis-public/main/README.md. Context: METR's own summary of its time-horizon method. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "τ³-bench added voice full-duplex evaluation, the `banking_knowledge` domain and 75+ task fixes" (Sierra Research, tau2-bench repository README, accessed 2026-11-06) https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md. Context: the benchmark's maintainers repaired dozens of their own task definitions. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "2.0 English runs are **not** input-identical to v1.0.1, so 2.0 scores are not comparable with the 1.x leaderboard." (Sierra Research, tau2-bench README, accessed 2026-11-06) https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md. Context: agent scores are tied to benchmark versions, so cross-time comparisons need care. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "But how performant are AI agents at helping to accelerate or even autonomously perform work-related tasks?" (TheAgentCompany authors, repository README, accessed 2026-11-06) https://raw.githubusercontent.com/TheAgentCompany/TheAgentCompany/main/README.md. Context: framing question of the benchmark. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Working with agents: oversight, deference and division of labour (from track file 07)

- "Across the four analytical choices, analysts went with their agent’s lean 86% of the time." (Amaç Herdağdelen, Hailey Joren, Alex Tamkin, Anthropic, "How AI agents shape experts’ judgment calls", 27 Oct 2026) https://www.anthropic.com/research/agents-judgment-calls. Context: RCT where agent preferences were set by coin flip. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "a human reviewer is only an independent check if their judgment is their own" (same authors, same page, 27 Oct 2026) https://www.anthropic.com/research/agents-judgment-calls. Context: why deference matters for oversight. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Manual prompts sit in the middle, and in practice users accept 93% of them anyway." (Anthropic engineering, "How we built Claude Code auto mode", 25 Mar 2026) https://www.anthropic.com/research/claude-code-auto-mode. Context: approval fatigue in step-by-step oversight. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The 17% false-negative rate on real overeager actions is the honest number." (same page, 25 Mar 2026) https://www.anthropic.com/research/claude-code-auto-mode. Context: residual risk of classifier-based approvals. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "effective oversight doesn’t require approving every action but being in a position to intervene when it matters" (Anthropic, "Measuring AI agent autonomy in practice", 18 Feb 2026) https://www.anthropic.com/research/measuring-agent-autonomy. Context: Anthropic's interpretation of rising auto-approve and interrupt rates. Verified: re-fetched and matched in the verification pass, 2026-11-06.

- "existing models are capable of more autonomy than they exercise in practice" (same page, 18 Feb 2026) https://www.anthropic.com/research/measuring-agent-autonomy. Context: the "deployment overhang" argument. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "On average, people make about 70% of the planning decisions but only 20% of the execution decisions" (Anthropic, "Agentic coding and persistent returns to expertise", 16 Jun 2026) https://www.anthropic.com/research/claude-code-expertise. Context: division of labour in ~400k sessions. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "77% of API transcripts show automation patterns (especially full task delegation) versus just 12% for augmentation" (Anthropic Economic Index, 15 Sept 2025) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report. Context: enterprise API use. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "primarily impacted behaviors that could be changed independently and not behaviors that required coordination to change" (Dillon, Jaffe, Immorlica, Stanton, Microsoft Research abstract, Apr 2025) https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/. Context: Copilot RCT, email vs meetings. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "someone who builds, delegates to, and manages agents to amplify their impact" (Microsoft Work Trend Index 2025, definition of "agent boss", 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: vendor framing (OPINION). Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Some may blend HR and IT or create new leadership roles" (same, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: org-design forecast. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "So far, 1 million skills have been published and nearly 280 million installed." (Anthropic, "Donating Agent Skills to the Agentic AI Foundation", 22 Oct 2026) https://www.anthropic.com/news/donating-agent-skills-to-the-agentic-ai-foundation. Context: vendor adoption claim. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Displacement and expectations (from track file 08)

- "We find no systematic increase in unemployment for highly exposed workers since late 2022, though we find suggestive evidence that hiring of younger workers has slowed in exposed occupations" (Maxim Massenkoff and Peter McCrory, Anthropic economists, "Labor market impacts of AI: A new measure and early evidence", 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: key-findings summary of CPS analysis. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The averaged estimate in the post-ChatGPT era is a 14% drop in the job finding rate compared to that in 2022 in the exposed occupations, although this is just barely statistically significant. (There is no such decrease for workers older than 25.)" (same source, 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: young-worker section. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The young workers who are not hired may be remaining at their existing jobs, taking different jobs, or returning to school." (same source, 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: authors' caveat on the hiring result. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Following adoption, junior employment declines sharply in adopting firms relative to non-adopters, while senior employment trends remain largely unchanged." (Seyed M. Hosseini

and Guy Lichtinger, Harvard, "Generative AI as Seniority-Biased Technological Change", version 6 May 2026, abstract) https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf. Context: headline result. Verified: re-fetched and matched in the verification pass, 2026-11-06.

- "The decline is driven primarily by slower hiring rather than increased separations." (same, abstract) Same URL. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "One potential implication is that if firms overestimate the near-term automation capabilities of GenAI, the resulting decline in junior employment may ultimately prove temporary." (same, introduction) Same URL. Context: authors' forward-looking-hiring interpretation. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Notably, respondents were on average more worried about job loss for others than for themselves." (Anthropic Economic Index report "Cadences", 26 Jun 2026) https://www.anthropic.com/research/economic-index-june-2026-report. Context: survey of about 9,700 Claude users. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The Economic Index Survey is not representative of the general population." (same) Same URL. Context: authors' own sample caveat. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Across US workers, ten more percentage points of work affected is associated with 0.4 points more forecasted unemployment." (Sacher, Cotter, Jones, Korinek, McCrory, "The AI expectations gap", Anthropic, 16 Oct 2026) https://www.anthropic.com/research/ai-expectations-gap. Context: Morning Consult survey, n=10,980. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Candidate quotes NOT verified (do not use in quotation marks yet)

These were seen only in search extracts or remembered from secondary coverage. They are paraphrased here. Fetch the source and copy the exact wording before quoting.

- **Moffatt v. Air Canada, 2024 BCCRT 149 (14 Feb 2024):** the tribunal described Air Canada's argument that its chatbot was a separate legal entity responsible for its own actions as a "remarkable submission". Source to fetch: https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html
- **Sebastian Siemiatkowski (Klarna CEO), 2025:** said the cost focus had produced lower-quality customer service. Exact wording not verified.
- **Commonwealth Bank of Australia, Aug 2025:** reported as conceding it had not adequately considered all relevant business factors when cutting 45 roles. Exact wording not verified.
- **Amjad Masad (Replit CEO), July 2025:** called the agent's deletion of a production database unacceptable and something that should never be possible. Exact wording not verified.
- **Satya Nadella, Microsoft FY26 Q4 call (29 Jul 2026):** said Microsoft has over 30 million paid Microsoft 365 Copilot seats with net seat adds more than doubling quarter over quarter. Wording appears in search extracts only; check the transcript.
- **Gartner (Anushree Verma), 25 Jun 2025:** described most agentic AI projects as early-stage experiments driven by hype. Check the press release.
- **Jim Covello (Goldman Sachs), Top of Mind 129, June 2024:** framed AI spending as a large investment problem that needs AI to solve complex problems to pay back. Check the PDF.

- **Hammer and Champy (1993):** described their 50 to 70% re-engineering failure figure as an unscientific estimate (as reported by Hughes 2011).
- **John Maynard Keynes (1930):** suggested three-hour shifts or a fifteen-hour week would satisfy most people's need to work. Check the essay text.
- **Will Stronge (Autonomy), July 2021:** called the Iceland trials an overwhelming success. Advocacy source.
- **Hosseini and Lichtinger cite Jerome Powell (Sept 2025 FOMC press conference)** on AI and young graduates. Check the Fed transcript.
- **Ron DeSantis (Florida Governor), 2025-26:** said in substance that an executive order cannot pre-empt state legislation. Check a primary transcript.