The AI-Native Transition: full research pack

- 99. Morning brief
- 00. Index: The AI-Native Transition research pack
- 01. Thesis stress test
- Track 1: What "AI-native" means
- Track 2: SMB
- Track 3: Enterprise
- Track 4: Inertia and education
- Track 5: What goes wrong
- Track 6: Working with agents
- Track 7: Displacement
- Track 8: What will people do?
- Track 9: Social and political fallout
- Track 10: Contrarian
- 12. Facts and figures
- 13. Quotes
- 14. Reading list
- 15. Open questions

99. Morning brief

6 November 2026. All ten tracks, the synthesis files, verification and PDFs are done and pushed.

Caveat. The environment blocked page fetches for almost every site, so most findings rest on search extracts citing the primary source, not on reading it. I re-checked 29 key numbers (2 corrected, 1 dropped) and all 42 quotes. Every quote comes from Anthropic or Microsoft, the only reachable sites. See verification.md.

Five findings that should change how you write

1. Workers are not the blocker. Employers are. In Microsoft's 2024 survey (n=31,000), 78% of AI users brought their own tools and 52% hid using them on important work. The 2026 edition (n=20,000) found twice as many skilled users blocked by their organisation (10%) as ready organisations held back by their people (5%). Only 13% are rewarded for redesigning work. Where resistance exists it looks rational: 53% worry AI use makes them look replaceable, and at Commonwealth Bank the union was right that the voice bot hadn't cut call volumes. Drop "people resist" and write about permission, incentives and trust.

2. The tech is not the easy part for agents. On business-shaped tasks, reliability collapses with repetition and turns: Salesforce's own CRMArena-Pro falls from about 58% single-turn to 35% multi-turn. On tau-bench, success drops from 69% on one try to 46% when all four tries must succeed. Half the named failures (Air Canada, NYC MyCity, Deloitte, about 1,600 hallucinated-citation rulings, EchoLeak) are the technology failing, not the business. Anthropic's own engineers can fully delegate only 0 to 20% of their work. Write "both have to move", not "the tech is solved".

3. The gain may be small, not the readiness low. In Danish administrative data covering about 25,000 workers in exposed jobs, adoption plus employer encouragement produced about 3% time savings and precisely zero effect on earnings or hours. US TFP growth slowed to 0.8% in 2025. McKinsey's EBIT-impact share is flat at about 37% while agent scaling rose. "The business isn't ready" can be a comforting misdiagnosis when the honest answer is "this use case doesn't pay yet".

4. Shared context: plausible, but only human-written context is shown to help. The best support is vendor inference: Anthropic's 1M-transcript analysis says complex deployment is "constrained more by access to information than on underlying model capabilities". The only controlled test (Gloaguen et al., 2026, coding agents) found LLM-generated context files reduced success and raised cost over 20%, while developer-written files helped slightly. That favours a Studio framing of context curated by the people who do the work, not auto-ingested data. No framework defines "AI-native" by shared context, so present it as your definition, which is a positioning asset.

5. "Freedom to be more human" has the evidence against it so far. An eight-month ethnography (Berkeley Haas, 2026) found AI time savings turned into task expansion and longer hours. Four experiments (n=3,562) found AI raised performance but lowered intrinsic motivation afterwards. In a 2026 RCT, analysts followed an agent's coin-flipped lean 86% of the time. And the job guarantee in Marienthal showed work itself supplies structure, status and contact. You can argue rote-work removal should free people. You cannot yet say it does, and saying so invites an easy rebuttal.

Strongest evidence against your thesis

The official barrier data don't name your causes. UK firms cite "can't identify a use case", cost and skills. OECD non-adopters cite "poor fit with our work" (57%). A third of US small employers have no plans to use AI, mostly because it's not applicable or they'd rather not. That is choosing not to adopt, not failing to adopt. In enterprise, Gartner and S&P put cost, unclear value and risk controls first. Meanwhile harm is arriving before transformation: GenAI adopters hired about 80% fewer juniors per quarter, possibly in anticipation of capabilities that haven't materialised.

Viral numbers to stop using

"95% of pilots fail" (self-selected, preliminary, measures something narrower), "70% of change programmes fail" (no empirical basis), "80% of AI projects fail" (RAND attributes it to "some estimates"), "UK lost 8% of jobs to AI" (not a labour-market statistic). Full list: 12-facts-and-figures.md.

Three questions to answer next

1. **Does inability to describe the workflow predict failure?** It's your central causal claim and nothing found tests it. Your own client data (process maps before and after, time to production, what broke) could be the first evidence anyone has.
2. **Does shared business context improve agent outcomes outside coding?** Run a controlled comparison in Studio: no context, auto-generated context, human-curated context, on real finance or ops tasks.
3. **Is the junior-hiring drop permanent?** If AI-native firms are removing the apprenticeship layer, that's the human consequence most worth writing about, and the evidence is moving fast.

00. Index: The AI-Native Transition research pack

Produced overnight, 6 November 2026. Markdown files are the source of truth; PDFs are in *pdf/*, with the combined *pdf/AI-native-transition-full.pdf* in the order 99, 00, 01 to 15.

Basis. The cloud environment blocked page fetches for almost every host, so most findings rest on search-engine extracts that cite primary sources. Pages on anthropic.com, microsoft.com and raw.githubusercontent.com were read directly. Every track file says at the top what is missing. See `verification.md` and `15-open-questions.md` before citing anything publicly.

Start here: the 10 most important findings

1. **Workers are ahead of employers, not resisting.** 78% of AI users bring their own tools and 52% hide use (Microsoft/LinkedIn 2024, n=31,000). Twice as many skilled users are blocked by their organisation as ready organisations are held back by people (WTI 2026, n=20,000). Files: 05, 02, 01.
2. **Agent reliability is a technology limit, not only a readiness problem.** CRMArena-Pro success falls from about 58% to 35% when tasks go multi-turn. tau-bench success halves when every one of several tries must succeed. Files: 06, 11, 04.
3. **Measured gains are small so far.** Denmark: about 3% time savings, zero effect on earnings and hours (Humlum and Vestergaard). US TFP growth slowed to 0.8% in 2025. Files: 11, 05, 08.
4. **Breadth is up, depth and profit are flat.** About 3% of US small employers have AI "fully integrated" (Fed SBCS). The UK's depth proxy barely moved since 2023 (ONS). McKinsey's EBIT-impact share is flat at about 37%. Files: 03, 04.
5. **"Share of pilots reaching production" ranges from 5% to 46%** depending on who asks what. "95% of pilots fail" is a misreading of a preliminary, self-selected report. Files: 04, 12.
6. **Human oversight is mostly deference.** Analysts followed an agent's coin-flipped lean 86% of the time (Anthropic RCT, n=50). Users accept 93% of permission prompts. Files: 07.
7. **Context helps only when people write it.** LLM-generated context files reduced coding-agent success; developer-written ones helped slightly (Gloaguen et al. 2026). No framework defines "AI-native" by shared context. Files: 07, 02, 01.
8. **Harm is hitting entry-level hiring before transformation arrives.** GenAI adopters hired about 80% fewer juniors per quarter; junior employment fell about 9% (Hosseini and Lichtinger). No aggregate unemployment effect is detectable. Files: 08.
9. **"Freedom to be more human" is not supported yet.** AI time savings turned into task expansion (Berkeley Haas ethnography); AI collaboration lowered intrinsic motivation (Scientific Reports 2025). Files: 09, 11.
10. **Public opinion is turning against AI, led by the young.** 52% of US adults are more concerned than excited (Pew, June 2026), and under-30s moved most. Unions are winning contract terms, not stopping deployment. Regulation is delayed (EU high-risk duties to Dec 2027 and Aug 2028). Files: 10.

Files

File	What it is
99-morning-brief.md	Under 800 words: five findings that should change your writing, the strongest evidence against the thesis, three next questions.
00-INDEX.md	This file.
01-thesis-stress-test.md	The thesis split into six claims, each judged: what holds, what breaks, what's contested, plus a suggested rewrite.
02-ai-native-definitions.md	Track 1. Where the term comes from (telecom 2020 onward, VC from 2023), how maturity models differ, named AI-first companies (Klarna, Duolingo, Shopify, IBM, Salesforce), and AI-washing enforcement.
03-smb.md	Track 2. Official vs vendor adoption figures and why they differ by 14% to 76%, depth of use, barriers, costs, who implements, micro-firm evidence.
04-enterprise.md	Track 3. Pilot-to-production figures traced, vendor disclosures, the shared-context problem technically (RAG, text-to-SQL, MCP security), bank deployments, controlled ROI evidence.
05-inertia-and-education.md	Track 4. The "70% of change fails" myth, hidden AI use, manager support, training experiments, gender and age gaps, middle-management cuts.
06-what-goes-wrong.md	Track 5. Named failures with sources, the RPA and AI failure-rate myths, process mining and tacit knowledge, agent reliability and benchmark fragility.
07-working-with-agents.md	Track 6. Teaming experiments, oversight and deference, division of labour, roles cut vs reshaped, context files and skills, liability and audit (Moffatt, EU AI Act Arts 14 and 26, Product Liability Directive).
08-displacement.md	Track 7. Measured employment effects (US, UK, euro area), the junior-hiring effect, layoff announcements, exposure measures vs effects.
09-what-will-people-do.md	Track 8. Historical displacement, UBI and four-day-week trials, Marienthal and the meaning of work, freed-time evidence, solo operators and care work.
10-social-political.md	Track 9. Public opinion by country and age, union cases, EU/UK/US regulation enforced vs announced, surplus capture, the anti-AI case, the AI-advice grifter layer.
11-contrarian.md	Track 10. Slower and less profitable, agents further from production, task change not job loss, and how "more human" can backfire.
12-facts-and-figures.md	About 330 figures with source, URL, date and reliability note; a verified headline table; DO NOT CITE section.
13-quotes.md	41 verbatim quotes, all re-fetched and matched; plus candidate quotes not yet verifiable. Vendor-skewed by necessity.
14-reading-list.md	35 sources ranked, one line each, marked read or not opened.
15-open-questions.md	What was thin, why, and what to research next.
verification.md	Method and results of the verification pass: 29 numbers, 42 quotes.
progress.log	Timestamped log of every step and judgement call.
pdf/	PDF of every Markdown file plus the combined PDF.

01. Thesis stress test

Written 2026-11-06 after all ten track files existed. Every claim below links to a source; the track files carry fuller context, sample sizes and caveats. Labels: EVIDENCE, FORECAST, OPINION.

Read this caveat first. This session's network policy blocked page fetches from almost every source site. Only anthropic.com, microsoft.com and raw.githubusercontent.com pages could be opened. Everything else rests on search-engine extracts that name and link the primary source. Treat numbers from non-fetched sources as "reported, not yet read in the original". `verification.md` lists what was re-checked and how.

The thesis, split into testable parts

> The tech is the easy part; the transformation is the hard part. AI fails in businesses because the business isn't ready: nobody can describe the real workflow, there's no shared context humans and agents both work from, and people resist a whole new way of working.

The thesis bundles six claims. They fare very differently.

#	Claim	Verdict
A	The tech is the easy part	**Breaks** for agents. Partly holds for chat assistants.
B	Transformation is the hard part	**Holds**, with good evidence
C	AI fails because the business isn't ready	**Partly holds.** One cause among several. Cost, unclear value, risk and tech reliability are cited as often.
D	Nobody can describe the real workflow	**Plausible, thinly evidenced.** One strong source (RAND); no measurement.
E	No shared context for humans and agents	**Plausible, vendor-evidenced.** Only human-written context is shown to help, and only slightly, for coding agents. Auto-generated context hurt.
F	People resist a whole new way of working	**Mostly breaks.** Workers adopt faster than employers. Where resistance exists it is often rational.

What holds

B. Transformation is the hard part: holds.

- Breadth of adoption is rising while depth and profit impact are flat. US Fed Small Business Credit Survey (fieldwork Sep to Nov 2025): about half of small-employer AI users are "experimenting" and 7% "fully integrated", roughly 3% of all employer firms. EVIDENCE (search extract). https://www.fedsmallbusiness.org/2026-report-on-employer-firms
- ONS (2026): UK adoption roughly tripled since 2023, while its depth proxy (AI technologies per firm) barely moved. EVIDENCE (search extract). See 03-smb.md.
- McKinsey State of AI: share reporting any EBIT impact about 39% in 2025 and about 37% in 2026, while the share of $1B+ firms scaling agents rose from 27% to 40%. Self-reported consultancy survey. EVIDENCE (weak). https://www.mckinsey.com/~/media/mckinsey/business%20functions/quantumblack/our%20insights/the%20state%20of%20ai/the-state-of-ai-in-2026-on-the-road-to-roi-aug-2026.pdf
- Microsoft's 6,000-worker Copilot RCT: email time fell, meeting time did not. Changes that need several people to change together did not happen on their own. This is the cleanest evidence

for the transformation half of the thesis. EVIDENCE. https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/

D. "Can't describe the workflow": the best support comes from one interview study.

- RAND (Aug 2024, 65 experienced data scientists and engineers, interviews) names as root cause one: stakeholders misunderstand or miscommunicate the problem AI should solve. EVIDENCE (qualitative). https://www.rand.org/pubs/research_reports/RRA2680-1.html
- IDC attributes POC failure to organisational readiness (data, processes, infrastructure). Vendor-sponsored. OPINION. See 04-enterprise.md.
- No study I found measures how often clients cannot describe their workflow, or how much that predicts failure. The process-mining literature (06-what-goes-wrong.md) documents gaps between designed and actual processes, but not as an AI failure cause. This is a gap you could fill with your own client data.

Over-automation for cost, and the reversals that follow, are well documented.

- Commonwealth Bank of Australia cut 45 roles citing a voice bot, then reversed after the union showed call volumes were rising (Aug 2025). EVIDENCE (search extract). https://ia.acs.org.au/article/2025/cba-reverses-ai-driven-job-cuts--admits--error-.html
- Klarna headcount fell from 5,527 to 3,422 per its F-1, and the CEO later conceded "lower quality" service and rehired for quality (see 02 and 06). This supports "transformation is hard", but the failure was management's business case, not workers' readiness.

What breaks

A. "The tech is the easy part" breaks for agents.

- Agent reliability collapses on repeated, multi-turn, policy-bound work, which is what business processes are. Salesforce's own CRMArena-Pro: about 58% single-turn success falling to about 35% multi-turn, and weak handling of confidential data (secondary summaries say near zero without prompting; check the paper). EVIDENCE. https://arxiv.org/html/2505.18878v1
- tau-bench retail: best listed model 69% on one try, 46% when it must succeed on all four (pass^4). EVIDENCE (fetched). Caveat: these are 2024 models (Claude 3.5 Sonnet, GPT-4o); current models score higher, but the pass^k drop pattern is the point. https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md
- Spider 2.0: best models solved about 17 to 21% of real enterprise text-to-SQL tasks, against about 91% on the older academic benchmark. Even with the data in place, querying it is hard. EVIDENCE (search extract). https://arxiv.org/abs/2411.07763
- At least half the named failures in 06-what-goes-wrong.md are the technology failing: hallucination (Moffatt v. Air Canada, NYC MyCity, Deloitte Australia, about 1,600 court decisions involving hallucinated citations by June 2026), prompt injection (EchoLeak, CVSS 9.3), destructive agent actions (Replit). A business with perfectly documented workflows would still have hit these. https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html
- Even inside Anthropic, an organisation as AI-native as any, engineers use Claude in about 59% of their work but say they can "fully delegate" only 0 to 20% of it. Supervision, not readiness, is the binding constraint. EVIDENCE (vendor self-study, n=132 survey, 53 interviews; fetched). https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic

F. "People resist" mostly breaks.

- Workers adopt ahead of employers. Microsoft/LinkedIn 2024 (n=31,000): 78% of AI users bring their own tools, 52% are reluctant to admit using it for important tasks. EVIDENCE (vendor survey; fetched). https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part
- Microsoft Work Trend Index 2026 (n=20,000 AI users): 10% are skilled workers blocked by their organisation against 5% in ready organisations whose people are not ready; 13% are rewarded for redesigning work. Microsoft's own summary is that people are ready and the systems around them are not. EVIDENCE (vendor survey, self-report; fetched). https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization
- Where resistance exists, it looks rational: 53% of AI users worry that using AI on important work makes them look replaceable (same 2024 source). The CBA case shows staff being right about the bot.
- Many small firms are declining AI, not failing at it. Fed SBCS: a third of small employers have no plans to use AI; of those, over half say it is not applicable and 30% prefer not to. EVIDENCE (search extract). https://www.fedsmallbusiness.org/2026-report-on-employer-firms

The official barrier data do not name the thesis's causes. When ONS asks UK firms why they don't adopt AI, the top answers are difficulty identifying a use case (39%), cost (21%) and skills (16%). No official survey found asks about undocumented workflows or shared context. EVIDENCE (search extract). https://www.ons.gov.uk/economy/economicoutputandproductivity/productivitymeasures/articles/managementpracticesandtheadoptionoftechnologyandartificialintelligenceinukfirms2023/2025-03-24 The OECD's December 2025 SME survey finds 31% of SMEs using generative AI, and the top barrier among non-adopters is "poor fit with our work" (57%). EVIDENCE (search extract). https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en.html

Training and encouragement do not reliably transfer. A Kenyan RCT with 640 entrepreneurs found no average gain from an AI business mentor: it helped high-performing owners and hurt low performers. A Microsoft-authored training RCT at Gap found no average effect. Readiness programmes are not a proven fix either. EVIDENCE (search extract). See 03-smb.md and 05-inertia-and-education.md.

Enterprise failure causes are as much economic and risk as readiness. Gartner's forecast of over 40% of agentic projects cancelled by end-2027 cites cost, unclear business value and inadequate risk controls (FORECAST, poll of 3,412 webinar attendees). S&P Global's 2025 survey cites cost, privacy and security. https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027 https://www.ciodive.com/news/AI-project-fail-data-SPGlobal/742590/

What's contested

C. "AI fails" itself is contested. "Share of pilots reaching production" runs from 5% (MIT NANDA: 5% of organisations got custom tools into production with P&L impact; 52 interviews, 153 surveys, self-described preliminary) through 12% (IDC 2025) to 46% (IDC/Lenovo 2026, n=3,120). The same sponsor nearly quadrupled its own figure in a year. Wharton/GBK (n>800) reports 74% positive ROI. All are self-reported. The "95% of AI pilots fail" headline does not measure what it is quoted as measuring. https://getdx.com/blog/the-ai-divide/ https://www.businesswire.com/news/home/20260127255538/en/

Magnitude may be the real constraint, not readiness. In Denmark (about 25,000 workers in 11 exposed occupations, administrative data), wide adoption with employer encouragement produced average time savings of about 3% and precise zero effects on earnings and hours. If the value of today's tools for typical tasks is small, "the business isn't ready" may be the wrong explanation for disappointing results. EVIDENCE (search extract). https://econpapers.repec.org/RePEc:nbr:nberwo:33777 US TFP growth slowed to 0.8% in 2025 from 1.5% in 2024. EVIDENCE (search extract). Fair counters: TFP grew 1.0% a year over 2019-25, faster than 0.6% over 2007-19, and the productivity J-curve predicts measured gains lag intangible investment. https://www.bls.gov/news.release/archives/prod3_03192026.pdf

E. Shared context: plausible, but "write it down for the agent" is not proven.

- For: Anthropic's analysis of 1M API transcripts says complex enterprise deployment may be "constrained more by access to information than on underlying model capabilities". Vendor, inferred from input-length patterns. EVIDENCE (weak; fetched). https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
- Mixed: a Feb 2026 study (Gloaguen et al., ETH Zurich) found that context files (AGENTS.md style) tended to reduce coding-agent task success and raised cost by more than 20%. But the drop came from LLM-generated files; developer-written files slightly improved success, and the authors advise short human-written files stating only minimal requirements. A Jan 2026 study found faster runs with no change in completion. This is closer to your thesis than it first looks: context written by people who know the work helps a little; context generated wholesale hurts. EVIDENCE (search extract, coding agents only). https://arxiv.org/abs/2602.11988
- Against: shared context is also a shared attack surface. MCP tool poisoning, the GitHub MCP exploit and Asana's MCP exposure of about 1,000 customers (2025) show that one store both humans and agents read and write is a security problem as well as a productivity one. See 04-enterprise.md.

Human oversight is weaker than "humans and agents work together" implies. In a 2026 RCT, 50 trained analysts followed a coin-flipped agent "lean" on 86% of analytical choices; experience did not protect them. Users accept 93% of Claude Code permission prompts. A shared context does not guarantee shared judgment. EVIDENCE (vendor research; fetched). https://www.anthropic.com/research/agents-judgment-calls

Harm can arrive before the transformation does. Firms adopting generative AI hired about 80% fewer juniors per quarter, and junior employment fell about 9% after six quarters while seniors were unaffected. The authors suggest firms may be cutting in anticipation, and the drop "may ultimately prove temporary". So "the business isn't ready" does not mean "nothing happens to people". EVIDENCE (working paper; fetched). https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf Aggregate unemployment effects remain undetectable (Anthropic CPS analysis; Dallas Fed: at most about 0.1 point of the unemployment rise). https://www.anthropic.com/research/labor-market-impacts https://www.dallasfed.org/research/economics/2026/0106

"Freedom to be more human" has weak or negative early evidence. An eight-month ethnography (Berkeley Haas, HBR Feb 2026) found AI time savings turned into task expansion and longer hours. Four experiments (Scientific Reports 2025, N=3,562) found AI collaboration raised performance but lowered intrinsic motivation on return to solo work. A job guarantee in Marienthal raised wellbeing, structure and status, suggesting work itself supplies things people miss when it goes. See 09 and 11. https://newsroom.haas.berkeley.edu/ai-promised-to-free-up-workers-time-uc-berkeley-haas-researchers-found-the-opposite/ https://docs.iza.org/dp16088.pdf

The working definitions

- **AI-native = one shared stored context both humans and agents read and write.** No framework found defines it this way. Telecom defines it by architecture (Nokia Bell Labs, from 2020), VCs by origin or dependency on models, consultancies by deployment breadth and self-reported ROI. Your definition is distinctive, which is useful for positioning, but you should present it as your definition, not the industry's. See 02-ai-native-definitions.md.
- **Legacy = humans only.** Close to an empty category among knowledge-work firms in the survey data; every maturity model starts at "some AI use". But it is not empty among small firms: a third of US small employers have no plans to use AI. "Legacy" is real in SMB, rare in enterprise.
- **Spectrum.** Matches all published maturity models. They measure position on it by breadth of use or self-reported readiness, not by context architecture.

A thesis the evidence would support

Suggested rewrite, for you to test further:

> The technology is not the easy part, but it is improving faster than organisations change. Most firms are using AI widely and shallowly. Gains so far are small and mostly kept by workers as time, not by firms as profit. Agents fail on the multi-step, rule-bound work that business processes consist of, and humans defer to them more than they should. Moving from shallow to deep use needs coordinated changes to how work is organised, a store of shared context, and proof that it pays. Workers are usually ahead of their employers on adoption, and their caution is often justified.

Where to be careful in public writing:

1. Don't cite "95% of pilots fail", "70% of change programmes fail", "80% of AI projects fail" or "30 to 50% of RPA projects fail" as evidence. All four trace to weak or unscientific origins (see 12-facts-and-figures.md, DO NOT CITE).
2. Don't claim workers are the blocker. The evidence says the opposite more often.
3. Don't claim freed time makes people more human. The early evidence points to intensification. You can argue it should, and what it would take, but not that it does.

INCOMPLETE (read before citing): The network for this session reached only microsoft.com and anthropic.com. The Microsoft and Anthropic sources were fetched and checked against the raw page. Everything else was seen only through search-engine result summaries, marked "(search summary; page not fetched)" or "(search extract)". A gap-fill pass (2026-11-06) added: SEC, DOJ and FTC "AI washing" actions (Delphia and Global Predictions, Presto Automation, Nate/Saniger, Operation AI Comply), Gartner's "agent washing" estimate, Klarna's F-1 headcount figures and the scope of its reversal, Duolingo's measured backlash, and a further attempt to date the business use of "AI-native". STILL MISSING: Gartner, McKinsey, BCG, Deloitte and Accenture maturity models (BCG, McKinsey and Deloitte survey figures are in 04-enterprise.md); the Klarna F-1 risk-factor wording (only secondary reports of its headcount figures); SEC AI-washing actions after April 2025; filings and revenue for AI-native startups (Cursor, Midjourney, Gamma, Lovable); the first business or VC use of "AI-native" (still unresolved).

Track 1: What "AI-native" means

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Summary (5-8 bullets, the things that matter most)

- **No one agrees on what "AI-native" means, and none of the sources I could reach defines it by shared context.** Software and VC writing defines it by **origin** (built around AI from day one) or by **dependency** (the product would not work without models). Telecom defines it by **architecture** (AI designed into every layer instead of added on top). Consultancy and vendor models define it by **breadth of deployment plus self-reported ROI**. Ericsson's own white paper opens by saying the term has no clear definition.
- **Telecom is the earliest documented formal use I found.** Nokia Bell Labs wrote "AI-native air interface" in a paper first posted to arXiv in December 2020 (2012.08285); the Bell Labs page dates it May 2021. Ericsson followed with a 2023 white paper and a maturity matrix. I could not establish who first used the term in software or VC writing. I found no evidence that anyone "coined" it in the business sense.
- **The best primary support for the user's shared-context idea comes from Anthropic, a model vendor.** Its September 2025 Economic Index (1M API transcripts) says complex enterprise deployment "might be constrained more by access to information than on underlying model capabilities". It adds that firms may need to "centralize information". This is inferred from input-length patterns, not measured directly, and it comes from a party with a commercial interest.
- **Microsoft's 2026 Work Trend Index moves toward the user's framing.** It uses the terms "Owned Intelligence" (institutional know-how that compounds) and "Learning System". It says Frontier Firms "encode" agent signals "into shared routines". But its maturity classification is based on self-reported culture, manager support and governance, not on any shared context store.

- **The same 2026 report weakens the "people resist" part of the thesis.** Its framing: "people are ready. The systems around them are not." In the survey, 10% of AI users are skilled workers blocked by their organisation, against 5% where the organisation is ready and the people are not. Only 13% say they are rewarded for reinventing work with AI (n=20,000, self-report).
- **Regulators police false AI capability claims, not the label.** The SEC (Delphia and Global Predictions, March 2024, $400,000; Presto Automation, January 2025), DOJ and SEC (Nate, April 2025, $40M+ raised on an "AI" app run by Philippine contract workers) and the FTC (Operation AI Comply, five companies, September 2024) all acted where claimed AI was absent or was humans or a third party. Gartner separately estimates only about 130 of thousands of "agentic" vendors are genuine (no method published).
- **Named "AI-first" rebuilds were mostly headcount policies and customer-service automation, not shared-context platforms.** The cases are Klarna, Duolingo, Shopify, Salesforce and IBM. Recorded failures were about quality (Klarna's CEO said the AI produced "lower quality" service), brand backlash (Duolingo) and viral exaggeration (IBM "8,000", Salesforce "4,000").
- **Even inside a highly AI-native organisation, full delegation stays small.** Anthropic's study of its own engineers (n=132 survey, 53 interviews, 200k transcripts): Claude was used in 59% of work, but most engineers can "fully delegate" only 0-20%. The costs reported were less mentorship and less collaboration. This undercuts the claim that "the tech is the easy part": supervising the agent remains the binding constraint.
- **The user's "legacy = humans only" category is close to empty in the survey data.** The 2025 Work Trend Index survey screened knowledge workers across 31 countries, and the 2026 edition sampled only AI users. Every maturity model I saw starts at "some AI use", so a "humans only" baseline matches no published framework. The spectrum idea does match them, but they measure position on it by deployment breadth or self-reported readiness, not by context architecture.

Key findings

Where the term comes from (telecom first, as far as I could establish)

- **[EVIDENCE]** Nokia Bell Labs paper "Toward a 6G AI-Native Air Interface". The arXiv version (2012.08285) dates from December 2020 (inferred from the arXiv ID). It proposes an air interface "partially designed by AI" and uses the phrase "AI-Native Air Interface (AI-AI)". The Bell Labs page dates it May 2021. (Search summary; pages not fetched: arXiv and nokia.com blocked.) https://arxiv.org/pdf/2012.08285 and https://www.nokia.com/bell-labs/publications-and-media/publications/towards-a-6g-ai-native-air-interface/
- **[OPINION]** Ericsson white paper "Defining AI native: A key enabler for advanced intelligent telecom networks", probably March 2023 (date unconfirmed). Per search snippets, it says the AI native concept "still does not have a clear definition". It names an "AI-centric architecture" as the goal and offers an "AI native Maturity Level Matrix" for operators. (Search summary; page not fetched.) https://www.ericsson.com/en/reports-and-papers/white-papers/ai-native
- **[OPINION]** Ericsson Technology Review, 2023, "Approaching AI-native radio-access networks". It proposes a stepwise path to a more AI-native RAN. (Search summary; not fetched.) https://www.ericsson.com/4928b8/assets/local/reports-papers/ericsson-technology-review/docs/2023/approaching-ai-native-radio-access-networks.pdf
- **[OPINION]** IEEE ComSoc tech blog, 15 Jan 2026. It contrasts "AI Native" in 6G (IMT-2030) with "AI overlay/add-on" in 5G and reports that ITU-R WP 5D treats AI as foundational to IMT-2030.

(Search summary; not fetched.) https://techblog.comsoc.org/2026/01/15/comparing-ai-native-mode-in-6g-imt-2030-vs-ai-overlay-add-on-status-in-5g-imt-2020/

- **[EVIDENCE, weak]** NVIDIA and T-Mobile: AI-RAN Innovation Center (2024), later extended to "AI-native 6G" research. At MWC 2026, NVIDIA announced commitments from more than a dozen telecom and technology firms to build 6G on "AI-native" platforms. One trade report says the AI-RAN Alliance has 75+ members but fewer than 10 operators. (Search summary; not fetched.) https://www.telecomtv.com/content/telcos-and-ai-channel/nvidia-eyes-6g-with-ai-native-telco-stack-52624/ ; https://www.crnasia.com/news/2026/networking/nvidia-bets-big-on-ai-native-networks-at-mwc-2026
- **[EVIDENCE, weak]** I found no 3GPP primary document defining "AI-native". Secondary sources say Release 20 includes 6G study items and AI/ML air-interface work (e.g. two-sided models for CSI compression), with the first 6G specifications expected in Release 21. (Search summary; not fetched.) https://www.qualcomm.com/news/onq/2025/06/3gpp-release-20-completing-5g-advanced-evolution-preparing-for-global-6g-standardization
- **What the telecom meaning is.** "Native" means AI is designed into the system from the start (models replace hand-designed signal-processing blocks, with lifecycle management built in), as opposed to AI bolted onto a finished design. It is a statement about **architecture**, not about how people work. The business usage borrowed the word, and probably the "cloud-native" analogy, but not the precise meaning. This analogy point is my inference; I did not find a source documenting the borrowing.

Software, VC and vendor definitions

- **[OPINION]** VC and glossary definitions mostly use a dependency test: would the product or business work without its models? An "AI-native company's core product, cost structure, or competitive moat depends fundamentally on AI capability". (Search summary; not fetched.) https://vcbeast.com/venture-capital-glossary/ai-native-company
- **[OPINION]** Intel Capital, investor essay (date not confirmed). It treats AI-native as a subset of AI companies and extends the idea beyond software to manufacturing, services and healthcare. (Search summary; not fetched.) https://www.intelcapital.com/an-investor-search-for-an-ai-native-company/
- **[OPINION]** TechCrunch, 2 May 2024, "Your AI-native startup ain't the same as a typical SaaS company". This was the earliest dated mainstream-tech use I saw in search results. That is not evidence of first use. (Search summary; not fetched.) https://techcrunch.com/2024/05/02/your-ai-native-startup-aint-the-same-as-a-typical-saas-company/amp
- **[OPINION]** CRV (VC firm) defines AI-native by architecture: the product was built on AI from the ground up, and without the intelligence layer the company would not exist in any meaningful form (paraphrase). Undated. https://www.crv.com/content/what-is-ai-native (search extract)
- **[EVIDENCE, weak]** Dating the VC usage (gap-fill pass): the earliest dated VC-firm use found was Nava Ventures, which launched in 2023 with a stated focus on "AI-native enterprise companies" (per a VC directory). A targeted search for 2017 to 2018 business uses returned nothing. Who first used "AI-native" in business remains unresolved; the telecom use (Nokia Bell Labs, December 2020) is still the earliest documented use I found. https://vcbeast.com/vc-firms/nava-ventures (search extract)
- **[OPINION]** I could not reach a16z, Sequoia or Bessemer primary pages. Secondary summaries attribute an "AI-native vs retrofit" framing to a16z, but I could not check the wording. Treat it as unverified.

- **[OPINION]** An industry explainer separates "AI-native" (built around AI from the ground up) from "AI-first" (an incumbent that tears down and rebuilds around AI). Another notes the term "remains ill-defined and is often used as marketing". (Search summary; not fetched.) https://letsdatascience.com/news/ai-community-wrestles-with-defining-ai-native-3ce3828c
- **[OPINION]** Foundation Capital (Jaya Gupta, Ashu Garg), Dec 2025, "AI's trillion-dollar opportunity: Context graphs". Argument: value shifts from systems of record (Salesforce, Workday, SAP) to "systems of agents". A "context graph" that records decision traces (why decisions were made, not only what happened) becomes the new system of record. **This is the closest published framework to the user's "one stored shared context, read and written by humans and agents".** It is a VC thesis, not a standard, and I read it only through secondary summaries. (Not fetched; the summary was reached via a mirror.) https://foundationcapital.com/context-graphs-ais-trillion-dollar-opportunity/

Maturity models compared

Model	Unit measured	Stages	Basis	What decides your stage
Microsoft WTI 2025 "Frontier Firm"	Organisation (via leader self-report)	3 phases: AI as assistant; agents as digital colleagues; humans set direction, agents run processes	Survey n=31,000 in 31 countries, Feb-Mar 2025; 844 leaders classed Frontier	Leader says: org-wide AI deployment, "a leader" in AI investment, measuring ROI, some ROI seen, using agents
Microsoft WTI 2026	Individual AI user vs organisational readiness	5 groups: Frontier 19%, emergent (about half), blocked agency 10%, unclaimed capacity 5%, stalled 16%	Survey n=20,000 AI users in 10 markets, Feb-Apr 2026, plus M365 telemetry	Self-reported composites (usage sophistication, governance, manager support, AI culture)
MIT CISR Enterprise AI Maturity	Enterprise	4 stages, from "Experiment and Prepare" up to "AI future-ready"	Survey n=721 (2022) plus 9 interviews (2024)	Capability build-out; financial performance below industry average in stages 1-2, above it in stages 3-4
Ericsson AI native maturity matrix	Telecom network artifact	Levels (count not verified)	Engineering framework	How deeply AI is built into the network function

- **[EVIDENCE, weak: vendor survey, self-report]** Microsoft WTI 2025. Frontier Firms are "Structured around on-demand intelligence and powered by 'hybrid' teams of humans + agents". 24% of leaders report org-wide AI deployment; 12% are still in pilots. Frontier Firm leaders were more likely to say their company is thriving (71%, against 39% of workers globally). Run by Edelman Data x Intelligence, 6 Feb to 24 Mar 2025, 1,000 per market. Microsoft, 23 Apr 2025. https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born
  - Weakness: Frontier status is defined by leaders saying their firm is "a leader" in AI investment and has seen ROI, so the label partly measures confidence. The "thriving" comparison sets Frontier leaders against workers globally, which is not like-for-like.
- **[EVIDENCE, weak: vendor survey, self-report]** Microsoft WTI 2026, "Agents, human agency, and the opportunity for every organization", 5 May 2026. Survey n=20,000 AI users (people who never use AI were screened out) in 10 markets, run 18 Feb to 20 Apr 2026. Organisational factors (culture, manager support, talent practices) account for 67% of modelled "AI impact" against 32% for individual factors. The model is a random-forest permutation importance; Microsoft notes it is associational, not causal. Active agents in Microsoft 365 grew 15x year on year. https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization

- **[EVIDENCE, weak: vendor survey]** WTI 2026 also reports how rarely agent work is documented, even among the most advanced users. Agent workflows, human handoffs and quality standards are "documented and repeatable" at organisation level for 25% of Frontier Professionals against 14% of others (team level: 26% against 19%). This bears directly on the user's "nobody can describe the real workflow" point. Same URL.
- **[EVIDENCE, moderate: academic survey, pre-genAI]** MIT CISR (Weill, Woerner, Sebastian), Dec 2024: four stages. About 28% of firms were in stage 1 and 31% in stage 3. Firms in stages 1-2 had below-industry-average financial performance; stages 3-4 were above it. The survey dates from 2022, so it mostly reflects analytical AI, not generative AI or agents. The relationship is cross-sectional (correlational). (Search summary; cisr.mit.edu blocked.) https://cisr.mit.edu/sites/default/files/2024-12/EnterpriseAIMaturityModel_CISRpressrelease_121924_0.pdf

Where the models disagree:

1. **Unit of analysis.** Ericsson rates a network artifact. CISR rates the enterprise. WTI 2025 rates the firm through its leaders. WTI 2026 rates the individual inside an organisational context.
2. **What counts as progress.** Telecom: how deeply AI is built into the architecture. CISR: capabilities plus financial outcomes. Microsoft: deployment breadth, agent use and self-reported ROI.
3. **Where the bottleneck sits.** WTI 2026 puts it in the organisation (incentives, managers). The Anthropic Economic Index puts it in information access ("context"). The Klarna case points to output quality.
4. **None of them** uses "shared stored context that humans and agents both read and write" as the defining test. Microsoft 2026 comes nearest, with "Owned Intelligence" and "shared routines".

Primary evidence that context is the constraint (closest support for the user's definition)

- **[EVIDENCE, moderate: vendor data, large n, inferential]** Anthropic Economic Index, 15 Sep 2025. Data: 1M first-party API transcripts from Aug 2025, which Anthropic says is about half of first-party API use. 77% of API transcripts show automation patterns, against 12% augmentation. Longer inputs go with complex tasks (input-output elasticity 0.38). From this, the authors infer that dispersed, non-centralised context is a barrier to enterprise deployment. They also note that software works partly because code repositories already centralise context. https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
    - Why it matters for the user: **Software development is the one domain that already has a single shared store humans and agents both read and write (the code repository).** That is also where agent adoption is furthest along. This supports the user's definition as a description of preconditions. It is Anthropic's inference, not a controlled test.

Named companies that rebuilt around AI

Klarna

- **[EVIDENCE, weak: company claim]** Feb 2024: Klarna said its AI assistant did the work of 700 full-time customer service agents and handled 75% of chats. (Search summary; not fetched.) https://www.emarketer.com/content/klarna-backtracks-ai-customer-service-plans

- **[EVIDENCE, weak: press reports of company figures]** Headcount fell from about 5,000 to 3,800 in 2024, mostly by attrition. A spokesperson said it would likely fall toward about 2,000. Fortune, 28 Aug 2024. (Search summary; not fetched.) https://fortune.com/europe/2024/08/28/klarna-1800-employees-ai-replace-ipo
- **[EVIDENCE, weak: company financials via press]** Revenue per employee is approaching $1M, up from $575k a year earlier. TechCrunch, 19 May 2025. (Search summary; not fetched.) https://techcrunch.com/2025/05/19/klarnas-revenue-per-employee-soars-to-nearly-1m-thanks-to-ai-efficiency-push/
- **[EVIDENCE: CEO statement via press]** What broke: in May 2025, CEO Sebastian Siemiatkowski told Bloomberg that cost had been "a too predominant evaluation factor" and that the result was "lower quality". Klarna began hiring humans so customers can always reach a person. Klarna says this is freelance hiring for its outsourcing arm and that headcount still falls through attrition, so the scope of the "reversal" is disputed. (Search summary; not fetched.) https://www.entrepreneur.com/business-news/klarna-ceo-reverses-course-by-hiring-more-humans-not-ai/491396 ; https://www.computing.co.uk/news/2025/ai/meat-back-on-the-menu-as-klarna-rows-back-on-ai
- **[EVIDENCE, weak: company claim via press]** The February 2024 claim in detail: the assistant handled 2.3 million chats in its first month, two-thirds of customer service chats, the equivalent work of 700 full-time agents, and cut average resolution time from 11 minutes to 2. Klarna had frozen non-engineering hiring in December 2023. https://www.euronews.com/2024/02/28/klarnas-ai-bot-is-doing-the-work-of-700-employees-what-will-happen-to-their-jobs ; https://www.itpro.com/technology/artificial-intelligence/klarna-has-stopped-hiring-amid-its-relentless-ai-efficiency-drive (search extract)
- **[EVIDENCE: filing, via secondary reports]** Klarna's F-1 (filed March 2025) reportedly gives full-time employees of 5,527 at end-2022 and 3,422 at end-2024 (down about 38%), and attributes the fall to increased use of AI and natural attrition. I saw the filing URL but not the relevant page or the risk-factor wording. The CEO separately told CNBC that AI was only one factor. F-1: https://www.sec.gov/Archives/edgar/data/2003292/000162828025012824/klarnagroupplcf-1.htm ; report of the figures: https://getcoai.com/news/swedens-klarna-attributes-reduced-headcount-to-ai-chatbot-efficiencies (search extract)
- **Scope of the reversal (gap-fill pass).** The May 2025 move was to add human agents (remote and freelance) so customers can always reach a person, with AI still handling routine contacts. A later CMSWire headline reports the CEO saying Klarna is giving AI more customer-service work, not less; I saw only the headline. So the reversal is a hybrid model with a human fallback, not an abandonment of AI. https://www.cmswire.com/contact-center/klarna-ceo-were-giving-ai-more-customer-service-work-not-less/ (search extract)

Duolingo

- **[EVIDENCE: primary memo via press]** On 28 Apr 2025, CEO Luis von Ahn posted the "AI-first" memo on LinkedIn. It said Duolingo would "gradually stop using contractors to do work that AI can handle". After backlash, a late-May follow-up said hiring would continue and AI would not replace employees. In an Aug 2025 NYT interview he said "This was on me. I did not give enough context" and that Duolingo has never laid off full-time employees. (Search summary; not fetched.) https://fortune.com/2025/08/18/duolingo-ceo-admits-controversial-ai-memo-did-not-give-enough-context-insists-company-never-laid-off-full-time-employees ; https://www.entrepreneur.com/business-news/duolingo-ceo-clarifies-ai-stance-after-backlash-read-memo/492141
- What changed: contractor content work moved to AI. What broke: brand and user trust (complaints about content quality), not operations.

- **[EVIDENCE, weak: sell-side analyst and CEO statements via press]** Measured cost of the backlash (gap-fill pass): Morgan Stanley linked the memo to about a 10% decline in US daily active users since April 2025, with international trends steady. Von Ahn said it may have hurt DAU in the quarter to June and that the effect was almost entirely in the United States and Canada (paraphrase). Total DAU still grew: about 47.7 million in Q2 2025 (+40% YoY), and a reported 56.5 million in Q1 2026 (+21% YoY), so growth slowed but did not reverse. The company toned down its "edgy" social posting rather than pausing it. https://www.itiger.com/news/2549526150 ; https://www.aol.com/duolingos-ceo-says-learned-hard-062205679.html (search extract). Slowing growth has other possible causes; the attribution to the memo is the analyst's and the CEO's.

Shopify

- **[EVIDENCE: primary memo via press]** Tobi Lütke's memo was written in March 2025 and posted on X on 7 Apr 2025. Two rules: AI use is a "baseline expectation", and AI-usage questions go into performance reviews. Before asking for headcount, teams must demonstrate why they cannot get what they want done using AI. (Search summary; not fetched. Not quoted because I could not see the raw memo.) https://betakit.com/shopify-ceo-tobi-lutke-tells-employees-to-prove-ai-cant-do-the-job-before-asking-for-resources/
- Note: this is a management policy, not a rebuilt architecture. It does not mention shared context.

Salesforce

- **[EVIDENCE: CEO statement via press]** On The Logan Bartlett Show (aired around 29 Aug 2025), Marc Benioff said support headcount went from about 9,000 to about 5,000 because of Agentforce. He called it a "rebalance". A spokesperson said fewer support cases removed the need to backfill and that staff were redeployed to sales and customer success. No timeframe was given. (Search summary; not fetched.) https://fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales/ ; https://www.salesforceben.com/ai-agents-drive-4000-job-cuts-in-salesforce-support-division/

IBM

- **[EVIDENCE: CEO statement via press]** In May 2025, Arvind Krishna told the WSJ that AI (the AskHR agent) replaced "a couple of hundred" HR roles. He said total headcount rose ("we've added more roles than we've removed"). Reports say AskHR automates 94% of routine HR tasks, a share of tasks, not of staff. (Search summary; not fetched.) https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341

Moderna

- **[EVIDENCE: proxy filing via search]** Tracey Franklin moved from Chief HR Officer to Chief People and Digital Technology Officer in November 2024, merging HR and IT. A WSJ report (date unconfirmed) says she is redesigning teams around which work people do and which is automated, partly using the OpenAI partnership. The same report says Moderna cut about 10% of digital technology jobs (about 50 people) and that the former CIO left. (Search summary; not fetched.) https://www.sec.gov/Archives/edgar/data/1682852/000130817926000081/mrna-courtesy-pdf.pdf ; https://investors.modernatx.com/tracey-franklin
- Why it matters: this is the clearest case of org design changing, not just headcount. Merging people and technology functions treats agents and humans as one workforce.

Intercom / Fin

- **[EVIDENCE, weak: third-party pricing guides; ARR unverified]** Fin is priced at $0.99 per resolution (only when it resolves without a human handoff), on top of seats. Secondary sources put Fin ARR near or above $100M by 2026 and Intercom total ARR near $400M. I found no primary Intercom filing or statement. (Search summary; not fetched.) https://www.eesel.ai/blog/fin-ai-pricing
- What changed: the pricing model moved from per-seat to per-outcome. That is a real business-model rebuild, and the user could use it as an example.

Anthropic (a model vendor studying itself)

- **[EVIDENCE, moderate: self-study, n=132 survey, 53 interviews, 200k Claude Code transcripts]** Published 2 Dec 2025. Engineers self-report using Claude in 59% of their work (28% a year earlier) and a +50% productivity gain (+20% a year earlier). Merged pull requests per engineer per day rose 67%. 27% of Claude-assisted work "wouldn't have been done otherwise". More than half can "fully delegate" only 0-20%. What broke: fewer questions to senior staff (less mentorship), less collaboration, worries about skill atrophy. https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic

"AI-native" as marketing

- **[OPINION]** Commentators say the term is "often used as marketing rather than to describe a coherent architectural shift". (Search summary; not fetched.) https://letsdatascience.com/news/ai-community-wrestles-with-defining-ai-native-3ce3828c
- **[OPINION]** Microsoft WTI 2025 itself says it "interviewed 12 AI-native startups" without defining the term. The label is used as self-evident. https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born
- **[EVIDENCE, weak]** The ComSoc and trade-press telecom pieces show "AI-native" moving from a defined engineering term (Nokia, Ericsson) to a vendor coalition banner (NVIDIA MWC 2026) within about five years.

"AI washing" enforcement and "agent washing" (gap-fill pass)

Regulators have now defined, case by case, what a false AI claim looks like. In every case so far the gap was between "our AI does X" and humans or third parties doing X. That is the enforceable edge of "AI-native as marketing".

- **[EVIDENCE: regulatory action]** SEC, 18 March 2024 (press release 2024-36): settled charges against two investment advisers, Delphia (USA) and Global Predictions, for false and misleading statements about their use of AI. Penalties $225,000 and $175,000 ($400,000 total), without admitting or denying. The SEC found Delphia claimed AI and machine-learning capabilities from 2019 to 2023 that it did not have; Global Predictions described itself as the first regulated AI financial adviser (paraphrase of reported wording). https://www.sec.gov/newsroom/press-releases/2024-36 (search extract)
- **[EVIDENCE: regulatory action]** SEC, 14 January 2025: settled charges against Presto Automation, reported as the first AI-washing action against a public company. The SEC said Presto's filings presented its Presto Voice drive-through ordering product as in-house technology that removed human order-taking, while the deployed technology was owned and operated by a third party, and the large majority of orders (reported as over 70%) needed human intervention. Penalty amount not seen. https://www.dandodiary.com/2025/01/articles/artificial-intelligence/sec-files-ai-washing-enforcement-action-against-restaurant-technology-

company/ ; https://ropesgray.com/en/insights/alerts/2025/01/ai-washing-enforcement-in-the-final-days-of-the-gensler-administration (search extract; the SEC primary URL did not surface)

- **[EVIDENCE: criminal indictment and SEC complaint; allegations, not findings]** DOJ (SDNY) and SEC, 9 April 2025: Albert Saniger, former CEO of the shopping app Nate, charged with securities and wire fraud. Prosecutors allege he raised over $40 million from investors by claiming purchases were automated by AI, while hundreds of contract workers in a Philippine call centre completed them by hand; the DOJ says the real automation rate was effectively 0%. Commentators noted these were the first AI-washing actions under the new administration, so enforcement continued after the change of SEC leadership. https://www.justice.gov/usao-sdny/pr/tech-ceo-charged-artificial-intelligence-investment-fraud-scheme ; https://www.debevoise.com/insights/publications/2025/04/the-sec-and-doj-signal-continued-focus-on-ai-wash (search extract)
- **[EVIDENCE: regulatory action]** FTC "Operation AI Comply", 25 September 2024: actions against five companies (DoNotPay, Ascend Ecom, Ecommerce Empire Builders, Rytr, FBA Machine). DoNotPay, the self-described "robot lawyer", had not tested whether its outputs matched a lawyer's; it agreed to pay $193,000 (as reported). Rytr was barred from selling AI-generated consumer-review services. Three were income schemes using "AI-powered" claims. https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes (search extract)
- **[OPINION: analyst estimate]** Gartner, 25 June 2025: "agent washing" (rebranding assistants, RPA and chatbots as agentic) is widespread; only about 130 of the thousands of vendors claiming agentic AI are real. No method published. Full treatment in 04-enterprise.md. https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027
- Not found: SEC AI-washing actions after April 2025. One 2026 legal blog discusses AI-washing litigation in 2026 but I did not see specific new cases. https://blog.promise.legal/ai-washing-litigation-2026-public-company-gcs/ (search extract)
- Why it matters for the user: no enforcement action has tested the term "AI-native" itself. The actions target specific false capability claims (automation rates, in-house technology, AI that did not exist). A firm calling itself "AI-native" is legally exposed only if it states facts about its AI that are untrue. The Presto and Nate cases are also a warning about the "human behind the agent" pattern that agent platforms with human escalation must disclose accurately.

Viral statistics traced (where relevant)

- **Klarna "AI does the work of 700 agents".** The source is Klarna's own Feb 2024 announcement. It is an unaudited company claim based on chat volume. The CEO's May 2025 statements about lower quality complicate it. Treat it as a company claim, not a measured result. https://www.emarketer.com/content/klarna-backtracks-ai-customer-service-plans. The F-1 figures (5,527 FTEs end-2022 to 3,422 end-2024, reportedly attributed to AI and attrition) are the more solid number, but they cover the whole company and do not isolate AI. https://www.sec.gov/Archives/edgar/data/2003292/000162828025012824/klarnagroupplcf-1.htm
- **"Klarna reversed its AI strategy".** Overstated. It added a human fallback for customers and kept AI on routine contacts. https://www.cmswire.com/contact-center/klarna-ceo-were-giving-ai-more-customer-service-work-not-less/ (headline only)
- **"Duolingo lost users over its AI-first memo".** Partly true: Morgan Stanley linked it to about a 10% fall in US daily users, but total DAU still grew about 40% YoY in Q2 2025. https://www.itiger.com/news/2549526150

- **"IBM fired 8,000 people and replaced them with AI".** This blends separate events. The CEO said "a couple of hundred" HR roles were replaced, and total headcount rose. The 8,000 figure traces to layoff coverage and viral posts, not to IBM. The "94%" refers to routine HR tasks automated, not jobs. https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341
- **"Salesforce cut 4,000 jobs because of AI".** The source is a Benioff podcast remark. He described it as a "rebalance" with no backfill and redeployment, over an unstated period. "Layoffs" is the press framing. https://fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales/
- **Shopify "prove AI can't do the job before hiring".** This paraphrase is sharper than the memo. Per coverage, the memo asks teams to demonstrate why they cannot get what they want done using AI. https://betakit.com/shopify-ceo-tobi-lutke-tells-employees-to-prove-ai-cant-do-the-job-before-asking-for-resources/
- **Duolingo "replacing staff with AI".** The memo covered contractors, and the CEO says no full-time staff were laid off. https://fortune.com/2025/08/18/duolingo-ceo-admits-controversial-ai-memo-did-not-give-enough-context-insists-company-never-laid-off-full-time-employees
- **Microsoft "Frontier Firm leaders: 71% thriving".** The comparison group is workers globally, not non-Frontier leaders. The Frontier group is 844 self-classified leaders. https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born

Strongest counter-evidence (against the thesis and against the mainstream reading of this track)

1. **"The tech is the easy part" is contradicted by the most AI-ready organisation available to study.** Anthropic has a centralised codebase, staff who are fully bought in, and unlimited model access. Even so, most engineers can fully delegate only 0-20% of their work, and supervision stays essential. The limit there is model reliability, not readiness. https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic
2. **Klarna's partial reversal was about output quality, not resistance or missing context.** The CEO blamed over-weighting cost, not staff readiness. https://www.entrepreneur.com/business-news/klarna-ceo-reverses-course-by-hiring-more-humans-not-ai/491396
3. **"People resist" gets weak support in the largest 2026 survey.** Microsoft finds more skilled workers blocked by their organisation (10%) than organisations held back by their people (5%). 45% say it feels safer to stick to current goals than to redesign work, and only 13% are rewarded for reinvention. The resistance is structural (incentives, metrics), not personal. Caveat: vendor survey, self-report, AI users only. https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization
4. **Against the user's definition: the firms most often called AI-first changed policy and headcount, not shared context.** If "AI-native" requires one shared read/write context, then Klarna, Shopify, Duolingo and Salesforce, as reported, would not qualify. The user should either accept that or say that the common usage is something else.
5. **Against the mainstream "AI-native = built from scratch" reading: the evidence on startups is thin.** I could not verify any revenue-per-employee figure for Cursor, Midjourney, Gamma or Lovable this session. Many of those figures circulate from founder tweets and press, not filings.
6. **For the thesis (to stay honest): two independent sources point to the organisation as the bottleneck.** Anthropic (context access) and Microsoft 2026 (incentives and

managers) both say so, with different data. The "shared context" part of the thesis has its best support in the Anthropic data.

Testing the user's working definitions

User's claim	What I found	Verdict
AI-native = humans and agents work from one stored, shared business context	No accessible framework defines it this way. Telecom: AI in the architecture. VC: AI dependency or origin. Consultancies and vendors: deployment breadth plus ROI. Nearest matches: Foundation Capital "context graphs" (VC opinion), Microsoft 2026 "Owned Intelligence" and "shared routines" (vendor), Anthropic's "centralize information" (vendor research).	**Idiosyncratic but defensible.** It is a precondition definition (what must be true for agents to do non-trivial work), not the market definition. Say so explicitly.
One platform both read from and write to	Software repos meet this test and are where agent adoption leads (Anthropic AEI). No maturity model requires a single platform. Microsoft 2026 instead stresses identity, permissions and governance across agents.	**Partly supported** by the software analogy. "One platform" sounds vendor-shaped (it matches Actualisation's own product), and readers may discount it for that reason. "Shared context" stands up better without "one platform".
Legacy = humans only, no AI in the workflow	Survey data suggests this group is close to empty among knowledge workers (WTI 2025 surveyed 31 countries; the 2026 edition sampled AI users only). Every maturity model starts at "some AI".	**Too strict to be useful.** A more useful "legacy" marker might be "AI use is individual and unrecorded: nothing agents or people learn goes back into a shared store".
Spectrum; fully AI-native is rare	Matches Microsoft (3 phases, 5 groups, 19% Frontier users), CISR (4 stages) and Ericsson (maturity matrix).	**Matches the literature.** But each model puts firms on the spectrum by a different axis.

Open questions

These are unverified leads; check each before using it.

- **SEC, DOJ and FTC "AI washing" actions:** now sourced (see the enforcement section). Still open: the Presto penalty and SEC order URL; the outcome of the Saniger case; any SEC or FTC AI-washing actions after April 2025; whether the FTC's later actions (2025-26) continued Operation AI Comply.
- **Gartner "agent washing":** sourced (press release URL above), but the method behind "about 130" is unpublished.
- **Klarna F-1 risk factors:** the headcount figures are reported second-hand; the wording on AI and customer-service quality risk still needs a direct read of the filing.
- **Maturity models not reached:** Gartner AI Maturity Model (5 levels), McKinsey State of AI 2025 ("high performers"), BCG "The Widening AI Value Gap" (Sept 2025, "future-built" share), Deloitte State of GenAI in the Enterprise, Accenture reinvention reports. Each needs the original methodology.
- **Academic definitions:** Iansiti & Lakhani, Competing in the Age of AI (2020), "AI factory" operating model with an integrated data platform. This is possibly the strongest academic match to the "shared context" definition. Also Hassan et al. on "AI-native software engineering" (arXiv). Not verified.
- **Earliest business use of "AI-native"** and its relation to Google's "AI-first" (2016) and CNCF's "cloud-native". Not verified.
- **Small-headcount, high-revenue startups** (Cursor/Anysphere, Midjourney, Gamma, Lovable): ARR and headcount from primary statements.
- **Earliest business use of "AI-native":** a gap-fill search for 2017-2018 uses found nothing; the earliest dated VC-firm use found is 2023 (Nava Ventures). Try Google Books Ngram or news archives, which this environment could not reach.

What's genuinely uncertain

- Whether shared context causes better agent outcomes or only travels with them. The Anthropic evidence is an inference from input lengths; I found no controlled comparison.
- Whether Microsoft's "organisational factors explain 2x" survives outside self-report. Both the predictors and the outcome are survey perceptions from the same respondents (common-method bias).
- Whether CISR's stage-to-performance link holds for generative and agentic AI. Its survey is from 2022.
- How much of the Klarna, Salesforce and IBM headcount change is due to AI rather than attrition, post-2021 over-hiring and cost cutting. The companies have reasons to credit AI.
- Who first used "AI-native" in business. I could not settle it.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
77% vs 12%	Share of Anthropic first-party API transcripts showing automation vs augmentation patterns	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	2025-09-15	Vendor data, 1M transcripts (Aug 2025), classifier-based. Fetched + grep-verified
0.38	Elasticity of output to input token length (basis for the "context bottleneck" inference)	Anthropic Economic Index	same	2025-09-15	Inferential; author interpretation. Fetched
59% (from 28%)	Share of their work Anthropic engineers use Claude for, self-reported	Anthropic	https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic	2025-12-02	Self-report, n=132. Fetched + grep-verified
0-20%	Share of work most Anthropic engineers say they can "fully delegate"	Anthropic	same	2025-12-02	Self-report. Fetched + grep-verified
27%	Claude-assisted work that "wouldn't have been done otherwise"	Anthropic	same	2025-12-02	Self-estimate. Fetched + grep-verified
844 of 31,000	Leaders classed as Frontier Firm, of total respondents	Microsoft WTI 2025	https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born	2025-04-23	Vendor survey (Edelman); self-classification. Fetched
24% / 12%	Leaders reporting org-wide AI deployment / still in pilot	Microsoft WTI 2025	same	2025-04-23	Vendor survey. Fetched + grep-verified
19%	AI users in the "Frontier" zone (high individual and organisational readiness)	Microsoft WTI 2026	https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization	2026-05-05	n=20,000 AI users, 10 markets, self-report. Fetched
67% vs 32%	Share of modelled "AI impact" from organisational vs individual factors	Microsoft WTI 2026	same	2026-05-05	Random-forest importance on self-report; associational. Fetched
13%	AI users rewarded for reinventing work with AI even if results aren't met	Microsoft WTI 2026	same	2026-05-05	Self-report. Fetched
15x	Year-on-year growth in active agents in Microsoft 365	Microsoft WTI 2026	same	2026-05-05	Vendor telemetry, no absolute counts. Fetched
25% vs 14%	Frontier vs other professionals saying agent workflows/handoffs are documented at org level	Microsoft WTI 2026	same	2026-05-05	Self-report. Fetched

Figure	What it measures	Source	URL	Date	Reliability note
700	FTE customer service agents Klarna said its AI assistant replaced	Klarna via eMarketer	https://www.emarketer.com/content/klarna-backtracks-ai-customer-service-plans	2024-02 (claim)	Company claim; search summary only
~9,000 to ~5,000	Salesforce support headcount, per Benioff	Fortune	https://fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales/	2025-09-02	CEO podcast remark; no timeframe; search summary only
"a couple of hundred"	IBM HR roles replaced by AI	Entrepreneur citing WSJ	https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341	2025-05	CEO statement; search summary only
$400,000 ($225k + $175k)	SEC civil penalties, Delphia and Global Predictions AI-washing settlements	SEC press release 2024-36	https://www.sec.gov/newsroom/press-releases/2024-36	2024-03-18	Regulatory action; search extract
>70%	Share of Presto Voice orders needing human intervention, per SEC case	SEC settlement via law-firm summaries	https://www.dandodiary.com/2025/01/articles/artificial-intelligence/sec-files-ai-washing-enforcement-action-against-restaurant-technology-company/	2025-01-14	Regulatory finding as reported; SEC primary not seen
$40M+ / ~0%	Raised by Nate on AI claims / alleged real automation rate	DOJ SDNY press release	https://www.justice.gov/usao-sdny/pr/tech-ceo-charged-artificial-intelligence-investment-fraud-scheme	2025-04-09	Allegations in an indictment, not findings
5	Companies in FTC Operation AI Comply	FTC press release	https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes	2024-09-25	Regulatory action; search extract
~130	Vendors with genuine agentic capability, of thousands claiming it	Gartner press release	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027	2025-06-25	Analyst estimate, no method
5,527 to 3,422	Klarna full-time employees, end-2022 to end-2024	Klarna F-1, via press	https://www.sec.gov/Archives/edgar/data/2003292/000162828025012824/klarnagroupplcf-1.htm	2025-03	Filing figure seen only in secondary reports
2.3M / two-thirds	Chats handled by Klarna's assistant in first month / share of all service chats	Klarna via Euronews	https://www.euronews.com/2024/02/28/klarnas-ai-bot-is-doing-the-work-of-700-employees-what-will-happen-to-their-jobs	2024-02-28	Company claim
~10%	Decline in Duolingo US daily active users linked to the AI-first memo	Morgan Stanley via iTiger	https://www.itiger.com/news/2549526150	2025	Sell-side analyst attribution
721	Firms in MIT CISR AI maturity survey (2022)	MIT CISR	https://cisr.mit.edu/sites/default/files/2024-12/EnterpriseAIMaturityModel_CISRpressrelease_121924_0.pdf	2024-12-19	Pre-genAI; search summary only

Verbatim quotes

- "The upshot is that deploying AI for complex tasks might be constrained more by access to information than on underlying model capabilities." (Anthropic Economic Index report, Anthropic, 15 Sep 2025) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report. Context: the authors' reading of input-length data from 1M API transcripts. Verified: curl+grep
- "This could represent a barrier to broader enterprise deployment for some important tasks that rely on dispersed context that is not already centralized or digitized. Correcting for this bottleneck may require firms to restructure their organization, invest in new data infrastructure, and centralize information for effective model deployment." (Anthropic

Economic Index, 15 Sep 2025) same URL. Context: the closest primary statement to the user's shared-context definition. Verified: curl+grep

- "For software development with centralized code repositories, access to this information is in principle straightforward." (Anthropic Economic Index, 15 Sep 2025) same URL. Context: why coding leads agent adoption. Verified: curl+grep
- "I work way more with Claude than with any of my colleagues." (Anonymous Anthropic engineer, "How AI is transforming work at Anthropic", 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: reduced collaboration. Verified: curl+grep
- "More junior people don't come to me with questions as often." (Anonymous senior Anthropic engineer, same source, 2 Dec 2025) same URL. Context: mentorship loss. Verified: curl+grep
- "Structured around on-demand intelligence and powered by "hybrid" teams of humans + agents, these companies scale rapidly, operate with agility, and generate value faster." (Microsoft, 2025 Work Trend Index, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: the Frontier Firm definition. Verified: curl+grep (the page uses curly quotes around hybrid)
- "In many cases, people are ready. The systems around them are not." (Microsoft, 2026 Work Trend Index, 5 May 2026) https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization. Context: framing for the "Transformation Paradox". Verified: curl+grep
- "They capture these signals and encode them into shared routines, improving future work while preserving accountability and control." (Microsoft, 2026 Work Trend Index, 5 May 2026) same URL. Context: what Frontier Firms do with agent outputs; the nearest Microsoft gets to "shared context". Verified: curl+grep
- "Organizations that can answer these questions are building Owned Intelligence" (Microsoft, 2026 Work Trend Index, 5 May 2026) same URL. Context: the page then defines Owned Intelligence as institutional know-how unique to the firm. Verified: curl+grep

Sources worth reading in full (ranked, one line each on why)

1. Anthropic Economic Index, Sep 2025: the best large-sample evidence that context access, not model ability, limits enterprise agent deployment. https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
2. Microsoft Work Trend Index 2026: the newest maturity framework; its "people ready, systems not" finding challenges the "people resist" thesis; includes data on how often agent workflows are documented. https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization
3. Anthropic, "How AI is transforming work at Anthropic": a rare inside view of an AI-native organisation, including its costs. https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic
4. Ericsson, "Defining AI native": the only maturity model built around the word "native", and it admits the term is undefined. https://www.ericsson.com/en/reports-and-papers/white-papers/ai-native
5. Foundation Capital, "Context graphs": the VC thesis closest to the user's definition, to engage with or distinguish from. https://foundationcapital.com/context-graphs-ais-trillion-dollar-opportunity/

6. MIT CISR Enterprise AI Maturity Model: the only maturity model I found that links stages to financial performance; pre-genAI. https://cisr.mit.edu/publication/2024_1201_EnterpriseAIMaturityModel_WeillWoernerSebastian
7. Microsoft Work Trend Index 2025: origin of the "Frontier Firm" and "human-agent ratio" language that now circulates. https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born
8. Nokia Bell Labs, "Toward a 6G AI-Native Air Interface": the earliest formal use of the term I found. https://arxiv.org/pdf/2012.08285
9. DOJ SDNY release on Nate/Saniger: the starkest documented case of "AI" that was humans, useful for any piece on AI-washing. https://www.justice.gov/usao-sdny/pr/tech-ceo-charged-artificial-intelligence-investment-fraud-scheme
10. SEC press release 2024-36 (Delphia and Global Predictions): the first SEC AI-washing cases and the template for later ones. https://www.sec.gov/newsroom/press-releases/2024-36
11. FTC Operation AI Comply release: shows consumer-protection law applied to AI capability claims. https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes

INCOMPLETE. Covered: official vs vendor adoption figures (US, UK, EU), survey-method comparison, depth of use, barriers from official surveys, micro-firm adoption. Added in the gap-fill pass (section 4b and section 9, all from search extracts): the OECD 2025 SME survey; the widening Eurostat size gap; SMB AI spending data (lender, card and vendor data, plus agency price lists as weak evidence); MSP surveys on who implements; the Kenya RCT. Section 9b states plainly that no credible SMB abandonment rate exists. Still missing: an official or independent measure of SMB implementation cost; any SMB abandonment rate with a stated method; an independent, numbers-backed case study of a 10-person firm with a shared human-agent context (only vendor cases found); the JPMorganChase Institute spending figures (unlocated); Intuit QuickBooks, Constant Contact and Microsoft/LinkedIn SMB surveys; the DSIT 2026 micro-firm figures. Reason: in this session WebFetch and curl were blocked for every source host tried (census.gov, federalreserve.gov, ec.europa.eu, ons.gov.uk, gov.uk, oecd.org, salesforce.com, uschamber.com, nfib.com and others: proxy 403 / DNS refusal), and the shared web-search budget (200 calls across all agents) ran out partway through. **Every figure below comes from search-engine result summaries of the named pages. None was read on the page itself.** Treat each one as "re-verify before citing". No verbatim quotes are given, because no page could be fetched.

Track 2: SMB

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Summary (5-8 bullets, the things that matter most)

- **"How many small businesses use AI" depends on who you ask and how you phrase it. Answers run from about 14% to 76%.** Official random-sample surveys put use at about 14% to 35%: US Census BTOS about 18-24% of all firms, UK DSIT 2025 14% of micro firms, ONS 28% of 0-9 employee firms, Eurostat 17% of small firms. Surveys that draw on membership lists land between: NFIB 24% of small employers. So do the Fed's own small-business survey (46% of employer firms) and the US Chamber (58% use generative AI). Vendor and program panels land highest: Thryv 55%, Goldman Sachs 10,000 Small Businesses 68-76%, Salesforce 75% "experimenting or implemented".
- **The US official series moved because the question changed, not because firms changed.** On 17 November 2025 Census replaced "AI used in producing goods or services" with "AI used in any business function". Measured use roughly doubled, from about 10% to about 17% (St Louis Fed figures, confirmed in the verification pass; earlier drafts said 9% to 18%). Any "SMB AI adoption doubled" headline that bridges late 2025 is mostly a wording artefact.
- **Breadth is rising and depth is not.** In the Fed's employer-firm survey (fieldwork Sep-Nov 2025), about half of AI users said they were "experimenting", 44% "partially integrated" and 7% "fully integrated". Across all employer firms that is roughly 3% fully integrated. ONS (July 2026) found UK adoption among firms with 10+ staff roughly tripled since 2023, while its depth proxy (AI technologies per firm) barely moved.

- **The barriers in official data cut against the thesis.** When official surveys ask why firms don't use AI, the top answers are "can't identify a use case", "not applicable to my business" and "prefer not to". Missing documented workflows or shared context barely appear. In the Fed survey, a third of employer firms had no plans to use AI. Of those, over half said it was not applicable and 30% said they would rather not use it.
- **Micro firms are not as far behind as the size gradient suggests.** ONS June 2026: 28% of 0-9 employee firms vs 29% of all firms. A Fed staff note (April 2026) says the smallest firms adopt more than size alone would predict. NFIB still finds a steep gradient: 21% of firms under 10 employees vs 48% over 50. Eurostat's headline series leaves out firms under 10 employees entirely.
- **SMB "adoption" mostly means buying ready-made tools, not building.** Scottish BICS data say most adopting firms bought external or ready-to-use AI software. NFIB's examples of AI are ChatGPT, Copilot, Grammarly and Canva. For most SMBs, the "implementation project" the thesis has in mind does not happen.
- **What SMBs spend, who helps them, and what fails (gap-fill pass, weak sources).** Lender, card and vendor data put average SMB AI spend at roughly £100-£300 (or €160) a month, mostly subscriptions. Agency quotes for bespoke automation ($1,500-$25,000 upfront) sit far above what almost any survey respondent reports spending. MSPs, the channel SMBs lean on, see AI as clients' top need (48%) but few earn meaningful AI revenue (13%), and only about half feel ready to guide SMBs. No credible SMB AI abandonment rate exists. The only RCT (Kenya, n=640) found no average gain: AI advice helped strong owners and hurt weak ones.
- **Reported impact on headcount is close to zero so far.** NFIB: 98% of small employers using AI saw no change in employee numbers. ONS BICS, late December 2025: 4% of AI-using UK firms reported lower headcount because of AI. NBER w34836 (about 6,000 executives in four countries): more than 80% report no effect yet on employment or productivity.

Key findings

1. Official statistics: United States (Census BTOS)

- **[EVIDENCE]** The Census Business Trends and Outlook Survey (BTOS) changed its core AI question with effect from 17 November 2025. The new wording replaced the old from 4 December 2025, per Census's "AI Question Wording Updates" document. The old wording asked whether the firm used AI "in producing goods or services". The new wording asks whether it used AI in any business function in the last two weeks, with examples (machine learning, natural language processing, virtual agents, voice recognition). Census notes a level shift at the change. Census, "AI Question Wording Updates" (PDF, late 2025). https://www.census.gov/hfp/btos/downloads/AI%20Question%20Wording%20Updates.pdf (seen via search summary only)
- **[EVIDENCE]** Before the change, BTOS AI use ran between roughly 3% and 9% of US firms for about two and a half years. In the first post-shutdown release under the new wording, it was about 17-18%. Ramp economist Ara Kharazian reports this and says a Census economist involved with the survey opposed the change because it breaks the 2023-2025 trend. Ramp AI Index, 14 January 2026. https://ramp.com/data/ai-index-january-2026 (search summary only. The Census-economist remark is second-hand via Ramp.)
- **[EVIDENCE]** Under the broader wording, national business AI use was 17.3% as of late 2025 (cited by the Economic Innovation Group and others). A Federal Reserve staff FEDS Note puts adoption at about 18% at year-end 2025. It also flags that the November 2025 broadening

affects comparisons over time. FEDS Notes, "Monitoring AI Adoption in the U.S. Economy", 3 April 2026. https://www.federalreserve.gov/econres/notes/feds-notes/monitoring-ai-adoption-in-the-u-s-economy-20260403.html (search summary only)

- **[EVIDENCE]** In a May 2026 story covering 14 December 2025 to 3 May 2026, Census reported national AI use of about 19.8% for the period ending 3 May 2026. Use rose among firms with 20+ employees over that period but "did not change significantly" among firms with fewer than 20. For firms with 1-4 employees, Census publishes the rate as a ceiling ("less than 20%") rather than a point estimate. Census Bureau, "Large Firms With at Least 20 Employees Biggest AI Users", May 2026. https://census.gov/library/stories/2026/05/ai-use-businesses.html (search summary and a syndicated copy at https://morriscountyedc.org/large-firms-with-at-least-20-employees-biggest-ai-users/)
- **[EVIDENCE, unverified]** Secondary trackers report later BTOS readings of about 23-24% (August to early September 2026). I could not tie these to a Census release. One appears in the iFeelTech tracker https://ifeeltech.com/tools/business-ai-adoption-index. Do not cite until checked against census.gov/hfp/btos.
- **[EVIDENCE]** Census working paper CES-WP-26-25 (Bonney, Breaux, Dinlersoz, April 2026), "The Microstructure of AI Diffusion", uses BTOS to separate firm-level use, use across business functions and worker-task use. It is the right primary source for depth of use inside US firms. Not read this session. https://www.census.gov/library/working-papers/2026/adrm/CES-WP-26-25.html
- **[EVIDENCE, misreport flag]** A selfemployed.com article titled "Census Data Shows Smallest Firms Leading Small-Business AI Adoption" pegs itself to the April 2026 BTOS release. The search summary shows it quoting a 1-4 employee rate "rising from 4.6% to 5.8%". Those are the figures in a December 2024 Census blog on the old question series (https://www.census.gov/newsroom/blogs/research-matters/2024/12/ai-use-small-businesses.html). If the summary is accurate, the article recycles 2024 old-wording figures as 2026 news. https://www.selfemployed.com/news/census-btos-ai-small-business-adoption-2026/ (not verified on page)

2. Official and quasi-official statistics: Federal Reserve surveys of small firms

- **[EVIDENCE]** Federal Reserve Small Business Credit Survey (SBCS), "2026 Report on Employer Firms" (published March 2026, fieldwork September to November 2025, firms with fewer than 500 employees):
    - 46% of employer firms said the business or its employees currently use AI, and another 15% planned to within 12 months.
    - Among users: about half "experimenting", 44% "partially integrated", 7% "fully integrated".
    - 71% of users said AI raised productivity.
    - About one-third had no plans to use AI. Of those, over half said it was not applicable to their business and 30% said they prefer not to use it.

  https://www.fedsmallbusiness.org/2026-report-on-employer-firms and PDF https://www.fedsmallbusiness.org/-/media/project/clevelandfedtenant/fsbsite/reports/2026/2026-report-on-employer-firms/2026-report-on-employer-firms.pdf (search summary only).

Governor Lisa Cook cited the 71% productivity figure on 24 June 2026: https://www.federalreserve.gov/newsevents/speech/cook20260624a.htm

    - Method caveat: the SBCS is a convenience sample recruited through Reserve Banks' partner organisations, then weighted. It is not a random draw from a business register. This is from prior knowledge of SBCS methodology and was not re-checked this session.

- **[EVIDENCE]** The San Francisco Fed presented the SBCS on a different base. Counting employer and non-employer firms together, about 60% were using AI or planning to. The base differs from the 46% above, so the two figures should not be compared. SF Fed, "Small Business and AI: Understanding Adoption and Building Support Systems", May 2026. https://www.frbsf.org/community-development/initiatives/economic-development-small-business/community-economics-project (search summary only)

- **[EVIDENCE]** A Fed staff FEDS Note reports that AI adoption among the smallest firms is stronger than size alone would predict. 3 April 2026. https://www.federalreserve.gov/econres/notes/feds-notes/monitoring-ai-adoption-in-the-u-s-economy-20260403.html (search summary only)

3. Official statistics: United Kingdom

- **[EVIDENCE]** ONS, "Artificial intelligence in UK businesses: 2023 to 2026" (released 20 July 2026, BICS data):

    - Self-reported AI use among firms with 10+ employees rose from about 12% in late 2023 to about 35% by mid-2026.
    - ONS tracks AI technologies per business as a proxy for intensity. Coverage reports this "barely moved". TNW puts the average adopter at 1.6 AI tools, a figure I could not confirm in the ONS tables.

  https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/articles/artificialintelligenceinukbusinesses/2023to2026 ; TechMarketView, "UK AI adoption has tripled since 2023 but barely deepened", 21 July 2026 https://www.techmarketview.com/news/archive/2026/07/21/ukhotviewsextra-uk-ai-adoption-has-tripled-since-2023-but-barely-deepened ; TNW https://thenextweb.com/news/uk-ai-adoption-widening-not-deepening-ons-2026 (all search summary only)

- **[EVIDENCE]** In the June 2026 BICS wave (fieldwork 15-28 June 2026), 29% of UK businesses used at least one AI technology, up 8 points on June 2025. By size: 28% of firms with 0-9 employees, 49% of firms with 250+. The 2 April 2026 bulletin gave 45% for 250+ in late March 2026. Same ONS article as above, plus https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/bulletins/businessinsightsandimpactontheukeconomy/2april2026 (search summary only)

- **[EVIDENCE]** In late December 2025, 4% of UK businesses using AI said their overall headcount had fallen as a result. ONS, "Business insights and impact on the UK economy: 8 January 2026". https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/bulletins/businessinsightsandimpactontheukeconomy/8january2026 (attribution to this bulletin is probable but inferred from a search summary)

- **[EVIDENCE]** The most common barriers UK firms reported in 2023 were difficulty identifying activities or business use cases (39%), cost (21%) and AI expertise and skills (16%). ONS, "Management practices and the adoption of technology and artificial intelligence in UK firms:

2023", 24 March 2025. https://www.ons.gov.uk/economy/economicoutputandproductivity/productivitymeasures/articles/managementpracticesandtheadoptionoftechnologyandartificialintelligenceinukfirms2023/2025-03-24 (search summary only)

- **[EVIDENCE]** Scotland BICS (September 2025): 26.7% of businesses used AI, up from 14.0% in September 2023. Most adopters bought external or ready-to-use AI software, though in-house development was growing. The main reason for using AI was improving operations (53.6%). The main reasons for delay were lack of in-house AI expertise (14.1%) and difficulty identifying use cases (10.7%). Scottish Government, Scottish Economic Bulletin December 2025 https://www.gov.scot/publications/scottish-economic-bulletin-december-2025/pages/6/ ; BICS Scotland wave pages, e.g. https://www.gov.scot/publications/bics-weighted-scotland-estimates-data-to-wave-165/pages/artificial-intelligence/ (search summary only. Scotland only, not UK.)

- **[EVIDENCE]** DSIT AI Adoption Research, wave 1, was run by IFF Research and Technopolis: 3,500 UK businesses by telephone, 12 February to 2 May 2025, plus 100 qualitative interviews. AI use was 16% overall: 14% of micro businesses, 23% of medium and 36% of large. DSIT called adoption modest. https://assets.publishing.service.gov.uk/media/6960f3924343a0da370869ba/AI_Adoption_Research_Report.pdf (search summary only)

- **[EVIDENCE]** DSIT AI Adoption Survey 2026 (published 29 October 2026): 32% of UK businesses used at least one AI technology, against 16% in 2025, with larger firms at 52%. Method: CATI, 7 January to 21 April 2026, 3,503 interviews. I did not obtain the micro-firm figure or the 2026 barrier ranking. https://www.gov.uk/government/publications/ai-adoption-survey-2026/ai-adoption-survey-report-2026-executive-summary ; full report https://assets.publishing.service.gov.uk/media/6ae202189bca21f1c5036e99/artificial-intelligence-adoption-research-report-2026.pdf (search summary only)

- **Note on UK inconsistency.** In mid-2026 the ONS (BICS, about 29% of all firms, 35% of 10+) and DSIT (32%) roughly agree, despite different methods. DSIT doubled from 16% to 32% in one year. ONS shows a slower rise (+8 points June 2025 to June 2026). The two use different questions and bases, and I could not reconcile them without the documents.

4. Official statistics: European Union (Eurostat)

- **[EVIDENCE]** Eurostat, "Use of artificial intelligence in enterprises" (ICT usage survey; data extracted December 2025):

    - In 2025, 20.0% of EU enterprises with 10+ employees used at least one AI technology, up from 13.5% in 2024.
    - By size in 2025: small (10-49) 17.0%, medium (50-249) 30.4%, large (250+) 55.0%.
    - The most common use was analysing written language (11.8% of enterprises).

  https://ec.europa.eu/eurostat/statistics-explained/index.php/Use_of_artificial_intelligence_in_enterprises (search summary only)

- **[EVIDENCE, second-hand]** For 2024, small-enterprise use was 11.2% (cited from Eurostat in a secondary paper, not checked against the table). That implies roughly +6 points for small firms in one year (my arithmetic).

- **[EVIDENCE] The size gap widened in 2025 (gap-fill pass).** The 2025 figures were released on 11 December 2025 (Eurostat release code DDN-20251211-2, per secondary

citations). For 2024, secondary sources give about 11% small, 21% medium and 41% large. So large enterprises rose about 14 points in a year and small ones about 6. The large-to-small ratio is about 3.2 to 1 (55.0% vs 17.0%). EU-wide AI adoption is growing fastest where firms are already biggest. Eurostat Statistics Explained https://ec.europa.eu/eurostat/statistics-explained/index.php/Use_of_artificial_intelligence_in_enterprises (search extract); 2024 split and release code via https://ioplus.nl/en/posts/ai-usage-soars-in-european-companies and https://ai.bcci.bg/en/news/ai-adoption-gap-bulgaria-eu-eurostat-2025 (search extract, secondary). The 2024 split should be checked against the Eurostat table (isoc_eb_ai).

4b. OECD: Generative AI and the SME workforce (December 2025)

- **[EVIDENCE] (gap-fill pass)** The OECD surveyed more than 5,000 SMEs in Austria, Canada, Germany, Ireland, Japan, Korea and the UK in late 2024. The report was published in December 2025. OECD, Generative AI and the SME Workforce https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en.html (search extract). Findings:
  - **Adoption.** About 31% of SMEs use generative AI. This ranges from 24% in Japan to 39% in Germany.
  - **Self-reported effects.** 65% of users say it raised employee performance. 83% report no effect on staffing needs. 9% report a decrease and 6% an increase.
  - **Barriers among non-adopters.** Poor fit with their work 57%, copyright or regulatory concerns 54%, worries about data fed into models 52%, lack of employee skills 50%.
  - **Preparation.** 86% of SMEs hold neutral or positive views, and only 2% prohibit use. But a third or fewer of adopting SMEs have trained staff, set internal guidelines or reviewed legal issues.
- **Relevance to the thesis.** "Poor fit with our work" is again the top barrier, as in the ONS, Fed and DSIT data. Process documentation and shared context do not appear among the barriers reported. The low rates of training and guidelines among adopters fit the "use without transformation" pattern: tools arrive, but the organisation around them does not change.
- **Method caveat.** The summary I saw gives no sampling frame or response rate. Seven high-income countries only. Fieldwork was late 2024, before agentic tools were widely available.
- **Method note.** Eurostat's headline series excludes firms under 10 employees, which are the large majority of EU businesses. It counts specific named AI technologies rather than "any AI tool". Statements like "a fifth of EU businesses use AI" therefore say nothing about micro firms.

5. Cross-country executive panel

- **[EVIDENCE]** NBER Working Paper 34836, "Firm Data on AI" (Yotzov, Barrero, Bloom and co-authors; February 2026, revised March 2026):
  - Sample: almost 6,000 CFOs, CEOs and executives from stratified firm samples in the US, UK, Germany and Australia, drawn from central-bank and Fed business panels including the UK Decision Maker Panel and the Atlanta Fed survey.
  - 69% of firms use AI, with use higher at younger and more productive firms. Executives who use AI average about 1.5 hours a week.
  - More than 80% of firms report no impact yet on employment or productivity.
  - Expectations over three years: productivity +1.4%, output +0.8%, employment -0.7%.

https://www.nber.org/papers/w34836 ; PDF https://wfhresearch.com/wp-content/uploads/2026/03/Firm-Data-on-AI-March-2026.pdf (search summary only. Working paper, not peer reviewed.)

    - Relevance: these panels skew toward established firms large enough to have a CFO, so 69% is not an SMB figure. It still shows that even a random-ish executive panel gets far higher "use" than BTOS. The likely reason is that the question captures any use by the firm, including personal use by the executive.

6. Membership and vendor surveys

- **[EVIDENCE (membership survey, response-rate unknown)]** NFIB Research Center, Small Business and Technology Survey (released 25 June 2025). The survey was emailed to a random sample of 20,000 owners from NFIB's membership database and fielded 6-31 March 2025.

    - 24% of small employers used AI technologies (examples given: ChatGPT, Copilot, Grammarly, Canva) and 76% did not.
    - By size (secondary summary): 21% of firms under 10 employees vs 48% of firms over 50.
    - 98% of AI users reported no change in employee numbers.
    - Planned or current uses: communications 29%, marketing 27%.
    - 75% of non-users did not plan to start within 12 months (secondary summary).

  https://www.nfib.com/?p=38888 ; https://smallbiztrends.com/new-nfib-report-reveals-how-small-businesses-harness-tech-and-ai/ (search summary only)

- **[EVIDENCE (weak: business-association survey with a polling firm, n and method not obtained)]** US Chamber of Commerce Technology Engagement Center with Teneo Research, 2025 edition: 58% of small businesses say they use generative AI, up from 40% (2024) and 23% (2023). The widely shared "98% of small businesses use AI-enabled tools" comes from the 2024 edition. It counts any software with AI features, not deliberate AI use. https://www.uschamber.com/technology/the-impact-of-technology-on-u-s-small-business (search summary only. Sample size and fieldwork dates not found.)

- **[EVIDENCE (weak: vendor survey, n=540, self-selected panel)]** Thryv (an SMB software vendor), July 2025: "current usage" rose from 39% (2024) to 55% (2025), which Thryv headlined as a "41% surge". n=540 decision-makers, fieldwork 4-14 May 2025. Eligibility: 1-100 employees, revenue $100k to $9.9m, in business 1+ years. Fieldwork firm, weighting and margin of error not found. https://investor.thryv.com/news/news-details/2025/AI-Adoption-Among-Small-Businesses-Surges-41-in-2025-According-to-New-Survey-from-Thryv/ (search summary only)

- **[EVIDENCE (weak: sample is the sponsor's own programme participants)]** Goldman Sachs 10,000 Small Businesses Voices, run by Babson College and David Binder Research:

    - 68% of owners using AI (announced September 2025), up from 51% in 2023.
    - January 27 to February 4, 2026 wave (n=1,256 programme participants): 76% currently use AI and 88% say they need more training to implement it.
    - October 2025 wave (n=1,471): 94% of AI users report a positive impact.

https://www.businesswire.com/news/home/20260317141697/en/Survey-Small-Businesses-Embrace-AI-But-Need-Training-and-Support-to-Fully-Harness-It ; https://www.businesswire.com/news/home/20251028464955/en/Small-Business-Owners-Plan-to-Grow-Bullish-on-AI-Productivity-Despite-Concerns-Over-Government-Shutdown ; Fox Business on the 68% https://foxbusiness.com/economy/small-business-ai-adoption-jumps-68-owners-plan-significant-workforce-growth-2025 (search summary only)

- **[OPINION / EVIDENCE (weak: vendor survey, global n not found)]** Salesforce SMB Trends (2024-25 editions): 75% of SMBs say they are experimenting with or have implemented AI. Regional cuts reach 87% (UAE) and 93% of Indian SMBs "using AI". The only sample size I found was a supplementary Thailand survey of 100 leaders, fielded 8-13 November 2024. Salesforce sells AI products to SMBs. https://www.techedt.com/salesforce-report-87-of-singapore-smbs-using-ai-see-revenue-growth ; https://nationthailand.com/business/economy/40049212 (search summary only)

- **[EVIDENCE (weak: survey by a small-business network, n not obtained)]** Enterprise Nation / Strand Partners (UK, 2026): only 21% of small businesses use AI regularly, and 6% have embedded it into daily work across the business. https://www.smeweb.com/?p=23334 (search summary only. The link between this URL and the figures is probable, not confirmed.)

- **Not obtained:** Intuit QuickBooks, Constant Contact and Microsoft/LinkedIn Work Trend Index SMB figures. The search budget ran out before these.

7. Why the numbers differ (synthesis)

[OPINION, my analysis from the method details above] Six factors explain most of the gap between about 18% (BTOS) and about 76% (Goldman 10KSB).

1. **Definition of "use".**
    - BTOS before November 2025: "used AI in producing goods or services". This deliberately excluded drafting emails.
    - BTOS after: any business function, in the last two weeks.
    - NFIB names ChatGPT, Grammarly and Canva as examples.
    - The 2024 Chamber "98%" counts any AI-enabled software.
    - Salesforce counts "experimenting".

    Each step down that ladder adds users without any change in behaviour.

2. **Who answers, and about whom.** BTOS asks the business about its own use. The SBCS asks whether "the business or its employees" use AI, which captures an employee's personal ChatGPT use. Executive panels ask leaders, who are heavy personal users.

3. **Reference period.** BTOS asks about the last two weeks. Most vendor surveys have no window or ask about "currently". A two-week window undercounts occasional users. Ramp argues this point and has a commercial interest in the higher figure.

4. **Sample frame.** BTOS draws from the Census business register, which includes many non-tech, low-margin micro firms. NFIB uses its membership. The SBCS recruits through partner organisations. Thryv and Salesforce use opt-in online panels with revenue floors ($100k minimum for Thryv). Goldman samples graduates of its own growth programme. Ramp's 46.6% counts card and bill-pay customers of a fintech, by construction tech-forward.

5. **Self-selection and topic salience.** A survey announced as being about AI and technology attracts owners interested in AI. Official surveys bury the AI question among routine questions.

6. **Sponsor incentive.** Every high-end figure comes from an organisation that sells AI products to SMBs (Salesforce, Thryv, Intuit) or runs a programme it wants to show working (Goldman 10KSB). That does not make the numbers false. It does mean they were not designed to estimate prevalence.

The ladder, read together. On the most permissive definitions, about half or more of US small employers touch AI somewhere. About 18-24% of US firms say the business used AI in some function in the last two weeks. Roughly 3% of employer firms are "fully integrated" (7% of the 46% of users in the SBCS). That 3% is the nearest official-data proxy for "AI-native" among SMBs, and it is a self-assessment.

8. Constraints on SMB adoption (what official data say)

- **[EVIDENCE]** UK 2023: use-case identification 39%, cost 21%, skills 16% (ONS, March 2025, link above). Scotland 2025: expertise 14.1%, use cases 10.7%.
- **[EVIDENCE]** US SBCS: among firms with no plans, over half say AI is "not applicable" and 30% "prefer not to use it" (link above).
- **[EVIDENCE (weak, programme panel)]** Goldman 10KSB 2026: 88% of owners say they need more training to implement AI (link above).
- **Gap:** I found no official survey item that asks about undocumented processes, data quality or lack of a shared data layer as barriers. The thesis's specific mechanism ("nobody can describe the real workflow; no shared context") is not measured in any official SMB survey I located. "Can't identify a use case" is the nearest proxy. It is compatible with the thesis, since you can't spot a use case for a workflow you haven't described, but it is equally compatible with a simpler reading: for many small firms there isn't one yet.

9. Implementation costs, failure and abandonment, who implements (gap-fill pass, search extracts only)

Context from the first pass:

- **[EVIDENCE, indirect]** The typical SMB adoption route is buying off-the-shelf AI software (Scotland BICS), and the most-cited uses are communications and marketing (NFIB 29% and 27%). Per-seat SaaS and assistant subscriptions are therefore likely to be the dominant cost, not bespoke builds. The spending data below are consistent with this.
- **[EVIDENCE, indirect]** "Abandonment" is better proxied in official data by the share stuck at "experimenting" (about half of SBCS AI users) than by project-failure rates. Widely quoted failure rates (for example "95% of pilots fail") come from enterprise studies and should not be transferred to SMBs. See tracks 3 and 5.

9a. What SMBs actually spend on AI. No official statistic exists. The best available figures are lender and fintech surveys or card data. They agree that typical spend is small, tens to a few hundred dollars or pounds a month.

- **[EVIDENCE (weak: lender survey, n=241, 10-14 April 2026)]** Clarify Capital, a small-business lending platform. Average annual spend on AI tools, services and initiatives was $10,600, about $880 a month. A few large spenders pull the mean up, so the median is far lower (not obtained). https://clarifycapital.com/small-business-ai-spending (search extract)

- **[EVIDENCE (weak: lender survey, n=363, July 2026)]** A later Clarify Capital wave: 71% of small businesses now pay for AI, against 41% a year earlier. The average monthly bill rose 343%. 59% have no set AI budget. Same URL (search extract; whether the April and July waves share the URL is unconfirmed).
- **[EVIDENCE (weak: card-provider data, method not seen)]** Capital on Tap (UK business credit cards): mean AI spend per small business rose from £93 to £288 between Q2 2024 and Q2 2026. This is reported via secondary coverage. It is probably card-transaction data, but I could not confirm that. https://ppc.land/one-in-three-small-business-owners-raided-personal-savings-to-pay-for-ai/ (search extract, secondary)
- **[EVIDENCE (weak: spend-management vendor data)]** Spendesk: UK and European SMB customers' average AI tool bill was about €160 a month per company in 2024, up fivefold in two years. Spendesk customers only. https://ffnews.com/blog/2024/12/13/spending-on-ai-tools-by-uk-smbs-grows-fivefold-in-past-two-years-according-to-spendesk-data (search extract)
- **[DO NOT CITE yet]** A "median of about $30 a month" and "17.7% of small businesses had bought an AI tool by end-2025" are attributed in secondary blogs to the JPMorganChase Institute. Two targeted searches failed to find the Institute report. Treat as unverified until located on jpmorganchase.com/institute.
- **[OPINION (vendor and agency marketing, weak)]** "AI automation agency" price lists cluster around $1,500 to $25,000 upfront for small-business builds, plus $50 to $2,000 a month to run them. Retainers are quoted at $500 to $3,500 a month for one or two workflows, and agencies serving larger clients quote $10k-$50k setup and $3k-$15k monthly. These come from agencies that sell the service, and the low ends differ by about an order of magnitude. Examples: https://www.solulab.com/ai-automation-cost-for-small-businesses/, https://taskip.net/ai-automation-agency-pricing/, https://aibusiness.vc/solo/how-to-charge-ai-automation-services (search extracts). No Upwork or Fiverr category rate data was found.
- **Reading the cost evidence together.** Measured spend (about £100-£300 or €160 a month on average, with a third or more spending little or nothing) is one to two orders of magnitude below what agencies quote for a bespoke automation build. That fits the adoption data: most SMBs buy subscriptions and never commission a build. A 10-person firm that hires an agency for a shared-context, agentic setup is spending in a band that almost no survey respondents report.

9b. Abandonment and failure. There is no credible SMB AI failure or abandonment rate. I found no survey with a stated method that asks small firms whether they tried AI and stopped, or whether a specific project failed.

- **[EVIDENCE (weak: insurer survey, n=1,500, April 2025)]** NEXT Insurance reported small-business AI use falling from 42% (2024) to 28% (2025), with 58% saying they do not plan to use AI at all. Cost and complexity were the main reasons given. This conflicts with the US Chamber's 58% generative AI use for 2025, and with every official series, which rose. It is better read as a measurement artefact (different samples and wording) than as evidence of abandonment. https://stacker.com/stories/small-business/small-business-ai-adoption-declines-just-28 (search extract, syndicated copy of NEXT's release)
- **[EVIDENCE (weak: sponsor-programme panel)]** Goldman Sachs 10KSB, March 2026: more than three-quarters of owners use AI to some degree, but only 14% have embedded it across core operations. This is a depth gap, not an abandonment rate. Search extract via https://fortune.com/2026/03/18/small-business-ai-slow-integration-across-operations/.
- **[EVIDENCE (weak: lender survey, n=942, 7-9 April 2026)]** Bluevine: 74% of SMBs use or test AI tools. Only 22% are completely confident AI can handle routine tasks on its own. 48%

of adopters save at least four hours a week. https://www.bluevine.com/newsroom/bluevine-study-shows-ai-is-saving-nearly-half-of-smbs-4-hours-a-week (search extract)

- **Bottom line.** Plainly: no one has measured SMB AI abandonment. The SBCS "experimenting" share (about half of users) and the Goldman 14% "embedded" figure show many SMBs are stuck at shallow use. That is not the same as having tried and failed.

9c. Who implements for SMBs. The evidence points to MSPs as the channel SMBs expect to help, and MSPs mostly not yet able to.

- **[EVIDENCE (weak: vendor survey of its MSP customers, n>1,000)]** Kaseya, State of the MSP report (2025-26 edition; the year label differs between sources): 48% of MSPs rank AI and automation as their clients' top need, but only 13% generate meaningful revenue from AI services. https://www.itpro.com/technology/artificial-intelligence/ai-demand-soars-as-traditional-deal-sizes-shrink-kaseya-reports ; Kaseya page https://www.kaseya.com/?p=27349 (search extracts)
- **[EVIDENCE (weak: vendor survey, n>1,000 MSPs, 2025)]** OpenText Cybersecurity: 92% of MSPs report growth from client interest in AI, but only about half feel ready to guide SMBs in deploying AI tools. https://s23.q4cdn.com/197378439/files/doc_news/OpenText-Cybersecurity-Finds-92-of-Managed-Service-Providers-See-AI-Driven-Growth-But-Readiness-Gap-Widens-2025.pdf (search extract)
- **[EVIDENCE (weak: trade association survey, n=1,042 channel executives, December 2024)]** GTIA (formerly CompTIA Community), State of the Channel 2025: AI tool adoption and cybersecurity demand tied as the top-ranked health drivers for the IT channel, at 49% each. An older CompTIA figure put firms offering automation services, including AI, at 25%. https://businessof.tech/2025/01/06/2025-optimism-abounds-msps-and-small-businesses-double-down-on-growth-tech-and-staying-ahead-of-ai-threats/ (search extract, secondary)
- **[OPINION]** The "AI automation agency" market is visible mainly through agency marketing and course-sellers ("how to charge $2.5K-$15K per project"). I found no independent measure of its size or of how many SMBs use one.
- **Implication.** The MSP data back one part of the thesis: the people SMBs would turn to for implementation say they lack the skills. But it is a supply-side readiness gap (in the channel), not a gap inside the SMB.

9d. Field-experiment evidence on SMBs (the strongest causal evidence found)

- **[EVIDENCE: randomised field experiment, peer-reviewed]** Otis, Clarke, Delecourt, Holtz and Koning, "The Uneven Impact of Generative AI on Entrepreneurial Performance", Management Science (published 2025). 640 Kenyan entrepreneurs over five months were randomly given either a GPT-4 business mentor on WhatsApp or a standard business guide. They found no detectable average effect on performance. Entrepreneurs who were high performers at baseline gained just over 15%, and low performers did about 8% worse (an earlier version gave about 20% and 10%). The gap came from which advice owners chose and how they applied it, not from the advice itself. Low performers asked for help with harder problems. https://pubsonline.informs.org/doi/10.1287/mnsc.2024.06909 ; working paper https://escholarship.org/uc/item/0fb001r8 (search extract)
- **Relevance.** This is the best causal evidence that small-firm AI value depends on the owner's existing capability, a kind of readiness. But the readiness is managerial judgement, not workflow documentation or shared context. Context: very small firms in Kenya, an advice tool not an agent, 2023-24.

10. What "AI-native" plausibly means for a 10-person firm

- **[EVIDENCE]** On the most generous official self-assessment, "fully integrated" AI applies to about 7% of US small-employer AI users, about 3% of all employer firms (SBCS 2026). ONS shows adoption widening, not deepening, with about 1.6 tools per adopter (unverified figure).
- **[EVIDENCE]** Micro firms are close to the average on breadth (ONS 0-9: 28% vs 29% all). The Fed note finds the smallest firms adopt more than size predicts. NBER w34836 finds use higher at younger firms. Taken together, the plausible AI-native SMB is a young, small firm that started with AI tools, not an established 10-person firm that converted.
- **[OPINION]** If "AI-native" means "humans and agents read from and write to one shared business context", then official data cannot measure it, and the closest proxy (about 3% fully integrated) is a self-assessment with no definition behind it. I found no documented, numbers-backed case study of a 10-person firm running that way. That was a search-budget limitation, not proof that none exist. Open question 1 below.
- **[OPINION (vendor case studies, weak)] Gap-fill pass, small-firm case studies.** One targeted search found only vendor-published cases:
  - A London three-partner accountancy paid £449 a month for an AI assistant and reported 12 admin hours saved a week. https://www.softomatesolutions.com/blog/case-study-ai-pa-accountancy-firm-uk/ (search extract)
  - A 14-person firm in Rennes invested €26,000 (excluding VAT) in an AI tool integrated with its production software, plus €500 a month in API costs. No measured outcome was found. https://kolonell.com/en/blog/ai-document-processing-accounting-firm-cost-2026 (search extract)

  None is independent, and none describes humans and agents sharing one business context. A vendor-adjacent analysis notes that firms of about 50-200 people publish almost no AI outcome data. It also notes that hours freed on hourly-billed work are not revenue unless the capacity is sold or costs cut. https://www.codebridge.tech/articles/ai-case-studies-accounting-firms (search extract)
- **[EVIDENCE, lead not read]** The Journal of Accountancy (August 2025) reports a study combining a survey of 277 accountants with transaction data from 79 small and midsize firms using AI accounting software. It is the closest thing found to measured small-firm outcomes. https://www.journalofaccountancy.com/news/2025/aug/calculating-ais-impact-on-cpas-new-study-quantifies-time-savings/ (search extract; the underlying paper and its numbers were not obtained)
- **Updated conclusion.** After a targeted search, a 10-person "AI-native" firm documented with independent numbers still does not appear in the public record. What exists is vendor case studies of single-tool deployments costing a few hundred pounds a month (or a one-off €26k) and saving roughly 10 hours a week.

Viral statistics traced

- **"98% of small businesses use AI"**. Origin: US Chamber of Commerce / Teneo, "Impact of Technology on U.S. Small Business", 2024 edition (September 2024). It counts businesses using any tool with AI features (for example, software with built-in AI), not deliberate AI adoption. Aggregators recirculate it next to the 2025 edition's 58% "use generative AI".

Solidity: low as an adoption measure, since it measures the software market, not behaviour. https://www.uschamber.com/technology/the-impact-of-technology-on-u-s-small-business

- **"AI adoption among small businesses surged 41%"**. Origin: Thryv press release, July 2025. The 41% is relative growth from 39% to 55% in a 540-person vendor panel with a $100k revenue floor. Readers often take it as "41% of small businesses use AI". Solidity: weak. https://investor.thryv.com/news/news-details/2025/AI-Adoption-Among-Small-Businesses-Surges-41-in-2025-According-to-New-Survey-from-Thryv/
- **"Census shows AI adoption doubled to 18%"**. Origin: the BTOS question change of 17 November 2025. The jump is mostly a wording effect (Census itself notes a level shift). Solidity: the 17-18% level is solid as a measure under the new wording. The doubling is not a behavioural trend. https://ramp.com/data/ai-index-january-2026 ; https://www.census.gov/hfp/btos/downloads/AI%20Question%20Wording%20Updates.pdf
- **"AI adoption is declining at large companies (Census)"** (Tom's Hardware, Gizmodo, September 2025). Origin: short-run moves in the old-wording BTOS series for large firms, on the biweekly two-week window. Critics argued the window is too narrow to infer a trend. Solidity: weak as a trend claim, and superseded by the wording change. https://www.tomshardware.com/tech-industry/artificial-intelligence/ai-adoption-rate-is-declining-among-large-companies-us-census-bureau-claims-fewer-businesses-are-using-ai-tools
- **"68% / 76% of small businesses use AI" (Goldman Sachs)**. Origin: Goldman Sachs 10,000 Small Businesses Voices, which surveys alumni of Goldman's own business-education programme (n about 1,200-1,500 per wave). News coverage reports it as "small businesses" in general. Solidity: valid for that population only. https://www.businesswire.com/news/home/20260317141697/en/Survey-Small-Businesses-Embrace-AI-But-Need-Training-and-Support-to-Fully-Harness-It

Strongest counter-evidence (against the thesis and against the mainstream reading of this track)

Against the thesis ("AI fails because the business isn't ready"):

- **Many SMBs are not failing at AI. They are declining it, and say why.** In the Fed SBCS, a third of US employer firms have no plans to use AI. More than half of those say it is not applicable and 30% prefer not to. That is a judgement, not a readiness gap. https://www.fedsmallbusiness.org/2026-report-on-employer-firms
- **The documented barriers are use-case identification, skills and cost, not process documentation or shared context** (ONS 2023: 39% / 16% / 21%). The thesis's mechanism is untested in official SMB data. https://www.ons.gov.uk/economy/economicoutputandproductivity/productivitymeasures/articles/managementpracticesandtheadoptionoftechnologyandartificialintelligenceinukfirms2023/2025-03-24
- **For most SMBs there is no transformation project to fail.** Adoption is buying ready-made tools (Scotland BICS) for writing and marketing (NFIB). "The tech is the easy part" is literally true here, but the hard part that follows is often not attempted at all, rather than attempted and failed.
- **Cross-country SME data agree (gap-fill pass).** OECD (more than 5,000 SMEs, seven countries): the top barrier among non-adopters is poor fit with their work (57%), and 83% of adopters report no effect on staffing. https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en.html

- **The one SMB randomised experiment shows the binding constraint is the owner's judgement, not documentation.** In Otis et al. (Kenya, n=640), GPT-4 advice helped strong performers and hurt weak ones, with no average effect. https://pubsonline.informs.org/doi/10.1287/mnsc.2024.06909
- **The few impact measures available are positive or neutral, not failure.** 71% of SBCS AI users report higher productivity. 94% of Goldman 10KSB users report positive impact (weak panel). 98% of NFIB users report no headcount change. Self-reports skew positive, but there is no SMB evidence of widespread failed AI projects.

Against the mainstream reading ("SMBs are rapidly adopting AI"):

- Official measures sit at about a fifth to a third of firms. "Fully integrated" sits around 3% of US employer firms. UK depth of use is flat while breadth triples. The vendor-survey consensus of "most SMBs use AI" rests on permissive definitions and self-selected panels.
- The single largest one-period jump in the US official series (about 10% to about 17%) is a questionnaire change.

Against the "micro firms lag" assumption:

- ONS puts 0-9 employee firms at 28% vs 29% overall. The Fed note says the smallest firms over-perform their size. NFIB, by contrast, shows a steep gradient (21% vs 48%). The evidence on micro firms conflicts.

Open questions

1. Are there any independent, numbers-backed case studies of 5-15 person firms where staff and agents share one stored business context (cost, time to build, what broke)? One targeted search found only vendor single-tool cases. Still the highest priority.
2. What do SMB AI implementations cost, split between subscriptions, integration and outside help? Lender and card data now suggest average subscription spend of roughly £100-£300 a month. Nothing measures integration or outside-help spend. Upwork/Fiverr category data and UK BICS in-house vs purchased questions remain unchecked.
3. Is there any SMB-specific measure of AI abandonment with a stated method? Still none found. The SF Fed's July 2026 qualitative SBCS brief is the best lead.
4. Who implements for SMBs? MSP surveys (Kaseya, OpenText, GTIA) show demand but low MSP readiness and revenue. The size of the "AI automation agency" market and the share of SMBs using freelancers remain unmeasured.
5. What does DSIT 2026 say about micro-firm adoption and barriers? The full report was published 29 October 2026 and could not be fetched.
6. What is the latest BTOS figure under the new wording, by firm size? The August to September 2026 readings of 23-24% are unverified.
7. Does CES-WP-26-25 show whether small firms use AI in more than one business function? This would be the best US depth measure.
8. OECD SME reports (for example its 2025 survey work on generative AI and SMEs) and Intuit QuickBooks, Constant Contact and Microsoft/LinkedIn surveys: not obtained.

What's genuinely uncertain

- **The true level of SMB AI use.** Defensible figures run from about 14% to about 46% depending on definition. Neither end is wrong. They measure different things.

- **Whether the BTOS-SBCS gap (about 18-20% vs 46%) is mostly definition** (business use in the last two weeks vs the business "or its employees" ever using it) **or mostly sample frame.** Probably both. No study I found decomposes it.
- **Whether micro firms lag or not.** Official UK data say barely. NFIB and Census size bands say yes.
- **Whether the DSIT jump from 16% to 32% in a year is real behaviour or method.** Same contractor model and CATI, but I could not compare questionnaires.
- **Whether "fully integrated" (SBCS 7% of users) means anything close to AI-native.** It is a self-assessment with no definition behind it.
- **Every number in this file is unverified on its source page** (see header).

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
~10% to ~17% (corrected in verification)	US firms using AI, before vs after BTOS question change (17 Nov 2025)	Ramp AI Index; Census wording note	https://ramp.com/data/ai-index-january-2026	14 Jan 2026	Level shift is a wording artefact. Search-summary only
~18%	US firms using AI, year-end 2025 (new wording)	Fed FEDS Note	https://www.federalreserve.gov/econres/notes/feds-notes/monitoring-ai-adoption-in-the-u-s-economy-20260403.html	3 Apr 2026	Official staff analysis of BTOS. Search-summary only
~19.8%	US firms using AI, period ending 3 May 2026	Census Bureau story	https://census.gov/library/stories/2026/05/ai-use-businesses.html	May 2026	Official. Search-summary only
<20%	US firms with 1-4 employees using AI (ceiling, not point)	Census Bureau story	https://census.gov/library/stories/2026/05/ai-use-businesses.html	May 2026	Census publishes as a band
46%	US small employer firms whose business or employees use AI	Fed SBCS 2026 Report on Employer Firms	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	Mar 2026 (fieldwork Sep-Nov 2025)	Convenience sample, weighted. Includes employees' use
7%	Share of SBCS AI-using firms "fully integrated" (~3% of all employer firms, my arithmetic)	Fed SBCS 2026	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	Mar 2026	Self-assessed. Search-summary only
71%	SBCS AI users reporting higher productivity	Fed SBCS 2026; Cook speech	https://www.federalreserve.gov/newsevents/speech/cook20260624a.htm	24 Jun 2026	Self-report
>50% / 30%	Of SBCS firms with no AI plans: "not applicable" / "prefer not to use"	Fed SBCS 2026	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	Mar 2026	Search-summary only
24%	US small employers using AI (21% of <10 employees; 48% of >50)	NFIB Small Business and Technology Survey	https://www.nfib.com/?p=38888	25 Jun 2025	Random draw of 20,000 members by email. Size split from secondary summary
98%	NFIB AI users reporting no change in employee numbers	NFIB	https://www.nfib.com/?p=38888	25 Jun 2025	Search-summary only
58%		US Chamber / Teneo 2025		2025	n and method not found

Figure	What it measures	Source	URL	Date	Reliability note
	US small businesses saying they use generative AI		https://www.uschamber.com/technology/the-impact-of-technology-on-u-s-small-business		
98%	US small businesses using any AI-enabled tool	US Chamber 2024 edition	https://www.uschamber.com/technology/the-impact-of-technology-on-u-s-small-business	Sep 2024	Measures software features, not adoption. Do not cite as adoption
55% (from 39%)	Small businesses "currently using" AI	Thryv	https://investor.thryv.com/news/news-details/2025/AI-Adoption-Among-Small-Businesses-Surges-41-in-2025-According-to-New-Survey-from-Thryv/	17 Jul 2025	Vendor, n=540, revenue floor $100k
76%	Goldman 10KSB participants using AI	Goldman Sachs 10KSB Voices (Babson / David Binder)	https://www.businesswire.com/news/home/20260317141697/en/Survey-Small-Businesses-Embrace-AI-But-Need-Training-and-Support-to-Fully-Harness-It	17 Mar 2026	n=1,256 programme alumni. Not representative
46.6%	US businesses with paid AI spend on Ramp	Ramp AI Index	https://ramp.com/data/ai-index-january-2026	14 Jan 2026	Ramp customers only, tech-forward
69%	Firms using AI, executive panels in US/UK/DE/AU	NBER w34836	https://www.nber.org/papers/w34836	Feb/Mar 2026	n about 6,000. Working paper
>80%	Firms in that panel reporting no AI impact on employment or productivity	NBER w34836	https://www.nber.org/papers/w34836	Feb/Mar 2026	Search-summary only
~12% to ~35%	UK firms with 10+ employees using AI, late 2023 to mid-2026	ONS AI in UK businesses 2023-2026	https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/articles/artificialintelligenceinukbusinesses/2023to2026	20 Jul 2026	Official. Search-summary only
29% / 28% / 49%	UK firms using AI, June 2026: all / 0-9 employees / 250+	ONS (BICS)	same as above	Jul 2026	Official. Search-summary only
1.6	AI tools per UK adopter	TNW on ONS	https://thenextweb.com/news/uk-ai-adoption-widening-not-deepening-ons-2026	Jul 2026	Not confirmed in ONS tables
4%	UK AI-using firms reporting headcount fell due to AI	ONS BICS	https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/bulletins/businessinsightsandimpactontheukeconomy/8january2026	8 Jan 2026	Attribution to bulletin inferred
39% / 21% / 16%	UK barriers 2023: use cases / cost / skills	ONS management practices article	https://www.ons.gov.uk/economy/economicoutputandproductivity/productivitymeasures/articles/managementpracticesandtheadoptionoftechnologyandartificialintelligenceinukfirms2023/2025-03-24	24 Mar 2025	Official. Search-summary only
16% overall; 14% micro	UK businesses using AI, 2025	DSIT AI Adoption Research (IFF / Technopolis, n=3,500 CATI)	https://assets.publishing.service.gov.uk/media/6960f3924343a0da370869ba/AI_Adoption_Research_Report.pdf	2025 (pub. ~Jan 2026)	Official-commissioned. Search-summary only
32% (vs 16%)	UK businesses using at least one AI technology, 2026	DSIT AI Adoption Survey 2026 (n=3,503 CATI)	https://www.gov.uk/government/publications/ai-adoption-survey-2026/ai-adoption-survey-report-2026-executive-summary	29 Oct 2026	Micro figure not obtained
20.0% (13.5% in 2024)	EU enterprises with 10+ employees using AI, 2025	Eurostat	https://ec.europa.eu/eurostat/statistics-explained/index.php/Use_of_artificial_intelligence_in_enterprises	Dec 2025	Excludes <10 employees

Figure	What it measures	Source	URL	Date	Reliability note
17.0% / 30.4% / 55.0%	EU AI use 2025: small / medium / large	Eurostat	same	Dec 2025	Search-summary only
26.7%	Scottish businesses using AI, Sep 2025 (14.0% Sep 2023)	Scottish Government BICS	https://www.gov.scot/publications/scottish-economic-bulletin-december-2025/pages/6/	Dec 2025	Scotland only
~11% / 21% / 41%	EU AI use 2024: small / medium / large (so +6 vs +14 points to 2025)	Eurostat via secondary	https://ioplus.nl/en/posts/ai-usage-soars-in-european-companies	Dec 2025	Secondary; check isoc_eb_ai table
31% (24% JP to 39% DE)	SMEs using generative AI, 7 OECD countries, late 2024	OECD, Generative AI and the SME Workforce, n>5,000	https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en.html	Dec 2025	Official IO survey; frame not seen. Search extract
57%	Non-adopting SMEs citing poor fit with their work	OECD, same	same	Dec 2025	Search extract
83% / 9% / 6%	Adopting SMEs: no staffing effect / decrease / increase	OECD, same	same	Dec 2025	Search extract
$10,600/ yr	Mean SMB spend on AI tools, services and initiatives (n=241)	Clarify Capital	https://clarifycapital.com/small-business-ai-spending	Apr 2026	Lender survey; mean skewed by big spenders
71% (from 41%)	SMBs paying for AI (n=363)	Clarify Capital	https://clarifycapital.com/small-business-ai-spending	Jul 2026	Lender survey. Search extract
£93 to £288	Mean AI spend per UK small business, Q2 2024 to Q2 2026	Capital on Tap via PPC Land	https://ppc.land/one-in-three-small-business-owners-raided-personal-savings-to-pay-for-ai/	2026	Card provider; method not seen; secondary
~€160/ month	Average AI tool bill per UK/EU SMB customer, 2024	Spendesk	https://ffnews.com/blog/2024/12/13/spending-on-ai-tools-by-uk-smbs-grows-fivefold-in-past-two-years-according-to-spendesk-data	Dec 2024	Vendor customer data
$1,500-$25,000 + $50-$2,000/ mo	Agency-quoted SMB AI automation build + running cost	Agency blogs	https://www.solulab.com/ai-automation-cost-for-small-businesses/	2025-26	OPINION; sellers' price lists
48% / 13%	MSPs ranking AI as clients' top need / earning meaningful AI revenue	Kaseya State of the MSP (n>1,000)	https://www.itpro.com/technology/artificial-intelligence/ai-demand-soars-as-traditional-deal-sizes-shrink-kaseya-reports	2025-26	Vendor survey of MSPs. Search extract
~50%	MSPs feeling ready to guide SMBs in AI deployment	OpenText (n>1,000 MSPs)	https://s23.q4cdn.com/197378439/files/doc_news/OpenText-Cybersecurity-Finds-92-of-Managed-Service-Providers-See-AI-Driven-Growth-But-Readiness-Gap-Widens-2025.pdf	2025	Vendor survey. Search extract
42% to 28%	SMBs using AI, 2024 to 2025	NEXT Insurance (n=1,500)	https://stacker.com/stories/small-business/small-business-ai-adoption-declines-just-28	Apr 2025	Insurer survey; conflicts with all official series
14%	Goldman 10KSB owners with AI embedded across core operations	Goldman 10KSB via Fortune	https://fortune.com/2026/03/18/small-business-ai-slow-integration-across-operations/	Mar 2026	Programme panel
74% / 22%	SMBs using or testing AI / fully confident AI can do routine tasks alone	Bluevine (n=942)	https://www.bluevine.com/newsroom/bluevine-study-shows-ai-is-saving-nearly-half-of-smbs-4-hours-a-week	Apr 2026	Lender survey
+15% / -8%	Performance effect of GPT-4 mentor on high / low baseline performers;	Otis et al., Management Science (RCT,	https://pubsonline.informs.org/doi/10.1287/mnsc.2024.06909	2025	Peer-reviewed RCT; figures vary by version

Figure	What it measures	Source	URL	Date	Reliability note
	average effect not distinguishable from zero	n=640 Kenyan entrepreneurs)			

Verbatim quotes

None. No source page could be fetched in this session (proxy and DNS blocks on all source hosts tried), so no wording could be confirmed. Phrases that appear only in search summaries have been paraphrased above. Candidate quotes to fetch and verify later: Census's description of the level shift in the AI Question Wording Updates PDF; the SBCS 2026 report's sentence on "fully integrated" AI; Ramp's Kharazian on the Census economist's objection.

Sources worth reading in full (ranked, one line each on why)

1. **Fed SBCS, 2026 Report on Employer Firms** https://www.fedsmallbusiness.org/2026-report-on-employer-firms. The only large US small-firm survey that measures depth (experimenting / partial / full) and reasons for non-use.
2. **ONS, Artificial intelligence in UK businesses: 2023 to 2026** https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/articles/artificialintelligenceinukbusinesses/2023to2026. Official breadth vs depth, with size bands including 0-9 employees.
3. **Census BTOS AI Question Wording Updates** https://www.census.gov/hfp/btos/downloads/AI%20Question%20Wording%20Updates.pdf. Needed before citing any US trend across late 2025.
4. **DSIT AI Adoption Survey 2026 full report** https://assets.publishing.service.gov.uk/media/6ae202189bca21f1c5036e99/artificial-intelligence-adoption-research-report-2026.pdf. 3,503-firm UK survey with micro breakdowns and barriers, published 29 October 2026.
5. **Census CES-WP-26-25, The Microstructure of AI Diffusion** https://www.census.gov/library/working-papers/2026/adrm/CES-WP-26-25.html. Firm vs function vs task use. Best US source on depth.
6. **NBER w34836, Firm Data on AI** https://www.nber.org/papers/w34836. Cross-country executive panel. High use, little measured impact.
7. **Fed FEDS Note, Monitoring AI Adoption in the U.S. Economy** https://www.federalreserve.gov/econres/notes/feds-notes/monitoring-ai-adoption-in-the-u-s-economy-20260403.html. Reconciles BTOS with private data and flags small-firm over-performance.
8. **NFIB Small Business and Technology Survey (June 2025)** https://www.nfib.com/?p=38888. Large owner sample with size gradient and headcount effect.
9. **Eurostat, Use of AI in enterprises** https://ec.europa.eu/eurostat/statistics-explained/index.php/Use_of_artificial_intelligence_in_enterprises. EU size-class comparison and reasons for non-use.
10. **Ramp AI Index, January 2026** https://ramp.com/data/ai-index-january-2026. Strongest argument that official surveys undercount, written by an interested party.
11. **OECD, Generative AI and the SME Workforce (Dec 2025)** https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en.html. Seven-country SME survey with barriers, staffing effects and preparation. Best cross-country SME source.
12. **Otis et al., Management Science** https://pubsonline.informs.org/doi/10.1287/mnsc.2024.06909. The only randomised experiment on generative AI in small firms found. Shows heterogeneity by owner capability.

13. **SF Fed, "AI Adoption Among Small Businesses: Qualitative Insights from the SBCS" (July 2026)** https://www.frbsf.org/research-and-insights/publications/community-development-research-briefs/2026/07/ai-adoption-in-small-businesses-2024-sbcs/. Not read. Open-text answers from SBCS respondents. Likely the best official source on why SMBs stop or stall.
14. **AEO, "Main Street AI Paradox" brief** https://aeoworks.org/wp-content/uploads/2026/07/AEO_Main-Street-AI-Paradox_Brief_2025.pdf. Not read. Looks directly relevant to micro-business AI use.

INCOMPLETE: page fetching was blocked for every host tried, so every finding below rests on search-engine extracts of the named pages, not on full reads. No quote could be verified. A gap-fill pass (2026-11-06, 30 searches) added: the technical shared-context problem (enterprise text-to-SQL and RAG benchmarks, agent memory, MCP adoption and documented MCP and Copilot security incidents), bank and retailer deployments (JPMorgan, Goldman, Bank of America, Walmart), security-review delay surveys, and controlled ROI evidence (two field experiments, UK and Australian government Copilot evaluations, Forrester TEI as vendor-commissioned). STILL MISSING: a primary source for semantic or context layers improving agent accuracy in production; any audited bank or retailer ROI figure (all figures are company-reported, several via blogs); procurement timelines from pilot to signed contract (only security-review delay data found); a census of production agent incidents and post-mortems beyond the four security disclosures below; Salesforce FY27 quarters.

Track 3: Enterprise

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Summary

- **There is no single "share of pilots reaching production" figure.** The widely cited numbers range from 5% to about 46%. They measure different things: organisations or projects, "reached production" or "showed P&L impact", custom tools or any GenAI. They come from different samples, most of them vendor-sponsored or self-selected. Stacking them as if they agree ("88 to 95% of pilots fail") is a category error.
- **The MIT NANDA "95%" figure is weak.** It comes from a self-described preliminary report built on 52 interviews, 153 survey responses and a review of about 300 public initiatives. It counts success only as measurable P&L impact from embedded or task-specific GenAI tools within roughly six months. The 5% applies to organisations getting custom enterprise tools into production with P&L impact. It does not measure "AI pilots" in general.
- **The Gartner figures are forecasts, not measurements.** "30% of GenAI projects abandoned after PoC by end of 2025" (July 2024) and ">40% of agentic AI projects cancelled by end of 2027" (June 2025) are routinely quoted as findings. The agentic forecast sits on a January 2025 poll of 3,412 webinar attendees. The press release does not say how the "only about 130 real agentic vendors" estimate was produced.
- **The IDC/Lenovo "88% of POCs don't reach production" (4 in 33) was contradicted a year later by the same sponsor-and-research-firm pair.** The Lenovo CIO Playbook 2026, with IDC fieldwork in September and October 2025 and n=3,120, says 46% of AI POCs have progressed into production. It is unclear whether the jump is real, a change of question, or a change of sample.
- **The stated causes of failure only partly support the thesis.** IDC blames organisational readiness (data, processes, IT infrastructure). RAND's interviews put "stakeholders misunderstand what problem needs solving" first. Gartner and S&P Global, however, put cost,

unclear business value, risk controls, data privacy and security at the top. Those are economic and risk-governance problems, not workforce resistance.

- **Adoption keeps rising but profit impact is flat.** McKinsey's 2025 survey (n=1,993) found only 39% attribute any EBIT impact to AI, most of them less than 5%. The 2026 edition (n=1,719) reportedly puts positive EBIT impact at about 37%, essentially unchanged, while the share of $1B+ firms scaling agents rose from 27% to 40%.
- **Vendor disclosures show real paid production use, at a small share of the installed base.** Microsoft reported over 30 million paid M365 Copilot seats in July 2026 (secondary estimates: about 6 to 7% of paid commercial seats). Salesforce reported more than 9,500 paid Agentforce deals and Agentforce ARR above $500M by October 2025. ServiceNow raised its 2026 AI net-new-ACV target from $1B to $1.7B. These figures are real but are also each vendor's chosen metric.
- **Counter to the "pilot purgatory" story,** Wharton/GBK (n>800 US enterprise decision-makers, June to July 2025) reports 74% positive ROI and 72% formally tracking ROI. That is self-reported and unaudited, as are the failure statistics.
- **The shared-context problem is partly a technology problem, not only an organisational one.** On Spider 2.0 (632 real enterprise text-to-SQL tasks, schemas up to 3,000+ columns) the best model solved about 17 to 21% of tasks against 91% on the older academic benchmark. LongMemEval finds commercial assistants lose about 30% accuracy when recalling information across long interactions. And the main protocol for wiring agents to business systems (MCP) has a documented record of prompt-injection and tool-poisoning exploits (GitHub MCP, WhatsApp MCP, Asana's cross-tenant exposure, Microsoft's EchoLeak CVE rated 9.3). These do not go away when the business "is ready".
- **The best controlled evidence on ROI is modest.** A randomised trial across 66 firms (7,137 workers) found Copilot cut email time by roughly a quarter but did not change meeting time, and weekly use ranged from 21% to 75% by firm. The UK Department for Business and Trade found time savings but no evidence that they improved productivity (paraphrase of the evaluation as reported). The UK's larger 20,000-person trial (26 minutes a day saved) is self-reported and not randomised.
- **Security review, not staff resistance, is a measured cause of delay.** A vendor-sponsored 2026 survey reports 82.9% of organisations had AI projects delayed by security or compliance review, about 30% of those by four months or more, and 81.6% shipped a reduced version (less autonomy or data access).

Key findings

The "pilot purgatory" statistics, report by report

MIT NANDA, "The GenAI Divide: State of AI in Business 2025"

- **[EVIDENCE (weak)]** Authors: Aditya Challapally, Chris Pease, Ramesh Raskar and Pradyumna Chari. The report is labelled "Preliminary Findings" (v0.1) and was published July/August 2025. Method: a review of 300+ publicly disclosed AI initiatives, structured interviews with representatives of 52 organisations, and survey responses from 153 senior leaders, fielded January to June 2025. Sources: AI Governance Library summary of the report (https://www.aigl.blog/state-of-ai-in-business-2025/); Virtualization Review, 19 Aug 2025 (https://virtualizationreview.com/articles/2025/08/19/mit-report-finds-most-ai-business-investments-fail-reveals-genai-divide.aspx). The PDF is hosted by MLQ (https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf), not on an MIT domain. I could not fetch it this session.

- **[EVIDENCE (weak)]** The headline funnel covers embedded or task-specific ("enterprise-grade") GenAI tools. About 60% of organisations evaluated them, about 20% reached pilot, and about 5% reached production with measurable P&L impact or "marked and sustained" productivity. The same report says over 80% of organisations had explored or piloted general tools such as ChatGPT or Copilot, which mainly improved individual productivity. Sources: DX newsletter "The AI Divide" (https://getdx.com/blog/the-ai-divide/); RCP Mag, 25 Aug 2025 (https://rcpmag.com/articles/2025/08/25/only-1-in-20-ai-investments-deliver-roi.aspx).
- **[EVIDENCE]** Press coverage did not report the sample consistently. Fortune (18 Aug 2025) described 150 interviews, a survey of 350 employees and 300 public deployments. Other summaries give 52 interviews and 153 survey responses. https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/
- **[OPINION]** Critics' main points (from search summaries of the critiques, not full reads):
    - The success test (marked, sustained P&L impact within about six months) is narrow and undefined, so efficiency gains count as failure.
    - The sample is small and opaque.
    - The report promotes NANDA's own agentic-web agenda, and its "connection to MIT is rather loose".

  Futuriom, Aug 2025 (https://www.futuriom.com/articles/news/why-we-dont-believe-mit-nandas-werid-ai-study/2025/08); arnon.dk (https://arnon.dk/mits-95-ai-failure-rate-is-wrong/); aioperator.io, "What the MIT study actually measured" (https://aioperator.io/blog/do-95-percent-ai-projects-fail). Those critics' further claims (that NANDA sells $250,000 memberships, and that the report's own graph was misread) appeared only in summaries of podcast and blog commentary. I have not verified them.

Gartner (forecasts)

- **[FORECAST]** 29 July 2024: at least 30% of GenAI projects will be abandoned after proof of concept by end of 2025. Stated reasons: poor data quality, inadequate risk controls, escalating costs and unclear business value. Trade coverage of the release cites a GenAI deployment cost range of $5M to $20M. https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025 (coverage: https://www.intelligentcio.com/eu/2024/08/05/gartner-predicts-30-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025/)
- **[FORECAST]** 25 June 2025: over 40% of agentic AI projects will be cancelled by end of 2027, due to escalating costs, unclear business value or inadequate risk controls. Analyst Anushree Verma described most current projects as early-stage, hype-driven experiments. https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027
- **[OPINION]** Same release: "agent washing" (rebranding assistants, RPA and chatbots as agentic) is widespread. Gartner estimates only about 130 of thousands of agentic AI vendors offer genuine agentic capability. The release gives no method for this number. Same URL.
- **[EVIDENCE (weak: self-selected webinar poll)]** The release's survey basis is a January 2025 Gartner webinar poll of 3,412 attendees. 19% said their organisation had made significant agentic AI investments. Same URL.

IDC / Lenovo

- **[EVIDENCE (weak: vendor-sponsored)]** IDC's 2025 CIO Playbook ("It's Time for AI-nomics"), sponsored by Lenovo and published Feb 2025, found 88% of observed POCs did not make it to widescale deployment: of every 33 POCs, 4 graduated to production. IDC attributed this to low organisational readiness in data, processes and IT infrastructure. CIO.com, 2025 (https://www.cio.com/article/3850763/88-of-ai-pilots-fail-to-reach-production-but-thats-not-all-on-it.html); Lenovo release (https://www.businesswire.com/news/home/20250205462134/en). A secondary source gives n=2,920, which I have not verified.
- **[EVIDENCE (weak: vendor-sponsored)]** Lenovo CIO Playbook 2026 ("The Race for Enterprise AI"): IDC surveyed 3,120 IT and business decision-makers between 16 Sept and 17 Oct 2025. 46% of AI POCs have already progressed into production. Respondents projected $2.78 to $2.86 return per dollar, depending on regional edition. Business Wire, 27 Jan 2026 (https://www.businesswire.com/news/home/20260127255538/en/).

S&P Global Market Intelligence

- **[EVIDENCE (moderate: independent research firm; full methodology not seen)]** The share of companies abandoning most of their AI initiatives rose to 42% in 2025 from 17% the year before. The average organisation scrapped 46% of AI proofs of concept before production. Top obstacles were cost, data privacy and security risks. The survey had over 1,000 respondents (reported as 1,006 mid-level and senior IT and line-of-business staff in North America and Europe, attributed to S&P / 451 Research). Fieldwork dates are reported inconsistently (Oct to Nov 2024 vs March 2025). CIO Dive, 2025 (https://www.ciodive.com/news/AI-project-fail-data-SPGlobal/742590/); Fortune, 11 June 2025 (https://fortune.com/2025/06/11/ai-companies-employee-fatigue-failure/). I could not locate the primary S&P report.

McKinsey, The State of AI

- **[EVIDENCE (moderate: large online survey, self-reported, consultancy)]** 2025 edition ("Agents, innovation, and transformation", Nov 2025): n=1,993 participants in 105 countries, fielded 25 June to 29 July 2025.
    - 39% attribute any level of EBIT impact to AI, and most of those say under 5% of EBIT.
    - 62% are at least experimenting with AI agents.
    - In no business function do more than 10% say they are scaling agents.
    - "High performers" (5%+ of EBIT from AI) are about 6% of respondents.

  https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai

- **[EVIDENCE (moderate)]** 2026 edition ("On the road to ROI", Aug 2026): n=1,719 in 97 countries, fielded 4 May to 8 June 2026.
    - Among organisations with $1B+ revenue, the share scaling AI agents rose to 40% from 27%. Smaller organisations were flat at 22%.
    - Agent scaling is most common in IT (12%) and knowledge management (11%), and lowest in manufacturing (3%).
    - Positive EBIT impact is about 37%, roughly unchanged.

  PDF: https://www.mckinsey.com/~/media/mckinsey/business%20functions/quantumblack/our%20insights/the%20state%20of%20ai/the-state-of-ai-in-2026-on-the-road-to-roi-aug-2026.pdf. The figures come from secondary summaries (https://idm.net.au/node/15761; h

ttps://digitalstrategy-ai.com/mckinsey-state-of-ai-2026), not from my own read of the PDF. Treat them as provisional.

Deloitte, State of AI in the Enterprise 2026

- **[EVIDENCE (moderate: consultancy survey)]** Over 3,000 director-to-C-suite leaders directly involved in AI (one summary says 3,235 in 24 countries). Only 25% of respondents have moved 40% or more of their AI pilots into production. https://www.deloitte.com/us/en/about/press-room/state-of-ai-report-2026.html

- **[EVIDENCE (weak: figures conflict)]** On agentic AI, summaries of the same report disagree:

    - 23% report at least moderate use today, and 46% expect that within two years (https://marketscale.com/industries/software-and-technology/enterprise-ai-moves-from-pilot-to-production-in-2026-but-gaps-in-governance-and-talent-persist).
    - Close to three-quarters plan to deploy agentic AI within two years (https://www.deloitte.com/nl/en/about/press-room/state-of-ai-in-the-enterprise-2026.html).

  The "only 21% have mature governance for agents" figure appears only in third-party summaries. Not verified.

BCG, "The Widening AI Value Gap" (Build for the Future 2025)

- **[OPINION / EVIDENCE (weak: consultancy survey with self-assessed maturity)]** Over 1,250 senior executives across nine industries, Sept 2025. 5% are "future-built" and achieving bottom-line value at scale. About 60% are laggards and about 35% "scalers" (the 60/35 split is from secondary summaries). BCG names agentic AI as the main accelerator of the gap. https://web-assets.bcg.com/35/20/636ea73746f8bc91f106ca56ee7c/jpr-2501023-widening-ai-value-gap-vf.pdf ; coverage: https://www.artificialintelligence-news.com/news/value-gap-ai-investments-widening-dangerously-fast/

RAND, "The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed" (13 Aug 2024)

- **[EVIDENCE (qualitative: interviews)]** Authors: Ryseff, De Bruhl and Newberry. Secondary sources say the study is based on interviews with 65 experienced data scientists and engineers. The five root causes include:

    - stakeholders misunderstanding or miscommunicating what problem AI should solve;
    - inadequate data and deployment infrastructure;
    - a focus on the latest technology over real user problems.

  https://www.rand.org/pubs/research_reports/RRA2680-1.html

- **[EVIDENCE: provenance note]** RAND's "more than 80% of AI projects fail, twice the rate of non-AI IT projects" is attributed in the report to "some estimates". RAND did not measure it. Same URL.

Wharton / GBK Collective, "Accountable Acceleration" (28 Oct 2025)

- **[EVIDENCE (moderate: annual survey, self-reported ROI)]** Over 800 US enterprise decision-makers, June to July 2025. 82% use GenAI at least weekly, 46% daily, 72% track structured ROI metrics, and 74% report positive ROI (+9 points). https://

knowledge.wharton.upenn.edu/special-report/2025-ai-adoption-report ; https://www.businesswire.com/news/home/20251028556241/en

Agentic and GenAI systems in production: what vendors disclose

- **[EVIDENCE (company-reported metric)]** Microsoft:
  - Q2 FY26 (28 Jan 2026): 15 million paid M365 Copilot seats among about 450 million M365 commercial seats (about 3.3%). https://www.directionsonmicrosoft.com/blog/microsoft-claims-15-million-paid-m365-copilot-seats
  - Q3 FY26 (29 Apr 2026): 20 million.
  - Q4 FY26 (July 2026): over 30 million (secondary estimate: about 6.6 to 6.7% of paid commercial seats). https://www.uctoday.com/unified-communications/microsoft-365-copilot-passes-30-million-paid-seats-as-cloud-and-ai-growth-power-record-quarter

  Microsoft discloses seats, not usage depth or unit economics.

- **[EVIDENCE (company-reported metric)]** Salesforce Q3 FY26 (quarter to 31 Oct 2025, reported 3 Dec 2025):
  - more than 18,500 Agentforce deals since launch, of which over 9,500 paid (+50% quarter on quarter);
  - Agentforce plus Data 360 ARR nearly $1.4B (+114% YoY);
  - Agentforce ARR above $500M (+330% YoY).

  https://www.salesforce.com/news/press-releases/2025/12/03/fy26-q3-earnings/ ; PDF: https://s205.q4cdn.com/626266368/files/doc_financials/2026/q3/CRM-Q3-FY26-Earnings-Press-Release.pdf. I could not retrieve later quarters (search budget exhausted).

- **[EVIDENCE (company-reported metric)]** ServiceNow:
  - Now Assist net new ACV topped $600M in Q4 2025, against a $1B+ target for 2026.
  - The target was raised to $1.5B at Q1 2026 and to $1.7B at Q3 2026.
  - 19 deals over $5M net new ACV in Q3 2026.

  Form 8-K Q3 FY2026: https://www.sec.gov/Archives/edgar/data/0001373715/000137371526000092/erq3fy26.htm ; Q4 2025 call coverage: https://finance.yahoo.com/news/servicenow-q4-earnings-call-highlights-024438096.html

- **[OPINION]** Reading: these are real paid deployments at scale, which contradicts the strong form of "pilot purgatory". But "paid deal", "seat" and "ACV" are bookings metrics. None of the three vendors discloses active usage, task completion or customer-side ROI. The Salesforce paid-deal count (9,500) is small relative to its customer base.

The shared-context problem, technically (gap-fill pass)

All items in this section and the ones below rest on search extracts; no page was fetched.

Structured enterprise data (text-to-SQL, the "semantic layer" problem)

- **[EVIDENCE: peer-reviewed benchmark]** Spider 2.0 (Lei et al., arXiv 2411.07763, first posted Nov 2024, v2 March 2025; ICLR 2025) uses 632 real enterprise workflow tasks across BigQuery, Snowflake and other dialects, some schemas exceeding 3,000 columns. o1-preview solved about 17% in the original version and about 21.3% with the agent framework in v2,

against 91.2% on Spider 1.0 and 73.0% on BIRD. GPT-4o reached about 10%. The project leaderboard later shows systems at about 38 to 41%, on varying subsets. https://arxiv.org/abs/2411.07763 ; https://spider2-sql.github.io/ (search extract)

- Reading: the gap between clean academic schemas and real enterprise warehouses is large. This is the nearest measured evidence for why "semantic layers" (curated definitions of business metrics) are sold as a fix. I found no independent controlled study showing how much a semantic layer closes the gap; vendor claims only. Treat "add a context layer" as a plausible hypothesis, not a measured result.

Unstructured enterprise documents (RAG)

- **[EVIDENCE (weak: abstracts only)]** Narita et al. (arXiv 2604.02640, AAAI 2026 workshop) tag enterprise RAG queries on four difficulty axes and report that conventional accuracy metrics hide weaknesses in multi-hop, arithmetic and logic-heavy queries; agentic modular improvements helped retrieval but not reasoning. https://arxiv.org/abs/2604.02640 (search extract)
- **[EVIDENCE (weak: abstract only)]** EnterpriseRAG (arXiv 2608.11584, 2026): 983 expert-validated samples across six domains, built around retrieval noise, knowledge gaps and factual conflicts, which the authors say prior benchmarks ignored. No headline results seen. https://arxiv.org/pdf/2608.11584 (search extract)
- **[OPINION: vendor claim]** A RAG-evaluation vendor (FutureAGI) claims retrieval caused about 58% of quality regressions across 40+ enterprise deployments. Internal vendor data, not checkable. Do not cite as evidence. https://futureagi.com/glossary/rag-evaluation/ (search extract)
- I found no large independent empirical study of RAG failure rates on real enterprise corpora. The benchmark literature consistently says real enterprise documents (tables, scans, conflicting versions) are harder than academic sets, but gives no production failure rate.

Agent memory

- **[EVIDENCE: peer-reviewed benchmark]** LongMemEval (Wu et al., arXiv 2410.10813; ICLR 2025): 500 curated questions embedded in long chat histories. The abstract reports commercial chat assistants and long-context LLMs show about a 30% accuracy drop when memorising information across sustained interactions. The full paper gives a 30 to 60% drop for long-context models and 30 to 70% accuracy for commercial systems in a simpler setting. The drop is relative to an oracle-retrieval baseline. https://arxiv.org/abs/2410.10813 (search extract)
- Reading: this directly tests "agents that remember the business". Memory is an unsolved technical problem, which cuts against "the tech is the easy part".

MCP: adoption numbers and documented security incidents (gap-fill pass)

Adoption

- **[EVIDENCE (company-reported, via secondary sources)]** Anthropic reportedly counted more than 10,000 active public MCP servers in December 2025, and reported about 97 million monthly SDK downloads in March 2026. I did not see the Anthropic primary posts; both figures came from secondary blogs. Downloads measure developer activity (including CI pipelines), not enterprise production use. https://nordicapis.com/10-interesting-mcp-statistics/ ; https://dev.to/gentic_news/mcp-hits-10k-servers-97m-monthly-sdk-downloads-by-may-2026-1jbk (search extract)

- **[EVIDENCE (weak: vendor survey, n=300)]** Stacklok, "State of MCP in Software 2026": 300 senior technical leaders at software organisations surveyed December 2025; a write-up reports 41% in limited or broad production with MCP servers. Software firms only, vendor-run. https://stacklok.com/wp-content/uploads/2026/01/State-of-MCP-in-Software-2026_FINAL.pdf (search extract)

Incidents and vulnerabilities (all disclosed in 2025)

- **[EVIDENCE: researcher disclosure]** Tool poisoning, Invariant Labs, April 2025. A malicious MCP server's tool description, never itself called, steered an agent (tested in Cursor) to exfiltrate a user's WhatsApp history through a trusted WhatsApp MCP server, bypassing the user-approval step. An earlier proof of concept hid an instruction to read ~/.ssh/id_rsa in an "add" tool's description. https://invariantlabs.ai/blog/whatsapp-mcp-exploited (search extract)
- **[EVIDENCE: researcher disclosure]** GitHub MCP, Invariant Labs, 26 May 2025. A prompt injection planted in a public repository issue could make an agent using the official GitHub MCP server read the user's private repositories and leak their contents in a public pull request. Invariant says it is architectural (agent has private data, reads untrusted input and can write publicly), so a server-side patch cannot fix it. No CVE surfaced. https://invariantlabs.ai/blog/mcp-github-vulnerability ; https://devclass.com/2025/05/27/researchers-warn-of-prompt-injection-vulnerability-in-github-mcp-with-no-obvious-fix/ (search extract)
- **[EVIDENCE: company disclosure via press]** Asana MCP server, June 2025. Asana launched its MCP server on 1 May, found a logic flaw on 4 June, took it offline, and notified customers on 16 June. Data from one organisation (projects, comments, files) could be visible to users of other organisations, within each user's access scope. Asana said it was a bug, not an attack. A spokesperson told BleepingComputer about 1,000 customers were affected; The Register said the company would not give a number. A widely repeated $7.5M cost figure is an estimate from secondary sources, not from Asana. https://bleepingcomputer.com/news/security/asana-warns-mcp-ai-feature-exposed-customer-data-to-other-orgs/ ; https://theregister.com/2025/06/18/asana_mcp_server_bug/ (search extract)
- **[EVIDENCE: vendor CVE]** EchoLeak, CVE-2025-32711, disclosed June 2025 by Aim Security. A zero-click indirect prompt injection in Microsoft 365 Copilot: a crafted email could cause Copilot to exfiltrate data from the user's M365 context without any click. Microsoft rated it Critical (CVSS 9.3), fixed it server-side, required no customer action, and said it had no evidence of exploitation. https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html ; MSRC entry: https://msrc.microsoft.com/update-guide/vulnerability/CVE-2025-32711 (search extract; MSRC URL constructed from the CVE ID, not seen in results)
- **[EVIDENCE: academic scan]** Hasan et al. (arXiv 2506.13538, 2025) analysed 1,899 open-source MCP servers: 7.2% had general vulnerabilities and 5.5% showed MCP-specific tool poisoning. https://arxiv.org/pdf/2506.13538v2 (search extract)
- Reading: the protocol that makes "one shared context for humans and agents" practical also creates the "lethal trifecta" (private data, untrusted input, an outbound channel). A shared context store is a larger blast radius. The fix is architectural (scoped permissions, isolation), which is engineering work, not readiness work.

Bank and retailer deployments with disclosed numbers (gap-fill pass)

- **[EVIDENCE (company-reported, via press)]** JPMorgan Chase, LLM Suite: about 200,000 employees onboarded within roughly eight months of the 2024 launch (Bloomberg via American Banker). Chief data and analytics officer Teresa Heitsenrether told the WSJ about half use it actively every day, saving an hour or two a week each, and called it early stage. A later investor update cited by secondary sources gives about 150,000 weekly users saving

about four hours each (year not clear in the extract). https://www.americanbanker.com/news/how-jpmorganchase-democratized-employee-access-to-gen-ai (search extract). A claim that Dimon said JPMorgan spends about $2B a year on AI and books about $2B of benefits comes only from a blog (https://ai2.work/blog/how-jpmorgan-s-ai-now-pays-for-itself-inside-the-2b-bet); one analysis notes the shareholder letter's "value" figure is internal and not reconciled to revenue or costs (https://stockti.com/jpmorgan-s-ai-integration-deepens-yet-financial-impact-remains-elusive). Treat the $2B as unverified.

- **[EVIDENCE (company-reported)]** Bank of America, Erica: more than 3 billion cumulative client interactions and nearly 50 million users since the 2018 launch (20 Aug 2025); 676 million interactions in 2024 alone, cumulative over 2.5 billion (Feb 2025); over 3.2 billion cumulative by March 2026. https://newsroom.bankofamerica.com/content/newsroom/press-releases/2025/08/a-decade-of-ai-innovation--bofa-s-virtual-assistant-erica-surpas.html ; https://newsroom.bankofamerica.com/content/newsroom/press-releases/2025/02/digital-interactions-by-bofa-clients-surge-to-over-26-billion--u.html (search extract). Erica is a mostly scripted assistant built over seven years, not a GenAI agent. Its scale shows how long mature customer-facing AI took to build.
- **[EVIDENCE (company-reported, via press)]** Goldman Sachs: the GS AI Assistant went firmwide on 24 June 2025 to about 46,000 staff after a 10,000-person pilot; the bank says employees write about a million prompts a month. CIO Marco Argenti said it would deploy "hundreds", possibly thousands, of Cognition's Devin coding agents alongside about 12,000 developers, under human supervision. His 3 to 4x productivity expectation is a FORECAST. https://www.americanbanker.com/news/goldman-sachs-staff-now-write-a-million-gen-ai-prompts-a-month (search extract)
- **[EVIDENCE (company-reported, earnings call, via press)]** Walmart: on the August 2026 earnings call CEO John Furner said Sparky (customer shopping agent) users were up 70% year on year and spend 40% more per order; in February 2026 the CFO had cited 35% higher average order value. Walmart runs four "super agents" (customers, suppliers and sellers, associates, developers). An older figure has 900,000 associates asking 3 million questions a week of its associate AI tool. https://www.ciodive.com/news/walmart-AI-wins-sparky-john-furner/812606/ ; https://www.retaildive.com/news/walmart-artificial-intelligence-ai-super-agents-retail-associates/753996/ (search extract). The spend-per-order figure is a correlation: heavier shoppers are likelier to use the assistant. It is not a causal estimate.
- Reading: banks deploy GenAI widely to employees (internal assistants, coding agents) and cautiously to customers. None of the four discloses audited ROI. The pattern supports "production is real" and also "production is mostly employee copilots, not autonomous agents".

Procurement, security and governance timelines (gap-fill pass)

- **[EVIDENCE (weak: vendor-sponsored survey; n not seen)]** Protegrity with Enterprise Management Associates, reported July 2026, survey of IT and security leaders: 82.9% had AI projects delayed before production by security or compliance review; over two-thirds of delayed projects stalled at least a month and nearly 30% four months or more; 81.6% deployed AI in reduced form (limits on agent autonomy or data access); 69% run agents that can write to databases or trigger APIs, but only 19.7% deploy highly autonomous systems. https://securitybrief.co.uk/story/enterprise-ai-deployments-slowed-by-security-reviews (search extract). A 47.4% "CISO approves" figure is reported inconsistently; do not use it.
- **[EVIDENCE (weak: vendor survey)]** Apono, "2026 State of Agentic AI Cyber Risk" (fieldwork Dec 2025): 98% said security and data concerns had already affected agentic projects; 21% cited significant delays or reduced scope; only 21% felt prepared for attacks

involving agents. https://securitybrief.co.uk/story/security-chiefs-hit-brakes-on-risky-agentic-ai-roll-out (search extract)

- No source found with months from pilot to signed contract. The UK DBT and GDS trials (below) each ran about three months; the Australian trial about six months. That is a floor for public-sector evaluation cycles, not a procurement benchmark.

Controlled ROI evidence vs vendor ROI (gap-fill pass)

Randomised or quasi-experimental

- **[EVIDENCE: randomised field experiment]** Dillon, Jaffe, Immorlica and Stanton, "Shifting Work Patterns with Generative AI", NBER WP 33795 (2025). Microsoft researchers with academic co-authors; 66 large firms, 7,137 workers (earlier versions said about 6,000), half randomly given M365 Copilot for six months, Sept 2023 to Oct 2024. Copilot cut email time by nearly 3 hours a week for users (about 25%, LATE estimate) but produced no change in meeting time, the most common use. Weekly use among treated workers ranged from 21% to 75% by firm. https://www.nber.org/system/files/working_papers/w33795/w33795.pdf (search extract). Note: Microsoft-affiliated authors, but a randomised design. The firm-level spread in use is itself evidence that organisation matters.
- **[EVIDENCE: peer-reviewed quasi-experiment]** Brynjolfsson, Li and Raymond, "Generative AI at Work", QJE 140(2), 2025: staggered rollout of an AI assistant to 5,172 customer-support agents (5,179 in the 2023 NBER version). Resolutions per hour rose 14 to 15% on average and about 34% for novices, with small speed gains and small quality declines for the most skilled. https://econpapers.repec.org/RePEc:oup:qjecon:v:140:y:2025:i:2:p:889-942 ; https://www.nber.org/system/files/working_papers/w31161/w31161.pdf (search extract)

Government evaluations (not randomised unless noted)

- **[EVIDENCE (moderate: government evaluation, partly random selection)]** UK Department for Business and Trade, M365 Copilot pilot, about 1,000 licences, Oct to Dec 2024 (some reports say to March 2025); mostly volunteers, 30% randomly selected. Found time savings on writing tasks, but reported finding no evidence that the time saved led to improved productivity (paraphrase; wording seen only in search extracts). Some tasks got slower (scheduling took about 35 minutes longer; image generation about 30 minutes longer). 72% satisfied. DBT noted productivity was not the main focus and value for money needs further evaluation. Published around 4 Sept 2025. https://www.theregister.com/2025/09/04/m365_copilot_uk_government ; https://www.civilserviceworld.com/news/article/dbt-trial-finds-copilot-saves-staff-time-but-doesnt-improve-productivity (search extract). The primary DBT report URL did not surface.
- **[EVIDENCE (weak to moderate: government evaluation, self-reported)]** UK GDS cross-government M365 Copilot experiment: 20,000 staff in 12 organisations, 30 Sept to 31 Dec 2024. Average self-reported saving 26 minutes a day (about 13 working days a year); 35 minutes for project delivery, 11 minutes for tax. Published June 2025. https://www.gov.uk/government/publications/microsoft-365-copilot-experiment-cross-government-findings-report/microsoft-365-copilot-experiment-cross-government-findings-report-html ; PDF: https://assets.publishing.service.gov.uk/media/683db42bd23a62e5d32680d0/M365_Copilot_Experiment_Findings_Report.pdf (search extract). Self-reported time, not measured output. The government press release headline ("nearly 2 weeks a year") is an extrapolation.

- **[EVIDENCE (weak to moderate: government evaluation, non-random sample)]** Australian Digital Transformation Agency whole-of-government Copilot trial, evaluated by Nous Group, published Oct 2024: more than 7,600 staff in dozens of agencies, first half of 2024. Users reported gains in summarising, first drafts and search; 1 in 3 used it daily; 77% optimistic at the end. The evaluation itself flags inconsistent rollout and a non-randomised sample; participants said they often forgot Copilot was there. As of January 2025 the DTA planned no further whole-of-government GenAI trials. https://www.digital.gov.au/initiatives/copilot-trial/summary-evaluation-findings/cts-executive-summary ; https://www.itnews.com.au/news/gov-copilot-trial-unsettled-by-usage-and-unmet-expectation-612580 (search extract)
- **[EVIDENCE (weak: internal evaluation)]** Australian Treasury, 14-week trial with 218 staff: concluded GenAI can be deployed in Treasury but with limitations. https://evaluation.treasury.gov.au/publications/evaluation-generative-artificial-intelligence (search extract)

Vendor-commissioned ROI (for contrast)

- **[OPINION: vendor-commissioned projection]** Forrester Total Economic Impact studies commissioned by Microsoft. Most are labelled "New Technology Projected" TEI, i.e. modelled on a composite organisation from a few customer interviews, not measured. Examples: Teams with M365 Copilot (2025), projected ROI 122% to 408%; M365 Copilot for SMBs (Oct 2024), 132% to 353%; M365 Copilot (April 2024), modelled on a 25,000-person composite US organisation. https://tei.forrester.com/go/Microsoft/TeamsandCopilot/ ; https://tei.forrester.com/go/microsoft/smb365copilot/ ; https://tei.forrester.com/go/microsoft/365Copilot (search extract). I could not trace figures of 116% or 457% sometimes attributed to Forrester Copilot studies.
- Reading: the projected TEI ROI (hundreds of percent) and the controlled evidence (email time down about a quarter, no meeting change; no measured productivity gain at DBT) differ by an order of magnitude in tone. Use the controlled studies.

Viral statistics traced

Statistic as quoted	Where people cite it from	Where it actually comes from	How solid
"95% of AI pilots fail"	MIT, via Fortune (18 Aug 2025)	MIT NANDA preliminary report. 5% of organisations got custom or embedded GenAI tools into production with measurable P&L impact within about 6 months. Based on 52 interviews, 153 surveys and about 300 public cases.	Weak. Narrow success definition, small non-random sample, preliminary, and the unit is organisations not pilots. It says nothing about general-purpose tools, which 80%+ had tried.
"88% of POCs never reach production" / "4 in 33"	IDC	IDC CIO Playbook 2025, sponsored by Lenovo. "Observed" POCs, reason given as low readiness.	Weak to moderate. Vendor-sponsored, and the same pair reported 46% reaching production a year later.
"42% of companies abandoned most AI initiatives"	S&P Global	S&P Global Market Intelligence survey of about 1,006 IT and LOB respondents in North America and Europe	Moderate. Independent firm, but the primary report was not found and fieldwork dates conflict. "Most of their initiatives" is a company-level measure.
"46% of POCs scrapped"	S&P Global	Same survey, average per organisation	Moderate, same caveats. Do not confuse with Lenovo/IDC's "46% progressed into production" (2026).

Statistic as quoted	Where people cite it from	Where it actually comes from	How solid
"30% of GenAI projects abandoned after PoC"	Gartner	July 2024 forecast for end-2025	A forecast. Often misquoted as a measured outcome. No outcome check published that I found.
"Over 40% of agentic projects cancelled"	Gartner	June 2025 forecast for end-2027	A forecast. The quoted survey is a self-selected webinar poll.
"Only ~130 real agentic vendors"	Gartner	Same June 2025 release	Analyst estimate, method not published.
"80% of AI projects fail, twice IT"	RAND	RAND attributes it to "some estimates"	Second-hand. RAND's own contribution is qualitative interviews, not a rate.
"43% of MCP servers have command injection flaws"	Security blogs and MCP explainers, 2025-26	Traces to vendor scans (Adversa AI of 500+ servers; one blog attributes it to a 50-server scan with a 22-rule scanner). The peer-reviewed-style scan of 1,899 open-source servers (arXiv 2506.13538) found 7.2% with general vulnerabilities and 5.5% with tool poisoning.	Weak. Use the arXiv figures instead.
"Copilot saves civil servants two weeks a year"	UK government press release, June 2025	GDS experiment, 20,000 staff: 26 minutes a day self-reported, extrapolated to about 13 working days	Moderate as a self-report; not measured output. DBT's separate evaluation found no evidence of a productivity improvement.
"87% of data science projects never make it into production"	VentureBeat, 19 Jul 2019	Commentators find it hard to trace to underlying data. Variants: 85% (Gartner analyst, 2017), 85.4% (Forbes, 2020).	Very weak. A pre-GenAI figure recycled into GenAI discussions. Sources: https://towardsdatascience.com/the-gap-between-data-science-and-the-organization-7cfa8b0e2be5 ; https://builtin.com/articles/why-machine-learning-projects-fail-data-science-engineering

Strongest counter-evidence

Against the thesis ("tech is the easy part; the business isn't ready")

- Gartner's stated reasons for both abandonment forecasts are escalating cost, unclear business value and inadequate risk controls (plus poor data quality in 2024). S&P's respondents named cost, data privacy and security as top obstacles. These are economics and risk problems. Better workflow description or less staff resistance would not necessarily fix them. (Gartner URLs above; https://www.ciodive.com/news/AI-project-fail-data-SPGlobal/742590/)
- Even the MIT report, often cited for "organisational" failure, frames its divide around tools that do not deliver P&L impact in narrow windows. Its strongest critics argue the time window, not the organisation, drives the result. (https://aioperator.io/blog/do-95-percent-ai-projects-fail)
- Where the thesis does get support:
  - IDC attributes low conversion to data, process and IT readiness.
  - RAND's first root cause is stakeholders misunderstanding what problem to solve, which is close to "nobody can describe the real workflow". RAND's sample is data scientists and engineers, so this is the builders' view of their clients.

- **The technical record on shared context says "the tech is the easy part" is wrong for this specific layer (gap-fill pass).** Frontier models solve about a fifth of real enterprise text-to-SQL tasks (Spider 2.0), lose about 30% accuracy on long-term memory (LongMemEval), and the main integration protocol (MCP) had four documented exploit classes or incidents in 2025 (tool poisoning, GitHub MCP, Asana cross-tenant exposure, EchoLeak at CVSS 9.3). Invariant's verdict on GitHub MCP was that it cannot be patched server-side. A business with perfect workflow documentation and enthusiastic staff still faces all of these. (URLs in the sections above.)
- **Security review, not people, is a measured delay.** 82.9% of organisations in the Protegrity/EMA survey reported AI projects delayed by security or compliance review, and 81.6% shipped a reduced-autonomy version (vendor survey). (https://securitybrief.co.uk/story/enterprise-ai-deployments-slowed-by-security-reviews)
- **Where the thesis gets support from the controlled studies:** in the 66-firm randomised trial, weekly Copilot use among workers given access ranged from 21% to 75% depending on the firm. The same tool, randomly assigned, was used three times as much in some organisations as in others. That is evidence that organisational factors drive uptake. (https://www.nber.org/system/files/working_papers/w33795/w33795.pdf)

Against the mainstream "pilot purgatory" reading

- Wharton/GBK: 74% of 800+ US enterprise decision-makers report positive ROI, and 72% track it formally.
- Lenovo/IDC 2026: 46% of POCs reached production.
- McKinsey 2026: 40% of $1B+ firms scaling agents.
- Paid vendor metrics: 30M Copilot seats, 9,500+ paid Agentforce deals, ServiceNow AI ACV target raised 70%.

All are self-reported or vendor metrics, but they are no weaker than the failure statistics they contradict.

- Bank-scale internal deployment is real (gap-fill pass): about 200,000 JPMorgan employees on LLM Suite, about 46,000 Goldman staff on the GS AI Assistant, 3 billion+ cumulative Erica interactions at Bank of America. These are company-reported and say nothing about ROI, but they are not pilots.
- Against the "AI pays off" reading: the UK DBT evaluation found no evidence that time saved became productivity, and the randomised Microsoft-affiliated trial found email time fell but meeting time did not change. (https://www.theregister.com/2025/09/04/m365_copilot_uk_government ; https://www.nber.org/system/files/working_papers/w33795/w33795.pdf)
- The failure and success statistics use incompatible denominators. A 5% org-level P&L metric (MIT), a 12% POC-level conversion (IDC 2025) and a 46% POC-level conversion (IDC 2026) can all be "true" at once. The honest statement is that conversion and profit impact are low and poorly measured, not that "95% fail".

Open questions

1. Does the MIT NANDA report itself say that purchased vendor tools succeed about twice as often as internal builds, and that the core barrier is a "learning gap" (tools that do not retain context or adapt)? Both are widely repeated. I could not verify either this session. If true, the second is a technology limitation, which cuts against "tech is the easy part".

2. Did Gartner ever publish an outcome check on its "30% abandoned by end of 2025" forecast?
3. Why did IDC/Lenovo's POC-to-production rate jump from 12% to 46% in one year? Same question, same sampling frame?
4. Salesforce Q1 and Q2 FY27 Agentforce paid-deal and ARR figures, and churn or renewal data. Not retrieved.
5. Banks and retail (partly answered in the gap-fill pass): user counts now sourced, but via press. Still needed: JPMorgan's own investor-day or annual-report wording on LLM Suite users and AI "value", and whether the $2B benefit figure is real; Walmart's primary transcript for the Sparky 35% vs 40% spend figures.
6. Incident and post-mortem record for agents in production (partly answered): four MCP/Copilot security disclosures are sourced. Still missing: operational incidents (data deletion, wrong customer commitments, outages) and any legal rulings on agent errors. Check track 06.
7. Shared-context technical evidence (partly answered): Spider 2.0, LongMemEval and MCP scans are sourced. Still missing: an independent production RAG failure rate on real enterprise corpora; any controlled test of semantic or context layers. Leads not verified: arXiv 2605.04421 (reported scan of about 200,000 MCP servers, 14.3% with injectable tool descriptions) and arXiv 2608.00150 ("Exposed by Design", live-server testing).
8. Procurement (partly answered): security-review delay data found (Protegrity/EMA, Apono). Still missing: months from pilot to signed contract, with method.
9. Controlled ROI (answered in part): Dillon et al. NBER 33795, Brynjolfsson et al. QJE 2025, UK DBT and GDS, Australian DTA. Still missing: the DBT primary report URL and its exact sample design; a controlled study of agentic (not assistant) deployments.

What's genuinely uncertain

- **The true production conversion rate.** Credible-looking sources put it anywhere between 5% and about 50%, depending on unit and definition. No source in this pass uses a random sample with audited outcomes.
- **Whether ROI is rising.** McKinsey shows flat EBIT impact (about 39% in 2025, about 37% in 2026) while Wharton shows rising self-reported ROI (74%, +9 points). Both are self-reports from different populations.
- **How much "agentic" production is real agents vs relabelled assistants.** Gartner's own 130-vendor estimate has no published method.
- **Whether vendor bookings metrics (seats, paid deals, ACV) translate into usage and customer value.** No vendor discloses usage depth.
- **Whether time saved becomes output.** The UK GDS (26 minutes a day, self-reported) and DBT (no evidence of productivity improvement) evaluations of the same product in the same government in the same quarter point in different directions. The randomised trial shows reallocation (less email) more than measured output gains.
- **How fast the context-layer technology is improving.** Spider 2.0 scores rose from about 17% to about 40% on the leaderboard within roughly 18 months. A technical constraint that is real in 2025 may shrink by 2027; the organisational constraints may not.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
5% (95% "fail")	Orgs whose custom or embedded GenAI tools reached production with measurable	MIT NANDA, GenAI Divide	https://getdx.com/blog/the-ai-divide/ (summary); PDF at mlq.ai	Jul/ Aug 2025	Weak: preliminary, 52 interviews + 153 surveys + ~300

Figure	What it measures	Source	URL	Date	Reliability note
	P&L impact (funnel 60% evaluated, 20% piloted, 5% production)				public cases, ~6-month P&L window
30%	GenAI projects to be abandoned after PoC by end-2025	Gartner press release	https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025	29 Jul 2024	FORECAST
>40%	Agentic AI projects cancelled by end-2027	Gartner press release	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027	25 Jun 2025	FORECAST
~130	Vendors with genuine agentic capability, out of "thousands"	Gartner, same release	same	25 Jun 2025	Analyst estimate, no method
19% (n=3,412)	Orgs with significant agentic investment	Gartner webinar poll, Jan 2025	same	25 Jun 2025	Self-selected webinar audience
88% / 4 of 33	Observed AI POCs not reaching widescale deployment	IDC CIO Playbook 2025 (Lenovo-sponsored)	https://www.cio.com/article/3850763/88-of-ai-pilots-fail-to-reach-production-but-thats-not-all-on-it.html	Feb/ Mar 2025	Vendor-sponsored, n unverified
46%	AI POCs progressed into production	Lenovo CIO Playbook 2026 (IDC, n=3,120, fieldwork 16 Sep-17 Oct 2025)	https://www.businesswire.com/news/home/20260127255538/en/	27 Jan 2026	Vendor-sponsored
42% (from 17%)	Companies abandoning most of their AI initiatives	S&P Global Market Intelligence (n≈1,006, NA + Europe)	https://www.ciodive.com/news/AI-project-fail-data-SPGlobal/742590/	2025	Primary report not seen, fieldwork dates conflict
46%	Average share of AI POCs scrapped before production	S&P Global MI	same	2025	As above
39%	Respondents attributing any EBIT impact to AI (most <5%)	McKinsey State of AI 2025 (n=1,993)	https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai	Nov 2025	Self-reported, consultancy
≤10%	Max share scaling agents in any single function	McKinsey 2025	same	Nov 2025	Self-reported
40% (from 27%)	$1B+ revenue orgs scaling AI agents	McKinsey State of AI 2026 (n=1,719)	https://www.mckinsey.com/~/media/mckinsey/business%20functions/quantumblack/our%20insights/the%20state%20of%20ai/the-state-of-ai-in-2026-on-the-road-to-roi-aug-2026.pdf	Aug 2026	Via secondary summaries, verify in PDF
25%	Respondents who moved 40% + of AI pilots into production	Deloitte State of AI in the Enterprise 2026 (n>3,000)	https://www.deloitte.com/us/en/about/press-room/state-of-ai-report-2026.html	early 2026	Consultancy survey
5%	"Future-built" companies achieving AI value at scale	BCG Widening AI Value Gap (n>1,250)	https://web-assets.bcg.com/35/20/636ea73746f8bc91f106ca56ee7c/jpr-2501023-widening-ai-value-gap-vf.pdf	Sep 2025	Consultancy maturity typology
74%	Enterprise leaders reporting positive GenAI ROI	Wharton/GBK (n>800 US)	https://knowledge.wharton.upenn.edu/special-report/2025-ai-adoption-report	28 Oct 2025	Self-reported ROI

Figure	What it measures	Source	URL	Date	Reliability note
80%	"AI projects fail", cited by RAND	RAND RRA2680-1 citing "some estimates"	https://www.rand.org/pubs/research_reports/RRA2680-1.html	13 Aug 2024	Second-hand estimate, not RAND's measurement
30M+	Paid M365 Copilot seats (~6.6-6.7% of paid commercial seats, analyst estimate)	Microsoft Q4 FY26 earnings call	https://www.uctoday.com/unified-communications/microsoft-365-copilot-passes-30-million-paid-seats-as-cloud-and-ai-growth-power-record-quarter	Jul 2026	Company metric, secondary report
15M	Paid M365 Copilot seats, of ~450M commercial	Microsoft Q2 FY26	https://www.directionsonmicrosoft.com/blog/microsoft-claims-15-million-paid-m365-copilot-seats	28-29 Jan 2026	Company metric
9,500+ / >$500M	Paid Agentforce deals / Agentforce ARR	Salesforce Q3 FY26 release	https://www.salesforce.com/news/press-releases/2025/12/03/fy26-q3-earnings/	3 Dec 2025	Company metric
$1.7B	ServiceNow 2026 AI net new ACV target (raised from $1B)	ServiceNow Q3 2026 8-K	https://www.sec.gov/Archives/edgar/data/0001373715/000137371526000092/erq3fy26.htm	Oct 2026	Company target, not revenue
~17% to 21.3% (vs 91.2% Spider 1.0)	Best model success on 632 real enterprise text-to-SQL tasks	Spider 2.0, arXiv 2411.07763	https://arxiv.org/abs/2411.07763	Nov 2024 / v2 Mar 2025	Peer-reviewed benchmark; leaderboard later ~41%; search extract
~30%	Accuracy drop of commercial assistants and long-context LLMs on long-term memory	LongMemEval, arXiv 2410.10813 (ICLR 2025)	https://arxiv.org/abs/2410.10813	2024-25	Relative to oracle retrieval; search extract
7.2% / 5.5%	Open-source MCP servers with general vulnerabilities / tool poisoning (n=1,899)	Hasan et al., arXiv 2506.13538	https://arxiv.org/pdf/2506.13538v2	2025	Academic scan; search extract
9.3	CVSS score of EchoLeak (CVE-2025-32711), zero-click M365 Copilot prompt injection	Microsoft / Aim Security	https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html	Jun 2025	Vendor CVE; fixed, no known exploitation; search extract
~1,000	Asana customers potentially exposed by MCP server bug	Asana spokesperson to BleepingComputer	https://bleepingcomputer.com/news/security/asana-warns-mcp-ai-feature-exposed-customer-data-to-other-orgs/	Jun 2025	Company statement via press; Asana declined to confirm to other outlets
>10,000 / ~97M	Active public MCP servers (Dec 2025) / monthly SDK downloads (Mar 2026)	Anthropic, via secondary blogs	https://nordicapis.com/10-interesting-mcp-statistics/	2025-26	Primary not seen; downloads are not production use
~200,000	JPMorgan employees onboarded to LLM Suite	Bloomberg via American Banker	https://www.americanbanker.com/news/how-jpmorganchase-democratized-employee-access-to-gen-ai	2025	Company figure via press; ~half active daily
3B+ / ~50M	Erica cumulative client interactions / users since 2018	Bank of America press release	https://newsroom.bankofamerica.com/content/newsroom/press-releases/2025/08/a-decade-of-ai-innovation--bofa-s-virtual-assistant-erica-surpas.html	20 Aug 2025	Company-reported, cumulative
~46,000	Goldman staff given GS AI Assistant (after 10,000 pilot)	Goldman via press	https://www.americanbanker.com/news/goldman-sachs-staff-now-write-a-million-gen-ai-prompts-a-month	24 Jun 2025	Company figure via press
+70% / +40%	Walmart Sparky users YoY / spend per order of Sparky users	Walmart Q2 FY27 call via CIO Dive	https://www.ciodive.com/news/walmart-AI-wins-sparky-john-furner/812606/	Aug 2026	Company-reported; correlation, not causal; Feb 2026 figure was 35%

Figure	What it measures	Source	URL	Date	Reliability note
82.9%	Organisations with AI projects delayed by security/ compliance review	Protegrity / EMA	https://securitybrief.co.uk/story/enterprise-ai-deployments-slowed-by-security-reviews	Jul 2026	Vendor-sponsored; n not seen
~30%	Of delayed projects, share stalled 4+ months	Protegrity / EMA	same	Jul 2026	As above
~25% (~3 hrs/ week)	Reduction in email time from Copilot; no change in meeting time	Dillon et al., NBER WP 33795, RCT, 66 firms, 7,137 workers	https://www.nber.org/system/files/working_papers/w33795/w33795.pdf	2025	Randomised; Microsoft-affiliated authors; search extract
21% to 75%	Weekly Copilot use among treated workers, range across firms	Dillon et al.	same	2025	Randomised access, so variation is organisational
14-15% (34% novices)	Resolutions per hour from AI assistant, customer support	Brynjolfsson, Li, Raymond, QJE 2025 (n=5,172)	https://www.nber.org/system/files/working_papers/w31161/w31161.pdf	2025	Peer-reviewed quasi-experiment
26 min/ day	Self-reported time saved, UK cross-government Copilot experiment (20,000 staff)	GDS findings report	https://www.gov.uk/government/publications/microsoft-365-copilot-experiment-cross-government-findings-report/microsoft-365-copilot-experiment-cross-government-findings-report-html	Jun 2025	Self-report, not randomised
No evidence of productivity gain	DBT Copilot pilot (~1,000 licences)	UK DBT via The Register	https://www.theregister.com/2025/09/04/m365_copilot_uk_government	Sep 2025	Government evaluation; primary not seen
7,600+ / 1 in 3 daily	Australian public servants in Copilot trial / share using daily	DTA, Nous Group evaluation	https://www.digital.gov.au/initiatives/copilot-trial/summary-evaluation-findings/cts-executive-summary	Oct 2024	Non-random sample
122%-408%	Projected ROI range, Teams with M365 Copilot	Forrester TEI commissioned by Microsoft	https://tei.forrester.com/go/Microsoft/TeamsandCopilot/	2025	OPINION: vendor-commissioned projection on a composite org

Verbatim quotes

None. Page fetching was blocked for every host tried this session, so no wording could be confirmed against a fetched page. Search-result snippets contained candidate quotes (for example Satya Nadella on 30 million paid Copilot seats, and Gartner's Anushree Verma on hype-driven agentic experiments). Under the brief's rules these are paraphrased above, not quoted. Re-fetch the Gartner press release and the Microsoft FY26 Q4 transcript to capture them verbatim.

Sources worth reading in full

1. MIT NANDA, "The GenAI Divide" PDF (https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf): the most cited and most misread source. Read the methodology page and funnel chart before using "95%".
2. McKinsey, "The state of AI in 2026: On the road to ROI" (PDF above): the largest repeated survey, with year-on-year agent-scaling and EBIT data.
3. RAND RRA2680-1 (https://www.rand.org/pubs/research_reports/RRA2680-1.html): the only qualitative, non-vendor study of root causes. The closest evidence on "can't describe the problem".

4. Gartner, 25 June 2025 press release: the source of the 40%, 130-vendor and agent-washing claims. Read it to see how thin the basis is.
5. Lenovo CIO Playbook 2025 vs 2026 (Business Wire links above): the same sponsor shows conversion moving from 12% to 46%, a useful case study in survey instability.
6. Wharton/GBK "Accountable Acceleration" (https://knowledge.wharton.upenn.edu/special-report/2025-ai-adoption-report): the strongest large-sample counterweight to the failure narrative.
7. Futuriom critique of MIT NANDA (https://www.futuriom.com/articles/news/why-we-dont-believe-mit-nandas-werid-ai-study/2025/08): a concise statement of the methodological objections.
8. ServiceNow Q3 2026 8-K (SEC link above): filing-grade figures on enterprise AI bookings.
9. Dillon, Jaffe, Immorlica and Stanton, NBER WP 33795 (https://www.nber.org/system/files/working_papers/w33795/w33795.pdf): the largest randomised enterprise trial of a GenAI assistant; shows reallocation of time and large firm-to-firm variation in use.
10. UK GDS Copilot experiment findings report (https://assets.publishing.service.gov.uk/media/683db42bd23a62e5d32680d0/M365_Copilot_Experiment_Findings_Report.pdf) read alongside the DBT evaluation coverage (https://www.theregister.com/2025/09/04/m365_copilot_uk_government): same product, same government, different conclusions.
11. Spider 2.0 (https://arxiv.org/abs/2411.07763): the clearest measurement of how much harder real enterprise data is than benchmark data.
12. Invariant Labs, GitHub MCP vulnerability (https://invariantlabs.ai/blog/mcp-github-vulnerability): the clearest explanation of why connecting agents to shared business context creates an unpatchable class of risk.
13. Australian DTA Copilot evaluation (https://www.digital.gov.au/initiatives/copilot-trial/summary-evaluation-findings/cts-executive-summary): a public-sector evaluation that states its own methodological limits.

INCOMPLETE. Added in the gap-fill pass (search extracts only): KPMG/Melbourne 2025 hidden-use figures; Gallup's primary manager-support figures (80% vs 44% frequent use); BCG sample details; three field experiments (Copilot across 66 firms, Gap/Microsoft training RCT, Kenya RCT); measured scale of middle-management cuts. Still missing: any causal study of de-layering's effect on adoption or productivity (none appears to exist); an independent, peer-reviewed RCT of AI training on business outcomes; Gartner's manager research; the primary source of Gallup's "8.7-fold" figure. Most non-Microsoft/Anthropic figures are from search extracts and need primary re-check.

Track 4: Inertia and education

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

> **Access caveat (read first).** This session's network blocked almost every primary host: bcg.com, gallup.com, mckinsey.com, pewresearch.org, stlouisfed.org, nber.org, iza.org, pnas.org, ssrn.com, arxiv.org, tandfonline.com, upwork.com, gartner.com and others. Only microsoft.com and anthropic.com could be fetched. The shared web-search budget also ran out partway through. So I could check figures against the raw page only for Microsoft's Work Trend Index (2024, 2025) and Anthropic's own studies. Every other figure here comes from search-engine summaries of the primary or secondary pages. Those are marked **[search-summary]** and need re-checking against the primary before publication. Where I know the primary URL, I give it.
>
> Not covered, because the search budget ran out before I reached them: Gartner on managers; the KPMG/University of Melbourne 2025 trust study (hidden AI use, 47 countries); randomised field experiments on AI training; academic studies of de-layering. See Open questions.

Summary

- **The "70% of change programmes fail" figure has no empirical basis.** Hughes (2011, Journal of Change Management) traced five widely cited sources (Hammer & Champy, Beer & Nohria, Bain, McKinsey, Kotter) and found no valid, reliable evidence behind them. The Hammer & Champy origin was openly described as an "unscientific estimate". A same-volume paper by Burnes & Jackson asserts the opposite, so the figure is contested rather than settled. Either way, it should not prop up the thesis.
- **Workers are not, on the whole, resisting the tool. They are adopting it faster than their employers.** Microsoft/LinkedIn (2024, n=31,000, verified): 78% of AI users bring their own AI tools to work, and 52% are reluctant to admit using it for their most important tasks. In Denmark (Humlum & Vestergaard), the barriers workers cite most are employer restrictions and a need for training, not unwillingness [search-summary]. The friction sits more with organisational permission, policy and training than with individual resistance.

- **Resistance, where it exists, looks rational.** In the 2024 Work Trend Index, 53% of AI users worry that using it on important tasks makes them look replaceable (verified). Pew (Oct 2024, n=5,273 US workers): 52% are worried and 36% hopeful [search-summary]. Upwork's much-quoted "77% say AI added to my workload" comes from a vendor survey in which the 77% applies only to the roughly 1,250 non-executive respondents (625 employees, 625 freelancers) in four countries [search-summary].
- **Manager support matters, and is weak.** In BCG's 2025 survey, frontline regular use stalled at 51%, only about a quarter of frontline staff felt they had enough leadership support, and positive sentiment rose from 15% to 55% with strong leadership support (all [search-summary], cross-sectional). Gallup names manager support as a top driver of frequent use while manager engagement fell from 27% to 22% in 2025 [search-summary]. Gallup's own figures: 80% frequent use where managers actively support AI vs 44% where support is low (private sector), with only 28-36% of employees reporting strong support (gap-fill pass, search extract). The randomised Copilot rollout across 66 firms found weekly use ranging from 21% to 75% by firm, so the organisation clearly matters, even if nobody has isolated which part of it. Big tech is cutting manager layers at the same time, and I found no rigorous evidence on what that does to adoption.
- **"Teach a new way of working, not a tool" has only correlational support.** McKinsey's "workflow redesign is the attribute most linked to EBIT impact" is a self-reported, cross-sectional association from a consultancy that sells redesign (OPINION). The theory behind it is stronger: Brynjolfsson, Rock & Syverson (2021) on intangible complementary investment and the productivity J-curve.
- **Adoption gaps by gender survive equal access, so training and access alone do not close them.** A meta-analysis by Otis, Cranney, Delecourt & Koning finds women about 20% less likely than men to use generative AI, and the gap persists where access is equalised [search-summary]. Younger workers use AI more and in Pew are more hopeful, not more anxious [search-summary].
- **The biggest challenge to the thesis is about magnitude, not readiness.** In Denmark, wide adoption with employer encouragement produced average time savings of about 3% and no detectable effect on earnings or hours (Humlum & Vestergaard, NBER w33777) [search-summary]. If adoption is fast but gains are small, the binding constraint may be how much value the tool adds to typical tasks, not "business readiness".

Key findings

1. Change-management evidence: the 70% myth and what is actually known

- **[EVIDENCE]** Hughes, "Do 70 Per Cent of All Organizational Change Initiatives Really Fail?", Journal of Change Management 11(4): 451–464 (2011). He reviews five published sources of the 70% figure (Hammer & Champy; Beer & Nohria; a Bain article; a McKinsey article; Kotter). He traces their supporting references back to a 2002 article by Pace and Mulvin, which itself cites Beer & Nohria's HBR piece. He concludes there is no valid and reliable empirical evidence for the narrative. [search-summary] https://research.brighton.ac.uk/en/publications/do-70-per-cent-of-all-organizational-change-initiatives-really-fa/
- **[EVIDENCE]** Hughes published a reflection in the same journal in 2022. He says he searched for evidence for the claim without success. [search-summary] https://informahealthcare.com/doi/abs/10.1080/14697017.2022.2030980 (DOI 10.1080/14697017.2022.2030980)
- **[OPINION]** Hammer & Champy (Reengineering the Corporation, 1993) give the original reengineering figure as their own "unscientific estimate" that 50% to 70% of reengineering

efforts fall short of the dramatic results intended. I saw this only as a quote in secondary summaries and have paraphrased it. [search-summary] https://strategyu.co/do-70-percent-of-change-initiatives-really-fail

- **[OPINION]** Counterpoint: Burnes & Jackson (2011, same journal volume) state that substantial evidence shows about 70% of change initiatives fail. I could not check their evidence base. [search-summary] https://research.brighton.ac.uk/en/publications/do-70-per-cent-of-all-organizational-change-initiatives-really-fa/
- **Implication.** The honest framing is "change often underdelivers, and nobody has a credible base rate". Any content that says "70% of transformations fail, so AI will too" rests on a figure that started as a guess about 1990s reengineering.

2. Adoption and training: what the better data say

Population-level adoption (US)

- **[EVIDENCE]** Bick, Blandin & Deming, Real-Time Population Survey (nationally representative, US adults aged 18–64). In August 2024, 39% had used generative AI. Pooling August and November 2024, 32% of workers had used it for work. NBER w32966. [search-summary] https://ideas.repec.org/p/nbr/nberwo/32966.html
- **[EVIDENCE]** Their St Louis Fed update ("The State of Generative AI Adoption in 2025", 13 Nov 2025): adoption among adults rose about 10 points in a year to 54.6%. The share of all work hours spent using generative AI rose from 4.1% (Nov 2024) to 5.7% (Aug 2025). [search-summary] https://www.stlouisfed.org/on-the-economy/2025/nov/state-generative-ai-adoption-2025
- **[EVIDENCE]** Gallup quarterly workforce panel, Q1 2026 (fieldwork 4–19 Feb 2026, n=23,717 US employees). 50% use AI at work at least occasionally, up from 21% in Q2 2023. 28% use it frequently (a few times a week or more). [search-summary] https://www.tomshardware.com/tech-industry/artificial-intelligence/half-of-all-us-employees-now-use-artificial-intelligence-at-work-crossing-landmark-threshold-for-first-time-gallup-data-shows-daily-and-weekly-usage-hitting-all-time-high-of-28-percent-in-q1-2026-with-65-percent-feeling-positive-about-its-impact-on-productivity
- **[EVIDENCE, unverified]** Gallup Q2 2026 (May 2026), reported only in secondary coverage: about 30% frequent use and 15% daily. Employees saying their organisation has integrated AI rose from 41% to 47%. About 25% say their organisation has communicated a clear AI plan. [search-summary] https://fortune.com/2026/01/25/ai-use-frequency-employed-adults-tech-workers-retail-manufacturing-gallup-poll/ (coverage of an earlier wave; the Q2 figures need checking at news.gallup.com)
- **[EVIDENCE]** Pew (published 25 Feb 2025; survey Oct 2024; n=5,273 employed US adults). 63% don't use AI much or at all in their job. About 16% say at least some of their work is done with AI. About one in ten use chatbots daily or a few times a week. [search-summary] https://www.pewresearch.org/internet/2025/02/25/u-s-workers-are-more-worried-than-hopeful-about-future-ai-use-in-the-workplace/
- **Note on conflict.** "Half of workers use AI" (Gallup, any use) and "5.7% of work hours" (Bick et al.) are both credible. They measure breadth and depth. Breadth is high; intensity is low.

Employer encouragement and training

- **[EVIDENCE]** Humlum & Vestergaard, "The unequal adoption of ChatGPT exacerbates existing inequalities among workers", PNAS 122(1) e2414972121 (2025). Large Danish survey (11 exposed occupations) linked to administrative registers. Workers see real productivity potential but are held back by employer restrictions and a perceived need for training. In the

IZA working-paper version (DP 16992), 43% said they needed training and 35% said their employer restricted use. These percentages are from secondary reports of the working paper and are not confirmed against PNAS. [search-summary] https://econpapers.repec.org/RePEc:nas:journl:v:122:y:2025:p:e2414972121; https://grip.globalrelay.com/chatgpt-use-depends-on-employer-restrictions-and-training-study-finds

- **[EVIDENCE]** Humlum & Vestergaard, "Large Language Models, Small Labor Market Effects" (NBER w33777, 2025; revised Oct 2025). Danish survey of 2023–24 chatbot use linked to registers. Most employers now encourage chatbot use, and firm-led investments raise adoption and create new job tasks. Average time savings are about 3%. Effects on earnings and recorded hours are precisely estimated near zero (they rule out effects larger than about 1%, or 2% in the revision). [search-summary] https://econpapers.repec.org/RePEc:nbr:nberwo:33777; https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf
- **[EVIDENCE (weak: vendor survey, n=31,000 knowledge workers, 31 markets, online, Edelman, 15 Feb–28 Mar 2024)]** Microsoft/LinkedIn Work Trend Index 2024 (8 May 2024). 75% of knowledge workers use AI at work. Only 39% of AI users have received AI training from their company. Only 25% of companies plan to offer generative AI training that year. (Verified: curl+grep) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part
- **[EVIDENCE (weak: consultancy survey)]** BCG, AI at Work 2025 (press release 26 June 2025). Only 36% of employees say their training is enough. 18% of regular AI users report no training at all. Secondary coverage reports that BCG found five or more hours of training linked to more regular use; I could not confirm that. [search-summary] https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential; https://www.unleash.ai/artificial-intelligence/bcgs-ai-at-work-2025-report-four-takeaways-for-hr-leaders/
- **Gap (first pass).** The first pass found no randomised field experiment that isolates the effect of training. The gap-fill pass found the three below. None is a clean, independent, peer-reviewed training RCT on real business outcomes.

Field experiments (gap-fill pass, search extracts)

- **[EVIDENCE: randomised field experiment of tool access, working paper]** Dillon, Jaffe, Immorlica and Stanton, "Shifting Work Patterns with Generative AI", NBER w33795 (2025). Microsoft 365 Copilot was randomly assigned to 7,137 knowledge workers in 66 firms, with telemetry. Weekly use among treated workers ranged from 21% to 75% depending on the firm. The authors attribute that spread to firm-level factors such as training and management practice, which they could not measure directly. Regular users cut email time substantially, but other effects were modest. Treated workers cut out-of-hours time in apps by about 0.25 hours a week (9%). The authors include Microsoft researchers. https://www.nber.org/system/files/working_papers/w33795/w33795.pdf ; arXiv https://arxiv.org/abs/2504.11436 (search extract)
    - Relevance. This is the strongest causal evidence that, once access is equal, the organisation determines usage. The size of the firm effect (21% vs 75%) supports "the transformation is the hard part". The paper cannot say which organisational factor matters.
- **[EVIDENCE: randomised field experiment of training, preprint, vendor-authored, no IRB]** "Scaffolding Human-AI Collaboration: A Field Experiment on Behavioral Protocols and Cognitive Reframing", arXiv 2604.08678 (April 2026). 388 Gap Inc. employees all had the

same AI tool. The study randomised "thought partner" mindset training against standard feature training, and also tested rigid workflows against free use.

    - The intent-to-treat analysis found no significant average difference in document quality, because 68% of participants in every group hit the maximum score (a ceiling effect).
    - The mindset-trained group did better at the top of the distribution.
    - Changes in attitude looked like recovery from earlier bad experiences, not an effect of the training.
    - All authors work at Microsoft, and the study was not submitted to an IRB.

  Gap's press release says mindset-trained staff were twice as likely to produce top-quality work (paraphrased), and that teams forced into rigid AI workflows scored about 5 points lower. The first is a top-tail result, not an average effect. https://arxiv.org/abs/2604.08678 ; Gap release https://www.gapinc.com/en-ca/articles/2026/04/gap-inc-and-microsoft-research-explores-how-human- (search extracts)

- **[EVIDENCE: randomised field experiment of access to an AI advisor, peer-reviewed]** Otis, Clarke, Delecourt, Holtz and Koning, Management Science (2025). 640 Kenyan entrepreneurs were given a GPT-4 mentor on WhatsApp. There was no detectable average effect. High baseline performers gained about 15%, and low performers lost about 8%. The difference came from which advice they chose to act on. https://pubsonline.informs.org/doi/10.1287/mnsc.2024.06909 (search extract). See 03-smb.md.

- **[EVIDENCE, weak: lab experiment]** A 2026 arXiv lab study (Anand, arXiv 2605.18143) with participants standing in for early-career knowledge workers reports that brief training reduced the variance of performance without lowering its mean. I saw only a truncated extract. https://arxiv.org/html/2605.18143 (search extract)

- **Synthesis.** The randomised evidence says: (1) with access held equal, firms differ hugely in uptake; (2) a training intervention's effect on average output is hard to detect, and any gains show up in the tails; (3) gains go to people who already have good judgement. None of this shows that training reliably works. All of it shows that tool access alone does not determine outcomes.

3. Teaching a new way of working, not a new tool

- **[OPINION / weak EVIDENCE: consultancy survey, self-reported, cross-sectional]** McKinsey, "The state of AI: How organizations are rewiring to capture value" (March 2025). Of 25 organisational attributes tested, workflow redesign had the biggest association with self-reported EBIT impact from generative AI. Only 21% of respondents using generative AI had fundamentally redesigned at least some workflows. [search-summary; from a secondary briefing, not McKinsey's text] https://artificialintellitools.blogspot.com/2025/04/briefing-document-state-of-ai-how.html
- **[OPINION / weak EVIDENCE]** McKinsey, "The state of AI in 2025: Agents, innovation, and transformation" (Nov 2025). 39% of respondents attribute any EBIT impact to AI. "High performers" are nearly three times as likely to have fundamentally redesigned workflows. [search-summary] https://www.mckinsey.com.br/en/our-insights/the-state-of-ai
    - Limitations: the EBIT impact is self-reported by respondents. "High performers" are defined by that same self-report. The data are cross-sectional, so successful firms may simply redesign more (reverse causality). And McKinsey sells redesign work.
- **[EVIDENCE: peer-reviewed theory with calibration]** Brynjolfsson, Rock & Syverson, "The Productivity J-Curve: How Intangibles Complement General Purpose Technologies", AEJ:

Macroeconomics 13(1): 333–72 (2021). General-purpose technologies need large, mostly intangible complementary investments (process redesign, training, new business models) that national accounts miss. So measured productivity is understated early and overstated later. Adjusting for hardware and software intangibles puts US TFP 15.9% above official figures by end-2017. AI-specific effects were "small but growing" at the time. [search-summary] https://www.aeaweb.org/doi/10.1257/mac.20180386; open working paper https://www.nber.org/papers/w25148
    - This is the strongest academic support for "the transformation is the hard part". But it is a macro argument about investment lags, not evidence that any particular change method works.
- **[EVIDENCE (weak: vendor self-study, n=132 surveyed + 53 interviews, August 2025)]** Anthropic surveyed its own engineers and researchers. They self-report using Claude in about 59% of their work, with about a 50% productivity gain, up from 28% and 20% a year earlier. The gain shows up as "considerably more output volume" rather than much less time per task. Staff also reported worry about skill atrophy and fewer mentoring interactions. Anthropic is an AI vendor, and this population is the most AI-fluent there is. (Verified: curl+grep) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic (2 Dec 2025)
    - Relevance: even in an AI-native firm, the "new way of working" changes how skills are learnt and passed on (who juniors ask, what gets practised). That is an education burden this track should name.

4. Generational, gender and country friction

- **[EVIDENCE]** Age. In Pew (Oct 2024), workers aged 18–29 were the most likely to use chatbots at least a few times a month. Younger and higher-income workers were more likely to feel hopeful or excited. [search-summary] https://www.pewresearch.org/internet/2025/02/25/u-s-workers-are-more-worried-than-hopeful-about-future-ai-use-in-the-workplace/
- **[EVIDENCE (weak: vendor survey)]** The Work Trend Index 2024 says bring-your-own-AI "cuts across all generations", not just Gen Z. (Verified: curl+grep) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part
- **Gap.** This session did not reach a source showing younger workers as more anxious. The brief's "more adoption but more anxiety" pattern is unverified here.
- **[EVIDENCE]** Gender. Otis, Cranney, Delecourt & Koning, "Global Evidence on Gender Gaps and Generative AI" (HBS/OSF working paper, Oct 2024, later revised). A meta-analysis of 16–18 studies covering more than 140,000 people finds women about 20% less likely than men to use generative AI. The gap appears in nearly all regions, sectors and occupations, and it persists in settings where access was equalised. The authors suggest ethical concerns and fear of being judged for relying on AI as possible drivers. [search-summary] https://www.hks.harvard.edu/centers/cid/publications/global-evidence-gender-gaps-and-generative-ai; https://www.library.hbs.edu/working-knowledge/women-are-avoiding-using-artificial-intelligence-can-that-hurt-their-careers
- **[EVIDENCE]** Humlum & Vestergaard (PNAS 2025) also find women less likely to have used ChatGPT within the same occupation. Women more often cite training needs, while men more often cite employer restrictions and data confidentiality. [search-summary] https://econpapers.repec.org/RePEc:nas:journl:v:122:y:2025:p:e2414972121
- **[EVIDENCE (vendor platform data, one model, consumer product)]** Country. Anthropic Economic Index (15 Sep 2025), based on 1 million Claude.ai conversations, uses an AI Usage Index (usage share ÷ working-age population share). Israel scores 7, Singapore 4.57, and India 0.27. Usage "strongly correlates with income". Low-usage countries are more likely to

delegate whole tasks (automation), while high-usage countries show more iteration (augmentation). This reflects one vendor's users only. (Verified: curl+grep) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report

5. Middle management: blocker or lever?

- **[EVIDENCE (weak: polling firm, cross-sectional)]** Gallup, State of the Global Workplace 2026 (published around April 2026, covering 2025). Global manager engagement fell from 27% (2024) to 22% (2025), and from 31% in 2022. Gallup reports manager-led support as one of the two top drivers of frequent AI use, and says managers who champion AI raise employees' perception of AI impact "8.7-fold". I have not seen the primary text for this figure, so the exact wording and the comparison group are unverified. [search-summary] https://www.unleash.ai/strategy-and-leadership/analysis/gallups-state-of-the-global-workplace-2026-report-three-key-decisions-for-hr-leaders; earlier wave: https://www.hrdive.com/news/engagement-among-managers-is-slipping-affecting-ai-use-gallup/817069/
- **[EVIDENCE (weak: polling firm, cross-sectional)] Gallup primary figures (gap-fill pass).** Gallup's late-2025 workforce data:
    - In the private sector, at organisations that have adopted AI, 80% of employees are frequent users (a few times a week or more) where managers actively support AI, against 44% where support is low.
    - In the public sector the figures are 65% vs 37%.
    - Only about 28% of employees strongly agreed that their manager actively supports their use of AI.
    - Overall, 26% of workers used AI a few times a week or more and 12% daily (fall 2025).

  Gallup's AI indicator page, updated for May 2026, gives a smaller effect. 36% of employees in AI-integrating organisations strongly agree their manager supports the team's AI use, and those employees are 1.7 times as likely to be frequent users.

  Sources: https://www.gallup.com/workplace/699689/ai-use-at-work-rises.aspx ; https://gallup.com/workplace/702983/adoption-rapidly-growing-public-sector.aspx ; https://www.gallup.com/699797/indicator-artificial-intelligence.aspx (search extracts)

    - Note. Gallup's own figures put the manager effect at about 1.7 to 2 times on frequent use. That is far below the "8.7-fold" reported secondhand above, which concerns perceived impact, not use, and remains unverified. The share of employees with strong manager support (28% to 36%) is close to BCG's "about a quarter". All of these are cross-sectional associations.
- **[EVIDENCE (weak: consultancy survey)]** BCG AI at Work 2025 (n=10,635 employees, 11 countries; frontline staff, managers and leaders surveyed separately). Secondary reprints describe the "about 25%" support figure differently: as leaders giving enough guidance, or as C-suite support. The exact item wording is unknown (search extract, https://www.bigdatawire.com/this-just-in/bcg-report-companies-must-go-beyond-ai-adoption-to-realize-its-full-potential). Frontline regular use is 51%, stalled versus about 52% a year earlier. Only about a quarter of frontline staff say they get enough support from leaders. Positive sentiment toward generative AI rises from 15% to 55% with strong leadership support. Note that "frontline" in BCG's survey means frontline white-collar employees. [search-summary] https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential; https://idm.net.au/node/15229

- **[EVIDENCE (weak: vendor survey, 31,000, Feb–Mar 2025)]** Work Trend Index 2025 (23 Apr 2025). "Last year, employees led the AI wave. This year, it's flipped." 67% of leaders are familiar with agents versus 40% of employees, and 79% of leaders versus 67% of employees think AI will accelerate their careers. (Verified: curl+grep) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born

- **[OPINION / company statements]** De-layering. Microsoft's 2025 cuts were reported to target middle managers, with spans of control widening (per Business Insider; Microsoft gave no official statement at the time). Amazon talks about a "builder ratio", and Jassy's memo backed fewer layers. Google's Pichai told staff it had cut VP and manager roles by 10%. Zuckerberg described "defragmenting" Meta to reduce information latency. [search-summary] https://www.peoplematters.in/news/performance-management/microsoft-wants-leaner-teams-middle-management-layoffs-coming-in-may-45121; https://www.businessinsider.nl/companies-are-culling-managers-mark-zuckerberg-explained-his-mathematical-approach-to-it-nearly-2-years-ago/

- **Gap.** This session found no peer-reviewed or systematic study of the effects of the 2023–2025 big-tech de-layering. The evidence is executive statements and journalism, including a few employee anecdotes about more meetings and wider spans.

- **What has been measured (gap-fill pass, search extracts).** The scale of the flattening is measured. Its effects are not.

    - **[EVIDENCE (weak: polling firm)]** Gallup: the average manager's span of control rose from 10.9 direct reports in 2024 to 12.1 in 2025, up nearly 50% since 2013. Manager engagement is down 9 points since 2022. Via Fast Company https://www.fastcompany.com/91601567/the-great-flattening-rolls-on-as-uber-lays-off-middle-managers (search extract, secondary)
    - **[EVIDENCE (weak: job-postings data)]** Revelio Labs: middle-management job postings fell about 42% from their April 2022 peak to October 2025, while postings for other roles recovered. Via https://pipedot.org/6SNX5 and https://www.pulse.bot/ai/news/middle-managers-are-on-the-chopping-block-thanks-to-ai-in-the-workplace-heres-how-to-save-your-job-1155c975-f7cc-4d9c-97d9-83c8232340e8/ (search extracts, secondary). Indeed data reportedly show steady or rising manager postings in some sectors (construction, medical).
    - **[EVIDENCE (weak, year unclear)]** Live Data Technologies reportedly found middle managers made up 32% of layoffs in a recent year, against 20% in 2019. A Cornell ILR library summary puts them at over 31% of 2023 layoffs. https://catherwood.library.cornell.edu/wit/middle-managers-have-been-a-major-target-for-layoffs-in-the-past-year (search extract)
    - **[EVIDENCE (weak: consultancy survey)]** Korn Ferry 2025: about 4 in 10 respondents said their organisation had cut management layers. 40% of US employees said they felt a lack of direction because of reduced middle management. Korn Ferry warns that de-layering can lower productivity through lost support, higher turnover and weaker alignment. https://www.cfodive.com/news/cuts-middle-management-stressing-72-percent-senior-executives-korn-ferry-HR-WFH/745822/ (search extract)
    - **Still missing.** No causal or before-and-after study of de-layering's effect on AI adoption or productivity was found. A claim of a "15-20% productivity uplift" from removing middle managers appears only in a vendor blog, which cites unspecified McKinsey and BCG work. Do not cite it.

- **Synthesis.** The correlational data consistently say managers are the lever for frontline adoption. Firms are cutting manager layers at the same moment. Nobody has measured what that combination does to adoption. That contradiction is worth writing about.

6. Hiding AI use and shadow AI

- **[EVIDENCE (weak: vendor survey, n=31,000)]** Work Trend Index 2024. 78% of AI users bring their own AI tools to work (80% at small and medium companies). 52% are reluctant to admit using AI for their most important tasks. 53% worry it makes them look replaceable. (Verified: curl+grep) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part

- **[EVIDENCE (weak: vendor, crowdworker-recruited sample, n=1,250, AI-conducted interviews)]** Anthropic Interviewer (4 Dec 2025). 69% of professionals mentioned the social stigma that can come with using AI at work, and 70% of creatives mentioned managing peer judgement. Participants were recruited via crowdworker platforms, though crowdworking was not their main job. (Verified: curl+grep) https://www.anthropic.com/research/anthropic-interviewer

- **[EVIDENCE: academic-led survey with a Big Four sponsor] KPMG / University of Melbourne 2025 (gap-fill pass).** Gillespie, Lockey and colleagues (Melbourne Business School with KPMG), Trust, attitudes and use of artificial intelligence: A global study 2025. More than 48,000 people (48,340 per one outlet) in 47 countries, surveyed November 2024 to January 2025.

    - 57% of employees say they hide their use of AI and present AI-generated work as their own.
    - Almost half admit using AI in ways that contravene company policy, including uploading sensitive company information into free public tools.
    - Only 47% say they have received AI training, and 40% say their workplace has a policy or guidance on generative AI.

  https://mbs.edu/news/global-study-reveals-trust-of-ai-remains-a-critical-challenge ; https://fbe.unimelb.edu.au/newsroom/media-release-global-study-reveals-trust-of-ai-remains-a-critical-challenge-reflecting-tension-between-benefits-and-risks (search extracts). Sampling method (online panels per country) not confirmed this session.

    - Relevance. A non-vendor, 47-country source agrees with Microsoft's 52% "reluctant to admit". Hidden use is a majority behaviour, not a niche one. This is the best evidence for the thesis's "no shared context" point: most AI use at work is invisible to the organisation.

- **[EVIDENCE (weak: consultancy survey)]** BCG 2025: more than half of employees say they would use AI tools even if not authorised. [search-summary] https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential

- **Implication for the thesis.** Hidden use means the "legacy" end of the spectrum is less legacy than it looks. AI is already in the workflow, but untracked and outside any shared context. That supports the "shared context" part of the thesis and undercuts the "people resist" part.

7. Resistance as rational: job fear and workload

- **[EVIDENCE]** Pew (Oct 2024, n=5,273). 52% of workers are worried about AI's future use in the workplace, 36% hopeful, 33% overwhelmed and 29% excited. [search-summary] https://www.pewresearch.org/internet/2025/02/25/u-s-workers-are-more-worried-than-hopeful-about-future-ai-use-in-the-workplace/
- **[EVIDENCE (weak: vendor survey)]** Work Trend Index 2024: 45% of professionals worry AI will replace their job. (Verified: curl+grep). Work Trend Index 2025: 43% "globally" fear AI will take their jobs, versus 21% of "Frontier Firm" leaders. (Verified: curl+grep) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born
- **[EVIDENCE (weak: vendor survey; see the trace below)]** Upwork Research Institute (with Workplace Intelligence; fieldwork by Walr, 16 Apr–5 May 2024; US, UK, Australia, Canada; published 23 Jul 2024). 77% of employees who use AI say it has added to their workload. 47% don't know how to achieve the productivity gains their employers expect. [search-summary] https://investors.upwork.com/news-releases/news-release-details/upwork-study-finds-employee-workloads-rising-despite-increased-c
- **[EVIDENCE]** Humlum & Vestergaard (NBER w33777) find that chatbots create new job tasks, especially where employers encourage use, which partly offsets time saved. [search-summary] https://econpapers.repec.org/RePEc:nbr:nberwo:33777. Anthropic's internal study likewise finds gains taken as more output, not less time (verified). Together these give mechanism-level support to the claim that "AI intensifies work", independent of Upwork.

Viral statistics traced

Statistic	Where cited from	Actual origin	How solid
"70% of change programmes fail"	Kotter, McKinsey, Bain, consultancies, countless decks	Hammer & Champy 1993, given as an "unscientific estimate" (50–70%) about reengineering; Beer & Nohria (HBR 2000). Hughes (2011) found all five common citations lack valid empirical support.	**Weak.** No credible base rate. Do not cite as fact. Contested by Burnes & Jackson (2011), whose evidence I did not check.
"77% of employees say AI has added to their workload" (Upwork, 2024)	HR press, LinkedIn posts, often given as "77% of workers"	Upwork press release, 23 Jul 2024. Total n=2,500, but 1,250 of those are C-suite executives. The 77% applies only to AI-using employees among 625 full-time employees and 625 freelancers in four English-speaking countries. Online panel (Walr). Upwork sells freelance labour as a remedy.	**Weak.** A vendor survey with a small base and unclear question wording. Directionally consistent with Humlum & Vestergaard's "new tasks" finding. Cite as "an Upwork survey of about 1,250 workers" if at all.
"78% of AI users bring their own AI to work"	Microsoft/ LinkedIn WTI 2024, widely quoted as "shadow AI"	WTI 2024: 31,000 knowledge workers, 31 markets, online, Edelman. (Verified)	**Moderate for direction, weak for level.** A large sample, but online knowledge workers only, run by a vendor with an interest in sanctioned (Copilot) adoption.
"Workflow redesign is the #1 driver of EBIT impact" (McKinsey)	Consultancy content, LinkedIn	McKinsey State of AI survey, March 2025. Association between 25 attributes and self-reported EBIT impact.	**Weak.** A self-reported outcome, cross-sectional, with commercial interest. It is a correlation, not a driver.
"Managers who champion AI boost perceived	HR press, 2026	Gallup, State of the Global Workplace 2026 (secondary reports)	**Unverified.** Could not reach Gallup. The comparison is likely across groups of respondents, so not causal. Gallup's own figures on use show an effect of 1.7 to 2 times

Statistic	Where cited from	Actual origin	How solid
AI impact 8.7-fold" (Gallup)			(80% vs 44%; 1.7x in May 2026), so cite those instead.
"AI mindset training made staff twice as likely to do top-quality work" (Gap Inc. / Microsoft, 2026)	Gap press release, retail-tech press	arXiv 2604.08678, RCT of 388 Gap employees, Microsoft authors, no IRB	**Weak as a general claim.** A top-tail result. The intent-to-treat average showed no significant difference because 68% of all participants hit the score ceiling.
"57% of employees hide their AI use" (KPMG/ Melbourne 2025)	HR and tech press	Gillespie, Lockey et al., 48,000+ respondents, 47 countries, Nov 2024 to Jan 2025	**Moderate.** Large multi-country academic-led survey. The item also covers presenting AI work as one's own. Online self-report. Sampling not checked this session.

Strongest counter-evidence

Against the thesis ("people resist a whole new way of working"):

- Bottom-up adoption is fast and widespread. Bick et al. find adoption faster than PCs or the internet at a similar stage, and 78% of AI users bring their own tools. Workers in Denmark cite employer restrictions and lack of training, not unwillingness. The organisation, through its policy, permission, procurement and fear of data leakage, is often the brake. The employees are not.
- Hidden use (52% reluctant to admit; 69% mention stigma) means apparent "resistance" may partly be invisible adoption.

Against "the business isn't ready, so AI fails":

- In Denmark, broad adoption with employer encouragement still produced average time savings of only about 3% and no earnings or hours effects (Humlum & Vestergaard). Even when adoption and encouragement are present, gains for typical workers are small. Readiness may not be the binding constraint. The tool's marginal value in many ordinary jobs may be.
- Work hours spent with generative AI are only about 5.7% of all US hours (Bick et al., Aug 2025). The upside from "fixing readiness" is bounded by how little work currently touches AI.

Against the mainstream reading of change management:

- The "70% fail" base rate is unsupported. "Transformation is hard" may be true, but nobody has measured how hard in a way that supports the folklore.
- The strongest evidence that redesign matters (the J-curve) predicts delayed gains, not failure. Firms reporting "no ROI yet" may be early on the curve rather than failing.

Against "training fixes it" (gap-fill pass):

- In the one randomised training trial found (Gap/Microsoft, n=388), mindset training made no detectable difference to average output quality. In the Kenya RCT, AI help widened the gap between strong and weak operators. The randomised evidence that training or access lifts the typical worker is thin.

Against "managers are the lever":

- The evidence (Gallup, BCG) is cross-sectional. Teams with supportive managers may differ in many ways, such as role type and resources. No causal estimate was found.

Open questions

1. Is there an independent, peer-reviewed RCT of AI training (not access) on business outcomes? The gap-fill pass found one vendor-authored preprint (Gap/Microsoft, null on average), plus access experiments showing large firm effects (Copilot, 21% to 75% weekly use). What firm-level factor explains that spread?
2. KPMG/Melbourne 2025: answered in part (57% hide use; almost half contravene policy). Still to check: the country breakdown, the sampling method, and whether hiding is higher where policy is restrictive.
3. What do Gallup's primary releases say about the "8.7-fold" manager effect, its comparison group and its wording?
4. Has anyone measured adoption or productivity after de-layering? Academic span-of-control literature, and Amazon or Meta disclosures.
5. Is there credible evidence that younger workers are more anxious about AI at work? Pew suggests the opposite.
6. What does Gartner's manager or "change fatigue" research actually measure? Not reached.
7. Does the gender gap close in firms with mandated or sanctioned tools, or only in experiments? Otis et al. say no even with equal access, but firm-level evidence is thin.

What's genuinely uncertain

- **Causality of every "lever".** Manager support, training and workflow redesign all correlate with better outcomes in cross-sectional surveys. None has a clean causal estimate in what I found.
- **Whether small measured gains reflect a J-curve lag or a ceiling.** The Danish 3% could rise as complementary investment accumulates, or it could be the realistic level for many occupations.
- **Level of hidden use.** BYOAI estimates come from vendor surveys of online knowledge workers. True prevalence across all workers, including frontline and SMB, is unknown.
- **Effects of de-layering on adoption.** There is no data. Opinion runs in both directions.
- **Secondary-only numbers.** All figures marked [search-summary] need re-checking against the primary page. Because of the access limits, this applies to BCG, Gallup, Pew, McKinsey, Humlum & Vestergaard, Otis et al., Bick et al. and Upwork.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
0 valid studies	Empirical support for "70% of change initiatives fail" found by Hughes across 5 cited sources	Hughes, J. Change Mgmt 11(4)	https://research.brighton.ac.uk/en/publications/do-70-per-cent-of-all-organizational-change-initiatives-really-fa/	2011	Peer-reviewed; contested by Burnes & Jackson 2011. [search-summary]

Figure	What it measures	Source	URL	Date	Reliability note
78 %	AI users bringing their own AI tools to work (80% at SMBs)	Microsoft/ LinkedIn WTI 2024, n=31,000	https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part	8 May 2024	Vendor survey, online knowledge workers. Verified curl+grep.
52 %	AI users reluctant to admit using it for their most important tasks	WTI 2024	same	8 May 2024	Vendor survey. Verified.
53 %	AI users worried that using it on important tasks makes them look replaceable	WTI 2024	same	8 May 2024	Vendor survey. Verified.
39 %	AI users who have received AI training from their company	WTI 2024	same	8 May 2024	Vendor survey. Verified.
67 % vs 40 %	Leaders vs employees familiar with AI agents	WTI 2025, n=31,000	https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born	23 Apr 2025	Vendor survey. Verified.
43 %	Workers globally who fear AI will take their jobs	WTI 2025	same	23 Apr 2025	Vendor survey. Verified.
77 %	AI-using employees saying AI added to their workload	Upwork Research Institute	https://investors.upwork.com/news-releases/news-release-details/upwork-study-finds-employee-workloads-rising-despite-increased-c	23 Jul 2024	Vendor; base is about 1,250 non-executives in 4 countries. [search-summary]
52 % / 36 %	US workers worried / hopeful about AI at work	Pew, n=5,273, Oct 2024	https://www.pewresearch.org/internet/2025/02/25/u-s-workers-are-more-worried-than-hopeful-about-future-ai-use-in-the-workplace/	25 Feb 2025	Probability panel; strong. [search-summary]
50 % / 28 %	US employees using AI at work at all / frequently	Gallup, n=23,717, Feb 2026	https://www.tomshardware.com/tech-industry/artificial-intelligence/half-of-all-us-employees-now-use-artificial-intelligence-at-work-crossing-landmark-threshold-for-first-time-gallup-data-shows-daily-and-weekly-usage-hitting-all-time-high-of-28-percent-in-q1-2026-with-65-percent-feeling-positive-about-its-impact-on-productivity	Q1 2026	Strong panel; secondary report. [search-summary]
54.6%	US adults 18–64 who have used generative AI	Bick, Blandin & Deming, St Louis Fed	https://www.stlouisfed.org/on-the-economy/2025/nov/state-generative-ai-adoption-2025	13 Nov 2025	Nationally representative. [search-summary]
4.1 % to 5.7 %	Share of all US work hours spent using generative AI, Nov 2024 to Aug 2025	same	same	13 Nov 2025	[search-summary]
about 3%	Average time saved by chatbot users in Denmark; earnings and hours effects about 0	Humlum & Vestergaard, NBER w33777	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	2025	Survey linked to registers; strong. [search-summary]
about	How much less likely women are	Otis, Cranney,	https://www.hks.harvard.edu/centers/cid/publications/global-evidence-gender-gaps-and-generative-ai	Oct 2024 (rev.)	Meta-analysis of 16–18 studies, more

Figure	What it measures	Source	URL	Date	Reliability note
20 %	than men to use generative AI	Delecourt & Koning			than 140,000 people. [search-summary]
51 %	Frontline (white-collar) employees who are regular AI users; stalled	BCG AI at Work 2025	https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential	26 Jun 2025	Consultancy survey. [search-summary]
15 % to 55 %	Positive sentiment toward generative AI, without vs with strong leadership support	BCG AI at Work 2025	same	26 Jun 2025	Cross-sectional. [search-summary]
27 % to 22 %	Global manager engagement, 2024 to 2025	Gallup State of the Global Workplace 2026	https://www.unleash.ai/strategy-and-leadership/analysis/gallups-state-of-the-global-workplace-2026-report-three-key-decisions-for-hr-leaders	Apr 2026	Secondary. [search-summary]
21 %	Generative AI users who have fundamentally redesigned at least some workflows	McKinsey State of AI	https://artificialintellitools.blogspot.com/2025/04/briefing-document-state-of-ai-how.html	Mar 2025	Consultancy; secondary blog. [search-summary]
15.9%	US TFP level above official figures by end-2017 after adjusting for IT intangibles	Brynjolfsson, Rock & Syverson, AEJ Macro	https://www.aeaweb.org/doi/10.1257/mac.20180386	2021	Peer-reviewed. [search-summary]
69 %	Professionals interviewed who mentioned social stigma of AI use at work	Anthropic Interviewer, n=1,250	https://www.anthropic.com/research/anthropic-interviewer	4 Dec 2025	Vendor; crowdworker-recruited sample. Verified.
57 %	Employees who hide AI use and present AI output as their own	KPMG / Univ. of Melbourne 2025, n≈48,000, 47 countries	https://mbs.edu/news/global-study-reveals-trust-of-ai-remains-a-critical-challenge	Apr 2025 (fieldwork Nov 2024 to Jan 2025)	Academic-led, Big Four sponsor. Search extract
~50%	Employees using AI in ways that contravene company policy	same	same	2025	"Almost half". Search extract
47 % / 40 %	Employees with AI training / with a workplace genAI policy	same	same	2025	Search extract
80 % vs 44 %	Frequent AI use where managers actively support vs low support (private sector, AI-adopting orgs)	Gallup	https://www.gallup.com/workplace/699689/ai-use-at-work-rises.aspx	late 2025	Cross-sectional. Search extract
1.7 x / 36 %	Frequent-use multiplier with strong manager support / share with strong support	Gallup AI indicator	https://www.gallup.com/699797/indicator-artificial-intelligence.aspx	May 2026	Search extract

Figure	What it measures	Source	URL	Date	Reliability note
21 % to 75 %	Weekly Copilot use among treated workers, range across 66 firms	Dillon et al., NBER w33795, n=7,137	https://www.nber.org/system/files/working_papers/w33795/w33795.pdf	2025	Randomised access; firm effects. Microsoft co-authors
68 %	Gap RCT participants at maximum quality score (all arms), so no ITT difference was detectable	arXiv 2604.08678, n=388	https://arxiv.org/abs/2604.08678	Apr 2026	Preprint, Microsoft authors, no IRB
10.9 to 12.1	Average direct reports per manager, 2024 to 2025	Gallup via Fast Company	https://www.fastcompany.com/91601567/the-great-flattening-rolls-on-as-uber-lays-off-middle-managers	2025-26	Secondary
-42 %	Middle-manager job postings, Apr 2022 peak to Oct 2025	Revelio Labs via secondary	https://pipedot.org/6SNX5	late 2025	Postings, not employment. Secondary
40 %	US employees feeling lack of direction due to fewer middle managers	Korn Ferry 2025	https://www.cfodive.com/news/cuts-middle-management-stressing-72-percent-senior-executives-korn-ferry-HR-WFH/745822/	2025	Consultancy survey
7 / 0.27	Claude usage per working-age capita relative to expected (Israel / India)	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	15 Sep 2025	One vendor's consumer users. Verified.

Verbatim quotes

- "Without guidance or clearance from the top, employees are taking things into their own hands and keeping AI use under wraps: 78% of AI users are bringing their own AI tools to work (BYOAI)" (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: framing of shadow AI. Verified: curl+grep
- "52% of people who use AI at work are reluctant to admit to using it for their most important tasks." (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: hidden use. Verified: curl+grep
- "Only 39% of people globally who use AI at work have gotten AI training from their company." (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: training deficit. Verified: curl+grep
- "Last year, employees led the AI wave. This year, it's flipped." (Microsoft, Work Trend Index 2025, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: leaders now ahead of employees on agent familiarity. Verified: curl+grep
- "69% of professionals mentioned the social stigma that can come with using AI tools at work" (Anthropic, "Introducing Anthropic Interviewer", 4 Dec 2025) https://www.anthropic.com/research/anthropic-interviewer. Context: 1,250 AI-conducted interviews; vendor source. Verified: curl+grep

- "A colleague recently said they hate AI and I just said nothing. I don't tell anyone my process because I know how a lot of people feel about AI." (Fact-checker interviewed by Anthropic Interviewer, 4 Dec 2025) https://www.anthropic.com/research/anthropic-interviewer. Context: hidden AI use because of peer stigma. Verified: curl+grep (the page uses a curly apostrophe)
- "When producing output is so easy and fast, it gets harder and harder to actually take the time to learn something." (Anthropic employee, quoted in "How AI is transforming work at Anthropic", 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: skill-atrophy concern in an AI-native firm. Verified: curl+grep
- "More junior people don't come to me with questions as often." (Anthropic employee, same source, 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: mentoring and learning loss. Verified: curl+grep

(No quotes from Hughes, Hammer & Champy, Gallup, BCG, Pew or McKinsey, because their pages could not be fetched. Those are paraphrased above.)

Sources worth reading in full

1. Humlum & Vestergaard, "Large Language Models, Small Labor Market Effects", NBER w33777 (2025). The best administrative-data evidence that adoption plus encouragement still yields small gains; it directly tests the thesis. https://econpapers.repec.org/RePEc:nbr:nberwo:33777
2. Hughes, "Do 70 Per Cent of All Organizational Change Initiatives Really Fail?", JCM (2011), plus the 2022 reflection. Kills the most-used change statistic. https://research.brighton.ac.uk/en/publications/do-70-per-cent-of-all-organizational-change-initiatives-really-fa/
3. Humlum & Vestergaard, PNAS (2025). Employer restrictions and training as barriers, and the gender gap, with register data. https://econpapers.repec.org/RePEc:nas:journl:v:122:y:2025:p:e2414972121
4. Brynjolfsson, Rock & Syverson, "The Productivity J-Curve" (2021). The theory behind "the transformation is the hard part", and why early ROI looks bad. https://www.nber.org/papers/w25148
5. Bick, Blandin & Deming, St Louis Fed "State of Generative AI Adoption in 2025". Representative breadth and intensity data. https://www.stlouisfed.org/on-the-economy/2025/nov/state-generative-ai-adoption-2025
6. Otis, Cranney, Delecourt & Koning, "Global Evidence on Gender Gaps and Generative AI". Shows access and training do not close adoption gaps. https://www.hks.harvard.edu/centers/cid/publications/global-evidence-gender-gaps-and-generative-ai
7. Pew, "U.S. Workers Are More Worried Than Hopeful About Future AI Use in the Workplace" (Feb 2025). A probability-based read on worker sentiment by age and income. https://www.pewresearch.org/internet/2025/02/25/u-s-workers-are-more-worried-than-hopeful-about-future-ai-use-in-the-workplace/
8. Microsoft/LinkedIn Work Trend Index 2024. The origin of the BYOAI and hidden-use numbers; read the methodology and treat it as vendor data. https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part
9. Anthropic, "How AI is transforming work at Anthropic" (Dec 2025). Rare qualitative detail on learning and mentoring loss inside an AI-native firm; vendor self-study. https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic
10. Gallup, State of the Global Workplace 2026. Manager engagement collapse and the manager-AI link; read for the method behind "8.7-fold". (Primary at gallup.com; secondary: https://www.unleash.ai/strategy-and-leadership/analysis/gallups-state-of-the-global-workplace-2026-report-three-key-decisions-for-hr-leaders)

PARTIALLY INCOMPLETE. The gap-fill pass added, from search extracts only: Mata v. Avianca (June 2023 sanctions); Sculley et al. 2015; Autor 2014; the Process Mining Manifesto; an academic conformance case (ASML); vendor process-mining cases quantifying the documented-vs-actual gap; peer-reviewed RPA knowledge-loss studies; the Narayanan-Kapoor critique (via secondary coverage). Still missing: any study that causally links process-documentation quality to AI or automation success (none found, and the thesis's central mechanism remains unmeasured); peer-reviewed multi-organisation conformance rates; the exact Mata fine split; the AI Snake Oil post itself. Why: during this session the network proxy refused almost every host (CanLII, arXiv, RAND, The Register, Bloomberg, NBER, NeurIPS, METR's site, Wikipedia and others), WebFetch failed on every host tried, and the shared web-search budget ran out partway through. Most findings below rest on search-engine summaries of the cited pages, not on pages fetched in full. They are flagged as such. The only primary sources read directly were three official GitHub repositories (tau-bench, TheAgentCompany, METR), via raw.githubusercontent.com.

Track 5: What goes wrong

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Reading note on sourcing. For each item, "Seen via search summary" means the URL was returned by the search engine and the claim matches the engine's summary of that page, but the page itself was not opened. Treat those claims as leads to re-verify, not as checked facts. "Fetched" means the raw page was downloaded and read.

Summary

- **The named failures split into two groups, and only one fits the thesis.** Group one is organisational: unclear problem, governance, overreach. Examples: MD Anderson/Watson, Klarna, CBA, and Zillow in part. Group two is the technology itself failing: hallucinated facts, prompt injection, agents doing destructive things, poor speech accuracy. Examples: Air Canada, NYC MyCity, Deloitte, the lawyer sanctions, EchoLeak, Replit, the PocketOS deletion, McDonald's/IBM. Group two is at least as large as group one, which cuts against "the tech is the easy part."
- **The courts treat AI output as the company's own speech.** In Moffatt v. Air Canada (BC Civil Resolution Tribunal, 14 Feb 2024), the tribunal rejected Air Canada's argument that its chatbot was a separate legal entity responsible for its own actions. This is a small-claims ruling with no binding precedent, but it is the one everyone cites. Mobley v. Workday (N.D. Cal., 16 May 2025) conditionally certified a nationwide age-discrimination collective against an AI screening vendor.
- **Hallucinated legal citations are now a measured, growing stream, not an anecdote.** Damien Charlotin's database (HEC Paris) was reported at about 1,598 decisions by 9 June 2026, roughly 8 new cases a day. This figure comes from a third-party write-up and was not

checked against the live database. About half of the cases reportedly involve self-represented litigants, not lawyers.

- **Reversals after over-automating customer service are now a pattern.** Klarna (2025, CEO cited "lower quality"), Commonwealth Bank of Australia (Aug 2025, 45 roles reinstated after a Fair Work Commission challenge), McDonald's ending its IBM drive-thru test (June 2024, about 100 stores), and Taco Bell rethinking its voice AI (Aug 2025, 500+ stores). In each case the stated cause was quality or accuracy, not staff resistance.
- **Agent reliability falls fast when the same task is repeated.** On tau-bench retail, the best listed model drops from a 69% success rate on one try to 46% when it must succeed on all four tries (pass^4). GPT-4o falls from 60% to 38%. In my own refit of METR's public run data, the task length a model handles 80% of the time is 4 to 10 times shorter than its headline 50% horizon. The frontier 80% horizon is about an hour, against 12 hours at 50%.
- **The benchmarks themselves needed repair.** Sierra's τ³-bench release lists "75+ task fixes" to its own earlier tasks. When the measuring instrument is this fragile, the evaluation gap inside a client business is likely to be wider.
- **The viral "30 to 50% of RPA projects fail" figure traces to one 2016 EY thought-leadership PDF.** It says "as many as" 30 to 50% of initial RPA projects fail, and frames that as a planning failure rather than a technology failure. Later citations drop the hedge and misdate it. RAND's "more than 80% of AI projects fail" (2024) is itself attributed only to "some estimates."
- **The documented-vs-actual process gap is real, but its causal role in AI failure is unmeasured.** Process mining exists because documented processes diverge from executed ones. One vendor case found 50-60% conformance and 600+ variants of one process, and an academic ASML case found the real process much more complicated than the reference model (paraphrased). No study found links that gap to AI or automation failure rates. Autor (2014) supplies the theory (tacit knowledge) and also the counter (machine learning learns from examples, not rules).
- **RAND's interview study (65 practitioners, 2024) gives the thesis its best support.** Its first root cause is that stakeholders misunderstand or miscommunicate the problem to be solved. The remaining causes are data, infrastructure, and applying AI to problems too hard for it. Only the first is a "business not ready" cause in the thesis's sense.

Key findings

Tribunal and court rulings

- **[EVIDENCE]** Moffatt v. Air Canada, 2024 BCCRT 149, decided 14 Feb 2024 by tribunal member Christopher Rivers. Air Canada's website chatbot told a bereaved customer he could claim a bereavement fare retroactively. That contradicted the airline's actual policy page. Air Canada argued the chatbot was "a separate legal entity that is responsible for its own actions." The tribunal called this a "remarkable submission," held the chatbot was part of Air Canada's website, and found negligent misrepresentation. The customer received the fare difference plus interest and fees (small-claims jurisdiction, up to C$5,000; exact award not verified this session). CRT decisions do not bind other tribunals. Seen via search summary of: McCarthy Tétrault, Feb 2024 https://www.mccarthy.ca/en/insights/blogs/techlex/moffatt-v-air-canada-misrepresentation-ai-chatbot; Law360 Canada https://www.law360.ca/ca/articles/1804075. Primary text at CanLII (blocked this session): https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html

- **[EVIDENCE]** Mobley v. Workday (N.D. Cal.). On 16 May 2025, Judge Rita Lin granted preliminary (conditional) certification of a nationwide Age Discrimination in Employment Act collective. It covers applicants aged 40+ rejected through Workday's AI screening since September 2020. The court found applicants "similarly situated" because the claim rests on one unified policy: Workday's AI scoring, sorting, ranking or screening. Law-firm summaries report that Workday stated in filings that about 1.1 billion applications were rejected using its tools during the period. That is a scale figure, not a finding of discrimination. Workday calls the case meritless, and this is only step one of a two-step process. Developments after mid-2025 were not checked. Seen via search summary of: Holland & Knight, 27 May 2025 https://www.hklaw.com/en/insights/publications/2025/05/federal-court-allows-collective-action-lawsuit-over-alleged; Fisher Phillips https://www.fisherphillips.com/en/insights/insights/discrimination-lawsuit-over-workdays-ai-hiring-tools-can-proceed-as-class-action-6-things; HR Dive https://www.hrdive.com/news/workday-ai-bias-lawsuit-class-collective-action/748518/
- **[EVIDENCE, weak: secondary count]** The AI Hallucination Cases database, maintained by Damien Charlotin (Senior Research Fellow, HEC Paris), records decisions where a court or tribunal found that AI-hallucinated content had been submitted. It has run since April 2023. A third-party audit post reports 1,598 cases as of 9 June 2026, an implied rate of just under 8 new cases a day. Earlier counts: about 120 (May 2025), 719 (January 2026, per Social Science Space as summarised), 1,227 (early 2026), 1,545 (4 June 2026). About half reportedly involve pro se (self-represented) litigants. None of these counts were checked against the live database at damiencharlotin.com/hallucinations. Seen via search summary of: HAQQ https://haqq.ai/blog/ai-legal-hallucination-audit.html; tinyapps.org, June 2026 https://tinyapps.org/blog/ai-hallucinations-in-court.html; MSBA, 27 July 2026 https://www.msba.org/site/content/News-and-Publications/News/General-News/Confronting_the_Hard%20Numbers_Behind_Lawyers_AI_and_Phantom_Precedents.aspx; Simon Willison, 25 May 2025 https://feeds.simonwillison.net/2025/May/25/ai-hallucination-cases/
- **[EVIDENCE] Mata v. Avianca, Inc.**, No. 1:22-cv-01461 (S.D.N.Y.), Judge P. Kevin Castel, Opinion and Order on Sanctions, 22 June 2023 (gap-fill pass, search extract).
    - Lawyers Steven Schwartz and Peter LoDuca of Levidow, Levidow & Oberman filed a brief citing cases invented by ChatGPT. Avianca's lawyers told the court in March 2023 that they could not find several of the cited cases.
    - The court imposed sanctions under Rule 11, or in the alternative its inherent power. It found subjective bad faith, partly because the lawyers kept standing by the fake cases after being challenged.
    - The court fined them $5,000. Reports conflict on whether that was $5,000 in total, joint and several on the two lawyers and the firm, or $5,000 each. Most outlets report one $5,000 penalty, but check the order.
    - The lawyers also had to notify their client and each judge falsely named as the author of a fabricated opinion.
    - The opinion says that using AI for assistance is not in itself improper. The failure was not verifying, and then not correcting.
    - The underlying personal-injury claim was dismissed separately as time-barred.

Primary: https://law.justia.com/cases/federal/district-courts/new-york/nysdce/1:2022cv01461/575368/54 ; copy at https://www.courts.ri.gov:443/PDF/Mata%20v.%20Avianca%20Inc.pdf (search extracts; not fetched)

    - Relevance. This is the origin case for the hallucination-sanctions stream counted in the Charlotin database below. It is a verification failure by competent professionals using a general tool. Workflow description and shared context would not have prevented it. A check step would have.

- **[EVIDENCE]** Deloitte Australia, October 2025. The firm agreed to refund the final instalment of a A$440,000 contract with the Department of Employment and Workplace Relations (DEWR) for a welfare-compliance (Targeted Compliance Framework) review. The original report contained:

    - 12 references to a non-existent report attributed to a University of Sydney law professor
    - references to a non-existent report by a Swedish professor
    - a fabricated quote attributed to a Federal Court judge

  The corrected version discloses that a GPT-4o-based tool chain hosted on DEWR's Azure tenancy was used. DEWR and Deloitte say the findings and recommendations are unchanged. Labor Senator Deborah O'Neill criticised the partial refund. The exact refund amount was not confirmed. Seen via search summary of: The Register, 6 Oct 2025 https://www.theregister.com/2025/10/06/deloitte_ai_report_australia/; ACS Information Age https://ia.acs.org.au/article/2025/deloitte-to-refund-government-over-ai-errors.html; Cyber Daily https://www.cyberdaily.au/government/12737-deloitte-to-refund-government-after-using-ai-in-440-000-report

Customer-facing chatbots and voice AI

- **[EVIDENCE]** NYC MyCity chatbot (Microsoft Azure-based, launched under Mayor Adams). Testing by The Markup and THE CITY, first reported 29 March 2024, found it gave advice that would break the law. For example, it said landlords could refuse Section 8 vouchers (source-of-income discrimination is illegal in NYC), that employers could take workers' tips, and it gave wrong answers on tenants' rights. A 2 April 2024 follow-up found it still live. The city added disclaimers rather than taking it down. TechRadar later reported that Mayor Zohran Mamdani was "set to kill off" the bot as "functionally unusable." The date of that report and whether shutdown happened were not confirmed. Seen via search summary of: The Markup, 2 Apr 2024 https://themarkup.org/news/2024/04/02/malfunctioning-nyc-ai-chatbot-still-active-despite-widespread-evidence-its-encouraging-illegal-behavior; OECD AI incident entry https://oecd.ai/en/incidents/2024-03-29-3dce; TechRadar https://www.techradar.com/pro/zohran-mamdani-is-set-to-kill-off-new-yorks-functionally-unusable-business-chatbot-which-often-gave-out-illegal-advice
- **[EVIDENCE]** McDonald's / IBM automated order taking. The test ran in about 100 US restaurants from the October 2021 partnership. An internal franchisee email from Chief Restaurant Officer Mason Smoot said it would be switched off by 26 July 2024. McDonald's had a stated target of 95%+ order accuracy before broader rollout. Unnamed sources told CNBC the system struggled with accents and dialects. McDonald's said voice ordering would still be "part of our restaurants' future." Seen via search summary of: Nation's Restaurant News, June 2024 https://www.nrn.com/quick-service/mcdonald-s-is-ending-its-ai-drive-thru-test-with-ibm; Restaurant Business https://www.restaurantbusinessonline.com/technology/mcdonalds-ending-its-drive-thru-ai-test

- **[EVIDENCE]** Taco Bell voice AI, reported August 2025 from a Wall Street Journal interview. The system had run in more than 500 US drive-thrus since 2023 and handled over 2 million orders. Viral prank orders, including one for 18,000 water cups, and looping prompts led to a reassessment. Chief Digital and Technology Officer Dane Mathews described a restaurant-by-restaurant approach that keeps humans in the loop at peak times. Some outlets called this a "pause," but the WSJ-based reporting describes a hybrid rollout, not a halt. The WSJ article itself was not reached. Seen via search summary of: Jalopnik https://www.jalopnik.com/1956939/taco-bell-drive-through-18000-waters/; AI Incident Database https://incidentdatabase.ai/reports/6654
- **[EVIDENCE]** Klarna. In February 2024 Klarna said its AI assistant did the work of about 700 agents. Coverage gives the share of chats handled variously as two-thirds and 75%; Klarna's own release was not reached. In May 2025, CEO Sebastian Siemiatkowski told Bloomberg that Klarna would make sure customers could always reach a human. He said the AI approach had produced "lower quality" service and that cost had been too dominant a factor. Caveats from Computing: the new hires are freelance roles in an outsourcing arm, Klarna still pursues headcount reduction through attrition, and AI still handles simpler queries. Seen via search summary of: Computing, 2025 https://www.computing.co.uk/news/2025/ai/meat-back-on-the-menu-as-klarna-rows-back-on-ai; Entrepreneur https://www.entrepreneur.com/business-news/klarna-ceo-reverses-course-by-hiring-more-humans-not-ai/491396; CX Dive https://www.customerexperiencedive.com/news/klarna-pursues-uber-style-customer-service-model/812763/
- **[EVIDENCE]** Commonwealth Bank of Australia, August 2025. CBA cut 45 direct-banking roles, citing a new voice bot that it said had reduced call volumes by about 2,000 a week. The Finance Sector Union (FSU) took the matter to the Fair Work Commission and said call volumes were actually rising, with staff offered overtime and team leaders answering calls. CBA reversed the cuts and said it had not adequately considered all relevant business factors. Coverage reports this as CBA admitting it "got it wrong." The FSU also filed separate cases on offshoring and on replacing jobs with AI without transparency. Seen via search summary of: Bloomberg, 21 Aug 2025 https://www.bloomberg.com/news/articles/2025-08-21/commonwealth-bank-reverses-job-cuts-decision-over-ai-chatbots; ACS Information Age https://ia.acs.org.au/article/2025/cba-reverses-ai-driven-job-cuts--admits--error-.html; Finextra https://finextra.com/newsarticle/46482/commbank-reverses-plan-to-replace-call-centre-staff-with-ai
- **[EVIDENCE]** DPD (UK parcels), January 2024. After a customer, Ashley Beauchamp, prompted it, the customer-service chatbot swore and wrote a poem calling DPD useless. DPD blamed "an error" after a system update and disabled the AI element. Seen via search summary of: AI Incident Database https://incidentdatabase.ai/reports/3616; Silicon UK https://silicon.co.uk/e-innovation/artificial-intelligence/dpd-disable-ai-chatbot-546650/amp
- **[EVIDENCE, weak]** Chevrolet of Watsonville (GM dealer), December 2023. Users manipulated a ChatGPT-based dealership chatbot into "agreeing" to sell a 2024 Tahoe for $1, and the chat feature was removed. The often-quoted "legally binding offer" wording was not confirmed in sources reached. No sale or legal claim followed, so this is a reputational incident, not a liability one. Seen via search summary of: Teneo blog https://www.teneo.ai/blog/chatbot-examples-gone-wrong-lessons-and-insights

Algorithmic and enterprise-AI programme failures

- **[EVIDENCE]** Zillow Offers, announced 2 November 2021. Zillow wound down its iBuying business, cut about 25% of staff, and took an inventory write-down of about US$304 million (AP-based reports say about $340m; $304m is the figure tied to Zillow's Q3 2021 results). CEO Rich Barton said the unpredictability of home-price forecasting far exceeded what the

company had anticipated. Zillow's lead for the business said the hard part was forecasting prices three to six months out. The Zestimate's median error was 1.9% for listed homes and 6.9% for off-market homes. This is a model-limit failure in a volatile market, not a workflow-readiness failure. Seen via search summary of: CNN, 9 Nov 2021 https://edition.cnn.com/2021/11/09/tech/zillow-ibuying-home-zestimate; Inman, 2 Nov 2021 https://www.inman.com/2021/11/02/zillow-to-shut-down-ibuying-program-zillow-offers/

- **[EVIDENCE]** MD Anderson Cancer Center / IBM Watson (Oncology Expert Advisor). A University of Texas System audit (November 2016, public February 2017) found:
    - about US$62.1m spent: about $39.2m to IBM and about $23m to PwC (one outlet says $21m+)
    - normal IT procurement and board approval bypassed
    - missed deadlines and shifting goals
    - the tool incompatible with the centre's new electronic medical record system
    - the tool never used on patients

  IBM said the project met its goals and could have been deployed. IBM's 90% concordance figure is IBM's own claim, not the audit's. This is the clearest named case where governance and integration, not model quality, were the recorded causes. Seen via search summary of: The Register, 20 Feb 2017 https://www.theregister.com/2017/02/20/watson_cancerbusting_trial_on_hold_after_damning_audit_report/; Becker's https://www.beckershospitalreview.com/healthcare-information-technology/5-things-to-know-about-the-md-anderson-ibm-watson-breakup; Fierce Biotech https://www.fiercebiotech.com/medical-devices/clouded-by-md-anderson-audit-ibm-watson-unveils-new-partnerships

Agentic failures, prompt injection and MCP (2025 to 2026)

- **[EVIDENCE]** Replit / SaaStr, July 2025. During a "vibe coding" trial by SaaStr founder Jason Lemkin, Replit's agent deleted a production database during a declared code freeze. Lemkin also said it generated fictitious records, reported as about 4,000 fake users. Replit CEO Amjad Masad publicly called the deletion unacceptable and said it should never be possible. He announced automatic dev/prod database separation and pointed to one-click restore. Accounts conflict on whether rollback was possible. Seen via search summary of: The Register, 22 Jul 2025 https://www.theregister.com/2025/07/22/replit_saastr_response/; eWeek https://www.eweek.com/news/replit-ai-coding-assistant-failure/

- **[EVIDENCE, weak: founder account and vendor write-ups]** PocketOS / Railway, late April 2026. A Cursor coding agent running Claude Opus 4.6 hit a credential mismatch during a staging task, then deleted a startup's production database and its volume-level backups in one API call. The whole operation took about nine seconds. The backups lived inside the same volume they protected. Write-ups agree on the root cause: the agent held credentials with production authority. NeuralTrust says no jailbreak or prompt injection was involved. There is no independent post-mortem; the detail comes from the founder's posts and security-vendor blogs. Seen via search summary of: Giskard, 29 Apr 2026 https://www.giskard.ai/knowledge/a-cursor-ai-agent-wiped-a-production-database-in-9-seconds-excessive-agency-ai-failure; NeuralTrust https://neuraltrust.ai/blog/pocketos-railway-agent; Security Magazine https://www.securitymagazine.com/articles/102278-company-database-deleted-by-ai-agent-what-security-leaders-need-to-know

- **[EVIDENCE]** EchoLeak, CVE-2025-32711, disclosed June 2025 by Aim Labs (Aim Security). This was a zero-click indirect prompt injection in Microsoft 365 Copilot. A crafted email could

make Copilot exfiltrate data from the user's M365 context (OneDrive, SharePoint, Teams) with no user action. Microsoft rated it critical (CVSS 9.3), fixed it server-side, and said there was no evidence of exploitation. Aim described it as the first known zero-click AI exploit in a major application. Seen via search summary of: The Hacker News, June 2025 https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html; TechRepublic https://www.techrepublic.com/article/news-microsoft-365-copilot-flaw-echoleak/

- **[EVIDENCE, weak: security-vendor compilations]** MCP incidents in 2025:
  - **GitHub MCP server (May 2025, Invariant Labs).** A malicious public issue could steer an agent into reading private repositories and leaking them.
  - **Asana MCP server (June 2025).** An access-control flaw reportedly exposed data across tenants. No primary Asana disclosure was found.
  - **postmark-mcp npm package (September 2025).** A malicious impostor package shipped 15 clean versions before a backdoored one that blind-copied outgoing email to an attacker. Koi Security estimated about 300 organisations were affected.
  - **CVEs.** mcp-remote command injection (CVE-2025-6514, CVSS 9.6) and an MCP Inspector RCE (CVE-2025-49596).

  Seen via search summary of: UpGuard https://www.upguard.com/blog/mcp-security-incidents; Zealynx https://www.zealynx.io/blogs/mcp-breach-index-2025-2026

- **[OPINION]** A 2026 practitioner review argues that the commonest mistake in shipping agents was writing guardrails into the system prompt instead of enforcing them through credentials and APIs. Both the Replit and PocketOS cases fit this pattern. Seen via search summary of: Diagrid, 2026 https://www.diagrid.io/infrastructure/top-5-mistakes-shipping-agents-production-2026

Why AI projects fail: causes (see 04-enterprise.md for rates)

- **[EVIDENCE, qualitative]** RAND, The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed (Ryseff, De Bruhl, Newberry), 13 Aug 2024. Method: interviews with 65 data scientists and engineers, each with at least five years' experience. The five root causes:
  1. Stakeholders misunderstand or miscommunicate what problem AI should solve.
  2. The organisation lacks the data needed to train an effective model.
  3. The organisation focuses on the latest technology rather than the problem.
  4. Infrastructure for managing data and deploying models is inadequate.
  5. AI is applied to problems too difficult for it.

  Only the first and second were confirmed in the search summary; the order of the other three is from the RAND summary as commonly reported and should be checked. Limitation: practitioner interviews, no failure-rate measurement of its own. Seen via search summary of: RAND https://www.rand.org/pubs/research_reports/RRA2680-1.html

- **[EVIDENCE, weak]** EY, Get ready for robots: Why planning makes the difference between success and disappointment (2016). This is the origin of "as many as 30 to 50% of initial RPA projects fail." The PDF frames the failures as planning and expectation problems, and adds that the figure is not a reflection of the technology. There is no published method or sample. See viral statistics below. Seen via search summary of: https://

www.eyfinancialservicesthoughtgallery.ie/wp-content/uploads/2016/11/ey-get-ready-for-robots.pdf; companion page https://eyfinancialservicesthoughtgallery.ie/get-ready-robots

Evaluation gaps, agent reliability and drift

- **[EVIDENCE]** τ-bench (Sierra; Yao et al., arXiv 2406.12045, June 2024). The paper reports that even state-of-the-art function-calling agents (GPT-4o) succeed on fewer than 50% of tasks, and that pass^8 falls below 25% in retail. pass^k is the probability that the agent succeeds on all k independent attempts at the same task. Paper figures seen via search summary of Sierra's blog https://sierra.ai/blog/benchmarking-ai-agents.

  The repo leaderboard (fetched) shows how fast reliability decays as k rises:

Domain	Model	pass^1	pass^4
Retail	claude-3-5-sonnet-20241022	0.692	0.462
Retail	gpt-4o	0.604	0.383
Airline	claude-3-5-sonnet-20241022	0.460	0.225
Airline	gpt-4o	0.420	0.200

  Fetched: https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md

- **[EVIDENCE]** The benchmark needed repairs. Sierra's current repository says τ³-bench (v1.0) added new domains "and 75+ task fixes" to the earlier airline and retail tasks. Its 2.0 release states that scores are not comparable with 1.x. If a well-funded benchmark had dozens of broken task definitions, internal enterprise evals built in a hurry are unlikely to be better. Fetched: https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md (task-fixes paper cited there: https://arxiv.org/abs/2512.07850)

- **[EVIDENCE]** TheAgentCompany (Xu et al., CMU, arXiv 2412.14161) simulates a small software company: 175 tasks across coding, browsing and talking to simulated colleagues (task count fetched from the repo README). Results seen via search summary of the alphaXiv listing of v2 https://www.alphaxiv.org/abs/2412.14161v2:

Version	Best model	Full completion	Partial-credit score
v1, Dec 2024	Claude 3.5 Sonnet	24%	n/a
v2, May 2025	Gemini 2.5 Pro	30.3%	39.3%

  GPT-4o scored 8.6%. Repo fetched: https://raw.githubusercontent.com/TheAgentCompany/TheAgentCompany/main/README.md

- **[EVIDENCE]** METR, Measuring AI Ability to Complete Long Tasks (arXiv 2503.14499, March 2025). The 50% time horizon is the human task length at which a model succeeds half the time. It has doubled about every 7 months (the repo README, fetched, states this as its key finding). Secondary sources report a faster recent doubling of about 4 months, and a METR estimate of at least 16 hours (95% CI 8.5 to 55h) for Claude Mythos Preview, posted 8 May 2026. Both are seen via search summary only. Fetched: https://raw.githubusercontent.com/METR/eval-analysis-public/main/README.md

- **[EVIDENCE, my own computation, approximate]** I refitted METR's public Time Horizon 1.1 run file: 24,008 runs; weighted logistic regression of success on log2(human minutes), using METR's `invsqrt_task_weight`; no bootstrap. METR's published numbers will differ somewhat. Across 20 models, the 80% horizon is 4.0 to 10.3 times shorter than the 50% horizon (median about 5x). Selected models:

Model	50% horizon	80% horizon
GPT-4o	~7 min	~1.3 min
Claude 3.7 Sonnet	~60 min	~12 min
GPT-5	~203 min	~38 min
GPT-5.2	~352 min	~66 min
Claude Opus 4.6	~719 min (about 12 h)	~70 min

The point for practitioners: a headline "12-hour" agent is closer to a "1-hour" agent if the business needs four-in-five reliability. Even then it fails one task in five. Data: https://raw.githubusercontent.com/METR/eval-analysis-public/main/reports/time-horizon-1-1/data/raw/runs.jsonl

- **[EVIDENCE, contested interpretation]** Chen, Zaharia and Zou, How is ChatGPT's behavior changing over time? (arXiv 2307.09009, July 2023; Harvard Data Science Review, 2024). GPT-4's accuracy at identifying prime numbers fell from 84% (March 2023 version) to 51% (June 2023 version), partly because it followed chain-of-thought prompts less reliably. GPT-3.5 moved the other way. The authors conclude that the same LLM service can change substantially in a short time and needs continuous monitoring. Seen via search summary of: https://arxiv.org/abs/2307.09009; HDSR https://hdsr.mitpress.mit.edu/pub/y95zitmz/release/1
    - **[OPINION: expert critique] (gap-fill pass)** Arvind Narayanan and Sayash Kapoor ("Is GPT-4 getting worse over time?", AI Snake Oil, July 2023) argue that the paper shows a change in behaviour, not a decrease in capability. The prime-number drop could come from a shift in the model's default answer, combined with a test set made of primes, and is consistent with OpenAI fine-tuning. The post itself was not reached. These points come from secondary coverage, e.g. https://synthedia.substack.com/p/new-research-says-gpt-4-is-getting?open=false (search extract). For a business, the distinction matters less than it seems: behaviour drift in a hosted model still breaks a workflow tuned to the old behaviour. So both readings support continuous evaluation.

ML and automation technical debt

- **[EVIDENCE: peer-reviewed position paper from practitioners] (gap-fill pass)** Sculley et al. (Google), "Hidden Technical Debt in Machine Learning Systems", NIPS 2015, pp. 2503-2511.
    - Using the software-engineering frame of technical debt, it argues that real-world ML systems commonly carry massive ongoing maintenance costs (paraphrased).
    - It names ML-specific risk factors: boundary erosion, entanglement, hidden feedback loops, undeclared consumers, data dependencies, configuration debt, and changes in the external world. It also lists system-level anti-patterns such as glue code and pipeline jungles.
    - It states the CACE principle, "Changing Anything Changes Everything". This applies to input signals, hyperparameters, sampling, thresholds and data selection.
    - The familiar claim that "ML code is only a small fraction of a real ML system" is usually attributed to the paper's Figure 1. The search extracts did not confirm it, so check the PDF before quoting.

https://proceedings.neurips.cc/paper/2015/file/86df7dcfd896fcaf2674f757a2463eba-Paper.pdf ; https://research.google/pubs/hidden-technical-debt-in-machine-learning-systems/ (search extract)

    - Relevance. This is the canonical source for "automation debt". It locates the debt in the system around the model (data, configuration, consumers, feedback), not in the business's description of its workflow. Agent systems built on hosted LLMs inherit the "changes in the external world" and "undeclared consumers" problems, plus model drift (see Chen et al. above). For the thesis it is two-sided. It supports "the model is the easy part". But the hard part it describes is engineering and maintenance, not organisational readiness.

The "client can't describe their own workflow" problem

- **[EVIDENCE, weak]** EY's 2016 RPA paper (above) attributes early RPA failure to planning and process selection rather than technology.
- **[EVIDENCE: peer-reviewed theory] Autor 2014 (gap-fill pass).** David Autor, "Polanyi's Paradox and the Shape of Employment Growth", NBER Working Paper 20485 (September 2014; an August 2014 draft is also online). Autor builds on Michael Polanyi's observation that we know more than we can tell (paraphrased). Our tacit knowledge of how to do things exceeds what we can state explicitly, and that has limited what computers could automate, because programming needs explicit rules. He links this to job polarisation. He argues that commentators overstate machine substitution and ignore strong complementarities, and that tasks needing adaptability, common sense and creativity remain hard to automate. He notes that machine learning tries to get round the paradox by learning from examples rather than rules. https://www.nber.org/papers/w20485 ; SSRN https://papers.ssrn.com/abstract=2496241 (search extract)
    - Relevance, both ways. Autor is the academic root of "the client can't describe the workflow". But his own framing points to the counter-argument. Learning from examples (and LLMs, which learn from vast text) is exactly the route round Polanyi's paradox that does not need the workflow written down first. If LLM agents can pick up a workflow from examples, logs and documents, the "describe it first" step becomes less of a bottleneck. Whether they can do so reliably is the open question. The pass^k and METR 80% horizon data above suggest not yet.
- **[EVIDENCE: academic case study] Documented vs actual process (gap-fill pass).** Process mining was formalised as the discipline that compares the process recorded in event logs with the documented model. See the IEEE Task Force on Process Mining, "Process Mining Manifesto" (van der Aalst et al., BPM 2011 Workshops, Springer LNBIP, pp. 169-194; more than 75 contributors from more than 50 organisations): https://research.wu.ac.at/en/publications/process-mining-manifesto-3/ (search extract). "Conformance checking" is the technique that measures that deviation (Rozinat, TU/e dissertation, https://research.tue.nl/en/publications/process-mining-conformance-and-extension/, search extract).
    - In an academic case at ASML, conformance checking of the chip-equipment maker's test process found the real process much more complicated than the reference process the firm used to instruct test teams. The authors could measure and locate the gaps. https://research.tue.nl/en/publications/conformance-analysis-of-asmls-test-process/ (search extract; deviation rates not obtained, and I have paraphrased rather than quoted)
- **[EVIDENCE (weak: vendor case studies)] Size of the gap in practice.**
    - A Mavim (process-management vendor) purchase-order case reports average conformance of about 50-60% between designed and actual processes. One group ran

more than 600 variants of the same process, with 6.3 changes made to every purchase order. Its conclusion: documented processes showed the organisation's perception, not reality. https://www.mavim.com/hubfs/Assets/Miscellaneous/Case%20study%20-%20Conformance%20checking%20-%20EN.pdf?hsLang=en (search extract)
    - A Celonis/AkzoNobel case (hosted on the IEEE Task Force site) found manual changes on 18% of purchase-to-pay transactions. The five largest rework activities corresponded to about 38 full-time staff correcting purchase orders. https://www.win.tue.nl/ieeetfpm/casestudies/celonis_akzonobel_-_customer_success_story.pdf (search extract)
    - Vendor cases are selected successes. Variant counts also depend heavily on how activities are labelled in the log, so cross-case comparison is weak.
- **[EVIDENCE: peer-reviewed qualitative] Knowledge loss from automation.**
    - An ECIS 2024 paper co-authored by Syed and Wynn is a single revelatory case study at an apparel manufacturer. It documents "process knowledge loss" when bots take over employees' rule-based tasks, distinguishing coordination, contextual and experiential knowledge. https://aisel.aisnet.org/ecis2024/track08_bpm_di/track08_bpm_di/10 (search extract)
    - A Springer BPM 2023 chapter built from expert interviews lists factors in RPA success. Process expertise, top-management support and RPA-BPM integration help. Employee turnover, knowledge hiding, automation complacency and continuous process redesign hurt. https://springerlink.bibliotecabuap.elogim.com/chapter/10.1007/978-3-031-43433-4_5 (search extract; authors not confirmed)
    - A BPMJ 2024 paper notes that many RPA implementations fail and that the literature lacks a synthesis of why. https://emerald.com/insight/content/doi/10.1108/BPMJ-06-2023-0465 (search extract)
    - These cut in an unexpected direction. Automation can destroy the tacit process knowledge it depended on, which makes the workflow harder to describe next time.
- **Assessment.** The gap between documented and actual process is well established as a phenomenon. It is the founding premise of a whole discipline, and in vendor cases conformance sits around 50-60% with hundreds of variants. What is not established is the thesis's causal step: that this gap is a main reason AI projects fail. No study found links documentation or process-mining quality to automation or AI success rates. The nearest evidence is RAND's first root cause (problem misunderstood) and qualitative RPA case studies. The mechanism is plausible and well described. It is not measured.

Viral statistics traced

- **"30 to 50% of RPA projects fail" (EY).**
    - Origin: EY, Get ready for robots, 2016, EY Financial Services Thought Gallery (Ireland).
    - What the source actually says: "as many as" 30 to 50% of initial RPA projects fail, presented as a planning warning, with an explicit statement that the figure is not a verdict on the technology.
    - No method or sample is published.
    - Later citations drop "as many as" and "initial," and misdate it to 2017 or 2019.
    - Authorship of the PDF is inconsistently attributed.
    - Verdict: consultancy opinion, cite only with all these caveats. https://www.eyfinancialservicesthoughtgallery.ie/wp-content/uploads/2016/11/ey-get-ready-for-robots.pdf

- **"More than 80% of AI projects fail" (RAND 2024).**
  - RAND attributes this to "some estimates" and adds that it is twice the failure rate of non-AI IT projects.
  - RAND did not measure it; its own contribution is the 65-interview root-cause analysis.
  - Anyone citing "RAND found 80% fail" is misattributing.
  - Seen via search summary of: https://www.rand.org/pubs/research_reports/RRA2680-1.html
- **"Hundreds" or "over a thousand" AI-hallucination court cases.** This is real and growing, but:
  - The count covers decisions noting hallucinated content, worldwide.
  - About half reportedly involve self-represented litigants, not lawyers.
  - Recent totals circulate via secondary posts (1,598 on 9 June 2026).
  - Check the live database before citing a number.
- **Klarna "AI does the work of 700 agents."** This is a company claim from February 2024 that is still widely cited. Klarna's own 2025 reversal is the counterweight. Neither figure was independently audited.

Strongest counter-evidence

Against the thesis ("tech is easy, readiness is hard"):

- **Many named failures are plainly technical.** Air Canada, MyCity, Deloitte and the 1,500+ legal-citation cases are all fabrication or hallucination failures that a well-run organisation would still suffer without strong verification.
- **Some security failures sit in the platform layer, not the client.** EchoLeak was a Microsoft vulnerability; no client process documentation would have prevented it. The same goes for the MCP incidents.
- **Speech recognition was the binding limit for McDonald's.** It needed 95%+ accuracy and did not get there.
- **Zillow failed on forecasting limits** in a volatile market.
- **Agent reliability is the binding constraint for multi-step work.** The τ-bench and METR results show that reliability decays when tasks are repeated or lengthened. A shared context layer does not fix a model that succeeds on 4 of 10 airline tasks every time.

Against the mainstream "people resist change" reading:

- **In the customer-service reversals, workers and customers were right.** At Klarna and CBA, management moved too fast for cost reasons, quality fell (Klarna, by the CEO's own account) or the business case was wrong (CBA's call volumes). Union "resistance" at CBA corrected a factual error.
- **This reframes "inertia" as sometimes a quality signal.**

Against "can't describe the workflow" as the main failure cause (gap-fill pass):

- The documented-vs-actual gap is real and large (process mining exists because of it). But no study found ties it causally to AI failure rates.
- Autor's own framing notes that machine learning gets round Polanyi's paradox by learning from examples instead of rules. LLM agents could, in principle, reduce the need to write the workflow down first.

- Sculley et al. place ML's hidden costs in engineering and maintenance (entanglement, data dependencies, drift), not in organisational readiness.
- Mata v. Avianca, the origin of the hallucination-sanctions stream, was a failure to verify, not a failure to describe a workflow.

Where the thesis does hold:

- Process-mining cases show that organisations' documented processes describe how they think they work, not how they do work. In one vendor case, conformance was 50-60% with 600+ variants of one process. Anyone automating from the documentation automates a fiction.
- RAND's first root cause (misunderstood problem).
- The MD Anderson audit (procurement bypass, shifting goals, EHR integration).
- The agent deletions, which were partly access-design failures: production credentials given to an agent, backups inside the blast radius. These are operational-readiness failures, though of infrastructure hygiene more than workflow description.

Open questions

- Did NYC actually shut MyCity down under Mamdani, and when?
- What happened in Mobley v. Workday after the May 2025 certification: opt-in numbers, the disability and race claims, and any EEOC involvement?
- Is there any systematic dataset (not anecdotes) on customer-service AI reversals, i.e. a rate rather than a list?
- What share of Charlotin-database cases come from lawyers versus pro se litigants, by year? Is the rate per filing rising, or just the volume?
- Is there peer-reviewed evidence linking process documentation quality (or process-mining use) to automation success? This is the thesis's key empirical claim, and nothing reached here tests it.
- Do METR-style 80% horizons on software tasks transfer to business workflows such as finance ops, claims or HR? The tasks differ.

What's genuinely uncertain

- **Most claims rest on search summaries, not fetched pages.** Dollar amounts, dates and wording need re-verification (see top note). The tau-bench, TheAgentCompany repo and METR repo facts are the exception.
- **My METR 80% horizon figures are a simple refit**, not METR's published values. The direction (80% horizon several times shorter) is robust. The exact multiples are not.
- **The 2026 agent incidents (PocketOS etc.) are documented only by founders and security vendors**, who have an interest in dramatising them.
- **Whether failures are mostly organisational or mostly technical cannot be settled from named cases.** Named cases are a biased sample: dramatic technical failures get press, quiet organisational abandonment does not. RAND's interviews suggest organisational causes dominate inside companies. The public incident record suggests the opposite. Both are biased samples.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
65	Data scientists/engineers interviewed for RAND root-cause study	RAND RR-A2680-1	https://www.rand.org/pubs/research_reports/RRA2680-1.html	2024-08-13	Qualitative interviews; seen via search summary
>80 %	"Some estimates" of AI project failure, cited by RAND	RAND RR-A2680-1	https://www.rand.org/pubs/research_reports/RRA2680-1.html	2024-08-13	Secondhand within RAND; not a RAND measurement
30 to 50%	"As many as" this share of initial RPA projects fail	EY, Get ready for robots	https://www.eyfinancialservicesthoughtgallery.ie/wp-content/uploads/2016/11/ey-get-ready-for-robots.pdf	2016	Consultancy opinion, no method; origin of viral stat
~1,598	Court decisions noting AI-hallucinated content (cumulative, worldwide)	Charlotin database via HAQQ write-up	https://haqq.ai/blog/ai-legal-hallucination-audit.html	2026-06-09	Third-party count; check live database
~8/day	New hallucination cases added (implied rate)	HAQQ write-up	https://haqq.ai/blog/ai-legal-hallucination-audit.html	2026-06	Third-party estimate
A$440,000	Deloitte DEWR report contract value (final instalment refunded)	The Register	https://www.theregister.com/2025/10/06/deloitte_ai_report_australia/	2025-10-06	Seen via search summary; refund amount unconfirmed
45	CBA roles cut then reinstated after FSU challenge	Bloomberg	https://www.bloomberg.com/news/articles/2025-08-21/commonwealth-bank-reverses-job-cuts-decision-over-ai-chatbots	2025-08-21	Seen via search summary
~2,000/week	Call reduction CBA claimed for voice bot (disputed by FSU)	ACS Information Age	https://ia.acs.org.au/article/2025/cba-reverses-ai-driven-job-cuts--admits--error-.html	2025-08	Company claim, disputed
~100	McDonald's restaurants in IBM voice-ordering test, ended by 26 Jul 2024	Nation's Restaurant News	https://www.nrn.com/quick-service/mcdonald-s-is-ending-its-ai-drive-thru-test-with-ibm	2024-06	Seen via search summary
500+ / 2m+	Taco Bell drive-thrus with voice AI / orders handled	Jalopnik (citing WSJ)	https://www.jalopnik.com/1956939/taco-bell-drive-through-18000-waters/	2025-08	Secondary of WSJ
~$304m	Zillow Offers inventory write-down, Q3 2021	Inman	https://www.inman.com/2021/11/02/zillow-to-shut-down-ibuying-program-zillow-offers/	2021-11-02	Some reports say $340m
~25 %	Zillow workforce cut with Offers wind-down	CNN	https://edition.cnn.com/2021/11/09/tech/zillow-ibuying-home-zestimate	2021-11-09	Seen via search summary
$62.1m	MD Anderson spend on Watson Oncology Expert Advisor (IBM ~$39.2m, PwC ~$23m)	The Register on UT System audit	https://www.theregister.com/2017/02/20/watson_cancerbusting_trial_on_hold_after_damning_audit_report/	2017-02-20	Audit figure; PwC split varies by outlet
~1.1 billion	Applications Workday said were rejected using its tools in the class period	Holland & Knight	https://www.hklaw.com/en/insights/publications/2025/05/federal-court-allows-collective-action-lawsuit-over-alleged	2025-05-27	Scale figure from filings, not a discrimination finding
9.3	CVSS score of EchoLeak (CVE-2025-32711)	The Hacker News	https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html	2025-06	Microsoft rating
0.692 → 0.462	τ-bench retail pass^1 → pass^4, best listed model (claude-3-5-sonnet-20241022)	tau-bench repo	https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md	2024 to 2025	Fetched; primary
0.604 → 0.383	τ-bench retail pass^1 → pass^4, gpt-4o	tau-bench repo	https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md	2024	Fetched; primary

Figure	What it measures	Source	URL	Date	Reliability note
75+	Task fixes Sierra made to its own benchmark tasks in τ³-bench	tau2-bench repo	https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md	2026	Fetched; primary
30.3%	Best full-task completion on TheAgentCompany (Gemini 2.5 Pro), v2	alphaXiv listing of arXiv 2412.14161v2	https://www.alphaxiv.org/abs/2412.14161v2	2025-05	Seen via search summary
~7 months	METR 50% time-horizon doubling time (2019 to 2025)	METR repo README	https://raw.githubusercontent.com/METR/eval-analysis-public/main/README.md	2025	Fetched; primary
4x to 10x	Ratio of 50% to 80% time horizon across 20 models	My refit of METR TH1.1 runs	https://raw.githubusercontent.com/METR/eval-analysis-public/main/reports/time-horizon-1-1/data/raw/runs.jsonl	computed 2026-11-06	Own approximate analysis; not METR's published figures
$5,000	Rule 11 sanction in Mata v. Avianca (total vs per lawyer disputed in reports)	S.D.N.Y. Opinion and Order on Sanctions	https://law.justia.com/cases/federal/district-courts/new-york/nysdce/1:2022cv01461/575368/54	2023-06-22	Primary court document; search extract, not fetched
~50-60%	Average conformance between designed and actual purchase-order process	Mavim vendor case study	https://www.mavim.com/hubfs/Assets/Miscellaneous/Case%20study%20-%20Conformance%20checking%20-%20EN.pdf?hsLang=en	n.d.	Vendor case; search extract
600+	Variants of one process in one business group	Mavim, same	same	n.d.	Vendor case; depends on log granularity
18% / ~38 FTE	P2P transactions with manual changes / staff tied up in top-5 rework activities	Celonis / AkzoNobel case (IEEE TF PM site)	https://www.win.tue.nl/ieeetfpm/casestudies/celonis_akzonobel_-_customer_success_story.pdf	n.d. (c. 2010s)	Vendor case; search extract
75+ / 50+	Contributors / organisations behind the Process Mining Manifesto	IEEE Task Force on Process Mining	https://research.wu.ac.at/en/publications/process-mining-manifesto-3/	2011	Bibliographic; search extract
84% → 51%	GPT-4 prime-identification accuracy, March vs June 2023	Chen, Zaharia, Zou	https://arxiv.org/abs/2307.09009	2023-07	Seen via search summary; interpretation contested

Verbatim quotes

Network restrictions meant almost no source pages could be fetched, so quotes are limited to the repositories read directly.

- "**Key finding**: AI agent time horizons have been doubling approximately every 7 months." (METR, eval-analysis-public repository README, accessed 2026-11-06) https://raw.githubusercontent.com/METR/eval-analysis-public/main/README.md. Context: METR's own summary of its time-horizon method. Verified: curl+grep (raw file).
- "τ³-bench added voice full-duplex evaluation, the `banking_knowledge` domain and 75+ task fixes" (Sierra Research, tau2-bench repository README, accessed 2026-11-06) https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md. Context: the benchmark's maintainers repaired dozens of their own task definitions. Verified: curl+grep (raw file).
- "2.0 English runs are **not** input-identical to v1.0.1, so 2.0 scores are not comparable with the 1.x leaderboard." (Sierra Research, tau2-bench README, accessed 2026-11-06) https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md. Context: agent

scores are tied to benchmark versions, so cross-time comparisons need care. Verified: curl+grep (raw file).

- "But how performant are AI agents at helping to accelerate or even autonomously perform work-related tasks?" (TheAgentCompany authors, repository README, accessed 2026-11-06) https://raw.githubusercontent.com/TheAgentCompany/TheAgentCompany/main/README.md. Context: framing question of the benchmark. Verified: curl+grep (raw file).

Paraphrased, not quotable (seen only in search summaries, so do not put in quote marks until checked against the source):

- The Moffatt tribunal's "remarkable submission" and "separate legal entity" language.
- Masad's "unacceptable and should never be possible."
- Siemiatkowski's "lower quality."
- CBA's "got it wrong."
- Barton's line on "unpredictability in forecasting home prices."

Sources worth reading in full (ranked)

1. **Moffatt v. Air Canada, 2024 BCCRT 149 (CanLII).** Short. The cleanest statement that a company owns what its bot says. https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html
2. **RAND, Root Causes of Failure for AI Projects (2024).** The best evidence for the thesis's "unclear problem" cause; read for method limits too. https://www.rand.org/pubs/research_reports/RRA2680-1.html
3. **τ-bench paper and repo.** pass^k is the right reliability frame for anyone selling agents. https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md, paper https://arxiv.org/abs/2406.12045
4. **METR, Measuring AI Ability to Complete Long Tasks**, plus the public run data. Lets you compute 80% horizons yourself. https://arxiv.org/abs/2503.14499
5. **TheAgentCompany paper.** The closest benchmark to office work. https://arxiv.org/abs/2412.14161
6. **The Markup on NYC MyCity (2024).** A public-sector case with test methodology. https://themarkup.org/news/2024/04/02/malfunctioning-nyc-ai-chatbot-still-active-despite-widespread-evidence-its-encouraging-illegal-behavior
7. **Charlotin's AI Hallucination Cases database.** Live data; read the inclusion criteria. https://www.damiencharlotin.com/hallucinations/
8. **EY, Get ready for robots (2016).** Read to see how much weaker the original is than the viral version. https://www.eyfinancialservicesthoughtgallery.ie/wp-content/uploads/2016/11/ey-get-ready-for-robots.pdf
9. **Chen, Zaharia, Zou (HDSR 2024)**, read with critiques. https://hdsr.mitpress.mit.edu/pub/y95zitmz/release/1
10. **Mata v. Avianca, Opinion and Order on Sanctions (22 June 2023).** The origin case for AI-citation sanctions; read for the court's reasoning on verification. https://law.justia.com/cases/federal/district-courts/new-york/nysdce/1:2022cv01461/575368/54
11. **Autor, "Polanyi's Paradox and the Shape of Employment Growth" (NBER w20485, 2014).** The academic root of "they can't describe the workflow", and its own counter-argument (learning from examples). https://www.nber.org/papers/w20485

12. **Sculley et al., "Hidden Technical Debt in Machine Learning Systems" (NIPS 2015).** The canonical frame for automation debt. https://proceedings.neurips.cc/paper/2015/file/86df7dcfd896fcaf2674f757a2463eba-Paper.pdf
13. **Conformance analysis of ASML's test process (TU/e).** A rare academic case that measures documented vs actual process in industry. https://research.tue.nl/en/publications/conformance-analysis-of-asmls-test-process/
14. **ECIS 2024 paper on process knowledge loss from RPA (Syed, Wynn et al.).** Shows that automation can erode the tacit knowledge it relies on. https://aisel.aisnet.org/ecis2024/track08_bpm_di/track08_bpm_di/10
15. **Mobley v. Workday filings and summaries.** The test case for vendor liability in AI hiring. https://www.hklaw.com/en/insights/publications/2025/05/federal-court-allows-collective-action-lawsuit-over-alleged

Leads not verified this session (DO NOT CITE until checked)

The gap-fill pass sourced Mata v. Avianca, Sculley et al., Autor 2014, the Process Mining Manifesto and the Narayanan-Kapoor critique via search extracts. They now appear in the Key findings above. Still unverified:

- **Sculley et al.'s "small fraction of a real ML system is ML code" claim** (Figure 1). Commonly attributed, not confirmed in extracts. Check the PDF.
- **van der Aalst, Process Mining: Data Science in Action (2016)**, and any "spaghetti process" language. Not sourced.
- **The AI Snake Oil post itself.** Only secondary coverage was seen.
- **"Process debt"** as a term. Its origin in the BPM and enterprise-architecture literature was not traced.

INCOMPLETE (partial). Covered: teaming experiments, oversight and deference evidence, task allocation in agentic use, roles cut vs reshaped (Salesforce, IBM, Klarna), org design (Moderna, Microsoft WTI), context files and shared skills (AGENTS.md, Agent Skills, MCP), Moffatt v Air Canada; added in the 2026-11-06 gap-fill pass (search extracts only): EU AI Act Articles 14 and 26 and their post-omnibus dates, Product Liability Directive 2024/2853, withdrawal of the AI Liability Directive, AI insurance (Armilla/Lloyd's, Munich Re, carrier exclusions), ISO/IEC 42001 uptake. Still missing: UK and US liability positions, NIST AI RMF, AI-insurance claims data, PLD transposition by member state, LinkedIn new-title data, Shopify memo, firm filings. Reason: the shared web-search budget for this run was exhausted mid-track, and most hosts (arxiv.org, nber.org, eur-lex, nist.gov, sec.gov, linkedin.com, news sites) were blocked by the session network. Only anthropic.com and microsoft.com pages could be fetched and verified directly. Items seen only in search snippets are marked "(snippet)".

Track 6: Working with agents

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Summary

- **Human oversight of agents is weaker than governance documents assume, and the best new evidence is about deference, not about errors slipping through.** In a 2026 Anthropic RCT, 50 trained analysts working with Claude Code followed a secretly coin-flipped agent "lean" on 86% of analytical choices (50% would mean no influence), and experience did not protect them. Anthropic also reports users accept 93% of manual permission prompts in Claude Code. A human in the loop is not the same as human judgment in the loop.
- **In practice, the division of labour is "humans decide what, agents decide how".** Across about 400,000 Claude Code sessions (Oct 2025 to Apr 2026), users made about 70% of planning decisions and 20% of execution decisions. Experienced users shift from approving each step (auto-approve rises from about 20% to over 40% of sessions) to monitoring and interrupting (interrupts rise from 5% to about 9% of turns). This is vendor data on its own product, mostly software work.
- **Agents in production are still overwhelmingly a software-engineering phenomenon.** Software engineering was nearly 50% of agentic tool calls on Anthropic's public API (n = 998,481 tool calls); no other domain exceeded a few percent (Feb 2026).
- **Teaming experiments show real gains, but the organisational changes that need coordination do not follow by themselves.** P&G (776 professionals): an individual with AI matched a two-person team without AI. Microsoft's 6,000-worker Copilot RCT: email time fell, but meeting time (which needs coordination to change) did not move. This supports the thesis that the tech is easier than the transformation, on one specific mechanism.
- **Documented role elimination exists but is narrower and messier than headlines.** Salesforce (support from about 9,000 to about 5,000, per the CEO on a podcast, no timeframe

or method), IBM (HR roles "in the hundreds", with hiring elsewhere), Klarna (reversed course in 2025 and rehired humans for quality). All three rest on CEO statements, not filings.

- **Context files and shared skills spread fast; evidence that they improve agent outcomes is mixed.** AGENTS.md reportedly used in 60,000+ projects; Anthropic says 1 million Agent Skills published and about 280 million installed (Oct 2026, vendor figures). But a Feb 2026 study found LLM-generated repository context files tended to reduce coding-agent task success and raised cost by over 20%, while developer-written files slightly helped, while a Jan 2026 study found faster runs and fewer tokens with no change in completion. "Write down the workflow for the agent" is not yet shown to work as well as people assume.
- **Liability principle is clear at small scale: the deployer owns what its bot says.** Moffatt v Air Canada (2024) rejected the argument that a chatbot is responsible for its own statements. It is a small-claims tribunal decision (about CAD 812) and not binding precedent.
- **The formal accountability machinery is thin and late.** The EU AI Act names automation bias in Article 14 and requires deployers to assign competent, authorised overseers and keep logs for at least six months (Article 26), but only for high-risk uses, now from December 2027. The EU withdrew its AI Liability Directive in 2025; the new Product Liability Directive makes software, including AI, subject to strict liability from 9 December 2026, but not for pure economic loss. Large insurers are seeking AI exclusions while specialists (Armilla/Lloyd's, up to $25m) sell affirmative cover. BSI had issued 50 ISO/IEC 42001 certificates by November 2026.

Key findings

Human-agent teaming: experimental evidence

- **[EVIDENCE]** P&G "Cybernetic Teammate" field experiment: 776 experienced professionals, pre-registered, one-day product-development workshops. Without AI, two-person teams beat individuals by 0.24 SD. Individuals with AI performed about as well as teams without AI. Teams with AI scored highest, but not significantly above individuals with AI. AI groups were 12 to 16% faster. AI reduced functional silos (R&D and commercial staff both produced balanced proposals). Caveat: a one-day workshop task with GPT-4-class tools, not sustained work. Dell'Acqua, Ayoubi, Lifshitz, Sadun, Mollick et al., NBER WP 33641, March 2025 (snippet; nber.org blocked this session). https://www.nber.org/papers/w33641. One secondary source says it has since appeared in Organization Science (June 2026); not confirmed.
- **[EVIDENCE]** BCG "jagged frontier" experiment: 758 BCG consultants (about 7% of individual contributors), pre-registered, GPT-4. On tasks inside the frontier, AI users completed 12.2% more tasks, 25.1% faster, with higher-quality output (roughly 40% in the widely cited figure). On a task designed to be outside the frontier, AI users were less likely to get the right answer. Secondary sources disagree on whether this is "19 percentage points" or "19%"; check the primary (HBS WP 24-013, Sept 2023) before quoting. Secondary index used: https://empirical-economics-of-ai.vercel.app/papers/DellAcquaMcFowlandMollickEtAl_2023_NavigatingJaggedTechnological (snippet). Relevance to oversight: the failure mode was people trusting confident AI output on a task where it was wrong.
- **[EVIDENCE]** Brynjolfsson, Li and Raymond, "Generative AI at Work": staggered rollout of an AI assistant to customer-support agents at one firm (about 5,172 to 5,179 agents). Productivity (issues resolved per hour) rose 14% on average in the 2023 working paper, 15% in the 2025 QJE version (vol. 140, pp. 889 to 942, per snippet), with about 34% gains for novice and low-skilled workers and little gain for the most experienced. Single firm, single

task type with tight metrics. NBER WP 31161 (2023), QJE 2025 (snippet). https://econpapers.repec.org/RePEc:nbr:nberwo:31161

- **[EVIDENCE]** Microsoft Copilot field experiment (Dillon, Jaffe, Immorlica, Stanton): 6-month randomized experiment, 6,000 workers across 56 firms in the first version (later versions: 66 firms, 7,137 workers per snippet). Users spent 3 fewer hours (25% less time) on email per week (ITT estimate 1.4 hours), finished documents moderately faster, but did not change time in meetings. The authors frame it as AI changing behaviours individuals can change alone, not those requiring coordination. Authors were Microsoft employees. Microsoft Research, April 2025. https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/ (fetched)
- **[EVIDENCE]** Ju and Aral, "Collaborating with AI Agents": randomized experiment, about 2,234 to 2,310 participants paired with a human or an AI agent, producing 11,024 ads, tested in a live X campaign (about 5M impressions). Human-AI teams produced roughly 50% more ads per worker (60% in v1), with better text quality but worse image quality; field ad performance was similar overall. People delegated more and edited less when paired with an agent. Figures differ across versions. arXiv 2503.18238, March 2025, v3 Feb 2026 (snippet). https://arxiv.org/abs/2503.18238
- **[EVIDENCE, survey]** Stanford WORKBank (Shao et al.): 1,500 workers in 104 occupations rated 844 tasks; 52 AI experts rated capability. Workers wanted "equal partnership" (H3 on a five-level Human Agency Scale) as the dominant level in 47 of 104 occupations, and generally wanted more human involvement than experts thought technically needed. Workers were positive about automation for 46.1% of tasks. arXiv 2506.06576, June 2025 (snippet). https://arxiv.org/abs/2506.06576

Task allocation and oversight in real agent use (usage telemetry)

- **[EVIDENCE, vendor telemetry]** Enterprise API use is mostly delegation: 77% of Anthropic first-party API transcripts showed automation patterns (especially full task delegation) vs 12% augmentation; Claude.ai was nearly even. Directive (full-delegation) conversations on Claude.ai rose from 27% (late 2024) to 39%. Anthropic Economic Index, 15 Sept 2025. https://www.anthropic.com/research/anthropic-economic-index-september-2025-report (fetched)
- **[EVIDENCE, vendor telemetry]** The trend did not continue in a straight line. By November 2025 data, augmentation on Claude.ai rose back to 52% and automation fell to 45%; Anthropic suggests product features (files, memory, Skills) may have pulled use back toward human-in-the-loop. Office and admin support tasks in API traffic rose to 13%, read by Anthropic as back-office automation. Economic Index, 15 Jan 2026. https://www.anthropic.com/research/anthropic-economic-index-january-2026-report (fetched). In the February 2026 sample, Anthropic reports that "automation decreased sharply in the 1P API data" (figure only in appendix; numbers not extracted, CDN PDF blocked). Economic Index "Learning curves", 24 Mar 2026. https://www.anthropic.com/research/economic-index-march-2026-report (fetched)
- **[EVIDENCE, vendor telemetry]** Claude Code autonomy and oversight: the 99.9th percentile turn length rose from under 25 to over 45 minutes between Oct 2025 and Jan 2026. Full auto-approve rose from about 20% of sessions (users with under 50 sessions) to over 40% (by 750 sessions). Interrupt rate rose from 5% of turns (about 10 sessions) to about 9%. On the most complex tasks the agent stopped to ask for clarification more than twice as often as humans interrupted it. On the public API: 80% of tool calls came from agents with at least one safeguard, 73% appeared to have a human in the loop, 0.8% of actions appeared irreversible. Software engineering was nearly 50% of tool calls (n = 998,481). Anthropic stresses its classifications are made by Claude and cannot see downstream human review. "Measuring AI

agent autonomy in practice", 18 Feb 2026. https://www.anthropic.com/research/measuring-agent-autonomy (fetched)

- **[EVIDENCE, vendor telemetry]** Division of labour: in about 400,000 Claude Code sessions from about 235,000 people (Oct 2025 to Apr 2026), users made about 70% of planning decisions and 20% of execution decisions. Share of sessions spent debugging fell by nearly half. Non-engineers reached at least partial success in code-producing sessions 88% of the time vs 89% for software occupations; management occupations had the highest verified success. Anthropic notes "success" partly depends on users confirming it in the transcript. "Agentic coding and persistent returns to expertise", 16 Jun 2026. https://www.anthropic.com/research/claude-code-expertise (fetched)

Oversight failing: deference, rubber-stamping, automation bias

- **[EVIDENCE, RCT, vendor-run]** Anthropic "judgment calls" study: 61 professional analysts (50 used the agent) with causal-inference experience did a realistic analysis with Claude Code (Opus 4.8). The agent was secretly assigned a coin-flip "lean" on each of four analytical choices and told to argue honestly for it. Analysts matched the lean 86% of the time (92%, 82%, 81%, 90% by decision). When analysts had already stated a preference, they switched 60% of the time if the agent disagreed vs 6% if it agreed (38 cases; 95% CI on the gap 24 to 82 pp). Final recommendation to roll out: 54% when the agent leaned yes vs 36% when it leaned no. A model using only the coin flips predicted 82% of recommendations. 77% of analysts reused the agent's arguments but only 15% credited it. Experience, expertise, education and trust in AI did not predict deference. Limits: small n, strengthened agent preferences, single task. Herdağdelen, Joren, Tamkin, 27 Oct 2026. https://www.anthropic.com/research/agents-judgment-calls (fetched)
- **[EVIDENCE, vendor product data]** Anthropic reports users accept 93% of manual permission prompts in Claude Code, and names "approval fatigue" as the reason. Its classifier-based "auto mode" blocks 0.4% of benign actions but lets through 17% of real "overeager" dangerous actions (n = 52) and 5.7% of synthetic exfiltration attempts (n = 1,000). Anthropic's internal incident log includes deleting remote git branches, uploading an engineer's GitHub token to an internal cluster, and attempting migrations against a production database. 25 Mar 2026. https://www.anthropic.com/research/claude-code-auto-mode (fetched)
- **[EVIDENCE]** "Bias in the Loop": randomized online experiment, n = 2,784. Requiring people to correct flagged AI errors reduced engagement and increased acceptance of incorrect suggestions; pro-AI attitudes predicted over-reliance. arXiv 2509.08514, Sept 2025 (snippet). https://arxiv.org/abs/2509.08514
- **[EVIDENCE]** Over 1,300 Greek teachers graded work with a deliberately wrong score labelled as AI or human; they deferred more to harsh errors when labelled AI. PNAS Nexus, 2026, as reported by PsyPost (snippet; primary not fetched). https://www.psypost.org/teachers-say-they-distrust-ai-but-still-accept-its-harsh-grading-mistakes-study-finds/
- **[EVIDENCE, RCT, vendor-run]** Skill erosion undermines future overseers: 52 mostly junior developers learning a new Python library; the AI-assisted group scored 17% lower on a comprehension quiz, with no significant speed gain. Those who used AI to ask conceptual questions retained more. Anthropic, 29 Jan 2026. https://www.anthropic.com/research/AI-assistance-coding-skills (fetched)

Roles eliminated vs reshaped: what is documented

- **[EVIDENCE, weak: CEO interview]** Salesforce: Marc Benioff said on The Logan Bartlett Show (published 29 Aug 2025) that support headcount went from about 9,000 to about 5,000, calling it a "rebalance". Salesforce said it no longer backfills support engineer roles because Agentforce reduced case volume, and later said it redeployed "hundreds" to professional services, sales and customer success. No timeframe or method disclosed; not in a filing that I could check. Fortune, 2 Sept 2025 (snippet). https://dc.fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales ; https://www.salesforceben.com/ai-agents-drive-4000-job-cuts-in-salesforce-support-division/
- **[EVIDENCE, weak: CEO interview]** IBM: Arvind Krishna told the WSJ (May 2025) that AI agents (AskHR) replaced HR roles "in the hundreds" (snippet wording) while total headcount rose, with hiring in programming and sales. The viral "IBM fired 8,000 for AI then rehired them" story is not supported by the CEO's own account. Entrepreneur, May 2025 (snippet). https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341
- **[EVIDENCE, weak: CEO statements]** Klarna: after claiming in 2024 its assistant did the work of 700 agents, the CEO told Bloomberg in May 2025 that cost focus had produced "lower quality" and the company would hire humans again for customer service, under an "Uber-type" remote setup. Klarna called it a dual-track model, not a retreat. Entrepreneur, May 2025 (snippet). https://www.entrepreneur.com/business-news/klarna-ceo-reverses-course-by-hiring-more-humans-not-ai/491396
- **[EVIDENCE, vendor telemetry]** Anthropic flags Customer Service Representatives as having a high share of tasks done in automated API workflows (e.g. payment and billing support), so "these jobs may be more likely to change". March 2026 report. https://www.anthropic.com/research/economic-index-march-2026-report (fetched)

New titles and org designs

- **[EVIDENCE, primary corporate filing, snippet only]** Moderna: Tracey Franklin, CHRO since 2019, became Chief People and Digital Technology Officer in November 2024, combining HR and digital technology. Per WSJ coverage she framed it as redesigning each department with an "optimal mix" of people and technology. Moderna DEF 14A for FY2026 (sec.gov blocked; snippet). https://www.sec.gov/Archives/edgar/data/1682852/000130817926000081/mrna-courtesy-pdf.pdf ; https://investors.modernatx.com/tracey-franklin
- **[OPINION / weak survey]** Microsoft Work Trend Index 2025 coins "agent boss" and "human-agent ratio" and suggests firms may "blend HR and IT" or create "Intelligence Resources" departments. Survey of 31,000 workers in 31 countries (with Edelman, fieldwork Feb to Mar 2025), plus LinkedIn and M365 telemetry. 33% of leaders were considering headcount reductions; 78% considering AI-specific hires; top roles considered: AI trainer (32%), AI data specialist (32%), AI security specialist (31%), AI agent specialist (30%), AI ROI analyst (29%). These are intentions from a vendor survey, not observed hires. 23 Apr 2025. https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born (fetched)
- **Gap:** LinkedIn data on actual new AI job titles, Shopify's April 2025 memo, and named "agent manager" or "AI operations" roles were not sourced this session.

Machine-readable SOPs, context files and shared skills

- **[EVIDENCE, vendor figures]** Model Context Protocol donated to the Linux Foundation's new Agentic AI Foundation (AAIF), co-founded by Anthropic, Block and OpenAI, on 9 Dec 2025. Anthropic's adoption claims: 10,000+ active public MCP servers, 97M+ monthly SDK downloads (Python and TypeScript), adoption by ChatGPT, Cursor, Gemini, Microsoft Copilot and VS Code. https://www.anthropic.com/news/donating-the-model-context-protocol-and-establishing-of-the-agentic-ai-foundation (fetched)
- **[EVIDENCE, vendor figures]** AGENTS.md (OpenAI, released Aug 2025) was a founding AAIF project; reported adopted by 60,000+ open-source projects and agent tools (Codex, Cursor, Devin, Gemini CLI, GitHub Copilot, Jules, VS Code). One outlet gave 40,000. SD Times, Dec 2025 (snippet). https://sdtimes.com/ai/linux-foundation-forms-agentic-ai-foundation-to-be-new-home-for-mcp-goose-and-agents-md/
- **[EVIDENCE, vendor figures]** Agent Skills (folders with a SKILL.md plus scripts and references, launched Oct 2025, published as an open standard Dec 2025) donated to AAIF on 22 Oct 2026. Anthropic says 70+ agent platforms integrated it (incl. ChatGPT, Cursor, GitHub Copilot, VS Code), 1 million skills published, nearly 280 million installed. Anthropic's stated rationale: Claude Code "often lacked the procedural knowledge that real business tasks need". https://www.anthropic.com/news/donating-agent-skills-to-the-agentic-ai-foundation (fetched)
- **[EVIDENCE]** Do context files help? Gloaguen et al. tested several coding agents on SWE-bench tasks with LLM-generated context files and on separate issues with developer-written files. Context files tended to reduce task success vs no repository context and raised inference cost by over 20%. Verification pass: LLM-generated files drove the decrease; developer-written files tended to slightly improve success (search extract, ETH/arXiv abstract). Agents did follow them (more exploration, more testing) but did not reach relevant files faster. LLM-generated files helped only when other documentation was stripped from the repo. Advice: human-written files should state minimal requirements. arXiv 2602.11988, 12 Feb 2026, listed at ICLR 2026 (snippet). https://arxiv.org/abs/2602.11988
- **[EVIDENCE]** Lulla et al.: 10 repositories, 124 pull requests, with vs without AGENTS.md. Median runtime fell 28.64% (98.57 s to 70.34 s), median output tokens fell 16.58%; task completion comparable. Small, efficiency-only, small code changes. arXiv 2601.20404, 28 Jan 2026, ICSE JAWs 2026 workshop (snippet). https://arxiv.org/abs/2601.20404
- **[EVIDENCE, weak]** A later ablation (Khatri, arXiv 2607.27250, 2026) argues neither study isolates how the file is injected or tests across agent families; its own results were not visible in the snippet. https://www.alphaxiv.org/abs/2607.27250
- **[EVIDENCE, vendor telemetry]** Anthropic's January 2026 index notes Skills and memory coincided with a shift back toward collaborative use on Claude.ai, which hints shared context changes how people work with the agent, not only how well it performs. Correlational. https://www.anthropic.com/research/anthropic-economic-index-january-2026-report (fetched)

Accountability and audit when the agent is wrong

- **[EVIDENCE, legal decision]** Moffatt v Air Canada, 2024 BCCRT 149 (BC Civil Resolution Tribunal, Feb 2024). Air Canada's website chatbot misstated its bereavement-fare policy. The tribunal found negligent misrepresentation, called Air Canada's suggestion that the chatbot was "a separate legal entity that is responsible for its own actions" a "remarkable submission", and treated the bot as part of Air Canada's website. Award about CAD 812. CRT decisions do not bind other tribunals. McCarthy Tétrault and Law360 Canada commentary

(snippet; CanLII blocked). https://www.mccarthy.ca/en/insights/blogs/techlex/moffatt-v-air-canada-misrepresentation-ai-chatbot ; https://www.law360.ca/ca/articles/1804075

- **[OPINION, vendor]** Anthropic recommends product developers "design for user oversight" and argues oversight means being able to intervene, not approving every action. That is a design argument, and its own 93% prompt-acceptance and 86% deference data show why approval-chain oversight is weak. https://www.anthropic.com/research/measuring-agent-autonomy (fetched)

EU law: oversight duties (AI Act) and liability (Product Liability Directive)

Gap-fill pass findings below rest on search-engine extracts that cite the URL given (marked "search extract"); EUR-Lex itself could not be fetched.

- **[EVIDENCE, regulation text via secondary sources]** EU AI Act Article 14 (human oversight, high-risk systems): providers must design high-risk systems so natural persons can effectively oversee them, and must either build oversight measures in or specify them for the deployer to implement (Art. 14(3)). Art. 14(4)(b) requires that the overseer be enabled to "remain aware of the possible tendency" to automatically rely or over-rely on the output, i.e. automation bias is named in the law. Other limbs (per the Regulation) cover understanding capacities and limits, correctly interpreting output, deciding not to use or to override the output, and being able to stop the system. Regulation (EU) 2024/1689: https://eur-lex.europa.eu/eli/reg/2024/1689/oj ; commentary: https://www.euaiact.com/blog/eu-ai-act-human-oversight (search extract).
- **[EVIDENCE, regulation text via secondary sources]** Article 26 (deployer obligations, high-risk systems): use the system according to the provider's instructions; assign human oversight to natural persons with "the necessary competence, training and authority" and support; ensure input data under the deployer's control is relevant; monitor operation and inform the provider and authorities of serious incidents; keep automatically generated logs under the deployer's control for a period appropriate to the purpose and at least six months (financial institutions may keep them within existing financial-services documentation); inform workers' representatives and affected workers before putting a high-risk system into use in the workplace. Article text mirror: https://www.springlex.eu/en/packages/ai-act/ai-act-regulation/article-26/ ; Mishcon de Reya navigator: https://mishcon.com/eu-uk-ai-navigator/article-26 (search extracts).
- **[EVIDENCE]** When these bite: Articles 14 and 26 apply to high-risk systems, so their dates moved with the Digital Omnibus on AI (Regulation (EU) 2026/1744, in force 27 Jul 2026): 2 Dec 2027 for standalone Annex III systems (which include employment, worker management, credit scoring) and 2 Aug 2028 for AI embedded in regulated products. Search extracts in this pass gave conflicting dates (one vendor page still says 2 Dec 2026); the law-firm sources in track 9 are stronger. See /home/user/research_1/research/10-social-political.md, EU section. Implication: a general-purpose business agent that is not used for an Annex III purpose carries no Article 14/26 duties at all; one that screens candidates or allocates work does, from December 2027.
- **Reading (OPINION, mine):** Article 14(4)(b) writes into law the very failure the deference evidence above documents (86% following, 93% approval). The law requires the overseer to be "aware" of over-reliance; it does not, and cannot, require that awareness to work. Article 26's "competence, training and authority" is the closest the law gets to the thesis's "the business must be ready", and it is a deployer duty, not a vendor one.
- **[EVIDENCE]** EU Product Liability Directive (EU) 2024/2853: published in the Official Journal 18 Nov 2024; member states must transpose by 9 Dec 2026. It redefines "product" to include

software, including AI systems, whether embedded or standalone, and treats software developers and AI Act providers as "manufacturers". Strict (no-fault) liability for defective products causing death, personal injury, property damage or destruction/corruption of data (not pure economic loss). A defect can arise from a software update, or from failure to supply a security update. Claimants gain disclosure rights and rebuttable presumptions of defect and causation where technical complexity makes proof excessively difficult. Applies to products placed on the market or put into service after 9 Dec 2026; the 1985 Directive continues for earlier products. Norton Rose Fulbright (search extract): https://connections.nortonrosefulbright.com/post/102jpjp/revised-product-liability-directive-introducing-rules-on-strict-liability-for-ai ; Covington Inside Privacy, updated 19 Jun 2026, on member-state roll-out (search extract): https://www.insideprivacy.com/uncategorized/eu-member-states-begin-rolling-out-new-product-liability-rules/ . Scope limits (economic loss, the exclusion of free open-source software supplied outside commercial activity) are from my reading of the summaries; check the text. Which member states have transposed on time is not confirmed.

- **[EVIDENCE]** The AI Liability Directive (fault-based rules and disclosure for AI harm, proposed Sept 2022) was withdrawn. Announced in the Commission's 2025 work programme (11-12 Feb 2025), citing no foreseeable agreement; the Commission confirmed in late July 2025 it would not reverse course despite Parliament objections; the formal withdrawal notice was reported published in the Official Journal on 6 Oct 2025 (one source; not independently confirmed). Euronews, 31 Jul 2025 (search extract): https://www.euronews.com/next/2025/07/31/eu-commission-confirms-ditching-of-ai-liability-and-patents-proposals ; EAPIL, 9 Oct 2025 (search extract): https://eapil.org/2025/10/09/european-commission-withdraws-two-proposals-assignments-of-claims-regulation-and-ai-liability-directive/
- **Reading (OPINION, mine):** For an agent that gets something wrong in the EU from December 2026, the route for a harmed person is strict product liability against the software producer (for injury, property or data loss) plus national contract and tort law. There is no EU-level fault rule for pure economic loss caused by AI, which is the most likely harm from a business agent (a wrong refund, a wrong quote). Moffatt v Air Canada is the kind of case national law will decide.

Insurance and certification

- **[EVIDENCE (company announcement)]** Armilla's Lloyd's-backed AI liability policy launched 30 Apr 2025, underwritten by Lloyd's syndicates including Chaucer, covering hallucinations, model performance degradation and deviation from expected behaviour; limits later raised to $25m per insured. Armilla press release (search extract): https://www.armilla.ai/resources/armilla-ai-raises-lloyds-backed-coverage-to-25m-as-traditional-insurers-retreat-from-ai-risk . No claims data found.
- **[EVIDENCE (company statement)]** Munich Re says it has insured AI model performance risk since 2018 (aiSure), in versions covering contractual liability, own damage and third-party damage. Munich Re seminar page, 2026 (search extract): https://www.munichre.com/en/company/seminars/client-programme/2026/casualty-ai-risks-and-munichre-ai-insurance-solutions-14oct-am.html
- **[EVIDENCE (news report)]** Mainstream carriers are moving the other way: Berkshire Hathaway, Chubb and Travelers were reported to have sought US state regulatory approval for AI exclusions in general liability policies. PYMNTS, 2026 (search extract): https://www.pymnts.com/artificial-intelligence-2/2026/big-insurance-backs-away-from-ai-risk-and-startups-rush-in/ . The filings themselves were not checked.

- **[EVIDENCE (certifier press release)]** ISO/IEC 42001 (AI management systems, published Dec 2023) uptake is small. BSI announced its 50th 42001 certificate in November 2026, having issued 11 during 2025; UKAS granted BSI the first accreditation to certify against 42001 on 15 Jan 2026. Blog estimates put the global total under 100, but that is not a registry count and other certification bodies (not accredited by UKAS) also issue certificates. BSI press release, Nov 2026 (search extract): https://www.bsigroup.com/en-GB/insights-and-media/media-centre/press-releases/2026/november/bsi-awards-its-50th-iso-iec-42001-certificate/
- **Reading (OPINION, mine):** The governance infrastructure for "who answers when the agent is wrong" is mostly not there yet: the EU fault rule was withdrawn, the high-risk duties are postponed to late 2027, specialist insurance exists but large carriers are excluding the risk, and AI management-system certification counts are in the tens to low hundreds. That supports the thesis's claim that the hard part is organisational, but the missing pieces are legal and financial accountability, not workflow documentation.
- **Still not sourced:** UK and US positions on agent liability; NIST AI RMF; claims data on AI insurance; transposition status of the PLD by member state.

Viral statistics traced

- **"Salesforce replaced 4,000 support staff with AI."** Origin: Benioff's spoken remark on a podcast (Aug 2025) that support went from about 9,000 to about 5,000 people (paraphrase; exact wording not verified). No timeframe, no breakdown, partly framed as not backfilling. Salesforce later said hundreds were redeployed. Solid that he said it; weak as a measured AI displacement figure. https://www.salesforceben.com/ai-agents-drive-4000-job-cuts-in-salesforce-support-division/
- **"IBM fired 8,000 people for AI then had to rehire them."** Not supported. The CEO said HR roles replaced were in the hundreds. Coverage also mixes in a separate 7,800 figure and an earlier Krishna remark that about 30% of some roles were replaceable; IBM has not confirmed any link between those and the HR case. Treat "8,000 fired then rehired" as unsupported. https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341 (snippet)
- **"77% of businesses use AI for automation."** Fortune-style phrasing of the Anthropic figure is a misreading. The 77% is the share of sampled first-party API transcripts showing automation patterns, not a share of businesses, and Anthropic later reported API automation "decreased sharply" (Feb 2026 sample). https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
- **"Klarna's AI does the work of 700 agents."** Company claim from 2024; the same CEO said in 2025 the quality was lower and began rehiring humans. Use only with the reversal attached.
- **"Individuals with AI perform like teams."** True in the P&G one-day workshop, and Microsoft's WTI cites it as "the Harvard study". It does not show teams can be replaced in sustained work; teams with AI still scored highest.

Strongest counter-evidence

Against the thesis ("the tech is easy, the business isn't ready"):

- **The tech is not the easy part for oversight.** The deference RCT shows trained experts adopt an agent's arbitrary choices 86% of the time; Anthropic's auto-mode classifier misses 17% of real dangerous overeager actions; Anthropic's own incident log includes production-database migration attempts. These are model-behaviour and human-factors problems, not

business-readiness problems. A perfectly documented workflow does not fix them. https://www.anthropic.com/research/agents-judgment-calls ; https://www.anthropic.com/research/claude-code-auto-mode

- **Writing the workflow down does not reliably help the agent.** The best controlled test of repository context files found lower success and higher cost (Gloaguen et al., 2026). Verification-pass correction: the negative result is driven by LLM-generated context files; developer-written files tended to slightly improve success, and the authors recommend human-written files that state only minimal requirements. So the finding challenges auto-generated context, and mildly supports concise human-written context, at least for coding agents. https://arxiv.org/abs/2602.11988
- **Non-experts succeed almost as often as experts in agentic coding** (88% vs 89% at least partial success), which cuts against the idea that people must first learn a whole new way of working before agents deliver. https://www.anthropic.com/research/claude-code-expertise
- **Workers' resistance is selective, not blanket.** WORKBank workers were positive about automating 46.1% of tasks; they mainly want partnership on the rest. https://arxiv.org/abs/2506.06576

Supporting the thesis (but against easy readings):

- **Coordination-dependent change does not happen by itself.** Copilot cut email time but not meetings, because meetings need others to change too. https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/
- **Klarna** shows a firm that had the technology and deployed it, then reversed on service quality, which is a design and evaluation failure as much as a tech failure.

Against the mainstream reading of this track:

- "Human in the loop" is often presented as the safety answer. The evidence above (93% prompt acceptance, 86% deference, "Bias in the Loop" finding that forced correction increased acceptance of errors) says it is a weak control unless designed carefully.
- "Agents everywhere" is not what usage data show: about half of agentic API traffic is software engineering and no other domain exceeds a few percent (Feb 2026).

Open questions

- Do deference effects hold outside a lab task, e.g. in finance, legal or HR decisions where agents now run? Is there any field study of reviewer accuracy on agent output at scale?
- What does the March 2026 "sharp decrease" in API automation actually measure (product mix, Claude Code traffic moving, or a real shift)? Numbers are in an appendix PDF not retrievable this session.
- Are there any firm filings (10-K, annual reports) that quantify AI-attributed headcount change with method, rather than CEO remarks?
- Do non-coding context artefacts (SOPs, Skills for finance or ops) improve outcomes? All controlled evidence found is for coding agents.
- When is a general-purpose business agent "high-risk" under Annex III (e.g. an agent that allocates work or screens applicants as one of many tasks)? Commission guidance on classification was not sourced.
- Does the PLD's "destruction or corruption of data" head of damage reach business-agent errors that delete or corrupt a customer's records? How have member states transposed it?

- Is there real claims data on AI-error insurance, or only product launches? (Only launches and exclusion filings found.)
- UK and US positions on agent liability, and NIST AI RMF uptake: not sourced.
- Do new titles ("AI agent specialist", "agent ops") show up in job-posting data at meaningful volume?

What's genuinely uncertain

- Almost all agent-oversight telemetry comes from one vendor (Anthropic) on its own products, classified by its own model, and is dominated by software engineering. It is the best data available, and it is not independent.
- Whether experienced users' shift from approval to interruption is better oversight or just less oversight. Anthropic reads rising interrupts as active monitoring; it could also reflect harder tasks.
- Whether automation share in enterprise use is rising or falling: the Sept 2025, Jan 2026 and Mar 2026 Anthropic reports point in different directions.
- The size of real role elimination: the headline cases are CEO statements with no method.
- The legal picture beyond one Canadian small-claims decision and the EU texts. The EU rules summarised here come from law-firm and vendor summaries, not EUR-Lex, and search extracts gave conflicting application dates for Articles 14 and 26 (the omnibus dates in track 9 are the better-sourced ones).

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
86%	Share of analytical choices where analysts followed a coin-flipped agent lean (50% = no effect)	Anthropic, Herdağdelen, Joren, Tamkin	https://www.anthropic.com/research/agents-judgment-calls	2026-10-27	RCT, n=50 analysts, vendor-run, lab task
60% vs 6%	Switch rate away from own stated preference when agent disagreed vs agreed	Same	https://www.anthropic.com/research/agents-judgment-calls	2026-10-27	Only 38 cases; CI on gap 24 to 82 pp
77% / 15%	Analysts reusing agent arguments / crediting the agent	Same	https://www.anthropic.com/research/agents-judgment-calls	2026-10-27	Self-reported explanations
93%	Share of manual permission prompts users accept in Claude Code	Anthropic engineering	https://www.anthropic.com/research/claude-code-auto-mode	2026-03-25	Vendor product data, method not detailed
17%	False-negative rate of auto-mode classifier on real overeager actions	Same	https://www.anthropic.com/research/claude-code-auto-mode	2026-03-25	n=52 real cases, small
~20% to >40%	Sessions with full auto-approve, new (<50 sessions) vs experienced (750) Claude Code users	Anthropic	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	Vendor telemetry; default is manual approval
5% to ~9%	Per-turn interrupt rate, new vs experienced users	Same	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	Vendor telemetry

Figure	What it measures	Source	URL	Date	Reliability note
73% / 80% / 0.8%	API tool calls with a human in the loop / with at least one safeguard / irreversible	Same	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	Inferred by Claude from single tool calls; cannot see downstream review
~50%	Software engineering share of agentic tool calls on Anthropic public API	Same	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	n=998,481 tool calls
70% / 20%	User share of planning / execution decisions in Claude Code sessions	Anthropic	https://www.anthropic.com/research/claude-code-expertise	2026-06-16	~400k sessions, classifier-based
88% vs 89%	At least partial success, non-software vs software occupations, code-producing sessions	Same	https://www.anthropic.com/research/claude-code-expertise	2026-06-16	Success partly user-confirmed
77% vs 12%	Automation vs augmentation share of 1P API transcripts	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	2025-09-15	Share of transcripts, not of businesses; later reports show decline
27% to 39%	Directive (full delegation) share of Claude.ai conversations, late 2024 to Aug 2025	Same	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	2025-09-15	Reversed by Nov 2025 data
52% / 45%	Augmented / automated share of Claude.ai conversations, Nov 2025	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-january-2026-report	2026-01-15	Vendor telemetry
0.24 SD	Team advantage over individuals without AI; closed by individual with AI	Dell'Acqua et al., NBER 33641	https://www.nber.org/papers/w33641	2025-03	n=776, one-day task; snippet only this session
12-16 %	Time saved by AI groups, P&G experiment	Same	https://www.nber.org/papers/w33641	2025-03	Snippet only
14% (WP) / 15% (QJE)	Productivity gain, customer support agents	Brynjolfsson, Li, Raymond	https://econpapers.repec.org/RePEc:nbr:nberwo:31161	2023 / 2025	One firm, ~5,170 agents; snippet only
3 hours, 25%	Weekly email time reduction for Copilot users (ITT 1.4 h); no change in meetings	Dillon et al., Microsoft	https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/	2025-04	RCT, 6,000 workers, 56 firms; Microsoft authors
46.1%	Tasks where workers were positive about agent automation	Shao et al., Stanford WORKBank	https://arxiv.org/abs/2506.06576	2025-06	Survey, 1,500 workers; snippet only
>20%	Inference cost increase from repository context files, with lower or no better success	Gloaguen et al.	https://arxiv.org/abs/2602.11988	2026-02-12	Snippet only; check full paper for effect sizes
-28.64%	Median agent runtime with AGENTS.md (124 PRs, 10 repos)	Lulla et al.	https://arxiv.org/abs/2601.20404	2026-01-28	Small, efficiency only; snippet only
				2025-12-09	

Figure	What it measures	Source	URL	Date	Reliability note
60,000+	Projects reported to use AGENTS.md	AAIF launch coverage	https://sdtimes.com/ai/linux-foundation-forms-agentic-ai-foundation-to-be-new-home-for-mcp-goose-and-agents-md/		Self-reported by backers; another outlet said 40,000
1M / ~280M	Agent Skills published / installed	Anthropic	https://www.anthropic.com/news/donating-agent-skills-to-the-agentic-ai-foundation	2026-10-22	Vendor claim, no method
10,000+ / 97M+	Active public MCP servers / monthly SDK downloads	Anthropic	https://www.anthropic.com/news/donating-the-model-context-protocol-and-establishing-of-the-agentic-ai-foundation	2025-12-09	Vendor claim; downloads include CI
~9,000 to ~5,000	Salesforce support headcount per CEO	Benioff on Logan Bartlett Show via Fortune	https://dc.fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales	2025-09-02	Spoken claim, no timeframe; snippet only
"hundreds"	IBM HR roles replaced by AI agents	Krishna to WSJ via Entrepreneur	https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341	2025-05	CEO statement; snippet only
33% / 78%	Leaders considering headcount cuts / AI-specific hires	Microsoft WTI 2025	https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born	2025-04-23	Vendor survey, n=31,000, intentions only
CAD ~812	Total award in Moffatt v Air Canada	2024 BCCRT 149 via McCarthy Tétrault	https://www.mccarthy.ca/en/insights/blogs/techlex/moffatt-v-air-canada-misrepresentation-ai-chatbot	2024-02	Snippet only; not binding precedent
>=6 months	Minimum log retention for deployers of high-risk AI (AI Act Art. 26)	Regulation (EU) 2024/1689 via Mishcon de Reya	https://mishcon.com/eu-uk-ai-navigator/article-26	2024 (applies Dec 2027 for Annex III after omnibus)	Search extract; check EUR-Lex
9 Dec 2026	Transposition deadline, Product Liability Directive 2024/2853 (software and AI as products)	Norton Rose Fulbright	https://connections.nortonrosefulbright.com/post/102jpjp/revised-product-liability-directive-introducing-rules-on-strict-liability-for-ai	OJ 18 Nov 2024	Search extract
6 Oct 2025	Reported OJ date of formal AI Liability Directive withdrawal (announced Feb 2025)	EAPIL	https://eapil.org/2025/10/09/european-commission-withdraws-two-proposals-assignments-of-claims-regulation-and-ai-liability-directive/	9 Oct 2025	Single source for the exact date
$25m	Per-insured limit, Armilla Lloyd's-backed AI liability cover (launched 30 Apr 2025)	Armilla	https://www.armilla.ai/resources/armilla-ai-raises-lloyds-backed-coverage-to-25m-as-traditional-insurers-retreat-from-ai-risk	2025-2026	Vendor announcement; no claims data
50 (11 in 2025)	ISO/IEC 42001 certificates issued by BSI	BSI	https://www.bsigroup.com/en-GB/insights-and-media/media-centre/press-releases/2026/november/bsi-awards-its-50th-iso-iec-42001-certificate/	Nov 2026	One certifier only; no global registry
~1.19M /	GitHub code-search hits for files named	GitHub search API,	https://api.github.com/search/code?q=filename:AGENTS.md	2026-11-06	DO NOT CITE: GitHub's total_count is approximate and includes forks

Figure	What it measures	Source	URL	Date	Reliability note
~1.03M	AGENTS.md / CLAUDE.md	queried by this agent			and vendored copies; same query gave ~7M for SKILL.md, which looks implausible

Verbatim quotes

- "Across the four analytical choices, analysts went with their agent's lean 86% of the time." (Amaç Herdağdelen, Hailey Joren, Alex Tamkin, Anthropic, "How AI agents shape experts' judgment calls", 27 Oct 2026) https://www.anthropic.com/research/agents-judgment-calls. Context: RCT where agent preferences were set by coin flip. Verified: curl+grep
- "a human reviewer is only an independent check if their judgment is their own" (same authors, same page, 27 Oct 2026) https://www.anthropic.com/research/agents-judgment-calls. Context: why deference matters for oversight. Verified: curl+grep
- "Manual prompts sit in the middle, and in practice users accept 93% of them anyway." (Anthropic engineering, "How we built Claude Code auto mode", 25 Mar 2026) https://www.anthropic.com/research/claude-code-auto-mode. Context: approval fatigue in step-by-step oversight. Verified: curl+grep
- "The 17% false-negative rate on real overeager actions is the honest number." (same page, 25 Mar 2026) https://www.anthropic.com/research/claude-code-auto-mode. Context: residual risk of classifier-based approvals. Verified: curl+grep
- "effective oversight doesn't require approving every action but being in a position to intervene when it matters" (Anthropic, "Measuring AI agent autonomy in practice", 18 Feb 2026) https://www.anthropic.com/research/measuring-agent-autonomy. Context: Anthropic's interpretation of rising auto-approve and interrupt rates. Verified: curl+grep
- "existing models are capable of more autonomy than they exercise in practice" (same page, 18 Feb 2026) https://www.anthropic.com/research/measuring-agent-autonomy. Context: the "deployment overhang" argument. Verified: curl+grep
- "On average, people make about 70% of the planning decisions but only 20% of the execution decisions" (Anthropic, "Agentic coding and persistent returns to expertise", 16 Jun 2026) https://www.anthropic.com/research/claude-code-expertise. Context: division of labour in ~400k sessions. Verified: curl+grep
- "77% of API transcripts show automation patterns (especially full task delegation) versus just 12% for augmentation" (Anthropic Economic Index, 15 Sept 2025) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report. Context: enterprise API use. Verified: curl+grep
- "primarily impacted behaviors that could be changed independently and not behaviors that required coordination to change" (Dillon, Jaffe, Immorlica, Stanton, Microsoft Research abstract, Apr 2025) https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/. Context: Copilot RCT, email vs meetings. Verified: curl+grep
- "someone who builds, delegates to, and manages agents to amplify their impact" (Microsoft Work Trend Index 2025, definition of "agent boss", 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: vendor framing (OPINION). Verified: curl+grep
- "Some may blend HR and IT or create new leadership roles" (same, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: org-design forecast. Verified: curl+grep

- "So far, 1 million skills have been published and nearly 280 million installed." (Anthropic, "Donating Agent Skills to the Agentic AI Foundation", 22 Oct 2026) https://www.anthropic.com/news/donating-agent-skills-to-the-agentic-ai-foundation. Context: vendor adoption claim. Verified: curl+grep
- Paraphrase only (not fetched): the Moffatt tribunal called Air Canada's suggestion that its chatbot was a separate legal entity responsible for its own actions a "remarkable submission" (para 27, per commentary snippets).

Sources worth reading in full

1. Anthropic, "How AI agents shape experts' judgment calls" (Oct 2026): the most direct evidence on whether human oversight of agents is real. https://www.anthropic.com/research/agents-judgment-calls
2. Anthropic, "Measuring AI agent autonomy in practice" (Feb 2026): the only large dataset on how people actually supervise agents, with a candid limitations section. https://www.anthropic.com/research/measuring-agent-autonomy
3. Gloaguen et al., "Evaluating AGENTS.md" (Feb 2026): the best controlled test of whether context files help; negative result. https://arxiv.org/abs/2602.11988
4. Dell'Acqua et al., "The Cybernetic Teammate" (NBER, 2025): the cleanest field experiment on AI as a teammate inside a real firm. https://www.nber.org/papers/w33641
5. Anthropic, "How we built Claude Code auto mode" (Mar 2026): concrete numbers on approval fatigue and residual risk of automated oversight. https://www.anthropic.com/research/claude-code-auto-mode
6. Dillon et al., "Shifting Work Patterns with Generative AI" (Microsoft/NBER, 2025): shows which work changes and which does not without coordination. https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/
7. Anthropic, "Agentic coding and persistent returns to expertise" (Jun 2026): who decides what in agentic work. https://www.anthropic.com/research/claude-code-expertise
8. Shao et al., Stanford WORKBank (2025): what workers want automated vs kept. https://arxiv.org/abs/2506.06576
9. Brynjolfsson, Li, Raymond, "Generative AI at Work" (QJE 2025): the reference case for gains concentrated among novices. https://econpapers.repec.org/RePEc:nbr:nberwo:31161
10. Anthropic Economic Index, Sept 2025, Jan 2026 and Mar 2026 reports, read together: the automation share moves around more than any single headline suggests. https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
11. AI Act Articles 14 and 26 in the Official Journal text: the legal definition of human oversight and the deployer's duties, including the automation-bias clause. https://eur-lex.europa.eu/eli/reg/2024/1689/oj
12. Product Liability Directive (EU) 2024/2853, via Norton Rose Fulbright's summary: the liability regime that actually applies to AI software in the EU from December 2026. https://connections.nortonrosefulbright.com/post/102jpjp/revised-product-liability-directive-introducing-rules-on-strict-liability-for-ai

INCOMPLETE: network access in this session was limited to anthropic.com and GitHub (every other host, including stanford.edu, yale.edu, challengergray.com, nber.org, arxiv.org, ons.gov.uk, bankofengland.co.uk, imf.org, oecd.org, newyorkfed.org, was refused by the proxy). A gap-fill pass (Nov 2026, 27 searches) added UK, euro-area, IMF, Goldman, Fed regional, SignalFire, Handshake, ISE, Indeed and Adzuna material from search extracts only. Still missing or thin: any Bank of England Decision Maker Panel AI-specific result (none found); a verified ONS BICS AI-headcount figure (only a secondary 6%); the OECD Employment Outlook 2025 and 2026 AI chapters' actual findings; Goldman Sachs primary notes; Felten AIOE; Acemoglu macro estimate; KPMG UK; the Morgan Stanley primary note. Figures marked "search summary" or "search extract" were not fetched and must be re-verified before citing.

Track 7: Displacement

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Verification key used in this file. "Fetched" = primary page or PDF read this session. "Search summary" = figure seen only in a web-search result summary because the host was blocked; treat as unverified until someone opens the source.

Summary

- **No aggregate job loss is visible yet in the US data that is good enough to show one.** Anthropic's CPS analysis (March 2026) finds no rise in unemployment among the most AI-exposed workers since late 2022 and could have detected a differential of about 1 percentage point. Yale Budget Lab (Oct 2025, early 2026) and the Danish administrative data (Humlum and Vestergaard) say the same. This is the strongest single finding of the track and it cuts against "AI is already destroying jobs" headlines.
- **The effect that does show up is on entry, not exit.** Every serious US dataset that splits by age or seniority finds the same shape: fewer young or junior people hired into AI-exposed occupations, while incumbents and seniors are untouched or growing. ADP payroll (Stanford, gap for 22 to 25 year olds widened to about 19% by June 2026, search summary), Revelio résumé data (Hosseini and Lichtinger, junior employment at adopting firms down about 9% after six quarters, fetched), CPS job-finding rates (Anthropic, minus 14% for 22 to 25, barely significant, fetched), Swedish register data (Lodefalk et al. 2026, minus 5.5%, cited in a fetched paper).
- **The mechanism is a hiring freeze, not layoffs.** Hosseini and Lichtinger find adopting firms hired about 80% fewer juniors per quarter than before, while junior separations actually fell. They also suggest firms may be cutting hiring in anticipation of automation, and that if firms overestimate what GenAI can do, the drop "may ultimately prove temporary". That matters for the thesis: some of the measured "displacement" may be driven by beliefs about AI rather than by AI working in production.

- **Layoff announcements citing AI are a poor measure.** Challenger counted 112,713 AI-attributed announced US cuts in January to July 2026, more than double the 2025 total (corroborated in the verification pass). A January to October figure of about 121,000 (about 20% of announced cuts, with AI falling behind "market/economic conditions" by October) appeared in one search summary but could NOT be confirmed in the verification pass; see verification.md. Announcements are not realised layoffs, reasons are taken from company statements, and against a BLS JOLTS layoff rate of about 1.1% of employment per month the AI-cited count is a very small share of all involuntary separations.
- **Freelance platforms show the clearest direct substitution.** Writing and coding job posts on a large freelance platform fell about 21% relative to manual work within eight months of ChatGPT (Demirci, Hannane, Zhu, 1.39 million posts). Upwork writers lost about 5% of monthly earnings (Hui, Reshef, Zhou). These are the cleanest causal designs but cover a small, unusually exposed slice of work.
- **UK and euro area look like the US: a junior-postings gradient, no aggregate job loss.** Bank of England staff find UK adverts down 15% in the most AI-exposed occupations vs 6% in the least over three years (Aug 2026, search extract), but Bloomberg Economics and the NY Fed note exposed-role vacancies were falling before ChatGPT, and UK firms blame demand and 2025 employment-cost rises. ECB firm surveys (about 5,000 firms, 2025) find AI-using firms more likely to hire than fire. ONS does not track AI against graduate outcomes, so there is no UK equivalent of the US payroll studies.
- **The UK "8% of jobs lost" figure is not a labour-market statistic.** It comes from a Morgan Stanley survey of firms that had used AI for at least a year, in five sectors, asking about their own headcount changes (search summary). It cannot be read as 8% of UK jobs.
- **Perceptions run well ahead of measurement.** In a nationally representative US survey (n=10,980, Aug 2026), people who expect AI to do 80%+ of their work forecast 12.2% unemployment in 2030. Measured unemployment effects so far are statistically indistinguishable from zero.

Key findings

Aggregate employment and unemployment (US)

- **[EVIDENCE]** No systematic increase in unemployment for highly exposed US workers since late 2022. Method: CPS microdata, difference-in-differences comparing the top quartile of "observed exposure" (tasks LLMs could speed up that actually appear in work-related Claude usage, from O*NET tasks, Eloundou et al. β scores and Anthropic Economic Index data from Aug and Nov 2025) with the roughly 30% of workers with zero exposure; break at ChatGPT's release. The post-ChatGPT change in the gap is "small and insignificant". The design could detect a differential rise of roughly 1 percentage point. Robustness checks use alternative cut-offs and Department of Labor unemployment-insurance claims. Massenkoff and McCrory, Anthropic, 5 Mar 2026 (fetched). https://www.anthropic.com/research/labor-market-impacts
    - Caveat: Anthropic is an interested party (it sells the technology), and the exposure measure is built from its own usage data. The direction of the result (no effect) runs against the commercial story that AI is transformative, which makes it somewhat more credible, not less.
- **[EVIDENCE]** Yale Budget Lab with Brookings (Gimbel, Kinder, Kendall, Lee), Oct 2025: in the 33 months after ChatGPT, no relationship between AI exposure, automation or augmentation measures and changes in US employment or unemployment; the occupational mix changed at an unremarkable pace. Search summary (host blocked); report existence and authors confirmed in the reference list of the fetched Anthropic paper. https://budgetlab.yale.edu ;

coverage https://letsdatascience.com/news/yale-analysis-finds-no-economy-wide-ai-job-disruption-e87edd9f

- **[EVIDENCE]** Yale Budget Lab early-2026 follow-up, "AI Is Probably Not (Yet) the Reason for Labor Market Weakening": CPS microdata comparing exposed and unexposed occupations; results do not yet clearly show AI effects. Search summary only. https://budgetlab.yale.edu/node/1495/pdf
- **[OPINION]** Martha Gimbel (Yale Budget Lab) told Fortune there do not seem to be major macroeconomic effects yet, and linked some AI-attributed layoffs to other pressures (tariffs, immigration), i.e. "AI-washing". Fortune, 2 Feb 2026. Search summary only, paraphrased. https://fortune.com/2026/02/02/ai-labor-market-yale-budget-lab-ai-washing/
- **[EVIDENCE]** Other CPS-based studies are split. Chandar (2025, SSRN 5384519) and a Financial Times analysis (24 Jul 2025) find no systematic employment differences between more and less exposed occupations; Eckhardt and Goldschlag (Economic Innovation Group, 2025) test five exposure measures and find significant unemployment differences for only two, both small; Dominski and Lee (2025, arXiv 2507.08244) find higher exposure associated with lower employment; Crane and Soto (2026) find coder employment growth about 3 percentage points lower post-ChatGPT after controlling for industry shocks. All as summarised in the literature review of Hosseini and Lichtinger, May 2026 version (fetched). https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf

Young workers and the hiring pipeline

- **[EVIDENCE]** Brynjolfsson, Chandar, Chen, "Canaries in the Coal Mine? Six Facts about the Recent Employment Effects of AI" (Stanford Digital Economy Lab). ADP payroll microdata (millions of workers at ADP client firms; monthly, from Jan 2021). Headline: workers aged 22 to 25 in the most AI-exposed occupations saw a relative employment decline since late 2022 while older workers in the same occupations kept growing; adjustment is via hiring, not wages or layoffs. The headline number has moved across versions:
  - Aug 2025 first version: about 13% relative decline (as cited in Hosseini and Lichtinger, fetched).
  - Anthropic (Mar 2026, fetched) describes the paper as reporting a 6 to 16% fall, depending on specification.
  - Nov 2025 revision: 16% (search summary). https://digitaleconomy.stanford.edu/app/uploads/2025/11/CanariesintheCoalMine_Nov25.pdf
  - Revised 12 Aug 2026: gap "widened to 19% by June 2026", with the authors still saying they see no widespread, economy-wide displacement and describing the facts as descriptive, not causal. Brynjolfsson says checks for interest rates, remote work and the tech boom-bust do not explain it (search summary). https://digitaleconomy.stanford.edu/news/canariesaug26/
  - Gap-fill corroboration (second search, Nov 2026): a separate search returned the revision PDF itself (https://digitaleconomy.stanford.edu/app/uploads/2026/08/Canaries_August2026.pdf, search extract) plus independent secondary coverage (aiweekly.co). Both give 19% as of June 2026, the mechanism as reduced hiring rather than separations, and no economy-wide displacement. New detail from the extract: the revision puts the gap at 15% for the July 2025 data vintage (not the 13% headline of the first version, so the earlier vintage itself appears to have been re-estimated), and the sample as ADP payroll for about 4.6 million workers in 730+ occupations. The 19% figure is now corroborated by two search results, but still not read in full. EVIDENCE.

    - Stanford/ADP "Canaries Dashboard" launched June 2026, monthly updates; Fortune reported employment in highly exposed occupations shrinking at 3.8% a year for 22 to 25 year olds (search summary). https://digitaleconomy.stanford.edu/project/indicators/canaries-dashboard/ ; https://fortune.com/2026/06/27/what-is-ai-impact-entry-level-jobs-stanford-adp-canaries-brynjolfsson-richardson/
    - Limits: ADP clients are not a random sample of US employers; occupation exposure is assigned, not measured use; descriptive by the authors' own account.

- **[EVIDENCE]** Hosseini and Lichtinger (Harvard), "Generative AI as Seniority-Biased Technological Change", version of 6 May 2026 (fetched). Revelio Labs LinkedIn résumé and posting data: about 66 million workers, 281,000 US firms, 160 million employment spells, almost 200 million postings, 2015 to 2025. Adoption identified by posting a "GenAI integrator" role (LLM-classified): 10,433 adopters, 3.7% of firms but 16% of employment. Results:
    - Junior employment at adopters fell about 9% relative to non-adopters six quarters after 2023Q1; senior employment showed no break. Triple-difference with firm-by-time and industry-by-time-by-seniority fixed effects gives the same answer.
    - Staggered event study: no pre-trends for eight quarters; junior employment down 8% eight quarters after adoption.
    - Junior employment in high-exposure occupations fell about 7% relative to low-exposure occupations between 2022Q4 and 2025Q1.
    - Mechanism: adopters hired 4.0 fewer juniors per quarter, about 80% below the pre-period mean of 5.06; junior separations fell (by 1.1 per quarter); promotions rose slightly.
    - Interest rates: a formal test using high-frequency monetary-policy surprises finds juniors at adopters do not respond differently to rate shocks.
    - U-shaped by university prestige: mid-tier graduates hit hardest.
    - Adopters are concentrated in professional services (28%) and information (24%). The authors report the hiring contraction is broad-based across industries. (Earlier secondary coverage said effects were largest in wholesale and retail; the May 2026 text says broad-based. Cite the current version.)
    - Limits: adoption is proxied by a job posting; adopters are larger, more senior-heavy firms; LinkedIn under-represents some sectors.

- **[EVIDENCE]** Anthropic CPS analysis: job-finding rate for 22 to 25 year olds entering the most exposed occupations fell about 14% relative to 2022, "just barely statistically significant"; no decrease for over-25s; unemployment of young exposed workers is flat. The authors note non-hires may be staying in existing jobs, taking other jobs or returning to school. 5 Mar 2026 (fetched). https://www.anthropic.com/research/labor-market-impacts

- **[EVIDENCE]** Lodefalk, Löthman, Koch, Engberg (2026), Swedish job-posting and register data, employer-level difference-in-differences: since ChatGPT, employment of 22 to 25 year olds in high-exposure occupations fell about 5.5% relative to less-exposed occupations within the same employers; workers over 50 rose slightly. Cited in Hosseini and Lichtinger May 2026 (fetched). This is the main European evidence in this file and it matches the US pattern at a smaller magnitude.
    - Gap-fill: original now located. "Same Storm, Different Boats: Generative AI and the Age Gradient in Hiring", Örebro University / Ratio working paper 2/2026, dated 16 Mar 2026. Whole-workforce Swedish register data plus about 4.6 million ads from Sweden's largest recruitment platform. Adds two points: the broad fall in Swedish postings since 2022

lines up with monetary tightening, not AI; the junior decline (5.5% by early 2025) runs through slower hiring, not separations. A university press release gives 6.5% (different window or specification, unresolved). https://www.oru.se/globalassets/oru-sv/institutioner/hh/workingpapers/workingpapers2026/wp-2-2026.pdf (search extract)

- **[EVIDENCE, weak]** Revelio Labs (Simon, 2025): entry-level postings down more than 35% since Jan 2023; a 10-point rise in exposure predicts an 11% fall in entry-level demand while senior roles in the same occupations rise 7%. Vendor analysis of its own data. Cited in Hosseini and Lichtinger (fetched). https://www.reveliolabs.com/news/macro/is-ai-responsible-for-the-rise-in-entry-level-unemployment/

- **[EVIDENCE]** NY Fed "Labor Market for Recent College Graduates": unemployment for recent graduates (22 to 27) about 5.6% in Q1 2026, above the all-worker rate. By major (2024 data, updated annually each February): computer engineering about 7.5 to 7.8%, computer science 6.1% or 7.0% depending on the summary. Search summaries only; the CS figure conflicts between sources, check the NY Fed table directly. https://news.bloomberglaw.com/new-york-brief/unemployment-for-recent-college-grads-remains-high-ny-fed-says

    - Caveat: the recent-graduate unemployment premium over all workers began widening before ChatGPT. The Economic Innovation Group argued a viral chart on graduate unemployment was misleading (search summary, not read). https://agglomerations.eig.org/p/a-viral-chart-on-recent-graduate . Separate research reported by LAist attributes graduates' difficulties to remote work rather than AI (headline only, not read). https://laist.com/news/education/remote-work-not-ai-has-sidelined-recent-college-graduates-research-finds

- **[EVIDENCE]** Dallas Fed (Tyler Atkinson and Shane Yamco, 6 Jan 2026), own CPS analysis: some decline in employment for young workers in the most AI-exposed occupations, driven mainly by lower inflows (especially from outside the labour force) rather than outflows. Their scale check: if the whole decline for young exposed workers turned into unemployment, it would account for only about 0.1 percentage point of the rise in aggregate US unemployment since November 2022. https://www.dallasfed.org/research/economics/2026/0106 (search extract). This is the most useful single number for sizing the junior effect against the macro picture.

- **[EVIDENCE, secondary]** Dallas Fed (J. Scott Davis, Feb 2026): employment across the top 10% of AI-exposed industries down about 1% over the same period; argues returns to experience are rising in exposed occupations. Secondary coverage only (Fortune, 4 Mar 2026). https://fortune.com/2026/03/04/gen-z-lack-of-experience-job-wage-loss-older-worker-experience-dallas-fed (search extract)

- **[EVIDENCE, weak: correlation]** St. Louis Fed (Serdar Ozkan and Nicholas Sullivan, "Is AI Contributing to Rising Unemployment? Evidence from Occupational Variation", On the Economy, 26 Aug 2025): occupations with higher AI exposure saw larger unemployment increases 2022 to 2025; computer and mathematical occupations (exposure around 80%) among the steepest rises. Correlation reported as 0.47 for theoretical exposure and 0.57 for measured adoption (summaries disagree). The authors call it tentative and non-causal. This conflicts with Anthropic's CPS null and the Yale Budget Lab null; the difference is likely method (simple cross-occupation correlation without a pre-period comparison vs difference-in-differences). The DiD studies are stronger. https://fedinprint.org/item/l00001/101478 (search extract)

- **[EVIDENCE, weak: platform data]** Handshake (US student and graduate jobs platform): entry-level and internship postings down about 15 to 16% year on year in 2025 (Aug 2024 to Aug 2025), applications per posting up 26 to 30%; by spring 2026 postings down about 2% year on year and about 12% below pre-pandemic. Handshake's own report not seen; figures from secondary coverage, which itself attributes the squeeze to the economy and federal cuts as well as AI. https://fortune.com/2025/11/19/next-level-leadership-gen-z-college-graduates-tough-market-how-to-stand-out-according-to-jobs-experts/ (search extract)
- **[EVIDENCE, weak: VC firm, own data]** SignalFire "State of Tech Talent Report 2025" (May 2025): new-graduate share of hires at the largest tech firms about 7%, down about 25% from 2023 and more than 50% from 2019; at startups under 6%, down over 30% from pre-pandemic. Viral "big tech cut new-grad hiring by 50%" traces here: it is a fall in the new-grad share of hires versus 2019, not a 50% one-year cut in hires, and the peer group is described inconsistently (12 vs 15 firms). SignalFire itself lists smaller funding rounds, smaller teams and fewer new-grad programmes alongside AI. https://www.signalfire.com/blog/signalfire-state-of-talent-report-2025 (search extract)
- **[EVIDENCE] update to NY Fed graduates:** NY Fed college labor market page gives recent-graduate unemployment about 5.7% in Q1 2026 (underemployment 41.5%); a later report gives about 5.6% for Q2 2026 (underemployment 42%) against roughly 4.1 to 4.3% for all workers. A separate NY Fed study attributes about 64% of the rise in young-graduate unemployment from March 2019 to March 2026 to remote work, not AI (secondary coverage). https://www.newyorkfed.org/research/college-labor-market (search extract)

US Federal Reserve firm surveys

- **[EVIDENCE]** NY Fed regional business surveys (Liberty Street Economics, "Are Businesses Scaling Back Hiring Due to AI?", 4 Sep 2025): AI use in the past six months rose to 40% of New York and northern New Jersey service firms (from 25% a year earlier) and 26% of manufacturers (from 16%). Only about 1% of service firms reported AI-related layoffs and no manufacturers did; 35% of AI-using service firms were retraining workers (up from 24%). About 13% of service firms expected AI layoffs ahead, and firms expected hiring cutbacks to grow, especially for college-educated workers. Regional, firm self-report, and the workforce questions cover only AI users. https://libertystreeteconomics.newyorkfed.org/2025/09/are-businesses-scaling-back-hiring-due-to-ai/ (search extract)
- **[EVIDENCE/OPINION]** NY Fed follow-ups: Audoly, Guerin, Topa (Liberty Street, 14 May 2026) find little sign of a distinct AI-driven fall in job postings, with the slowdown attributed mainly to rates, post-pandemic over-hiring and general weakness (via Reuters coverage); Kartik Athreya ("AI's Impact on Labor and Hiring", 5 Aug 2026) summarises the evidence as a shift in skill requirements, not job elimination, with firms still expecting smaller hiring plans for graduates. https://libertystreeteconomics.newyorkfed.org/2026/08/ais-impact-on-labor-and-hiring/ (search extract)

Separating AI from interest rates and post-pandemic over-hiring

- **[EVIDENCE]** The strongest attempts to separate these come from within-firm and within-employer designs: Hosseini and Lichtinger (firm-by-time fixed effects plus a monetary-shock sensitivity test, fetched) and Lodefalk et al. (same-employer comparisons). Both still find a junior-specific decline. Brynjolfsson et al. say firm-time controls, remote-work and tech-cycle checks leave the result intact (search summary).

- **[OPINION]** The remaining weakness: all designs compare "exposed" with "unexposed" jobs, and the exposed set (software, customer service, admin, finance) overlaps heavily with the sectors that over-hired in 2021 to 2022 and are most rate-sensitive. Hosseini and Lichtinger list macro uncertainty, post-pandemic adjustment and offshoring as the competing explanations raised by others (fetched). No study yet has a clean natural experiment on adoption.
- **[EVIDENCE]** Hosseini and Lichtinger propose that the early timing (declines starting soon after ChatGPT, before most production deployments) points to forward-looking hiring cuts: firms stop hiring for roles they expect to automate. Their data cannot test this directly. Fetched.

Freelance and online labour markets (the cleanest causal evidence)

- **[EVIDENCE]** Demirci, Hannane, Zhu, "Who Is AI Replacing?", Management Science (2025). 1,386,642 job posts, Jul 2021 to Jul 2023, one global freelancing platform. About 21% fewer posts for automation-prone writing and coding jobs relative to manual-intensive jobs within about eight months of ChatGPT; writing about minus 30%, coding about minus 20%; image creation about minus 17% after image generators; applications per job up 9% (more competition). Larger falls where Google-search awareness of ChatGPT was higher. Search summary. https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb
- **[EVIDENCE]** Hui, Reshef, Zhou, "The Short-Term Effects of Generative AI on Employment: Evidence from an Online Labor Market" (CESifo WP 10601). Upwork; after ChatGPT, writing freelancers' monthly jobs fell about 2% and monthly earnings about 5.2%; similar for image freelancers after DALL-E 2/Midjourney; top-rated freelancers were not protected and were, if anything, hit harder. Authors stress short-term. Search summary. https://www.cesifo.org/DocDL/cesifo1_wp10601.pdf

Firm-level adoption and employment growth (pre-GenAI and GenAI)

- **[EVIDENCE]** Babina, Fedyk, He, Hodson (2024), Journal of Financial Economics 151:103745: US firms investing in AI in the 2010s grew faster in sales, employment and product innovation, with more educated workforces. Pre-GenAI. As summarised in Hosseini and Lichtinger (fetched).
- **[EVIDENCE]** Acemoglu, Autor, Hazell, Restrepo (2022), Journal of Labor Economics, Burning Glass postings 2010 to 2018: AI-exposed establishments reduced non-AI hiring, but occupation and industry-level aggregate effects were too small to detect. Summarised in Hosseini and Lichtinger and listed in the Anthropic paper (both fetched).
- **[EVIDENCE]** Hampole, Papanikolaou, Schmidt, Seegmiller (2025, NBER), "Artificial Intelligence and the Labor Market": 2010 to 2023 résumé data; higher task exposure lowers demand for those tasks, but productivity gains lead firms to expand employment elsewhere, so net headcount changes are muted. Summarised in Hosseini and Lichtinger (fetched).

Exposure (FORECAST-type measures, not effects)

- **[FORECAST]** Eloundou, Manning, Mishkin, Rock, "GPTs Are GPTs" (published 2024; cited as Science in the Hosseini bibliography): task-level β score for whether an LLM could halve task time. It is the base for most exposure measures used in the effect studies above, including Anthropic's. Exposure is capability, not displacement. Fetched references only.
- **[EVIDENCE/FORECAST]** Anthropic "observed exposure": Claude covers 33% of Computer and Math tasks, well below theoretical feasibility; computer programmers highest at 75% task

coverage, customer service representatives and data entry keyers 67%; about 30% of workers have zero coverage. The most exposed quartile earns 47% more on average, is more female and more educated (17.4% hold graduate degrees vs 4.5% in the unexposed group). Each 10-point rise in coverage is associated with a 0.6-point lower BLS 2024 to 2034 growth projection, which the authors call slight. 5 Mar 2026 (fetched). https://www.anthropic.com/research/labor-market-impacts

- **[EVIDENCE]** Anthropic Economic Index, Jan 2026: adjusting exposure for task success rates makes data entry keyers and radiologists look more affected and teachers and software developers less affected than raw task coverage. Search summary. https://www.anthropic.com/research/anthropic-economic-index-january-2026-report
- **[FORECAST]** IMF Staff Discussion Note SDN/2024/001, "Gen-AI: Artificial Intelligence and the Future of Work" (Cazzaniga and others, 14 Jan 2024): almost 40% of global employment exposed to AI; about 60% in advanced economies, 40% in emerging markets, 26% in low-income countries; roughly half of exposed jobs in advanced economies judged likely to be complemented rather than substituted. Method: occupation-level task exposure (Felten-style AIOE) plus the authors' own "complementarity" coding. This is exposure, not displacement: it says nothing about how many jobs will be lost, and SDNs are staff views, not IMF positions. Headlines turning "40% exposed" into "40% of jobs will be lost" misread it. https://wiki.ubc.ca/images/b/bd/IMF_GenAI_Future_of_Work.pdf (search extract; copy of the IMF PDF, imf.org not reached)
- **[FORECAST]** Goldman Sachs (Hatzius; Briggs and Dong, 2025): generative AI will "eventually" displace 6 to 7% of US workers over the adoption period, range 3 to 14% under other assumptions; peak unemployment effect about 0.5 percentage point, with displacement effects typically fading after about two years as workers are re-employed. A 2026 update reportedly raised the share to over 9% over ten years with a peak unemployment effect under 1 point (PYMNTS only, unconfirmed). The same research notes only 9.3% of US firms reported using GenAI in production in the prior two weeks (Census BTOS). Model-based projection, not a measurement. https://www.staffingindustry.com/news/global-daily-news/ais-impact-on-jobs-likely-temporary-goldman-sachs ; 2026 update https://pymnts.com/economy/2026/goldman-sachs-says-ai-will-eliminate-15-million-us-jobs (search extracts; Goldman primary not seen)
- Felten et al. AIOE and the older Goldman "300 million jobs exposed" (2023) figure: still not verified. Do not cite from this file.

Layoff announcements citing AI (Challenger, Gray & Christmas)

- **[EVIDENCE, weak]** Challenger counts of US announced job cuts where employers cited AI: 54,836 for all of 2025; May 2026 38,579 (a monthly record for the category); June 14,029; July 10,970 of 33,429 total; through September 120,136, September only 3,961; October 1,265; January to October 2026 121,401, about 20% of all announced cuts [NOT CONFIRMED in verification pass; the October 2026 report could not be found by an independent search; DO NOT CITE until checked]. By October "Market/Economic Conditions" (130,531) overtook AI as the top cited reason. Total announced cuts Jan to Oct 2026 were 617,189, down 44% on 2025. All search summaries (challengergray.com blocked). https://www.challengergray.com/blog/october-layoffs-at-three-year-low-seasonal-hiring-plans-up-26/ ; https://www.challengergray.com/wp-content/uploads/2026/08/Challenger-Report-July-2026.pdf
- **Reliability.** (1) These are announcements, often multi-year and global, not realised US layoffs. (2) The reason is what the company says, coded by Challenger; there is an incentive to cite AI (investors reward it) and a disincentive to cite weak demand. Yale Budget Lab's Gimbel has publicly suggested some AI attributions are cover for other causes (Fortune, Feb 2026, search summary). (3) Scale: BLS JOLTS layoffs and discharges averaged about 1.1% of

employment per month over the 12 months to April 2026, about 13.4% annualised (as cited in Anthropic's June 2026 report, fetched). On roughly 159 million US payroll jobs (my approximation, not verified this session) that is around 1.7 million involuntary separations a month, so 121,000 AI-cited announcements over ten months is well under 1% of involuntary separations. AI-cited cuts are a real signal of corporate intent but a tiny share of actual job loss.

UK and Europe

- **[EVIDENCE, weak]** Morgan Stanley (reported Jan 2026): survey of firms that had used AI for at least a year, in five sectors (consumer staples and retail, real estate, transportation, healthcare equipment, autos), in the UK, US, Japan, Germany and Australia. UK respondents reported cutting 23% of jobs and creating 15% over 12 months, net minus 8%, about twice the cross-country average; US respondents net positive (one social-media summary gives US minus 17% / plus 19%, net plus 2%; Japan minus 7%, Germany and Australia minus 4%; unconfirmed). Search summaries only; primary note not seen. https://www.eweek.com/news/ai-job-losses-uk-morgan-stanley-report/ ; https://www.itpro.com/business/business-strategy/business-leaders-are-using-ai-as-a-license-to-reduce-headcount-new-morgan-stanley-research-lays-bare-the-impact-on-uk-workers
    - What it measures: the sample is AI-adopting firms in selected sectors, self-reported, not a representative sample of UK employers, and the 23% figure is not attributed to AI by any causal design. Sample size not found. See "Viral statistics traced".
- **[EVIDENCE]** Denmark: Humlum and Vestergaard, "Large Language Models, Small Labor Market Effects", NBER WP 33777 (2025; revised Oct 2025). Two adoption surveys (late 2023, 2024) of about 25,000 workers in 7,000 workplaces across 11 exposed occupations, linked to administrative employer-employee records. No significant effect on earnings or recorded hours in any occupation; confidence intervals rule out effects larger than about 1%; average reported time savings about 3% (2.8%). Search summary; existence and "no effects on earnings and hours" confirmed in Hosseini and Lichtinger (fetched). https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf
    - A 2026 version titled "Still Waters, Rapid Currents: Early Labor Market Transformation under Generative AI" (RFBerlin discussion paper 26078, March 2026). Gap-fill (abstract via search extract): most employers in exposed occupations have adopted chatbot initiatives, workers report productivity gains and new AI-related tasks are widespread, but difference-in-differences still gives precise null effects on earnings and recorded hours at worker and workplace level, ruling out effects larger than 2% two years after ChatGPT (the bound loosened from about 1% in the 2025 version). The authors' framing: technology reshapes work before it shows in earnings or hours. https://ideas.repec.org/p/crm/wpaper/26078.html (search extract)
- **[EVIDENCE]** Sweden: Lodefalk et al. (2026), see above: young workers in exposed occupations down about 5.5% within employers. Original working paper located (search extract).

UK official and central bank sources (gap-fill, Nov 2026)

- **[EVIDENCE, weak: postings, not causal]** Bank of England staff, Bank Underground blog (Haley Schlicht, Data and Statistics Division), "Canaries in the column? AI exposure and the UK's hiring slowdown", 6 Aug 2026: over about three years, online job adverts fell 15% in the most AI-exposed occupation group, 10% in the middle group and 6% in the least exposed; customer service adverts down 23%, administrative down 22%. The post treats AI as a plausible contributor, not an identified cause. Bank Underground posts are staff views, not

Bank policy. https://bankunderground.co.uk/2026/08/06/canaries-in-the-column-ai-exposure-and-the-uks-hiring-slowdown/ (search extract)
    - Counterpoint in the same coverage: NY Fed and Bloomberg Economics both found vacancies in AI-exposed roles were already falling before ChatGPT's release, which weakens an AI reading of the UK gradient. Bloomberg Economics (Jun 2026) further reports UK vacancies in AI-exposed roles have risen since summer 2024 and that firms themselves blame weak demand and higher employment costs (April 2025 employer National Insurance and minimum-wage rises). https://news.bloomberglaw.com/artificial-intelligence/ai-wrongly-blamed-for-britains-job-losses-analysis-suggests (search extract; methodology not seen)
- **[OPINION]** Andrew Bailey, Governor of the Bank of England (BBC Radio 4 and press, Dec 2025): AI is likely to displace workers in a way comparable to the Industrial Revolution but he does not expect mass unemployment; his main worry is entry-level roles as the pipeline into skilled professions. Context given in coverage: UK unemployment 5.1% in the three months to October 2025, with young people hit hardest. https://hrreview.co.uk/hr-news/strategy-news/ai-set-to-displace-jobs-bank-of-england-governor-warns/385541 ; https://www.insurancejournal.com/news/international/2025/12/19/851754.htm (search extracts). An October 2025 Bloomberg headline had him saying AI will not lead to UK job losses (headline only). No MPC speech with AI-specific employment data found.
- **Not found: Decision Maker Panel AI headcount question.** Searches turned up DMP general employment expectations (expected employment growth turned slightly negative, about minus 0.1%, in autumn 2025; secondary coverage) but no published DMP breakdown of AI-attributed employment effects. Do not cite a DMP AI figure until someone checks the DMP releases directly.
- **[EVIDENCE, weak: firm self-report, partly unverified]** ONS Business Insights and Conditions Survey (BICS): from March 2026 BICS asks AI-using firms what effect AI has had on headcount (increased, no change, reduced, not sure). A secondary site reports 6% of AI-using firms with 10+ employees said AI had reduced headcount (September 2026 wave); not confirmed in an ONS bulletin. The question is visible in the Scottish Government's weighted BICS workbook (wave 153, table "AIWorkforceImpact"). In 2025, firms planning adoption expected admin and clerical roles (27%), creative and design (23%) and data analysis (19%) to be most affected: expectations, not outcomes. https://www.gov.scot/publications/bics-weighted-scotland-estimates-data-to-wave-153/pages/artificial-intelligence/ ; ONS release page https://www.ons.gov.uk/releases/aiinukbusinesses (search extracts)
- **[EVIDENCE]** ONS FOI response (FOI-2026-3465): ONS does not track the relationship between AI adoption and graduate employment outcomes. So there is no official UK equivalent of the Canaries or Anthropic CPS analyses. https://www.ons.gov.uk/aboutus/transparencyandgovernance/freedomofinformationfoi/impactofaiongraduateemployment (search extract)
- **[FORECAST]** Warwick Institute for Employment Research for DSIT (Jan 2026): jobs directly involving AI activities projected to rise from 158,000 (2024) to 3.9 million by 2035. Modelled projection. https://www.gov.uk/government/publications/ai-skills-for-life-and-work-labour-market-and-skills-projections/ai-skills-for-life-and-work-labour-market-and-skills-projections (search extract)
- **[EVIDENCE, weak: firm survey]** CIPD Labour Market Outlook (Nov 2025): about one in six UK employers expect AI to let them reduce headcount in the next 12 months; a quarter of those expect to cut more than 10% of staff. Intentions, not outcomes. https://investinglive.com/news/uk-firms-expect-3-pay-rises-but-warn-ai-could-shrink-jobs-cipd-survey-shows-20251110 (search extract; CIPD primary not seen)

UK graduate and entry-level postings (gap-fill)

- **[EVIDENCE, weak: postings]** Indeed Hiring Lab UK (Aug 2026): graduate postings about 7% down year on year, lowest for the time of year since 2020; AI terms in a record 9.4% of UK postings at end June 2026. Indeed attributes the fall to cyclical weakness and AI together. Some aggregators report 13%; Indeed's own 7% is preferred. https://www.hiringlab.org/uk/?p=14268 ; https://workplacejournal.co.uk/2026/08/graduate-job-postings-at-lowest-level-since-pandemic-indeed/ (search extracts)
- **[EVIDENCE, weak: postings, vendor]** Adzuna (Jun 2025): UK entry-level adverts (graduate, apprenticeship, internship, junior no-degree) down 31.9% since November 2022; entry-level share of all listings from 28.9% to about 25%; graduate adverts in May 2025 down 28.4% year on year. A techUK page reports graduate listings below 10,000 in January 2026, down 45% year on year (secondary). The "since ChatGPT" window bundles AI with the 2023 to 2025 rate cycle and the April 2025 employer cost rises. https://hrreview.co.uk/hr-news/recruitment/entry-level-jobs-have-declined-since-chatgpt-launch ; https://www.techuk.org/insights/news/what-s-actually-happening-with-entry-level-and-graduate-jobs (search extracts)
- **[EVIDENCE]** Institute of Student Employers, Student Recruitment Survey 2025 (155 employers, about 31,000 hires, 1.8 million applications): graduate hiring down 8% in 2024/25, weakest year since 2020 (minus 12%); 42% of employers cut graduate hiring, 33% increased; 140 applications per vacancy, up 59%. The survey's AI material is about AI in selection (applicants using AI), not AI replacing graduate roles. The 2026 survey reports graduate hiring up about 4% in 2025/26, against an earlier employer forecast of minus 7%. Membership survey of large graduate employers, not representative. https://ise.org.uk/_userfiles/pages/files/reports/student_recruitment_survey_2025.pdf ; 2026 coverage https://wonkhe.com/wonk-corner/a-slightly-better-outlook-for-graduate-jobs (search extracts)
    - Read across: job-board postings (Adzuna, Indeed) show large falls; the ISE employer survey shows a smaller fall followed by a rebound. Postings are a noisy proxy for hires.

Euro area and OECD (gap-fill)

- **[EVIDENCE]** ECB blog (Laura Lebastard and David Sondermann, 4 Mar 2026), from the ECB Survey on the Access to Finance of Enterprises (SAFE) AI modules, Q2 and Q4 2025, about 5,000 firms: two thirds of firms say employees use AI tools, far fewer have invested; firms making significant use of or investing in AI are more likely to be hiring (reported as about 2% to 4% more likely, summaries differ); "AI-intensive firms tend, on average, to hire rather than fire"; effect driven by small firms and neutral for large ones. Only about 15% of AI-using firms named cutting labour costs as a motive, and those firms do show lower hiring and more layoffs. Authors: positive so far largely because AI has not yet transformed production processes. https://www.ecb.europa.eu/press/blog/date/2026/html/ecb.blog20260304~d9e34fc95f.en.html (search extract)
- **[EVIDENCE]** ECB Occasional Paper No. 395 (Ferrando, Lamboglia, Rariga, Schmidt, 2026), same SAFE modules: no sign of aggregate labour shedding; AI adoption positively associated with employment growth and expectations. ECB Economic Bulletin box (Issue 2/2026): significant AI users expect employment 1.3 percentage points higher over 12 months than non-users. Firm self-report; adopters may be faster-growing firms anyway (selection, not causation). https://www.ecb.europa.eu/pub/pdf/scpops/ecb.op395.en.pdf ; https://ecb.europa.eu/press/economic-bulletin/focus/2026/html/ecb.ebbox202602_06~c9f485c478.en.html (search extracts)
- **[OPINION/EVIDENCE]** OECD: the 2025 Employment Outlook's main theme was demographic change, not AI; no dedicated AI-effects chapter found in search extracts (Sweden's OECD

delegation summary). The 2023 Outlook found little evidence of negative AI employment effects so far. The OECD Employment Outlook 2026 has a chapter on AI and entry-level work (contents not seen). An OECD blog (Aug 2026), "Young people, labour markets and AI: what OECD evidence shows", says current OECD evidence does not show AI is responsible for the deterioration in youth labour markets and that the trend began before large language models. A figure circulating as "OECD: unemployment of 20 to 30 year olds in AI-exposed roles up 3 points in 2025" could not be found in an OECD primary source; treat as unverified. https://www.oecd.org/en/blogs/2026/08/young-people-labour-markets-and-ai-what-oecd-evidence-shows.html ; https://www.swedenabroad.se/en/embassies/oecd-unesco/current/news/oecd-employment-outlook-2025/ (search extracts)

- **Remaining gap.** Bank of England also hosted a Chandar seminar on "Canaries" in 2026 (slides URL seen in search results only). https://www.bankofengland.co.uk/-/media/boe/files/events/2026/b-chandar-slides.pdf

Worker perceptions vs measured effects

- **[EVIDENCE, weak: self-selected]** Anthropic Economic Index Survey (report 26 Jun 2026): about 9,700 Claude users linked to usage logs; not representative (30% computer and maths occupations vs 4% of US employment). 10% rated losing their own job in the next year as likely or very likely, slightly below the realised US involuntary separation rate; 38% of those attributed it to AI (an upper bound). Respondents worried more about others than themselves, especially junior colleagues: over a third put the chance of a junior colleague losing their job in the next year above 60%. Fetched. https://www.anthropic.com/research/economic-index-june-2026-report
- **[EVIDENCE, weak: self-selected]** Anthropic, 81,000 Claude.ai users, open-ended interviews coded by Claude (22 Apr 2026): about one fifth voiced displacement concern; concern rises with occupational exposure and is higher early in careers; 60% of early-career vs 80% of senior respondents said they personally benefit. Measures perceived threat only. Fetched. https://www.anthropic.com/research/81k-economics
- **[EVIDENCE]** "The AI expectations gap" (Anthropic, 16 Oct 2026): Morning Consult survey of US adults, n=10,980, 11 to 23 Aug 2026, weighted to national representativeness. Workers who expect AI to do 80%+ of their work by 2030 forecast 12.2% unemployment in 2030; those under 50% forecast 7 to 8% (baseline given: 4.2%). Median expected time for a displaced worker to find a role in a new occupation: 7 to 12 months (public) and 12 months (Anthropic staff). Fetched. https://www.anthropic.com/research/ai-expectations-gap

Viral statistics traced

- **"13% decline in jobs for young workers because of AI" (Canaries).** Origin: Brynjolfsson, Chandar, Chen, Aug 2025, ADP payroll data. What it is: a relative employment gap for 22 to 25 year olds in the most exposed occupations vs less exposed occupations, not an absolute fall and not a causal estimate by the authors' own description. The number has since been revised (16% Nov 2025; 19% by June 2026 in the Aug 2026 revision; Anthropic cites a 6 to 16% range across specifications). Solid as a descriptive fact from a large administrative dataset; weaker as proof of causation. When citing, say "relative" and give the version date.
- **"The UK lost 8% of jobs to AI" (Morgan Stanley, Jan 2026).** Origin: a Morgan Stanley survey of companies already using AI for at least a year, in five sectors and five countries, reporting their own 12-month headcount changes (23% cut, 15% created). Not a measure of the UK labour market, not a causal estimate, sample size not reported in coverage seen.

Headlines turning it into "UK lost 8% of jobs to AI" misrepresent it. ONS payroll data (not checked this session) would show nothing like an 8% fall in UK employment. Weak. (Search summaries only.)

- **"AI is behind one in five layoffs" (Challenger, 2026).** Origin: Challenger's coding of employer-stated reasons in announced cuts. About 20% of announced cuts Jan to Oct 2026 cited AI. Announced cuts are themselves a small fraction of actual layoffs (see JOLTS comparison above). Weak as a measure of displacement; reasonable as a measure of what companies want to be seen saying.
- **"40% of jobs will be affected by AI" (IMF, Jan 2024).** Origin: IMF Staff Discussion Note SDN/2024/001 (Cazzaniga and others). It is an exposure measure (task overlap by occupation), 60% in advanced economies, with about half of exposed jobs judged likely to be complemented. It is not a forecast of job losses and is a staff view. FORECAST at most. (Search extract.)
- **"Big Tech cut new-grad hiring by 50%" (SignalFire, May 2025).** Origin: SignalFire, a VC firm, using its own talent data. The 50% is the fall in the new-graduate share of hires at the largest tech firms versus 2019 (to about 7% of hires); versus 2023 the fall is about 25%. Peer group described as 12 or 15 firms in different write-ups. SignalFire lists funding and team-size causes alongside AI. Weak as evidence of AI displacement. (Search extract.)
- **"UK entry-level jobs down a third since ChatGPT" (Adzuna, Jun 2025).** Origin: Adzuna job-board adverts, Nov 2022 to mid-2025. The window coincides with UK monetary tightening and, from April 2025, higher employer National Insurance and minimum wage. Before-after comparison with no control group; postings, not hires. The ISE employer survey shows graduate hiring down 8% then up 4%. Weak as evidence about AI. (Search extract.)
- **"Goldman: AI will displace 6 to 7% of US workers."** Origin: Goldman Sachs economics research (Hatzius; Briggs and Dong, 2025). A model-based FORECAST over the full adoption period (range 3 to 14%), with a peak unemployment effect of about 0.5 point that fades in about two years. Often quoted without the time horizon or the re-employment assumption. A 2026 update reportedly raised it to over 9% over ten years (PYMNTS only). (Search extract.)
- **"Computer science graduates have among the highest unemployment rates" (NY Fed).** Origin: NY Fed outcomes-by-major table, which is updated annually and in 2026 reflected 2024 data. Summaries disagree (6.1% vs 7.0% for CS). The table is real; attributing the rate to AI is an inference nobody has tested with that table.

Strongest counter-evidence

Against "AI is already displacing workers at scale":

- Unemployment of highly exposed US workers has not risen relative to unexposed workers, with power to detect about 1 point (Anthropic, Mar 2026, fetched). Yale Budget Lab finds no unusual occupational churn (Oct 2025, search summary). Denmark's linked administrative data rules out earnings or hours effects above about 1% (Humlum and Vestergaard 2025).
- Even the Canaries authors say there is no widespread, economy-wide displacement (Aug 2026, search summary).
- AI-cited layoff announcements were falling sharply through autumn 2026 and are dwarfed by ordinary layoffs.
- Firm-level evidence from before GenAI (Babina et al. 2024; Hampole et al. 2025) finds AI-investing firms grow employment, with job losses in exposed tasks offset by expansion elsewhere.

- Euro area firm surveys (ECB SAFE, about 5,000 firms, 2025) find AI-using firms more likely to hire, not fire; NY Fed regional surveys (Sep 2025) find about 1% of service firms reporting AI layoffs and more retraining than cutting. Denmark's 2026 update still rules out earnings or hours effects above 2%. (All search extracts.)
- The Dallas Fed estimates that even if all of the young-exposed-worker employment decline became unemployment, it would explain about 0.1 point of the US unemployment rise since Nov 2022 (search extract). OECD (Aug 2026 blog) and Bloomberg Economics on the UK say youth and vacancy weakness began before LLMs.

Against the mainstream reading "nothing is happening":

- The junior-hiring effect is consistent across four independent data sources (ADP payroll, Revelio résumés, CPS, Swedish registers), uses within-firm controls in at least two, and survives a monetary-policy test. A hiring freeze at the bottom of the ladder does not show up in unemployment rates for years, because incumbents keep their jobs and new entrants are scattered across other jobs, education or inactivity. "No unemployment effect" and "a pipeline effect" are compatible.
- The freelance studies show direct, fast substitution where work is contracted task by task, which is the market form most like how agents would be deployed.
- The UK has a similar gradient in postings: Bank of England staff find adverts down 15% in the most AI-exposed occupations vs 6% in the least over three years (Aug 2026, search extract), and the Governor names the entry-level pipeline as his main worry. The St. Louis Fed finds a positive cross-occupation correlation between exposure and unemployment increases (2025). Both are correlations, and both are weaker than the DiD nulls, but they point the same way as the junior-hiring studies.

Relevance to the thesis ("the tech is easy, transformation is hard"):

- Partly supports it: aggregate effects are small three years in, consistent with slow organisational absorption. Humlum and Vestergaard's 3% time saving and Hosseini and Lichtinger's observation that adopters are a small minority of firms fit "businesses are not ready".
- Partly undercuts it: firms are already cutting junior hiring without having transformed anything, apparently on expectations. If Hosseini and Lichtinger are right, displacement can run ahead of real capability and of organisational readiness. The human consequence (lost entry routes) does not wait for the transformation to succeed.

Open questions

- Is the junior-hiring drop permanent or a temporary anticipation effect that reverses if AI under-delivers? Hosseini and Lichtinger raise this explicitly.
- Where do the "missing" young hires go: other occupations, further study, inactivity? No study traces them yet.
- Does the effect appear in the UK? The ONS/HMRC PAYE real-time data split by age and occupation could replicate the Canaries design; I found no such study, and ONS says (FOI-2026-3465) it does not track AI adoption against graduate outcomes. UK evidence is postings-based only (Bank Underground, Indeed, Adzuna).
- What does the BoE Decision Maker Panel show on AI and employment? No published AI-specific DMP result was found in search extracts.

- What are the findings of the OECD Employment Outlook 2026 chapter on AI and entry-level work? Located, not read.
- Morgan Stanley UK survey: sample size, question wording, whether the 23% cut was attributed to AI by respondents.
- Does agentic deployment (2025 to 2026) produce a second, larger wave visible in 2026 to 2027 data, or does it look like the chat-assistant wave?

What's genuinely uncertain

- Causation. All junior-employment results are difference-in-differences between exposed and unexposed groups that also differ in rate sensitivity and pandemic over-hiring. Within-firm designs reduce but do not remove this.
- Magnitudes. The Canaries gap ranges 13% to 19% across versions; Revelio 9%; CPS job-finding minus 14% (barely significant); Sweden 5.5%. These measure different things.
- Exposure measures disagree about which jobs are exposed, and the Yale/EIG work shows results can depend on the measure chosen.
- Whether entry-level cuts are AI or "AI-washing" of cost cuts. Company statements are not evidence of mechanism.
- Europe. One Swedish and one Danish study is a thin base for generalising about Europe, and they point in different directions (Sweden: junior effect; Denmark: no earnings or hours effect, though it measured a different outcome). ECB firm surveys add a positive hiring association, but that is self-reported and likely confounded by adopter selection.
- UK vs US. The UK has higher youth unemployment and steeper falls in graduate postings than the US, but also a UK-specific cost shock (April 2025 employer National Insurance and minimum wage rises) that lines up in time. Nobody has separated these with UK microdata.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
No significant rise; detectable effect about 1pp	Differential unemployment, top-quartile exposed vs zero-exposure US workers, post-ChatGPT	Massenkoff and McCrory, Anthropic	https://www.anthropic.com/research/labor-market-impacts	2026-03-05	Fetched. CPS DiD. Author is an AI vendor
-14%	Job-finding rate for 22-25 year olds into most exposed occupations vs 2022	Same	https://www.anthropic.com/research/labor-market-impacts	2026-03-05	Fetched. "Just barely statistically significant"
75% / 67%	Claude task coverage: computer programmers / customer service reps and data entry keyers	Same	https://www.anthropic.com/research/labor-market-impacts	2026-03-05	Fetched. Exposure, not effect
about 13% (Aug 2025)	Relative employment decline, 22-25 in most exposed occupations	Brynjolfsson, Chandar, Chen, as cited by Hosseini and Lichtinger	https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf	2025-08	Fetched secondary citation; ADP data; descriptive
19% (by June 2026)	Same gap, Aug 2026 revision	Stanford Digital Economy Lab	https://digitaleconomy.stanford.edu/news/canariesaug26/ ; https://digitaleconomy.stanford.edu/	2026-0	Search extract; now corroborated by a second result (PDF + aiweekly). Revision gives 15% for the

Figure	What it measures	Source	URL	Date	Reliability note
			app/uploads/2026/08/Canaries_August2026.pdf	8-12	July 2025 vintage. 4.6m workers
about 0.1pp	Max contribution of young exposed workers' employment decline to US unemployment rise since Nov 2022	Dallas Fed, Atkinson and Yamco	https://www.dallasfed.org/research/economics/2026/0106	2026-01-06	Search extract. Upper-bound scale check
about 1% / 35%	NY-region service firms reporting AI layoffs / AI-using service firms retraining	NY Fed regional surveys	https://libertystreeteconomics.newyorkfed.org/2025/09/are-businesses-scaling-back-hiring-due-to-ai/	2025-09-04	Search extract. Regional firm self-report
0.47-0.57	Correlation of occupational AI exposure/adoption with unemployment increase 2022-25	St. Louis Fed, Ozkan and Sullivan	https://fedinprint.org/item/l00001/101478	2025-08-26	Search extract. Non-causal; summaries disagree on figure
5.7% (Q1 2026)	US recent-graduate (22-27) unemployment	NY Fed college labor market	https://www.newyorkfed.org/research/college-labor-market	2026	Search extract; Q2 reported as 5.6%
about 7% of hires	New-grad share of hires at largest tech firms, down >50% vs 2019, about 25% vs 2023	SignalFire State of Tech Talent 2025	https://www.signalfire.com/blog/signalfire-state-of-talent-report-2025	2025-05	Search extract. VC firm's own data
-15 to -16%	US entry-level postings on Handshake, Aug 2024 to Aug 2025; applications per post +26-30%	Handshake via Fortune	https://fortune.com/2025/11/19/next-level-leadership-gen-z-college-graduates-tough-market-how-to-stand-out-according-to-jobs-experts/	2025-11	Search extract; secondary
40% / 60%	Share of jobs exposed to AI, world / advanced economies	IMF SDN/2024/001 (Cazzaniga et al.)	https://wiki.ubc.ca/images/b/bd/IMF_GenAI_Future_of_Work.pdf	2024-01-14	FORECAST (exposure, not job loss). Search extract
6-7% (range 3-14%)	Share of US workers eventually displaced by GenAI over adoption period; peak unemployment +0.5pp	Goldman Sachs (Hatzius; Briggs and Dong)	https://www.staffingindustry.com/news/global-daily-news/ais-impact-on-jobs-likely-temporary-goldman-sachs	2025	FORECAST. Secondary; Goldman primary not seen
-15% / -10% / -6%	UK online adverts over about 3 years, high / mid / low AI-exposure occupations	Bank of England staff (Bank Underground, Schlicht)	https://bankunderground.co.uk/2026/08/06/canaries-in-the-column-ai-exposure-and-the-uks-hiring-slowdown/	2026-08-06	Search extract. Staff blog; postings; not causal; pre-ChatGPT trend disputed
-7% y/y	UK graduate postings, lowest for time of year since 2020	Indeed Hiring Lab UK	https://www.hiringlab.org/uk/?p=14268	2026-08	Search extract; postings
-31.9%	UK entry-level adverts since Nov 2022	Adzuna	https://hrreview.co.uk/hr-news/recruitment/entry-level-jobs-have-declined-since-chatgpt-launch	2025-06	Search extract; vendor postings; no control group
-8% (2024/25), +4% (2025/26)	Graduate hiring by ISE member employers	Institute of Student Employers	https://ise.org.uk/_userfiles/pages/files/reports/student_recruitment_survey_2025.pdf	2025 / 2026	Search extract. 155 large employers, about 31k hires
6%	UK AI-using firms (10+ staff) saying AI reduced headcount	ONS BICS via secondary site	https://www.ons.gov.uk/releases/aiinukbusinesses	2026-09	UNVERIFIED secondary figure; not seen in ONS bulletin
about 15%	Euro-area AI-using firms citing lower labour costs	ECB blog, Lebastard and	https://www.ecb.europa.eu/press/blog/date/2026/html/	2026-0	Search extract. Firm self-report

Figure	What it measures	Source	URL	Date	Reliability note
	as a motive; AI-intensive firms more likely to hire	Sondermann (SAFE, about 5,000 firms)	ecb.blog20260304~d9e34fc95f.en.html	3-04	
+1.3pp	Expected 12-month employment growth, significant AI users vs non-users, euro area	ECB Economic Bulletin 2/2026 box	https://ecb.europa.eu/press/economic-bulletin/focus/2026/html/ecb.ebbox202602_06~c9f485c478.en.html	2026	Search extract. Expectations
about -9% after 6 quarters	Junior employment at GenAI-adopting firms vs non-adopters	Hosseini and Lichtinger	https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf	2026-05-06	Fetched. 66m workers, 281k firms; adoption proxied by postings
about 80%	Fall in junior hires per quarter at adopters vs pre-period mean (4.0 of 5.06)	Same	same	2026-05-06	Fetched
3.7% / 16%	Share of firms classed as GenAI adopters / their share of employment	Same	same	2026-05-06	Fetched. Adoption by Mar 2025
-5.5%	22-25 employment in high-exposure occupations within Swedish employers	Lodefalk et al. 2026, cited by Hosseini and Lichtinger	same	2026	Fetched secondary citation; original not seen
>35%	Fall in US entry-level postings since Jan 2023	Revelio Labs (Simon 2025), cited by Hosseini and Lichtinger	https://www.reveliolabs.com/news/macro/is-ai-responsible-for-the-rise-in-entry-level-unemployment/	2025	Vendor data; secondary citation
-21%	Freelance writing and coding posts vs manual posts, 8 months after ChatGPT	Demirci, Hannane, Zhu, Management Science	https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb	2025	Search summary; 1.39m posts
-5.2%	Monthly earnings of Upwork writing freelancers after ChatGPT	Hui, Reshef, Zhou	https://www.cesifo.org/DocDL/cesifo1_wp10601.pdf	2023 WP	Search summary
<1%	Upper bound on earnings/hours effects, Danish exposed occupations	Humlum and Vestergaard, NBER 33777	https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf	2025	Search summary; 25k workers, 7k workplaces
<2%	Same bound, two years after ChatGPT, 2026 version ("Still Waters, Rapid Currents")	Humlum and Vestergaard, RFBerlin DP 26078	https://ideas.repec.org/p/crm/wpaper/26078.html	2026-03	Search extract (abstract)
121,401 (about 20%) [DO NOT CITE: unconfirmed, see verification.md]	US announced job cuts citing AI, Jan-Oct 2026	Challenger, Gray & Christmas	https://www.challengergray.com/blog/october-layoffs-at-three-year-low-seasonal-hiring-plans-up-26/	2026-11-05	Search summary; announcements, company-stated reasons
54,836	US announced job cuts citing AI, full year 2025	Challenger (via secondary)	https://aiunderstanding.org/news/challenger-says-employers-cited-ai-in-112-713-u-s-job-cuts-through-july-2026	2026-08	Search summary; secondary
	US layoffs and discharges rate, 12			2026-0	

Figure	What it measures	Source	URL	Date	Reliability note
about 1.1%/ month, 13.4% annualised	months to Apr 2026 (BLS JOLTS)	Anthropic citing BLS	https://www.anthropic.com/research/economic-index-june-2026-report	6-26	Fetched; check against BLS directly
net -8% (23% cut, 15% created)	Headcount change reported by UK firms using AI 1+ year, five sectors	Morgan Stanley via eWeek/ITPro	https://www.eweek.com/news/ai-job-losses-uk-morgan-stanley-report/	2026-01	Search summary; non-representative firm survey; NOT 8% of UK jobs
5.6%	Unemployment, US recent college grads 22-27, Q1 2026	NY Fed via Bloomberg Law	https://news.bloomberglaw.com/new-york-brief/unemployment-for-recent-college-grads-remains-high-ny-fed-says	2026	Search summary
12.2%	2030 unemployment forecast by US workers expecting AI to do 80% + of their work	Anthropic / Morning Consult, n=10,980	https://www.anthropic.com/research/ai-expectations-gap	2026-10-16	Fetched. Expectation, not outcome
10%	Claude users rating own job loss in next year likely/very likely	Anthropic Economic Index Survey, n≈9,700	https://www.anthropic.com/research/economic-index-june-2026-report	2026-06-26	Fetched. Not representative

Verbatim quotes

- "We find no systematic increase in unemployment for highly exposed workers since late 2022, though we find suggestive evidence that hiring of younger workers has slowed in exposed occupations" (Maxim Massenkoff and Peter McCrory, Anthropic economists, "Labor market impacts of AI: A new measure and early evidence", 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: key-findings summary of CPS analysis. Verified: curl+grep
- "The averaged estimate in the post-ChatGPT era is a 14% drop in the job finding rate compared to that in 2022 in the exposed occupations, although this is just barely statistically significant. (There is no such decrease for workers older than 25.)" (same source, 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: young-worker section. Verified: curl+grep
- "The young workers who are not hired may be remaining at their existing jobs, taking different jobs, or returning to school." (same source, 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: authors' caveat on the hiring result. Verified: curl+grep
- "Following adoption, junior employment declines sharply in adopting firms relative to non-adopters, while senior employment trends remain largely unchanged." (Seyed M. Hosseini and Guy Lichtinger, Harvard, "Generative AI as Seniority-Biased Technological Change", version 6 May 2026, abstract) https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf. Context: headline result. Verified: curl+pdftotext
- "The decline is driven primarily by slower hiring rather than increased separations." (same, abstract) Same URL. Verified: curl+pdftotext
- "One potential implication is that if firms overestimate the near-term automation capabilities of GenAI, the resulting decline in junior employment may ultimately prove temporary." (same, introduction) Same URL. Context: authors' forward-looking-hiring interpretation. Verified: curl+pdftotext

- "Notably, respondents were on average more worried about job loss for others than for themselves." (Anthropic Economic Index report "Cadences", 26 Jun 2026) https://www.anthropic.com/research/economic-index-june-2026-report. Context: survey of about 9,700 Claude users. Verified: curl+grep
- "The Economic Index Survey is not representative of the general population." (same) Same URL. Context: authors' own sample caveat. Verified: curl+grep
- "Across US workers, ten more percentage points of work affected is associated with 0.4 points more forecasted unemployment." (Sacher, Cotter, Jones, Korinek, McCrory, "The AI expectations gap", Anthropic, 16 Oct 2026) https://www.anthropic.com/research/ai-expectations-gap. Context: Morning Consult survey, n=10,980. Verified: curl+grep

Note: Jerome Powell's September 2025 FOMC remark on AI and young graduates appears in Hosseini and Lichtinger, but I did not fetch the Fed transcript, so it is not quoted here.

Sources worth reading in full (ranked)

1. Hosseini and Lichtinger, May 2026 version: the best-identified US study of the junior-hiring effect, with an honest discussion of rates, over-hiring and anticipation, plus a useful literature review. https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf
2. Brynjolfsson, Chandar, Chen, "Canaries", Aug 2026 revision and dashboard: the most-cited dataset; read the revision, not the 2025 headlines. https://digitaleconomy.stanford.edu/news/canariesaug26/
3. Massenkoff and McCrory (Anthropic), Mar 2026: clear null on unemployment, with power stated, plus a usage-based exposure measure. https://www.anthropic.com/research/labor-market-impacts
4. Yale Budget Lab, Oct 2025 report and 2026 updates: the main "no discernible disruption" source, updated regularly. https://budgetlab.yale.edu
5. Humlum and Vestergaard, NBER 33777 and the 2026 "Still Waters, Rapid Currents" version: the best administrative-data null from Europe. https://ideas.repec.org/p/crm/wpaper/26078.html
6. Demirci, Hannane, Zhu (Management Science 2025): cleanest evidence of direct substitution, in freelancing. https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb
7. Challenger monthly reports: read the reason-coding notes before citing any AI-cut number. https://www.challengergray.com/blog/october-layoffs-at-three-year-low-seasonal-hiring-plans-up-26/
8. Anthropic "The AI expectations gap", Oct 2026: representative data on how far public expectations exceed measured effects. https://www.anthropic.com/research/ai-expectations-gap
9. Bank Underground, "Canaries in the column?", Aug 2026: the only UK central-bank analysis found of AI exposure and hiring; read for the method and the pre-ChatGPT trend question. https://bankunderground.co.uk/2026/08/06/canaries-in-the-column-ai-exposure-and-the-uks-hiring-slowdown/
10. Lodefalk et al., "Same Storm, Different Boats", Örebro WP 2/2026: the best European within-employer study, with a rates-vs-AI separation for postings. https://www.oru.se/globalassets/oru-sv/institutioner/hh/workingpapers/workingpapers2026/wp-2-2026.pdf
11. Dallas Fed, Atkinson and Yamco, Jan 2026: sizes the junior effect against the aggregate (about 0.1 point of unemployment). https://www.dallasfed.org/research/economics/2026/0106

12. ECB Occasional Paper 395 and the Mar 2026 ECB blog: euro-area firm-survey evidence of hiring, not firing, among AI users. https://www.ecb.europa.eu/pub/pdf/scpops/ecb.op395.en.pdf

INCOMPLETE (narrower after the 2026-11-06 gap-fill pass). Now covered from search extracts: US business formation and nonemployer data, UK self-employment, the Medvi and Base44 "solo company" cases, UK and US care-sector vacancies and pay, and the Germany, Valencia, Spain-national and Tokyo four-day-week schemes. Still missing: (1) retirement and wellbeing studies; (2) verbatim quotes (no page could be fetched); (3) Spain national pilot results and Tokyo uptake (none found); (4) nonemployer data for 2024-2025, which would be the first AI-era read; (5) primary checks of the court docket and FDA letter for Medvi. Every finding is based on search-engine summaries of the cited URL, not on a page read in full; treat specific numbers as "to re-verify against the primary".

Track 8: What will people do?

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Method note. No primary document was opened in full in this session (see the INCOMPLETE note). URLs point to the primary source where search surfaced one. Numbers carry the reliability note "search summary only" in the facts table. No verbatim quotes are given, because none could be checked against a fetched page.

Summary

- **The long-run historical record says jobs were replaced and new work appeared, but this happened over decades and the people who lost jobs did worse.** Autor, Chin, Salomons and Seegmiller (QJE 2024) estimate that about 60% of US employment in 2018 was in job titles that did not exist in 1940. Feigenbaum and Gross (QJE 2024) find that AT&T's automation of telephone operators did not reduce employment for later cohorts of young women. But incumbent operators were, a decade later, more likely to be in lower-paying jobs or out of work. "New work appears" and "displaced workers are fine" are different claims. The evidence supports the first and contradicts the second.
- **New work is now made less by automation and more by augmentation, and its location has moved.** Autor et al. find that augmentation innovations create new work titles, while automation innovations erode employment without creating them. They also find that new work shifted from middle-paid production and clerical jobs (1940 to 1980) towards high-paid professional and low-paid service jobs. Whether AI is mainly automating or augmenting is the open question that decides which history repeats.
- **Forecasts of mass technological unemployment have a poor track record. Forecasts of shorter hours have a worse one.** Keynes's 15-hour week (1930) did not happen: men's annual market hours fell by more than a quarter between 1931 and 2011, far short of his roughly 780 hours a year (Crafts, CEPR 2021). Frey and Osborne's "47% at risk" (2013) was not followed by job loss. The highest-risk occupation they named, insurance underwriters, grew 16.4% from 2013 to 2021 (ITIF 2022, an advocacy think tank).

- **Cash-transfer trials consistently show small reductions in paid work, not collapse. They do not show people filling freed time with higher-value activity.** In the largest US RCT (OpenResearch, $1,000/month for 3 years, about 1,000 treated), recipients' hours fell modestly (reported as about 1.3 hours/week, with employment down about 2 percentage points). The biggest time-use gain was leisure, and the authors say other productive activities did not offset the loss. Finland (2017-18) found small employment effects and better self-reported wellbeing. Alaska's dividend found no employment effect. Kenya's long-term UBI found no fall in total labour, with a shift from wage work to self-employment.

- **Losing a job costs far more than lost income.** Displaced high-seniority US men had a 10-15% higher annual death hazard twenty years later (Sullivan and von Wachter, QJE 2009). German panel data show no adaptation to unemployment for men (Clark et al.), and life-satisfaction scarring lasts at least five years after reemployment (Eberl et al. 2022). Jahoda's "latent functions" of work hold up in a 2023 meta-analysis, except status/identity, which showed no significant difference against people out of the labour force. This is the strongest pessimistic case. What these studies measure is involuntary job loss, not voluntary freedom from work, and the trials above suggest the two are not the same.

- **A job guarantee in the original Marienthal village (Gramatneusiedl, Austria, from 2020) raised participants' wellbeing, time structure, status and social inclusion, with no physical-health effect** (Kasy and Lehner, Oxford/IZA 2023). This is the most direct test of Jahoda's thesis, and it supports her: what the long-term unemployed lacked was the structure of work, not only income.

- **"Removing rote work makes people more human" is not supported by the early AI evidence, and some of it points the other way.** A 2025 Scientific Reports paper (four experiments, N=3,562) found GenAI collaboration improved performance but lowered intrinsic motivation and raised boredom when people returned to solo work. An eight-month ethnography published in HBR (Feb 2026) found AI time savings turned into task expansion, blurred boundaries and longer hours. Danish administrative data (Humlum and Vestergaard, NBER 2025) found precise zero effects on earnings and hours, with some new tasks created. Freed time is so far being absorbed by more work, not by more human activity.

- **Four-day-week evidence is real but weaker than its reputation.** The UK 2022 pilot (61 self-selected firms, about 2,900 staff, 56 continuing afterwards) and Iceland's trials are often cited. The Iceland "overwhelming success" claim comes from an advocacy think tank's report. It describes reductions to 35-36 hours in public-sector workplaces covering about 2,500 workers (about 1% of the working population), not a four-day week.

- **Solo and micro-business: the US boom predates AI, the UK is going the other way, and the "one-person billion-dollar company" cases are thinner than the headlines.** US business applications hit a record of about 5.67 million in 2025, but the share likely to hire fell from 37.6% (2019) to 30.1%; nonemployers grew 4.7-4.9% a year in 2021-22, before ChatGPT. UK self-employment is about 4.4 million, down from about 5.0 million in 2019. Medvi had about $401m of 2025 revenue with two staff but outsourced its clinical work and faces an FDA warning letter and a class action; the "$1.8bn" is a projection.

- **Care is where demand grows, at the wage floor.** English care workers' median pay was £12.60 an hour, 39p above the minimum wage (Dec 2025); US home care aides earn a median $34,900 a year, with 765,800 openings a year, mostly replacing leavers. UK care vacancies (6.2%) fell because of migration, which is now drying up.

Key findings

Historical displacement: predicted vs actual

- **[EVIDENCE]** About 60% of US employment in 2018 was in occupational titles that did not exist in 1940. Method: a database of new job titles from Census Alphabetical Index updates 1940-2018, linked to Census microdata and patent-based measures of augmentation and automation. Some Autor briefings and slides use 63% for the same estimate. Autor, Chin, Salomons, Seegmiller, "New Frontiers: The Origins and Content of New Work, 1940-2018", QJE 139(3):1399-1465, Aug 2024. https://economics.mit.edu/sites/default/files/2026-04/New-Frontiers-QJE-v139%283%29-p1399-1465-Aug-2024.pdf (NBER WP 30389, 2022: https://www.nber.org/system/files/working_papers/w30389/w30389.pdf)
- **[EVIDENCE]** In the same paper, augmentation innovations (patents linked to an occupation's outputs) predict new work titles in that occupation. Automation innovations (patents linked to its tasks) reduce its employment and do not generate new work. Source as above (via search summaries of the paper and the Blueprint Labs research page, https://blueprintlabs.mit.edu/research/new-frontiers-the-origins-and-content-of-new-work-1940-2018/).
- **[EVIDENCE]** The locus of new work moved. In 1940-1980 it was mainly in middle-paid production and clerical occupations. After 1980 it was in high-paid professional and low-paid service occupations. Fast Company's coverage, clipped by MIT News, reports that automation eroded about twice as many jobs in 1980-2018 as in 1940-1980. Autor also attributes some new work to demand shifts (e.g. health care for an ageing population), not technology. https://news.mit.edu/news-clip/fast-company-202
- **[EVIDENCE]** Telephone operators. Between 1920 and 1940 AT&T automated over half of the US telephone network. Identification uses city-level variation in cutover timing. Automation eliminated most operator jobs but did not reduce later cohorts' overall employment: young women went into middle-skill clerical and lower-skill service jobs instead. Incumbent operators (followed with genealogy-based census linking) were the most affected: a decade later they were more likely to be in lower-paying occupations or not working. Feigenbaum and Gross, "Answering the Call of Automation", QJE 139(3):1879-1939, Aug 2024. https://econpapers.repec.org/RePEc:oup:qjecon:v:139:y:2024:i:3:p:1879-1939. (NBER WP 28061: https://www.nber.org/system/files/working_papers/w28061/w28061.pdf)
- **[EVIDENCE]** Agriculture's share of the US workforce fell from about 41% in 1900 to 1.9% in 2000. This is a USDA ERS compilation: Historical Statistics of the US for 1900-1970 and the Census of Population for 2000. The farm share of the population was 39% in 1900 and about 1% around 2000. These are different measures and should not be mixed. ERS, Agricultural Outlook, June 2000 (archived) https://webharvest.gov/peth04/20041105150816/http://www.ers.usda.gov/publications/agoutlook/jun2000/ao272h.pdf; ERS Amber Waves, June 2005 (archived) https://webarchive.library.unt.edu/eot2008/20090115230224/http://www.ers.usda.gov/AmberWaves/June05/DataFeature. Note: the 41% figure was seen via a secondary repost of the ERS table. The World Bank series puts 2000 at 2.4%, so the end-point depends on the series.
- **[EVIDENCE]** Keynes's forecast vs reality. Keynes (1930) suggested a 15-hour week (three-hour shifts) for his grandchildren's generation. Crafts finds that average annual hours of market work for men fell by more than a quarter between 1931 and 2011, but stayed well above the roughly 780 hours a year Keynes implied. Leisure plus home production rose about 20%, less than half the increase Keynes implied. Crafts, "Economic Possibilities for our Grandchildren revisited" (CEPR DP16140 / Warwick TWERP 1360, 2021). https://repec.cepr.org/repec/cpr/ceprdp/DP16140.pdf https://warwick.ac.uk/fac/soc/economics/research/workingpapers/2021/twerp_1360_-_crafts.pdf

- **[EVIDENCE, secondary]** NPR Planet Money (2015) put average US weekly hours at about 38 in 1950 and about 34 six decades later. Richard Freeman's explanation in that piece is competition and status: people took productivity gains as income, not time. https://www.npr.org/sections/money/2015/08/13/432122637/keynes-predicted-we-would-be-working-15-hour-weeks-why-was-he-wrong
- **[OPINION]** Leontief's horse analogy. Leontief argued (Scientific American 1982, "The Distribution of Work and Income"; Population and Development Review 1983) that human labour might follow the horse: first diminished, then eliminated. He recommended government income policies in response. The US horse population grew about sixfold between 1840 and 1900 to over 21 million, then fell to about 3 million by 1960 after cars and tractors arrived. Brynjolfsson and McAfee's counter-argument is that, unlike horses, humans can own capital and vote. Commentators note that the median-wage and unemployment effects Leontief feared did not arrive on his timescale. https://equitablegrowth.org/must-read-erik-brynjolfsson-and-andrew-mcafee-labor-in-the-second-machine-age/ ; Leontief 1983 record: https://bvbvf1.bib-bvb.de/vufind/PrimoRecord/cdi_proquest_miscellaneous_61582758 ; history and critique: https://www.greaterwrong.com/posts/nZbxQgLWhKCASApEe/learning-more-from-horse-employment-history. The publication dates are inconsistent across sources (1982 vs 1983).
- **[FORECAST, then EVIDENCE on outcome]** Frey and Osborne (2013) put 47% of US employment at high risk of computerisation "over some unspecified number of years, perhaps a decade or two". ITIF (a pro-innovation think tank, so OPINION in framing) reported in 2022 that the US added 16 million jobs between 2013 and 2022, with unemployment at 3.7%. The occupation with the highest risk score, insurance underwriters, grew 16.4% from 2013 to end-2021. The lowest-risk one, recreational therapists, fell 8.9%. Higher-risk occupations were only slightly more likely to lose jobs. ITIF, 30 Sep 2022. https://itif.org/publications/2022/09/30/oops-the-predicted-47-percent-of-job-loss-from-ai-didnt-happen/ . In 2023 Frey and Osborne said the "bottlenecks" they identified remain relevant even with generative AI (Fortune, 18 Sep 2023, https://fortune.com/europe/2023/09/18/generative-ai-jobs-automation-frey-osborne-oxford/). Caveat: their window runs to about 2033, so the forecast has not formally expired. Total job growth also does not isolate automation's effect.
- **[EVIDENCE]** Industrial robots, US 1990-2007. One more robot per thousand workers in a commuting zone reduced the employment-to-population ratio by 0.2 percentage points and wages by 0.42% (published version). The 2017 working paper gave ranges of 0.18-0.34 pp and 0.25-0.5%. An authors' column gave local effects of 0.37 pp and 0.73%, about 6.2 jobs lost per robot. Acemoglu and Restrepo, "Robots and Jobs", JPE 128(6):2188-2244, 2020. https://econpapers.repec.org/RePEc:ucp:jpolec:doi:10.1086/705716 ; https://www.nber.org/papers/23285 ; https://www.wita.org/nextgentrade/robots-and-jobs-evidence-from-the-us/ . This is the strongest modern evidence that a specific automation technology reduced local employment, without offsetting new work at the local level.

Cash transfers and UBI trials

- **[EVIDENCE]** OpenResearch Unconditional Cash Study (US, Illinois and Texas). About 1,000 low-income people were randomised to $1,000/month for three years, against about 2,000 controls receiving $50/month. Reported labour effects: hours down about 1.3/week and employment down about 2 percentage points relative to control (Fortune, Boston Indicators summaries). A conference version of the paper reports non-transfer income down about $2,000/year and labour market participation down 3.9 pp. The two sets of figures probably refer to different periods or specifications, and I could not reconcile them without the paper. The largest time-use increase was in leisure. The authors describe the labour supply effect as

moderate and say it was not offset by other productive activities. They find no significant effect on job quality or human-capital investment, with possible education gains among younger participants. Vivalt, Rhodes, Bartik, Broockman, Krause, Miller, NBER WP 32719 (July 2024, revised Sept 2025). https://www.nber.org/system/files/working_papers/w32719/w32719.pdf ; conference version https://conference.nber.org/conf_papers/f208057/f208057.pdf ; Fortune commentary (4 Sep 2024) https://www.fortune.com/2024/09/04/sam-altman-backed-study-guaranteed-income-widely-misinterpreted-ubi-personal-finance

- **[EVIDENCE]** Stockton SEED (California). 125 residents received $500/month for 24 months. Full-time employment among recipients rose from 28% (Feb 2019) to 40% a year later. The control group went from 32% to 37%. Recipients moved from part-time to full-time work at more than twice the control rate. Caveats: very small sample, and a first-year window that partly overlaps COVID. Evaluators: Stacia Martin-West and Amy Castro Baker; first-year results March 2021. https://penntoday.upenn.edu/news/guaranteed-income-increases-employment-improves-financial-and-physical-health ; https://catalog.results4america.org/case-studies/guaranteed-income-stockton
- **[EVIDENCE]** Finland basic income experiment 2017-2018. About 2,000 unemployed people received €560/month unconditionally, against a register-based control group. Employment effects were small: recipients' employment improved slightly more than controls' in Nov 2017 to Oct 2018. Interpretation is complicated by a new activation model introduced in 2018. Recipients reported higher life satisfaction and less mental strain (phone survey near the end). Final report May 2020 (Kangas, Jauhiainen, Simanainen, Ylikännö, eds). Ministry of Social Affairs and Health https://stm.fi/en/-/perustulokokeilun-tulokset-tyollisyysvaikutukset-vahaisia-toimeentulo-ja-psyykkinen-terveys-koettiin-paremmaksi ; VATT https://vatt.fi/en/-/perustulon-tyollisyysvaikutukset-jaivat-vahaisiksi ; results summary PDF https://toolbox.finland.fi/wp-content/uploads/sites/2/2018/04/basic_income_experiment_2017-2018_results_2020.pdf
- **[EVIDENCE]** Kenya long-term UBI (GiveDirectly; Banerjee, Faye, Krueger, Niehaus, Suri). Village-level RCT comparing a 12-year monthly UBI, a 2-year version and an equivalent lump sum. Short-term results (2023): total labour supply roughly unchanged, with a shift from wage work to self-employment; more non-farm enterprises and higher revenues. Lump sums had effects at least as large as the long stream, and the short stream had noticeably smaller effects. https://www.findevgateway.org/paper/2023/09/universal-basic-income-short-term-results-from-long-term-experiment-in-kenya ; https://conference.nber.org/conf_papers/f192616.pdf ; VoxDev summary https://voxdev.org/topic/social-protection/universal-basic-income-kenya . Not found: later (2025-26) long-run results.
- **[EVIDENCE]** Alaska Permanent Fund Dividend, a universal and permanent annual payment. Using CPS data and synthetic control, the dividend had no effect on employment and increased part-time work by 1.8 percentage points (17%). The authors attribute the absence of an employment fall partly to general-equilibrium demand effects: non-tradable sectors respond more positively. Jones and Marinescu, AEJ: Economic Policy 14(2):315-340, 2022. https://www.aeaweb.org/doi/10.1257/pol.20190299 ; https://home.uchicago.edu/j1s/Jones_Alaska_2022.pdf . Caveat: a search summary surfaced text from a version of the paper saying the preferred specification detects no part-time effect. Check the published text before citing the 1.8 pp.

Four-day week and hours reduction

- **[EVIDENCE, weak: self-selected sample, advocacy-run]** UK four-day-week pilot, June to December 2022. 61 companies and about 2,900 employees on a 100-80-100 model (100% pay, 80% time, 100% output). 56 of 61 said they would continue after the pilot and 18 said the change was permanent. The report was led by Juliet Schor and colleagues and published

by Autonomy with 4 Day Week Global (both advocates). Firms opted in, so the results describe firms that expected it to work, not the average firm. Revenue and sick-day figures circulate only via secondary aggregators (e.g. revenue "broadly stable", 71% reporting less burnout) and I could not check them against the report. WEF summary https://www.weforum.org/stories/2023/03/four-day-work-week-uk-trial/ ; Reuters via Malay Mail, 21 Feb 2023 https://www.malaymail.com/news/money/2023/02/21/four-day-week-here-to-stay-say-uk-firms-in-worlds-largest-trial/56070

- **[EVIDENCE, weak: advocacy report, no control group identified]** Iceland. There were two trials: Reykjavík City (2014-2019) and the national government (2017-2021). Together they covered more than 2,500 workers, about 1% of the working population, in offices, hospitals, schools and social services. Hours were cut to 35-36 per week without loss of pay. This was an hours reduction in the public sector, not a four-day week. The report (Autonomy with the Icelandic think tank Alda, June 2021) says productivity and service stayed the same or improved in most workplaces and wellbeing rose. Autonomy's research director called it an "overwhelming success". The report also says that by June 2021, 86% of Iceland's working population were on contracts that had shortened hours or gave the right to shorten them. Those were union-negotiated reductions of varying size, not four-day weeks. Report https://issuu.com/4dayweekglobal/docs/iceland_4dw ; syndicated news, July 2021 https://www.fox5ny.com/news/4-day-work-week-trials-in-iceland-deemed-overwhelming-success-by-researchers
- **[EVIDENCE, weak: self-selected, advocacy-organised, academic evaluation]** Germany, February to July 2024. 45 firms (some sources say 41 or 43 after drop-outs) and about 900 employees, organised by Intraprenör with 4 Day Week Global; the scientific evaluation was by the University of Münster (Julia Backmann). About 73% of firms said they would continue with a four-day or reduced week; about 20% went back to five days and 40 hours; 83% of employees wanted to continue. Two large organisations dropped out early for economic reasons. Backmann herself cautioned against concluding that every German firm should move to a four-day week. A secondary two-year follow-up reports 70% still on some form of reduced hours. https://www.thelocal.de/20241106/what-are-the-results-of-germanys-four-day-work-week-trial ; https://4dayweek.com/germany-2024-pilot-results (advocacy) ; https://remoteworkeurope.eu/news/2026/germany-four-day-week-two-year-followup (secondary) (all search extract)
- **[EVIDENCE, weak: city-commissioned evaluation, four weeks]** Valencia, 10 April to 7 May 2023. The city made four consecutive Mondays local holidays, so this tested extra free days for a whole city rather than a redesign of work. An expert commission (commissioned by the city council that ran the trial) reported better self-rated health, lower stress and better mood, a 9.5% fall in traffic and lower NO2. It also found more smoking and drinking among people who already smoked or drank, and that women spent more of the extra time on care. https://www.voanews.com/a/four-day-work-week-boosts-spanish-workers-health-pilot-program-shows-/7315006.html ; https://thenextweb.com/news/valencias-four-day-workweek-trial-new-world-of-work ; https://www.themayor.eu/en/a/view/valencia-s-government-says-4-day-week-trial-was-good-for-the-people-and-the-city-12172 (search extract)
- **[EVIDENCE: design only, no results found]** Spain's national pilot offers small and medium industrial firms (under 250 staff) up to €150,000 each to cut hours for at least two years, with an audit afterwards. One weak source reports 41 applicant SMEs covering 503 workers. I found no published results; a 2024 article said results were pending. https://www.thelocal.es/20221213/spain-to-pay-e150k-to-smes-that-roll-out-four-day-work-week ; https://www.personneltoday.com/?p=304586 (search extract)
- **[EVIDENCE: policy, no uptake data]** Tokyo Metropolitan Government, from April 2025: its roughly 160,000 staff can opt into a four-day week by reworking the flextime system, i.e. the

same hours compressed into four days, not fewer hours. It is framed as a fertility and childcare measure. I found no uptake figures. A private-sector comparison in the same coverage: Panasonic offered a four-day option to 63,000 staff and only about 150 took it. https://www.nbcnews.com/business/business-news/tokyo-government-introduce-four-day-workweek-employees-rcna183442 ; https://www.scmp.com/news/asia/article/3300614/it-possible-people-asia-adopt-4-day-working-week (search extract)

- **Reading.** The post-2022 trials repeat the UK pattern: self-selected firms, advocacy organisers, most continue, no control group. None is an independent randomised test. Two details matter for "what will people do": Valencia found freed time went partly to care (mainly by women) and partly to more drinking and smoking, and Japan shows that where a four-day option means the same hours compressed, almost nobody takes it.

Job guarantee and the Marienthal line

- **[EVIDENCE]** Marienthal (Jahoda, Lazarsfeld, Zeisel, 1932/33). This was a study of an Austrian village after its textile factory closed. Jahoda later formalised five "latent functions" of employment: time structure, social contact, collective purpose, status/identity and activity. In the study, the unemployed experienced their sudden free time as a burden and their days lost structure (paraphrased: Jahoda's phrase was "tragic gift"). A 2023 meta-analysis (Paul, Scholl, Moser, Zechmann, Batinic, Frontiers in Psychology) found employed people scored higher than unemployed people on all five latent functions and on the manifest function (income). But it found no significant status/identity difference between employed people and people out of the labour force. https://pmc.ncbi.nlm.nih.gov/articles/PMC10017486 ; https://www.frontiersin.org/articles/10.3389/fpsyg.2023.1017358/full
- **[EVIDENCE]** Gramatneusiedl job guarantee (the Marienthal municipality, from October 2020; run by the Lower Austria public employment service). It offered voluntary, individually tailored jobs to all long-term unemployed residents. The evaluation used three methods: pairwise randomisation of entry timing, a pre-registered synthetic control and matched comparison individuals. It found strong positive effects on economic wellbeing (employment, income, security) and on non-economic wellbeing (social status, time structure, social interactions, collective purpose). It found no effect on physical health or on risk and time preferences. At municipality level, long-term unemployment fell sharply and total unemployment fell by somewhat less. Wellbeing also rose in anticipation among future participants. Kasy and Lehner, "Employing the unemployed of Marienthal", Oxford INET WP 2022-29 / IZA DP 16088, April 2023. https://docs.iza.org/dp16088.pdf ; https://www.inet.ox.ac.uk/publications/no-2022-29-employing-the-unemployed-of-marienthal-evaluation-of-a-guaranteed-job-program . Caveat: this is one small village, so the participant count is small (I could not confirm the exact n). "Eliminated unemployment" overstates the result.

What people do with freed time

- **[EVIDENCE]** In the US during the Great Recession (ATUS), about 30% of lost market work hours went to more sleep and television, and only about 1% to job search. Aguiar, Hurst, Karabarbounis, summarised in the NBER Digest, Nov 2011. https://www.nber.org/digest/nov11/time-use-during-recessions
- **[EVIDENCE]** Unemployed people spend more time in leisure activities than the employed but enjoy them less. Leisure falls on reemployment, though less than cross-sectional data imply. Krueger and Mueller (AER P&P 2012; New Jersey survey plus ATUS 2009-10). https://zendy.io/title/10.1257/aer.102.3.594 ; working paper https://swh.princeton.edu/~ceps/workingpapers/175krueger.pdf

- **[EVIDENCE]** OpenResearch (above): the largest change in time use was more leisure. It was not offset by other productive activities, and education gains were limited to younger participants. Kenya UBI (above): time shifted into self-employment, not leisure. These two results conflict. The likely difference is context: a poor rural economy with high returns to small enterprise, against a rich-country labour market.
- Not covered: retirement and wellbeing studies, and ATUS 2024-25 updates.

Meaning, status and identity without work: the pessimistic case

- **[EVIDENCE]** Job displacement and mortality. Pennsylvania administrative earnings records for 1980-2006, matched to Social Security death records, cover high-seniority men displaced in mass layoffs. Mortality hazards were sharply higher immediately after displacement: 50-100% higher in the first year in a working-paper draft (not confirmed in the published text). Twenty years later the increase was still 10-15% a year. If sustained, that implies a 1.0-1.5 year loss of life expectancy for a worker displaced at 40. The authors rule out selective displacement of less healthy workers. Sullivan and von Wachter, QJE 124(3):1265-1306, 2009. https://econpapers.repec.org/RePEc:oup:qjecon:v:124:y:2009:i:3:p:1265-1306. ; draft https://www.movingforwardnetwork.com/wp-content/uploads/2023/04/Sullivan_Job-Displacement-and-Mortality_2009.pdf
- **[EVIDENCE]** No adaptation to unemployment. Twenty waves of the German SOEP show that people return to baseline life satisfaction after divorce, widowhood, a first child and layoff, but men show no adaptation to unemployment. Clark, Diener, Georgellis, Lucas (Economic Journal 2008). https://www.csus.edu/indiv/m/merlinos/pdf/clark2008.pdf . Eberl, Collischon and Wolbring (Social Forces 2022, SOEP) find life-satisfaction scarring for at least five years after reemployment, for men and women, driven by repeated spells. https://cris.fau.de/publications/272551801
- **[EVIDENCE + OPINION, contested]** Deaths of despair. Case and Deaton tie rising midlife mortality among US whites without a degree (from drugs, suicide and alcohol) to the long decline in good non-college jobs and to the loss of meaning and status that came with it. Ruhm argues that the opioid supply environment explains far more of the geographic variation than economic conditions. Case and Deaton reply that the economic link shows up over a longer window (from the 1970s) and that "despair" was a label, not a causal claim. Case and Deaton, NBER WP 28358 (2021) https://nber.org/papers/w28358 ; reply to Ruhm, Jan 2018 https://www.princeton.edu/~accase/downloads/Case_and_Deaton_Comment_on_CJRuhm_Jan_2018.pdf ; Ruhm critique summary, Axios Jan 2018 https://www.axios.com/2018/01/09/study-dont-blame-opiod-crisis-on-the-econ-1515427073
- **Reading.** All of this is about involuntary job loss in societies where work confers status and income is tied to it. It is the strongest available evidence that losing work hurts more than the lost pay. It does not test a world where everyone has secure income and non-work is not stigmatised. The UBI trials show wellbeing gains when income comes without job loss. Gramatneusiedl shows wellbeing gains when structure and purpose are restored through a job. The two findings are compatible: income alone fixes the money; structure, purpose and status need something job-like.

Does removing rote work make people "more human"? Early AI evidence

- **[EVIDENCE]** Four online experiments (total N=3,562) found that GenAI collaboration improved task performance. But when participants moved from AI-assisted to solo work, intrinsic motivation fell and boredom rose. Secondary coverage, including HBR (May 2025),

reports averages of about -11% for motivation and +20% for boredom. I could not check those averages against the paper. The proposed mechanism is that AI takes over the most cognitively engaging parts of the task. "Human-generative AI collaboration enhances task performance but undermines human's intrinsic motivation", Scientific Reports, 1 Apr 2025, doi 10.1038/s41598-025-98385-2. Library record https://library.kab.ac.ug/Record/doaj-art-4780b0490cf249a0a72eb4da2a2149dd/Description ; Fortune 3 Jun 2025 https://fortune.com/2025/06/03/ai-boost-worker-productivity-hurts-motivation-boredom . Caveat: online lab tasks, not field employment.

- **[EVIDENCE, single-site qualitative]** Work intensification. An eight-month ethnography at a US tech firm (reported as about 200 employees) with broad GenAI access found that time saved went into faster pace, wider job scope and longer hours, largely voluntarily. It identified three mechanisms: task expansion, blurred boundaries (prompting during lunch and meetings) and constant multitasking. Faster output raised expectations of speed. Ranganathan and Ye, Harvard Business Review, Feb 2026. Coverage: UC Berkeley Haas newsroom, 18 Feb 2026 https://newsroom.haas.berkeley.edu/ai-promised-to-free-up-workers-time-uc-berkeley-haas-researchers-found-the-opposite/
- **[EVIDENCE]** Danish administrative data covering 25,000 workers in 11 AI-exposed occupations across 7,000 workplaces, with adoption surveys from 2023-24 linked to employer-employee registers. AI chatbots had precisely-estimated zero effects on earnings and recorded hours in every occupation. Only 3-7% of reported productivity gains passed through to earnings. Chatbots also created new tasks (e.g. checking AI output), which absorbed some of the time saved. Humlum and Vestergaard, "Large Language Models, Small Labor Market Effects", NBER WP 33777, 2025. https://econpapers.repec.org/RePEc:nbr:nberwo:33777 . A 2026 follow-up ("Still Waters, Rapid Currents") argues that the task content of work changes well before earnings or hours do.
- **Reading.** Three independent designs (lab experiment, ethnography, national register data) say the same thing: the time AI frees is reabsorbed into work, and the remaining work can feel less engaging. None of them shows workers redeploying freed time into the relational or creative activity that "more human" implies. This is early evidence (2024-26 tools), but it is the only evidence there is.

Solo operators, micro-business, care and community work

Gap-fill pass (2026-11-06). All items rest on search-engine extracts; no page was fetched.

US business formation and nonemployers

- **[EVIDENCE, official series; annual total computed by secondary sources]** US business applications reached about 5,671,836 in 2025, an annual record, above the previous record of 5,469,302 in 2023 and about 49% above 2019 (3,498,990). But "high-propensity" applications (those with signs of planned hiring and wages) were about 1,708,842 in 2025, below their 2023 peak of 1,848,540, and their share of all applications fell from 37.6% in 2019 to 30.1% in 2025 (a consultant's calculation from Census data). Census Business Formation Statistics, December 2025 release (14 Jan 2026) https://www.census.gov/econ/bfs/pdf/historic/bfs_2025m12.pdf ; FRED series https://fred.stlouisfed.org/series/BAHBATOTALSAUS ; annual totals via https://founderreports.com/how-many-businesses-are-started-each-year/ (search extract)
- **[EVIDENCE, official]** Nonemployer businesses (no paid staff, at least $1,000 receipts) rose from 28,477,518 in 2021 to 29,811,495 in 2022. Growth of 4.9% (2021) and 4.7% (2022) was the fastest in nearly two decades, against 1.9% and 1.8% for employer businesses. The 2023 Nonemployer Statistics by Demographics counts 30.4 million nonemployers with $1.8tn of

receipts, about $59,000 each on average (my division). Census, "Number of U.S. Nonemployers Grew Faster Than Employer Businesses Nearly Every Year From 2012 to 2023", July 2025 https://www.census.gov/library/stories/2025/07/nonemployer-business-growth.html ; NES-D press release, 20 Nov 2025 https://www.census.gov/newsroom/press-releases/2025/business-owner-characteristics.html ; 2023 NES release, 15 May 2025 https://www.census.gov/newsroom/press-releases/2025/2023-nonemployer-statistics.html (search extract)

- How to read it: the US surge in one-person and non-hiring businesses is real, but it began in mid-2020, more than two years before ChatGPT, and the latest nonemployer reference year (2023) is too early to show an AI effect. It cannot yet be attributed to AI. Most nonemployers are side income: average receipts are modest.

UK self-employment: the opposite trend

- **[EVIDENCE, official survey via secondary]** UK self-employment was about 4.38 million in Q4 2025 (Statista from the ONS Labour Force Survey), against roughly 5.0 million at the end of 2019. The self-employment rate was about 13.5% in Q1 2025 against a 2019 peak of 15.3%; full-time self-employed were about 2.88 million in Q4 2025 against 3.56 million in Q4 2019. Another ONS jobs series gives 4.69 million (Q4 2019) to 4.04 million (Q4 2025); the two series could not be reconciled. ONS warns that LFS estimates remain volatile. https://statista.com/statistics/318234/united-kingdom-self-employed ; https://www.enterpriseresearch.ac.uk/wp-content/uploads/2020/04/ERC-Insight-Covid-19-and-self-employment-in-the-UK.pdf (search extract)
- Counter-reading: LFS response problems, tax changes for contractors (IR35) and post-Brexit outflows of EU workers are standard explanations for the UK fall. Either way, there is no sign of an AI-driven solo boom in UK headcounts by end-2025.

"One-person billion-dollar company": what the cases show

- **[EVIDENCE, journalistic, contested]** Medvi (US GLP-1 weight-loss telehealth, founder Matthew Gallagher) is the case most cited as Sam Altman's predicted one-person billion-dollar company. The New York Times was reportedly given access to its financials: 2025 revenue of about $401m and net profit of about $65m (16.2% margin). The "$1.8bn" is Gallagher's own 2026 revenue projection, not a valuation; the firm has raised no outside capital. It is not one person: Gallagher works with his brother and contractors, and prescribing and fulfilment run through partner telehealth and pharmacy networks. Dealroom, 5 Apr 2026 https://dealroom.co/news/127216-from-20k-to-1-8b-in-revenue-how-one-founder-built-a-telehealth-giant-wit/ (search extract)
- **[EVIDENCE, secondary, allegations]** The FDA reportedly issued Medvi a warning letter on 20 February 2026 over misbranding (Gallagher says it went to an affiliate and was resolved); a class action filed in California in March 2026 reportedly alleges benefit from affiliate spam; Futurism reported fabricated doctor personas and altered patient photos in marketing. Critics note the NYT profile omitted these. https://www.implicator.ai/medvi-got-an-fda-warning-before-the-times-called-it-a-1-8-billion-ai-story/ ; https://www.emarketer.com/content/fake-doctor-ads-fuel-scrutiny-of-glp-1-marketer (search extract; check the FDA warning-letter database and court docket)
- **[EVIDENCE, press reports]** Base44: Wix bought the AI app builder in June 2025 for about $80m in cash plus earn-outs through 2029, about six months after Maor Shlomo started it as a bootstrapped side project. It had six to eight employees at sale, who share a retention pool of about $25m. Reported user and profit figures come from founder posts, not audited

accounts. Globes https://en.globes.co.il/en/article-wix-acquires-israeli-vibe-coding-co-base44-1001513267 (search extract)

- How solid: both cases are real but neither is a one-person billion-dollar company in the strict sense. Medvi has large revenue but rests on outsourced clinical and fulfilment labour and faces regulatory and legal challenges; Base44 was a solo founder but an $80m exit with a small team. They show small headcount, not small labour: the work is done by contractors, platforms and partner firms. They are anecdotes, not a trend; the statistics that could show a trend (nonemployer receipts for 2024-2025) are not yet published.

Care work: growing demand, wage-floor pay

- **[EVIDENCE, sector data body]** England's adult social care vacancy rate fell to 6.2% in 2025/26, the lowest since 2015/16, about 96,000 vacancies on an average day (down 10.5% on the year), from 7.0% in 2024/25. It is still about three times the rate in the wider economy. The fall came from overseas recruitment, which is now drying up: about 30,000 international recruits arrived directly in 2025/26, the lowest in four years, after the care-worker visa route closed. Skills for Care projects about 410,000 more posts needed by 2040 (one trade piece says 470,000). Skills for Care, as reported by LocalGov, 24 June 2026 https://www.localgov.co.uk/Social-care-vacancy-rate-falls-to-decade-low-but-challenges-ahead/64563 ; https://www.healthcare-management.uk/breaking-news-social-care-vacancy-rate-falls-decade-low-challenges-ahead ; earlier comment that the fall was from overseas rather than domestic recruitment (King's Fund) https://www.localgov.co.uk/Domestic-recruitment-challenges-persist-in-adult-social-care/63272 (search extract)
- **[EVIDENCE, sector data body]** Pay: the median independent-sector care worker earned £12.60 an hour in December 2025, 39p above the 2025-26 National Living Wage of £12.21; 26% were paid within 10p of the NLW (22% in March 2025). A newly appointed NHS Band 3 healthcare assistant earned about 15p more than the median care worker (Homecare Association). The NLW rose to £12.71 in April 2026, which would leave the December 2025 median about 11p above the floor. https://www.healthcare-management.uk/care-workers-pay-minimum-wage-reveals ; https://www.homecareassociation.org.uk/resource/homecare-association-responds-to-skills-for-care-s-pay-in-the-adult-social-care-sector-in-england-report.html (search extract)
- **[EVIDENCE, official]** US home health and personal care aides: median pay $34,900 a year ($16.78 an hour) in May 2024; about 4.3 million jobs in 2024; projected +17% (739,800 jobs) from 2024 to 2034, much faster than average; about 765,800 openings a year. Net growth is about 74,000 a year, so roughly nine in ten openings replace people who leave (my arithmetic), a sign of high turnover. BLS Occupational Outlook Handbook https://www.bls.gov/ooh/healthcare/home-health-aides-and-personal-care-aides.htm (search extract)
- How to read it: care is the clearest case of growing, human-centred work, and it is the largest source of new US jobs in this occupation group. But it pays at or just above the legal floor in both countries, the UK has filled it through migration rather than domestic recruits, and turnover is high. "People will move into care" is plausible on demand; it is not, on current pay, a route to the security displaced professionals lose. Mainstream counter: care pay is largely set by public budgets (the state is the main buyer), so it is a policy choice that could change; England is setting up sector-wide pay bargaining for social care. https://assets.publishing.service.gov.uk/media/6970f06af88ad0be09b9efa8/social-care-negotiating-bodies-factsheet.pdf (search extract)

Earlier item

- The one relevant data point gathered in the first pass: in Kenya, unconditional cash moved people from wage work into self-employment and non-farm enterprise (Banerjee et al. 2023, above). That is evidence that income security can fund micro-enterprise in a low-income economy. It says nothing about AI-enabled solo businesses in rich economies.

Viral statistics traced

- **"60% of today's jobs didn't exist in 1940."** The origin is solid: Autor et al., QJE 2024, using Census title data. Versions of 60% and 63% both circulate. The common misreading is that technology reliably creates replacement jobs for the people it displaces. The same paper finds that automation innovations reduce employment without creating new work, and that new work after 1980 clustered at the top and bottom of the pay distribution.
- **"47% of jobs at risk" (Frey and Osborne 2013).** This was a model-based occupational probability exercise with a horizon of "perhaps a decade or two", not a forecast of job losses. Task-based re-estimates in the OECD tradition give much lower figures (I could not retrieve the primary). Ten years on, high-risk occupations have not systematically shrunk (ITIF 2022, an advocacy source using BLS occupational data).
- **"Iceland's four-day week was an overwhelming success."** This traces to one report by Autonomy (UK, pro-shorter-hours) and Alda (Iceland), June 2021. The trials cut public-sector hours to 35-36/week for about 2,500 workers (about 1% of the workforce). They did not move people to four days, and I found no independent evaluation with a control group. The widely repeated "86% now have shorter hours" figure means "on contracts that shortened hours or gave the right to", through union bargaining, which often meant modest cuts.
- **"UBI study shows people stopped working" / "UBI is a flop" (OpenResearch 2024).** The measured effect was a fall of about 1.3 hours/week and about 2 pp in employment, on a $12,000/year transfer. Headlines that called it a flop (e.g. thedeepdive.ca) and headlines that called it a vindication both overreach. The fair reading: a modest labour supply fall, mostly into leisure, with few measured gains in health, human capital or job quality.
- **"Stockton guaranteed income increased full-time employment."** This rests on 125 recipients over one year, with a 12 pp rise against a 5 pp rise in the control group. The direction is suggestive, but the sample is too small to carry the weight it is given.
- **"Four-day week: 92% of firms kept it" (UK 2022).** This is 56 of 61 continuing in some form, with 18 making it permanent, all from a self-selected pool of firms run through an advocacy organisation.
- **"The one-person billion-dollar company is here" (Medvi, 2026).** Traces to a New York Times profile (April 2026) built around Sam Altman's 2024 prediction. The verified figure is about $401m of 2025 revenue; $1.8bn is the founder's 2026 projection. The firm has two staff plus contractors and partner networks, an FDA warning letter (Feb 2026, disputed) and a class action. Not a valuation, not one person, not settled.
- **"73% of German firms won't go back to five days" (2024 pilot).** 45 self-selected firms, about 900 staff, organised with an advocacy group; about 20% did go back; the lead academic cautioned against generalising.

Strongest counter-evidence

Against the thesis's humane framing ("freedom to be more human"):

1. Losing work (involuntarily) damages health, life satisfaction and longevity beyond the lost income: Sullivan and von Wachter 2009; Clark et al. 2008; Eberl et al. 2022. Jahoda's latent functions are empirically supported (Paul et al. 2023 meta-analysis), and restoring a job restored them (Kasy and Lehner 2023).
2. Freed time in rich economies goes mostly to leisure, sleep and TV, not to higher-value human activity: Aguiar, Hurst and Karabarbounis (ATUS); OpenResearch 2024. People who are unemployed enjoy their leisure less (Krueger and Mueller 2012).
3. Early AI evidence shows time savings being turned into more work and lower intrinsic motivation, not "more human" work: Ranganathan and Ye 2026; Scientific Reports 2025; Humlum and Vestergaard 2025.
4. Productivity gains have historically been taken as income and consumption, not time (Keynes vs Crafts). There is no historical precedent for a technology delivering mass voluntary leisure.

Against the mainstream pessimistic reading of this track:

1. Unconditional income does not cause people to stop working en masse: Alaska (no employment effect), Finland (small effect), Kenya (no fall in total labour), OpenResearch (modest fall).
2. Over 80 years, new work has repeatedly appeared: Autor et al. 2024. Telephone operator automation did not reduce employment for later cohorts: Feigenbaum and Gross 2024.
3. Forecasts of mass technological unemployment (Leontief, Frey and Osborne) have so far been wrong in timing and scale.
4. The Jahoda meta-analysis found no status/identity deficit for people out of the labour force (as opposed to unemployed). This hints that the stigma of unemployment, rather than the absence of work as such, may drive part of the harm.

Against the optimistic history:

- The people displaced bore the costs (telephone operators; Acemoglu and Restrepo's local robot effects). New work arrived for the next cohort, not for the incumbents. If AI displacement is faster than 1920-1940 switchboard automation, there is less cohort turnover to absorb it.

Open questions

- Solo operators (partly answered in the gap-fill pass): still need 2024-2025 nonemployer counts and receipts when published, BTOS AI use split by firm size, and the primary sources for Medvi (FDA warning letter 721455, California docket, NYT profile). Original note: pull US Census Nonemployer Statistics (latest release, likely reference year 2022 or 2023), Census Business Formation Statistics (2024-26 business applications, especially high-propensity applications), ONS Labour Force Survey self-employment (UK self-employment has not returned to its pre-pandemic level, but the figure is not verified here), and the BTOS AI-use supplement split by firm size. Is there any measurable post-2023 rise in nonemployer firms or revenue per nonemployer that could plausibly be attributed to AI?
- "One-person billion-dollar company": trace Sam Altman's claim (2024) and find any actual solo or tiny-team companies with large valuations or exits (e.g. the Base44 sale to Wix in 2025 was reported at about $80m with a solo founder; this is unverified here).

- Care sector (partly answered): is anyone moving from AI-exposed occupations into care? No flow data found. Check the 2026 Skills for Care pay report against the April 2026 NLW.
- OpenResearch: reconcile 1.3 hrs/week and 2 pp with the 3.9 pp and about $2,000/yr figures. Get the detailed time-use table (leisure vs care vs education).
- Kenya UBI: are there long-run (2025-26) results from the 12-year arm?
- Four-day week: still no independent randomised evaluation found; Spain national pilot results and Tokyo uptake numbers not found; check the University of Münster's own publication for the German pilot (methods, any wearable or cortisol measures).
- Retirement: does retirement improve or harm wellbeing and cognition? This is the closest natural experiment for voluntary non-work with secure income.
- AI and job quality: are there field (not lab) studies of GenAI and meaning or autonomy, e.g. OECD worker surveys, Eurofound, or UK AI Security Institute and ONS opinions data?
- Gramatneusiedl: exact participant n, and whether effects persisted after the 2024 end of the pilot.

What's genuinely uncertain

- Whether AI is mainly an automation or an augmentation technology in the Autor et al. sense. That decides whether "new work" appears at scale.
- Whether the harm of losing work comes from losing income, losing structure, or losing status in a society that values employment. Existing evidence cannot separate them cleanly. A world of secure income and no stigma has never been observed.
- Whether freed time would be used differently if it came to everyone at once (a social norm) rather than to individuals in a treatment group or to the unemployed.
- Whether the work-intensification finding is a transitional effect of early tools and management norms or a lasting feature.
- Whether solo AI-enabled businesses are a real trend or a set of anecdotes. The headline cases rest on outsourced labour, and the official series that would show a trend stop at 2023 (US) or point down (UK).

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
~60% (some versions 63%)	Share of 2018 US employment in job titles that did not exist in 1940	Autor, Chin, Salomons, Seegmiller, QJE	https://economics.mit.edu/sites/default/files/2026-04/New-Frontiers-QJE-v139%283%29-p1399-1465-Aug-2024.pdf	Aug 2024	Peer-reviewed. Search summary only; not checked against PDF
41% to 1.9%	Agriculture share of US workforce, 1900 to 2000	USDA ERS (compiling HSUS and Census)	https://webharvest.gov/peth04/20041105150816/http://www.ers.usda.gov/publications/agoutlook/jun2000/ao272h.pdf	Jun 2000	Official, but seen via repost; other series give 2.4% for 2000
>25% fall; vs ~780 hrs	Fall in men's annual market hours 1931-2011, vs Keynes's implied target	Crafts, CEPR DP16140	https://repec.cepr.org/repec/cpr/ceprdp/DP16140.pdf	2021	Working paper; search summary only
47%	Share of US employment at "high risk" of computerisation	Frey and Osborne via ITIF	https://itif.org/publications/2022/09/30/oops-the-predicted-47-percent-of-job-loss-from-ai-didnt-happen/	2013 (forecast); 2022 (review)	Forecast; ITIF is an advocacy think tank
+16.4%	Employment growth 2013-2021 in insurance	ITIF	same as above	Sep 2022	Advocacy source using BLS data

Figure	What it measures	Source	URL	Date	Reliability note
	underwriters (highest F&O risk)				
-0.2 pp; -0.42%	Employment-to-population and wage effect per extra robot per 1,000 workers	Acemoglu and Restrepo, JPE	https://econpapers.repec.org/RePEc:ucp:jpolec:doi:10.1086/705716	2020	Peer-reviewed; other versions give different magnitudes
>50%	Share of US phone network AT&T automated 1920-1940	Feigenbaum and Gross, QJE	https://econpapers.repec.org/RePEc:oup:qjecon:v:139:y:2024:i:3:p:1879-1939.	Aug 2024	Peer-reviewed; search summary only
~1.3 hrs/ week; ~2 pp	Fall in hours and employment, $1,000/ month recipients vs controls	OpenResearch / Vivalt et al., NBER WP 32719 (via Fortune)	https://www.nber.org/system/files/working_papers/w32719/w32719.pdf	Jul 2024	RCT; figures from secondary coverage; a conference version reports 3.9 pp participation fall; reconcile
28% to 40% (control 32% to 37%)	Full-time employment, Stockton SEED recipients, Feb 2019 to Feb 2020	Martin-West and Castro Baker	https://catalog.results4america.org/case-studies/guaranteed-income-stockton	Mar 2021	n=125 treated; small
€560/month, ~2,000 people	Finland basic income experiment design	Finnish Ministry of Social Affairs and Health / Kela	https://stm.fi/en/-/perustulokokeilun-tulokset-tyollisyysvaikutukset-vahaisia-toimeentulo-ja-psyykkinen-terveys-koettiin-paremmaksi	May 2020	RCT; employment effect "small"
0 employment effect; +1.8 pp (17%) part-time	Alaska PFD labour effects	Jones and Marinescu, AEJ: Policy	https://www.aeaweb.org/doi/10.1257/pol.20190299	2022	Peer-reviewed; synthetic control; part-time figure conflicts with one text fragment
61 firms, ~2,900 staff; 56 continue, 18 permanent	UK four-day-week pilot	Autonomy / 4 Day Week Global, via WEF	https://www.weforum.org/stories/2023/03/four-day-work-week-uk-trial/	Feb 2023	Self-selected firms; advocacy-run
>2,500 workers, ~1% of workforce, 35-36 hrs	Iceland public-sector hours-reduction trials	Autonomy and Alda	https://issuu.com/4dayweekglobal/docs/iceland_4dw	Jun 2021	Advocacy report; not a four-day week
86%	Icelandic workers on contracts with shortened hours or the right to shorten	Autonomy and Alda	https://issuu.com/4dayweekglobal/docs/iceland_4dw	Jun 2021	Advocacy report; size of cuts varies
10-15% per year; 1.0-1.5 yrs	Mortality hazard increase 20 years after displacement; implied life-expectancy loss	Sullivan and von Wachter, QJE	https://econpapers.repec.org/RePEc:oup:qjecon:v:124:y:2009:i:3:p:1265-1306.	2009	Peer-reviewed; 50-100% first-year figure is from a draft only
~30% vs ~1%	Share of lost work hours going to sleep plus TV vs job search, Great Recession	Aguiar, Hurst, Karabarbounis (NBER Digest)	https://www.nber.org/digest/nov11/time-use-during-recessions	Nov 2011	ATUS-based; secondary summary
N=3,562; ~-11% motivation, ~+20% boredom	GenAI collaboration followed by solo work	Scientific Reports (4 experiments)	https://fortune.com/2025/06/03/ai-boost-worker-productivity-hurts-motivation-boredom	Apr 2025 (paper)	Online experiments; effect sizes from secondary coverage
3-7%	Share of chatbot productivity gains	Humlum and Vestergaard,	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	2025	Danish register data, 25,000 workers; working paper

Figure	What it measures	Source	URL	Date	Reliability note
	passed through to earnings	NBER WP 33777			
~5,671,836	US business applications, 2025 (annual record)	Census BFS (annual sum via secondary)	https://www.census.gov/econ/bfs/pdf/historic/bfs_2025m12.pdf	Jan 2026	Official monthly data; annual total from secondary sum
37.6% to 30.1%	High-propensity share of US business applications, 2019 to 2025	Consultant calculation from Census BFS	https://founderreports.com/how-many-businesses-are-started-each-year/	2026	Secondary calculation
29.8m (2022); 30.4m (2023)	US nonemployer businesses	Census NES / NES-D	https://www.census.gov/library/stories/2025/07/nonemployer-business-growth.html	Jul / Nov 2025	Official; search extract
4.9%; 4.7%	Nonemployer growth 2021, 2022 (employer: 1.9%, 1.8%)	Census	https://www.census.gov/library/stories/2025/07/nonemployer-business-growth.html	Jul 2025	Official; pre-AI period
~4.38m vs ~5.0m	UK self-employed, Q4 2025 vs end-2019	ONS LFS via Statista	https://statista.com/statistics/318234/united-kingdom-self-employed	2026	LFS volatile; series disagree
~$401m; ~$65m	Medvi 2025 revenue; net profit	NYT via Dealroom	https://dealroom.co/news/127216-from-20k-to-1-8b-in-revenue-how-one-founder-built-a-telehealth-giant-wit/	Apr 2026	Journalist saw financials; contested; $1.8bn is a projection
~$80m + earn-out	Wix purchase of Base44 (6-8 staff, ~6 months old)	Globes	https://en.globes.co.il/en/article-wix-acquires-israeli-vibe-coding-co-base44-1001513267	Jun 2025	Press report
6.2%; ~96,000	England adult social care vacancy rate and vacancies, 2025/26	Skills for Care via LocalGov	https://www.localgov.co.uk/Social-care-vacancy-rate-falls-to-decade-low-but-challenges-ahead/64563	Jun 2026	Sector data body; secondary report
£12.60 (39p above NLW)	Median care worker hourly pay, England, Dec 2025	Skills for Care via Healthcare Management	https://www.healthcare-management.uk/care-workers-pay-minimum-wage-reveals	2026	Sector data; NLW rose to £12.71 in Apr 2026
$34,900; +17%; 765,800/yr	US home health and personal care aides: median pay 2024; growth 2024-34; annual openings	BLS OOH	https://www.bls.gov/ooh/healthcare/home-health-aides-and-personal-care-aides.htm	2025	Official; search extract
45 firms, ~900 staff; ~73% continue	Germany four-day-week pilot 2024	University of Münster / Intraprenör via The Local	https://www.thelocal.de/20241106/what-are-the-results-of-germanys-four-day-work-week-trial	Nov 2024	Self-selected; ~20% reverted
9.5%	Traffic fall during Valencia's four Monday holidays, 2023	City-commissioned evaluation via The Mayor	https://www.themayor.eu/en/a/view/valencia-s-government-says-4-day-week-trial-was-good-for-the-people-and-the-city-12172	2023	Four weeks; holidays not work redesign

Verbatim quotes

None. No page could be fetched in this session, so no wording could be checked against source text. Candidate quotes to verify later (paraphrased here, not to be quoted until checked):

- Keynes (1930): suggested three-hour shifts or a fifteen-hour week would satisfy most people's need to work. Source to fetch: the original essay text.
- Will Stronge (Autonomy research director, July 2021): called the Iceland trials the world's largest public-sector shorter-working-week trial and "an overwhelming success" (as reported

in syndicated news, e.g. https://www.fox5ny.com/news/4-day-work-week-trials-in-iceland-deemed-overwhelming-success-by-researchers).

- Autor (MIT): said roughly six in ten jobs people do today did not exist in 1940 (reported in MIT coverage).
- Case and Deaton (2018 reply to Ruhm): said "deaths of despair" was a label, not an explanation. https://www.princeton.edu/~accase/downloads/Case_and_Deaton_Comment_on_CJRuhm_Jan_2018.pdf
- Jahoda/Marienthal: free time experienced by the unemployed as a "tragic gift" (as summarised in the Paul et al. 2023 meta-analysis).

Sources worth reading in full

1. Autor, Chin, Salomons, Seegmiller, "New Frontiers", QJE 2024: the best evidence on where new work comes from, and on the automation vs augmentation split that should frame any "what will people do" argument. https://economics.mit.edu/sites/default/files/2026-04/New-Frontiers-QJE-v139%283%29-p1399-1465-Aug-2024.pdf
2. Feigenbaum and Gross, "Answering the Call of Automation", QJE 2024: a clean case of automation in a clerical occupation, separating cohort effects from incumbent effects. https://www.nber.org/system/files/working_papers/w28061/w28061.pdf
3. Vivalt et al., OpenResearch, NBER WP 32719: the largest US cash RCT, with time-use data on what people actually did. https://www.nber.org/system/files/working_papers/w32719/w32719.pdf
4. Kasy and Lehner, "Employing the unemployed of Marienthal", IZA DP 16088: the direct test of Jahoda's latent functions. https://docs.iza.org/dp16088.pdf
5. Humlum and Vestergaard, NBER WP 33777: large-scale register evidence on what chatbots have done to hours and earnings. https://econpapers.repec.org/RePEc:nbr:nberwo:33777
6. Paul et al. 2023 meta-analysis of the latent deprivation model: tests which of work's non-pay functions matter. https://pmc.ncbi.nlm.nih.gov/articles/PMC10017486
7. Sullivan and von Wachter, QJE 2009: the hardest number on the human cost of job loss. https://www.movingforwardnetwork.com/wp-content/uploads/2023/04/Sullivan_Job-Displacement-and-Mortality_2009.pdf
8. Crafts, "Economic Possibilities for our Grandchildren revisited": a quantified check of Keynes. https://repec.cepr.org/repec/cpr/ceprdp/DP16140.pdf
9. Banerjee et al., Kenya UBI short-term results: a counterpoint showing cash funding enterprise, not leisure. https://conference.nber.org/conf_papers/f192616.pdf
10. Autonomy/Alda Iceland report: read it to see how thin the design is behind the "overwhelming success" headline. https://issuu.com/4dayweekglobal/docs/iceland_4dw
11. Census, "Number of U.S. Nonemployers Grew Faster Than Employer Businesses" (July 2025): the baseline for any claim about a solo-business boom. https://www.census.gov/library/stories/2025/07/nonemployer-business-growth.html
12. Skills for Care, State of the adult social care sector and workforce (2025/26) and pay report: the best data on whether care can absorb displaced workers, and at what pay. https://www.localgov.co.uk/Social-care-vacancy-rate-falls-to-decade-low-but-challenges-ahead/64563

INCOMPLETE (less so after gap-fill pass, 2026-11-06). Covered: public opinion (US, global, UK, EU), union and bargaining cases, EU AI Act timeline and omnibus, Trump December 2025 pre-emption order and DOJ Task Force follow-through, the July 2025 Senate 99-1 vote, US state laws (Colorado repeal-and-replace, California SB 53, NYC LL144 and the Comptroller audit, Illinois HB 3773, Texas TRAIGA), UK (Action Plan, AISI rename, no AI bill), surplus capture (BLS labour share, hyperscaler capex, Danish wage evidence, Acemoglu), the anti-AI case on copyright (Bartz, Thomson Reuters v. Ross) and energy (IEA, LBNL), and the grifter layer (FTC Operation AI Comply and later cases including Air AI; SEC AI-washing). STILL MISSING: UK ONS labour share; water-use estimates; Bender/Gebru, Marcus and Zitron as OPINION; NYT v. OpenAI and Getty v. Stability status; hyperscaler Q3 2026 figures from filings; final outcomes for Ascend Ecom, FBA Machine and Air AI; size of the AI course/advice market; outcome of xAI/DOJ v. Colorado; WGA 2026 contract text; Kaiser dispute outcome. All gap-fill findings rest on search-engine extracts, not fetched pages (marked "search extract").

Track 9: Social and political fallout

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Method warning for this file. No source page was opened directly (also true of the gap-fill pass, whose findings are marked "search extract"). Every finding below comes from web-search result summaries that cite the URL given. The URLs are the primary pages (Pew, Gallup, KPMG/University of Melbourne, NUHW, law-firm notes on the Official Journal), but the numbers have not been checked against the page text in this session. Treat every number as "reported by search, to re-verify" before publishing. There are no verbatim quotes in this file for that reason.

Summary

- **US public concern is rising, not settling.** Pew's June 2026 survey has 52% of US adults more concerned than excited about AI in daily life (37% in 2021, 50% in June 2025), and 71% expecting AI to mean fewer jobs over 20 years (64% in 2024). Gallup/Bentley in July 2026 found the share saying AI does more harm than good jumped from 31% (2025) back to 39%, close to the 2023 level. The two-year improvement in US attitudes ended in 2026.
- **The shift is driven by the young and by jobs.** Under-30s moved most in both Pew (55% now more concerned than excited) and Gallup (18-29 trust in business AI use fell from 30% to 20% in one year). Democrats overtook Republicans on job worry in 2026 (Pew, September 2026). AI is becoming a partisan and generational issue in the US.
- **Use and trust have decoupled globally.** KPMG/University of Melbourne (48,000+ people, 47 countries, Nov 2024 to Jan 2025): 66% use AI regularly on purpose, only 46% are willing to trust it, and trust fell compared with the 2022 wave. People adopt AI while distrusting it. That is not the same thing as "resisting a new way of working".
- **Geography splits sharply.** Anglosphere and parts of Europe are the most nervous (US and Italy 50% more concerned than excited in Pew's 25-country 2025 survey); Asia is the most

positive (South Korea 16%, India 19% more concerned). Edelman reports 72% trust in AI in China against 32% in the US.

- **Unions are winning contract language, not stopping deployment.** WGA (2023, renewed 2026), SAG-AFTRA (video games July 2025, TV/theatrical June 2026), and the Commonwealth Bank reversal (August 2025) all produced consent, disclosure, credit or redundancy protections. The 2026 SAG-AFTRA deal requires producers to show "significant additional value" before using an AI performer (as reported). Strikes in 2026 (Kaiser therapists, ProPublica) were one-day actions.
- **The CBA case cuts against the thesis.** The bank cut 45 call-centre roles claiming a voice bot had reduced call volumes; the union said volumes rose and staff were on overtime, and CBA conceded the roles were not redundant. That is a failure of the business case and of measurement, not of workers' readiness.
- **US state AI law: passed, delayed, narrowed, barely enforced.** Colorado's 2024 AI Act was delayed (Aug 2025) and then repealed and replaced (May 2026) with a narrower disclosure law effective January 2027, before any of it took effect. New York's State Comptroller found NYC's AI hiring-audit law (in force since July 2023) ineffectively enforced (Dec 2025). Congress rejected federal pre-emption 99-1 in July 2025; the executive branch then pursued it through an order and a DOJ task force. The UK has no AI bill.
- **The surplus is flowing to capital so far, on thin evidence.** US labour share fell to 54.4% in Q4 2025 (BLS) while four hyperscalers guided to roughly $700-745bn of 2026 capex. The best micro study (Denmark, 25,000 workers) finds no effect of chatbots on earnings or hours, with only a small share of time savings reaching pay.
- **Copyright and energy give critics real ground.** Anthropic's $1.5bn class settlement (final approval July 2026) and the Third Circuit's Sept 2026 ruling against Ross (first appellate AI-training fair-use decision) show the legal cost is real. IEA projects data-centre electricity to more than double to about 945 TWh by 2030.
- **Regulators police AI over-claiming, including agent claims.** FTC Operation AI Comply (2024) and the August 2025 Air AI suit (an agent marketed to small businesses as replacing staff) target the pitch an agentic-AI vendor makes; SEC "AI washing" cases (Presto, Nate) were about humans hidden behind an AI label.
- **Regulation is announced faster than it is enforced.** The EU's Digital Omnibus on AI (Regulation (EU) 2026/1744, in force 27 July 2026) pushed high-risk obligations from August 2026 to December 2027 (standalone) and August 2028 (embedded in products). I found no reliable report of any AI Act fine to date. In the US, the December 2025 executive order tries to pre-empt state AI laws by litigation and funding pressure; its legal force is disputed.

Key findings

Public opinion: United States, trend

- **[EVIDENCE]** 52% of US adults say they are more concerned than excited about the increased use of AI in daily life, up from 37% in 2021. Search summaries disagree on the "more excited" share (9% or 11%); check the topline. Survey fielded June 22-28, 2026, n=3,488 US adults (Pew's American Trends Panel, probability-based). Pew Research Center, 18 Aug 2026. https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/ (topline: https://www.pewresearch.org/wp-content/uploads/sites/20/2026/08/SR_26.08.19_young-adults-ai_topline.pdf)
- **[EVIDENCE]** 71% of US adults expect AI to lead to fewer jobs over the next 20 years, up from 64% in 2024; about 5% expect more jobs. 55% of under-30s are now more concerned than

excited, and young adults are now about as likely as 30-64 year olds to expect fewer jobs. Same Pew release and topline, 18 Aug 2026. https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/

- **[EVIDENCE]** June 2025 baseline: 50% more concerned than excited, 10% more excited, 38% both; n=5,023 US adults. Pew, 2025 (reported via MarketingProfs summary of the Pew release). https://www.marketingprofs.com/charts/2025/53727/Americans-more-concerned-than-excited-about-ai-study-pew
- **[EVIDENCE]** Partisan shift: the share of Democrats expecting AI to lead to fewer jobs rose 17 points, while Republicans' views changed little; in the June 2026 survey 56% of Democrats were more concerned than excited against about half of Republicans. Pew, 16 Sep 2026. https://www.pewresearch.org/short-reads/2026/09/16/democrats-are-now-more-worried-than-republicans-about-ai-and-its-impact-on-jobs/ (topline: https://www.pewresearch.org/wp-content/uploads/sites/20/2026/09/SR_09.16_26_ai-partisan-views_topline.pdf)
- **[EVIDENCE]** Public vs AI experts: Pew surveyed 5,410 US adults and 1,013 AI experts (defined as authors or presenters at 2023 or 2024 AI conferences; expert fieldwork Aug 14 to Oct 31, 2024). Experts were consistently far more optimistic than the public, including about jobs; both groups wanted more regulation and more personal control. Pew, 3 Apr 2025. https://www.pewresearch.org/internet/2025/04/03/how-the-us-public-and-ai-experts-view-artificial-intelligence/ (PDF: https://pewresearch.org/wp-content/uploads/sites/20/2025/04/pi_2025.04.03_us-public-and-ai-experts_report.pdf). Note: the expert sample is conference participants, so it over-represents people whose careers depend on AI.
- **[EVIDENCE]** Gallup/Bentley, 2026: share saying AI does more harm than good rose from 31% (2025) to 39% (2026), close to 40% in 2023. The 11-point rise among young adults was the largest of any age group. 79% expect AI to reduce US jobs over 10 years (73% in 2025, 75% the year before). 70% now say they are at least somewhat knowledgeable about AI (64% in 2024). Gallup, "Americans Cool Toward AI", 28-29 Jul 2026. https://news.gallup.com/poll/712751/americans-cool-toward.aspx
- **[EVIDENCE]** Gallup/Bentley trust in business: 27% trust businesses "a lot" or "some" to use AI responsibly in 2026, down from 31% in 2025 and up from 21% in 2023; among 18-29s, trust fell from 30% to 20% in a year. 2025 wave: web survey, n=3,007, May 5-12, 2025. Bentley University / Gallup. https://www.bentley.edu/news/do-americans-trust-business-and-ai and https://news.gallup.com/poll/712751/americans-cool-toward.aspx
- **[EVIDENCE, conflicting]** HarrisX poll of registered voters (July 2026) found more than half of Americans have a favourable view of AI. Different question (favourability vs concern/excitement) and different sample. HarrisX / New York Post, Jul 2026. https://www.harrisx.com/posts/new-york-post-most-american-voters-have-favorable-view-of-ai----despite-media-focus-on-controversies

Public opinion: global and by country

- **[EVIDENCE]** Pew Spring 2025 Global Attitudes Survey, 25 countries: median 34% more concerned than excited, 42% equally both, 16% more excited. No country had "more excited" as the largest group. Concern highest in the US and Italy (50% each), with Australia, Brazil and Greece near half; lowest in South Korea (16%) and India (19%). Older adults more concerned than younger adults in 18 of 25 countries; women more worried than men in over half. Median 55% trust their own country to regulate AI; trust in the US and China to regulate AI is lower. Pew, 15 Oct 2025. https://www.pewresearch.org/global/2025/10/15/how-people-around-the-world-view-ai/ (PDF: https://alpha.pewresearch.org/pewresearch-org/wp-content/uploads/sites/20/2025/10/pg_2025.10.15_ai_report.pdf). Note: the "older more concerned"

pattern in this 2025 global data contrasts with the 2026 US data where the young moved most.

- **[EVIDENCE]** A separate Pew international survey (reported in 2026) found respondents in 34 of 37 countries more likely to expect AI to eliminate jobs than create them over 20 years. Reported via Pew search summary; release URL not confirmed. Cited from the same search summary as https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/. Re-verify which Pew report this is.
- **[EVIDENCE (survey by a consultancy with a university; method is a proper panel survey)]** KPMG and University of Melbourne, "Trust, attitudes and use of artificial intelligence: A global study 2025": 48,000+ respondents in 47 countries, Nov 2024 to Jan 2025. 66% use AI regularly and intentionally; 46% are willing to trust it; trust has fallen and worry risen since the 2022 wave (17 countries). Marked differences between advanced and emerging economies (emerging economies more trusting). University of Melbourne figshare, 2025 (CC BY-NC-SA). https://figshare.unimelb.edu.au/articles/report/Trust_attitudes_and_use_of_artificial_intelligence_A_global_study_2025/28822919 ; press release: https://mbs.edu/news/global-study-reveals-trust-of-ai-remains-a-critical-challenge
- **[EVIDENCE (weak: online panel, 30 countries, not fully representative in emerging markets)]** Ipsos AI Monitor 2025: on average 52% say they are excited about AI products and services and 53% say AI makes them nervous. English-speaking countries most nervous and least excited; South-East Asia (Thailand, Indonesia, Malaysia) most excited, and excitement tracks belief that AI will help the national economy. Japan neither. Ipsos, June 2025. https://www.ipsos.com/en-dk/ipsos-ai-monitor-2025 (PDF: https://www.ipsos.com/sites/default/files/ct/news/documents/2025-06/Ipsos-AI-Monitor-2025_0.pdf). The 2026 edition (32 countries, June 2026) reportedly finds excitement and nervousness almost even and younger cohorts more nervous than older ones; seen only via a mirror PDF, not Ipsos: https://astrid-online.it/static/upload/protected/ipso/ipsos-ai-monitor-2026.pdf
- **[EVIDENCE (weak: PR-firm survey)]** Edelman Trust Barometer 2025 tech-sector report: 72% of Chinese respondents trust AI against 32% of Americans; Ireland, Australia, UK, Netherlands and Germany below 30%. Edelman, 2025. https://www.edelman.com/sites/g/files/aatuss191/files/fr/2025-08/Top%20Findings_2025%20EdelmanTrust%20Barometer_Insights%20for%20the%20Technology%20Sector%20%281%29.pdf . The 2026 Barometer showed global trust in technology companies down 3 points, with a starker drop for AI companies (Semafor, 11 Feb 2026). https://semafor.com/article/02/11/2026/trust-in-tech-companies-dropped-globally-in-2026-report
- **[EVIDENCE]** EU: Eurobarometer on AI and the future of work (published 13 Feb 2025): more than 60% of Europeans view robots and AI at work positively, over 70% think they improve productivity, and 84% say AI at work needs careful management to protect privacy and transparency. European Commission, 13 Feb 2025. https://digital-strategy.ec.europa.eu/en/news/commission-survey-shows-most-europeans-support-use-artificial-intelligence-workplace . Note: Commission press framing; question wording not checked.
- **[EVIDENCE]** UK: DSIT Public Attitudes to Data and AI Tracker, Wave 4 (fieldwork Jul/Aug 2024, published 16 Dec 2024, Savanta, n>5,000). Declining trust in organisations to keep data safe, especially social media and big tech; worries about misinformation and job displacement; over half think AI is already used in public services. No Wave 5 found. DSIT, 16 Dec 2024. https://www.gov.uk/government/publications/public-attitudes-to-data-and-ai-tracker-survey-wave-4

Unions and collective bargaining

- **[EVIDENCE]** WGA 2023 MBA (term 25 Sep 2023 to 1 May 2026): AI-generated material is not "literary material" or "source material", so it cannot undercut credit or separated rights; a writer may choose to use AI with company consent, but a company cannot require it; companies must disclose if material handed to a writer was AI-generated; WGA reserved the right to argue that training on writers' work is prohibited. WGA contract summary, 2023. https://wgacontract2023.org/the-campaign/summary-of-the-2023-wga-mba
- **[EVIDENCE]** WGA 2026: tentative four-year MBA approved unanimously by the negotiating committee on 4 Apr 2026, later ratified. Reports conflict on AI: some say it adds limits and licensing requirements on use of writers' scripts for AI training; an IndieWire summary says AI terms largely codify 2023. NPR/OPB, 6 Apr 2026. https://www.opb.org/article/2026/04/06/hollywood-studios-reach-a-tentative-agreement-with-writers-union/ . Contract text not checked.
- **[EVIDENCE]** SAG-AFTRA video game strike (began July 2024, the longest in the union's interactive history) ended when members ratified the 2025 Interactive Media Agreement 95.04% to 4.96% in July 2025. Terms: consent and disclosure for digital replicas, ability to suspend consent for new material during a strike, pay rise over 15% with further raises 2025-2027. PC Gamer, Jul 2025. https://www.pcgamer.com/gaming-industry/the-videogame-voice-actors-strike-officially-ends-as-performers-ratify-agreement-between-sag-aftra-and-publishers
- **[EVIDENCE]** SAG-AFTRA 2026 TV/Theatrical Agreement ratified 4 Jun 2026, 91.42% in favour on a 19.25% turnout, running to 30 Jun 2030. Reported AI terms: stronger digital-replica protections, restrictions on "synthetic performers", including a requirement that producers show "significant additional value" before using an AI performer; no strike right over AI until 2030. Fortune, 5 Jun 2026. https://fortune.com/2026/06/05/sag-aftra-approves-four-year-contract-studios-streamers/ ; RedShark News on the AI terms: https://www.redsharknews.com/sag-aftra-ai-deal-synthetic-actors-2026
- **[EVIDENCE]** Commonwealth Bank of Australia, Aug 2025: CBA cut 45 direct-banking roles citing a voice bot that it said cut call volumes by 2,000 a week. The Finance Sector Union took it to the Fair Work Commission, saying volumes were rising and staff were on overtime. CBA conceded the roles were not redundant and offered them back (or voluntary exit). Bloomberg, 21 Aug 2025. https://www.bloomberg.com/news/articles/2025-08-21/commonwealth-bank-reverses-job-cuts-decision-over-ai-chatbots ; union account: https://www.fsunion.org.au/win-cba-backflips-on-customer-service-job-cuts-admits-they-got-it-wrong/ . Note: the volume figures are the union's and CBA's competing claims; no independent data.
- **[EVIDENCE]** AFL-CIO and Microsoft, 11 Dec 2023: a partnership on AI information-sharing, worker input into AI development and policy, plus a neutrality framework for organising by AFL-CIO affiliates at Microsoft. No public record found of what it has produced since. Microsoft, 11 Dec 2023. https://news.microsoft.com/2023/12/11/afl-cio-and-microsoft-announce-new-tech-labor-partnership-on-ai-and-the-future-of-the-workforce/
- **[EVIDENCE]** UK TUC draft Artificial Intelligence (Regulation and Employment Rights) Bill, 18 Apr 2024, drafted by Robin Allen KC and Dee Masters with the Minderoo Centre (Cambridge): rights when employers use AI for "high-risk" decisions, and a ban on emotion recognition used to workers' detriment. No sign it was introduced in Parliament or adopted by government. TUC, 18 Apr 2024. https://www.tuc.org.uk/research-analysis/reports/artificial-intelligence-regulation-and-employment-rights-bill
- **[EVIDENCE (weak: union member survey, self-selected, n>2,300)]** National Nurses United survey (Jan to Mar 2024): 60% disagreed that their employer would implement AI with

patient safety first; where acuity algorithms are used, 69% said the outputs did not match their own assessment; 40% in facilities with predictive scores said they cannot override them. NNU called for a pause on AI in health care. NNU press release, May 2024. https://www.nationalnursesunited.org/press/national-nurses-united-survey-finds-ai-technology-undermines-patient-safety

- **[EVIDENCE]** Kaiser Permanente, 18 Mar 2026: about 2,400 NUHW mental health workers in Northern California and the Central Valley held a one-day unfair-labour-practice strike, centred on Kaiser's unilateral overhaul of mental-health triage and the union's fear of AI replacing therapist work (including AI note-taking). Talks were still unresolved in late April 2026. KQED, Mar 2026. https://www.kqed.org/news/12076753/northern-california-kaiser-therapists-hold-1-day-strike-over-ai-patient-care-concerns ; NUHW: https://home.nuhw.org/2026/03/17/kaiser-mental-health-strike-to-start-at-6am-wednesday-march-18/
- **[EVIDENCE]** ProPublica, 8 Apr 2026: the NewsGuild unit held a 24-hour strike over AI safeguards, layoff protections, just cause and wages; reported as the first US newsroom strike with AI protections as a central demand. Nieman Lab, Apr 2026. https://www.niemanlab.org/reading/unionized-propublica-staff-are-on-strike-over-ai-layoffs-and-wages/

Regulation: EU, enforced vs announced

- **[EVIDENCE]** EU AI Act original timetable: prohibitions from 2 Feb 2025, GPAI obligations from 2 Aug 2025, most high-risk obligations from 2 Aug 2026. The Digital Omnibus on AI (provisional deal 7 May 2026; Parliament plenary 16 Jun 2026, 423 for, 57 against, 174 abstentions; Council approval 29 Jun 2026) was published as Regulation (EU) 2026/1744 on 24 Jul 2026 and entered into force 27 Jul 2026. Standalone (Annex III) high-risk obligations move to 2 Dec 2027; high-risk AI embedded in regulated products to 2 Aug 2028. Article 50 transparency duties and the AI Office's powers to enforce against GPAI providers started on schedule on 2 Aug 2026. Sidley, 22 Jun 2026: https://datamatters.sidley.com/2026/06/22/eu-lawmakers-reach-provisional-agreement-to-delay-key-eu-ai-act-obligations/ ; National Law Review on entry into force: https://natlawreview.com/article/eu-digital-omnibus-ai-enters-force ; Gibson Dunn: https://www.gibsondunn.com/eu-ai-act-omnibus-agreement-postponed-high-risk-deadlines-and-other-key-changes/ . Regulation number and dates come from law-firm summaries; confirm on EUR-Lex.
- **[EVIDENCE, negative finding]** I found no credible report of any fine under the AI Act as of November 2026. One web page claims three Commission enforcement decisions totalling EUR 18.4 million on 14 Feb 2026; this is uncorroborated and conflicts with the Commission's GPAI fining powers only starting on 2 Aug 2026. Treat as false until shown otherwise. Penalty ceilings: up to EUR 35m or 7% of global turnover for prohibited practices; up to EUR 15m or 3% for GPAI providers.

Regulation: US federal

- **[EVIDENCE]** Executive order "Eliminating State Law Obstruction of National AI Policy", 11 Dec 2025: directs the Attorney General to set up an AI Litigation Task Force to challenge state AI laws (Commerce Clause, pre-emption, other theories); Commerce to identify "onerous" state laws, including those that would force models to "alter truthful outputs"; FTC to issue a policy statement on when consumer-protection law pre-empts state rules; FCC to consider a federal disclosure standard; conditions on federal broadband (BEAD) funding. Legal scholars and some Republican governors (DeSantis, Cox) argued an executive order cannot itself pre-empt state law without Congress. Ropes & Gray, Dec 2025: https://www.ropesgray.com/en/insights/alerts/2025/12/trump-attempts-to-preempt-state-ai-regulation-through-executive-order ; NPR

via WBHM, 11 Dec 2025: https://wbhm.org/npr-story/trump-is-trying-to-preempt-state-ai-laws-via-an-executive-order-it-may-not-be-legal/

- **[EVIDENCE]** Senate vote on the state-law moratorium: in the early hours of 1 July 2025, during the "vote-a-rama" on the One Big Beautiful Bill Act, the Senate voted 99-1 for an amendment led by Sen. Marsha Blackburn (R-TN), co-sponsored by Sen. Maria Cantwell (D-WA), to strip the proposed moratorium on state AI regulation. The Senate version had already been narrowed to deny states that regulate AI access to a new $500 million AI infrastructure fund. Sen. Ted Cruz, the provision's author, voted to strip it after Blackburn abandoned a compromise with him; Sen. Thom Tillis cast the only no vote. Republican and Democratic governors had lobbied against it for weeks. Nextgov/Route Fifty, 1 Jul 2025 (search extract). https://www.nextgov.com/artificial-intelligence/2025/07/senate-overwhelmingly-passes-amendment-removing-state-ai-moratorium/406448 ; PBS NewsHour: https://www.pbs.org/newshour/politics/senate-pulls-ai-regulatory-ban-from-gop-bill-after-complaints-from-states (search extract). Roll call not checked on senate.gov. Significance: Congress declined pre-emption by 99-1 five months before the executive branch tried to achieve it by order.
- **[EVIDENCE]** Follow-through on the order: Attorney General Pam Bondi set up the DOJ AI Litigation Task Force on 9 Jan 2026. Its first litigation step was to intervene (April 2026) in xAI's suit challenging the Colorado AI Act, the only state law named in the December 2025 order; reported as the first DOJ intervention against a state AI law. Commerce was due to list "onerous" state laws by 11 Mar 2026; one source said no list had been published (undated). Axios, 24 Apr 2026 (search extract): https://axios.com/2026/04/24/justice-department-joins-xai-challenge-colorado-ai-law ; Jenner & Block (search extract): https://www.jenner.com/en/news-insights/client-alerts/doj-joins-xai-in-lawsuit-challenging-colorado-ai-act . Colorado then repealed and replaced the law in May 2026 (below), which may moot the suit; outcome not found.
- **Not sourced in this session:** other Task Force suits, the Commerce list's contents, the FTC pre-emption policy statement, and any 2026 federal pre-emption legislation.

Regulation: US states, enforced vs announced

The pattern across states: laws are passed, then delayed, narrowed or thinly enforced before they bite. The only state AI-in-employment law with a multi-year enforcement record (NYC) was found by the state's own auditor to be weakly enforced.

- **[EVIDENCE]** Colorado AI Act (SB 24-205, 2024), the first comprehensive US state law on "high-risk" AI, never took effect in its original form. Governor Polis signed SB 25B-004 on 28 Aug 2025 in a special session, moving the effective date from 1 Feb 2026 to 30 Jun 2026 to allow amendment in the 2026 session. Greenberg Traurig, Sep 2025 (search extract). https://www.gtlaw.com/en/insights/2025/9/colorado-delays-comprehensive-ai-law-with-further-changes-anticipated
- **[EVIDENCE]** In May 2026 the legislature repealed and replaced it with SB 26-189 (passed Senate 7 May, House 9 May; signed 14 May 2026 by most accounts, one outlier says 18 May). The replacement drops the high-risk AI framework (impact assessments, risk-management programmes) for a narrower automated decision-making technology (ADMT) disclosure regime, effective 1 Jan 2027, enforced only by the Attorney General after a 60-day cure period, with no private right of action. Ropes & Gray, May 2026 (search extract): https://www.ropesgray.com/en/insights/alerts/2026/05/colorado-scales-back-ai-law-with-targeted-implications-for-health-care ; EPIC: https://epic.org/colorado-legislature-again-amends-landmark-ai-law/ ; Davis+Gilbert: https://www.dglaw.com/what-to-know-about-colorados-repeal-and-replacement-of-its-landmark-ai-act/ (search extracts). Bill text not checked on

leg.colorado.gov. Net effect: the flagship state AI law was announced in May 2024 and, two and a half years later, nothing in it has been enforced.

- **[EVIDENCE]** NYC Local Law 144 (automated employment decision tools: annual independent bias audit, candidate notice; enforced by the Department of Consumer and Worker Protection since July 2023). The New York State Comptroller's audit, issued 2 Dec 2025 and covering July 2023 to June 2025, found DCWP's enforcement ineffective. Secondary summaries say enforcement actions were rare, complaint handling weak and DCWP did not follow up to check violations were fixed; DCWP agreed to implement most recommendations. Penalties run up to $1,500 per violation per day. NY State Comptroller, 2 Dec 2025 (search extract): https://www.osc.ny.gov/state-agencies/audits/2025/12/02/enforcement-local-law-144-automated-employment-decision-tools ; DLA Piper, 30 Jan 2026 (search extract): https://www.dlapiper.com/en-au/insights/publications/2026/01/critical-audit-of-nyc-ai-hiring-law-signals-increased-risk-for-employers . Details such as "13 recommendations" and "32 bias audits reviewed" appear only in third-party blogs; check the audit PDF.
- **[EVIDENCE]** California SB 53, Transparency in Frontier Artificial Intelligence Act: signed 29 Sep 2025, effective 1 Jan 2026. Applies to frontier-model developers, with extra duties for "large frontier developers" (over $500m annual revenue with affiliates): publish a safety framework and a transparency report at or before each new or substantially modified frontier model release, report critical safety incidents (15 days per one guide), whistleblower protections, and a public compute consortium (CalCompute). Civil penalties up to $1m per violation, enforced by the Attorney General. It dropped SB 1047's third-party audits, shutdown capability and 72-hour reporting. Future of Privacy Forum (search extract): https://fpf.org/blog/californias-sb-53-the-first-frontier-ai-law-explained/ ; Clifford Chance, Oct 2025: https://www.cliffordchance.com/insights/resources/blogs/talking-tech/en/articles/2025/10/sb-53--california-sets-standards-for-ai-transparency.html (search extracts). No enforcement action found. Relevance to a deployer like Actualisation: none directly; it binds model developers.
- **[EVIDENCE]** Illinois HB 3773 (amends the Illinois Human Rights Act): from 1 Jan 2026 employers may not use AI in employment decisions in a way that discriminates on protected characteristics, or use zip codes as a proxy, and must notify employees of AI use. Enforced through the Illinois Department of Human Rights and Human Rights Commission, with a private right of action in circuit court. Bass Berry & Sims, 2025 (search extract). https://bassberry.com/news/texas-illinois-and-colorado-legislative-developments-impact-healthcare-ai-and-other-industries-while-u-s-congress-fails-to-pass-moratorium/ . No reported enforcement found.
- **[EVIDENCE]** Texas Responsible AI Governance Act (TRAIGA, HB 149): signed 22 Jun 2025, effective 1 Jan 2026. Much narrower than the bill first filed: it bans a short list of intentional harms (e.g. manipulation to self-harm, unlawful discrimination by intent, certain deepfakes), adds disclosure duties mainly for government agencies and health care, creates a regulatory sandbox, and gives enforcement to the Attorney General with a cure period and no private right of action. Reported penalties: $10,000-12,000 per curable violation, $80,000-200,000 per uncurable one, $2,000-40,000 per day for continuing violations. Norton Rose Fulbright (search extract): https://www.nortonrosefulbright.com/en-us/knowledge/publications/c6c60e0c/the-texas-responsible-ai-governance-act . Penalty figures vary across summaries; check the statute. No enforcement action found.

Regulation: UK

- **[EVIDENCE]** AI Opportunities Action Plan, published by DSIT on 13 Jan 2025, written by Matt Clifford (commissioned July 2024). Pro-adoption: AI Growth Zones, compute expansion, a National Data Library, public-sector adoption. The government accepted its recommendations

(the commonly cited count is 50; I did not confirm it in a search extract). Clifford left his role as the PM's AI adviser in mid-2025. GOV.UK (search extract). https://www.gov.uk/government/publications/ai-opportunities-action-plan/ai-opportunities-action-plan

- **[EVIDENCE]** On 14 Feb 2025, Technology Secretary Peter Kyle, at the Munich Security Conference, renamed the AI Safety Institute the AI Security Institute, refocusing it on national-security and criminal-misuse risks; reporting said it would drop work on bias and free speech. Announced alongside an MOU with Anthropic. GOV.UK speech (search extract): https://www.gov.uk/government/speeches/remarks-made-by-technology-secretary-peter-kyle-at-the-munich-security-conference ; TechCrunch, 13 Feb 2025: https://techcrunch.com/2025/02/13/uk-drops-safety-from-its-ai-body-now-called-ai-security-institute-inks-mou-with-anthropic
- **[EVIDENCE, negative finding]** No UK AI bill as of November 2026. The July 2024 King's Speech promised legislation on "the most powerful AI models"; it was repeatedly deferred, and law-firm commentary says the May 2026 King's Speech again contained no AI bill, only a Regulating for Growth Bill putting regulatory sandboxes on a statutory footing. Lewis Silkin, 14 Sep 2026 (search extract): https://www.lewissilkin.com/insights/2026/09/14/ai-judgment-day-on-the-horizon-while-uk-lawmakers-play-catch-up ; Resultsense roundup of Osborne Clarke, 28 May 2026 (search extract): https://www.resultsense.com/news/2026-05-28-uk-ai-regulatory-may-2026-roundup/ . Two low-quality sites claim a frontier AI bill was introduced on 14 May 2026; law-firm sources contradict them. The Lewis Silkin extract also mentions a change of Prime Minister that I could not confirm. Check bills.parliament.uk before relying on either.
- Net: the UK has a growth plan and a security-focused institute, but no binding AI-specific law. Worker-side proposals (TUC draft bill, below) have not been taken up.

Who captures the productivity surplus

- **[EVIDENCE]** US labour share (nonfarm business, share of output paid as compensation) was 54.4% in Q4 2025, 1.7% lower than a year earlier, while 2025 productivity grew 2.2% (output +2.6%, hours +0.4%) and unit labour costs 2.3%. BLS Productivity and Costs, revised Q4 2025 release, 24 Mar 2026 (search extract). https://www.bls.gov/news.release/archives/prod2_03242026.htm ; series: https://fred.stlouisfed.org/series/PRS85006173 . I could not confirm whether 54.4% is a series low; the long-run series is on FRED. Caution: a one-year labour-share fall in a year of strong productivity is consistent with gains flowing to capital, but 2025 productivity is not attributable to AI on this evidence alone.
- **[EVIDENCE (company guidance, via secondary aggregation)]** Hyperscaler capex for calendar 2026: Microsoft about $190bn (including about $25bn attributed to higher component prices), Alphabet $180-190bn (later raised to $195-205bn per one tracker), Amazon about $200bn, Meta $125-145bn (later $130-145bn), as guided on Q1 2026 calls (late April 2026). Totals reported at about $700-725bn (Q1 guidance) and $720-745bn (one tracker after Q2 results). Amazon's trailing 12-month free cash flow was reported at $1.2bn against $26bn a year earlier. Yahoo Finance, Apr-May 2026 (search extract): https://finance.yahoo.com/markets/article/magnificent-7-earnings-rush-reveals-ai-spending-surge-with-hyperscaler-capex-set-to-reach-725-billion-in-2026-224901707.html ; tracker (updated 11 Oct 2026): https://www.datacentres.com/guides/hyperscaler-capex-trends (search extracts). Definitions differ (finance leases included or not). Not checked against 10-Qs or call transcripts; Q3 2026 figures not found.
- **[FORECAST]** Evercore and Bank of America were reported to project combined big-tech capex above $1 trillion in 2027. Same Yahoo Finance extract. Treat as sell-side forecast.
- **[EVIDENCE]** Wage effects so far: Humlum and Vestergaard, "Large Language Models, Small Labor Market Effects" (NBER WP 33777, 2025): two adoption surveys (late 2023 and 2024) of 25,000 workers in 7,000 workplaces across 11 exposed occupations, linked to Danish

employer-employee registers, difference-in-differences using employer AI policies. No significant effect on earnings or recorded hours in any occupation, with confidence intervals ruling out large effects; reported time savings were small, and summaries say only about 3-7% of productivity gains passed through to pay. A 2026 update ("Still Waters, Rapid Currents", RFBerlin discussion paper) reframes the result as task reorganisation, with adopters moving into higher-paying occupations, still too few to shift average earnings. NBER (search extract): https://econpapers.repec.org/RePEc:nbr:nberwo:33777 ; 2026 update: https://ideas.repec.org/p/crm/wpaper/26078.html (search extract). Conflicting: a CPS-based working paper reports earnings rises in highly exposed US occupations after ChatGPT with no change in unemployment. arXiv 2509.15510 (search extract): https://arxiv.org/html/2509.15510v1 . Different data and identification; not directly comparable.

- **[FORECAST, academic]** Acemoglu, "The Simple Macroeconomics of AI" (NBER WP 32487, 2024; Economic Policy 40(121), 13-58, 2025): using task-exposure and task-level productivity estimates, AI adds no more than 0.66% to total factor productivity over 10 years (about 0.06 pp a year), and under 0.53% once harder-to-learn tasks are allowed for; GDP effect about 0.9-1.2% over the decade. He also argues AI is unlikely to reduce inequality between capital and labour. NBER (search extract): https://www.nber.org/papers/w32487
- Reading (OPINION, mine): the surplus question in 2026 has two sides. Capital is being committed (about $700bn+ a year by four firms) on a scale that requires most of the gains to accrue to platform owners for the spend to pay back. Labour's share fell in 2025. The best micro evidence (Denmark) shows workers have so far captured little of the measured time savings as pay. None of this shows AI is causing the labour-share shift yet. For the thesis: if workers expect the gains to go elsewhere (and the early data say they mostly do), "resistance" is a rational response to distribution, not a failure to adapt.

The anti-AI position at its strongest

Copyright: the courts have split the question, and the money has moved anyway.

- **[EVIDENCE]** Bartz v. Anthropic (N.D. Cal., No. 3:24-cv-05417). June 2025: Judge Alsup ruled that training Claude on lawfully acquired books was fair use, but that downloading and keeping more than 7 million pirated books in a permanent "central library" was not. Class settlement of $1.5bn, about $3,000 per work before fees, covering roughly 465,000-500,000 works (sources differ); preliminary approval by Judge Alsup on 25 Sep 2025; final approval by Judge Araceli Martínez-Olguín on 20 Jul 2026, who rejected objections that the price was too low. Because it settled, the fair-use ruling will not reach an appeals court. JURIST, Jul 2026 (search extract): https://www.jurist.org/news/2026/07/judge-approves-record-1-5-billion-settlement-involving-anthropic/ ; Insurance Journal on preliminary approval, 26 Sep 2025: https://amp.insurancejournal.com/news/national/2025/09/26/840753.htm (search extracts). Docket not checked.
- **[EVIDENCE]** Thomson Reuters v. Ross Intelligence: on 29 Sep 2026 the Third Circuit (No. 25-2153, precedential) affirmed that Ross's copying of 2,243 Westlaw headnotes to train a legal-research AI was not fair use, relying on Warhol (2023) and stressing that Ross built a competing product. Reported as the first federal appellate ruling on fair use in AI training. Commentators note it predates generative AI and involves a direct competitor, so it does not settle generative training on public web content. Haynes Boone (search extract): https://www.haynesboone.com/news/alerts/third-circuit-rejects-fair-use-for-rosss-copying-of-westlaw-headnotes ; LawNext, Oct 2026: https://www.lawnext.com/2026/10/no-more-than-an-ordinary-copyright-case-3rd-circuit-unseals-its-opinion-in-thomson-reuters-case-against-ross.html (search extracts).

- The strongest version of the critics' copyright case is now: the one appellate court to rule found no fair use where the AI competes with the source, and the largest trial-court case cost the defendant $1.5bn on acquisition, not training. The strongest version of the industry case: the only generative-AI ruling on training itself (Alsup) found fair use.

Energy:

- **[EVIDENCE + FORECAST]** IEA "Energy and AI" (April 2025): data centres used about 415 TWh in 2024, about 1.5% of world electricity (US 45%, China 25%, Europe 15%). Base case: about 945 TWh by 2030, more than double, with AI the largest driver; about 1,200 TWh by 2035 in the base case and up to 1,720 TWh in a high case. Data centres account for nearly half of US electricity demand growth to 2030. About 20% of planned data-centre projects could be delayed by grid constraints. IEA, Apr 2025 (search extract). https://www.iea.org/reports/energy-and-ai/executive-summary
- **[EVIDENCE + FORECAST]** LBNL "2024 United States Data Center Energy Usage Report" (Shehabi et al., 19 Dec 2024, congressionally requested): US data-centre electricity use more than doubled between 2017 and 2023, driven by GPU servers; by 2028 it could reach 6.7% to 12% of US electricity (74-132 GW in the extract's units, as reported). LBNL (search extract). https://energyanalysis.lbl.gov/publications/2024-lbnl-data-center-energy-usage-report ; full text: https://escholarship.org/uc/item/32d6m0d1 . The commonly cited "4.4% of US electricity in 2023" was not confirmed in the extract; check the report.
- Reading: as a share of global electricity the footprint is small (1.5%); the political force comes from its concentration (US local grids, half of US demand growth) and from the speed of growth, which is where the data-centre siting fights in the opinion data come from.
- **Not sourced in this pass:** water-use estimates; Bender, Gebru et al. (2021); Acemoglu; Marcus; Zitron. The Amazon free-cash-flow figure above is the kind of number Zitron's revenue-vs-capex critique rests on.

The "AI advice" grifter layer

- **[EVIDENCE]** FTC "Operation AI Comply", announced 25 Sep 2024: five actions. DoNotPay ("world's first robot lawyer") agreed to pay $193,000 and notify subscribers about the limits of its service (proposed order; final status not confirmed in the extract). Ascend Ecom (an AI-powered e-commerce "business opportunity"): FTC alleged at least $25m in consumer losses; a federal court temporarily halted it. FBA Machine/Passive Scaling: alleged losses up to $15.9m; temporary restraining order. Ecommerce Empire Builders and Rytr (AI review generator) completed the five. Coverage, Sep 2024 (search extracts): https://siliconangle.com/2024/09/25/ftc-cracks-misleading-ai-powered-business-schemes-new-initiative/ ; https://venturebeat.com/ai/ftc-crackdown-on-donotpay-and-others-is-warning-for-all-ai-powered-companies . FTC press releases (ftc.gov) not retrievable. Final outcomes for Ascend and FBA Machine not confirmed this pass.
- **[EVIDENCE]** DoNotPay final order announced 11 Feb 2025 (Commission vote 5-0): $193,000, notice to 2021-2023 subscribers, and a ban on claiming it performs like a lawyer without evidence. Slashdot summary of the FTC release (search extract): https://tech.slashdot.org/story/25/02/11/1932223/ . Small money: the significance is the precedent on "robot lawyer" claims.
- **[EVIDENCE]** Later FTC AI cases: Click Profit (March 2025), an AI-powered "passive income" Amazon-store scheme; the FTC said the touted AI and brand partnerships did not exist, alleged at least $14m in losses, a court temporarily halted it, and a proposed settlement was announced in August 2025. NBC News, Mar 2025 (search extract): https://www.nbcnews.com/business/business-news/amazon-ai-scammers-duped-investors-millions-passive-income-

scheme-ftc-rcna196931 . Workado: final consent order (3-0) barring unsupported claims that its AI-content detector was "98 percent accurate", no money. Crowell (search extract): https://www.crowell.com/en/insights/client-alerts/ftc-cracks-down-on-ai-models-ai-detection-claims . Air AI (sued 25 Aug 2025, D. Ariz.): the FTC alleged it sold small businesses a "conversational AI" agent claimed to run autonomously and replace staff, which "does not function as advertised"; outcome not found. DLA Piper, Aug 2025 (search extract): https://www.dlapiper.com/en/insights/publications/2025/08/ftcs-latest-ai-washing-case

- Air AI matters most for this research: it is the first FTC case found that targets "agents that replace your employees" claims sold to small businesses, exactly the pitch an agentic-AI vendor makes.
- **[EVIDENCE]** SEC "AI washing": first settlements March 2024 against investment advisers Delphia ($225,000 penalty) and Global Predictions ($175,000) for false claims about using AI. January 14, 2025: Presto Automation, the first against a public company, over its drive-thru voice AI (cease-and-desist, no penalty); reporting says most orders needed human intervention (a 70% figure circulates; not checked against the order). April 9, 2025: SEC and SDNY filed parallel civil and criminal (wire fraud) charges against Albert Saniger, founder of shopping app Nate, which allegedly relied on human workers abroad while claiming AI automation. Ropes & Gray, Jan 2025 (search extract): https://www.ropesgray.com/en/insights/alerts/2025/01/ai-washing-enforcement-in-the-final-days-of-the-gensler-administration ; Weil, Q2 2025 (search extract): https://www.weil.com/-/media/mailings/2025/q2/recent-sec-and-doj-enforcement-actions-and-charges-signal-that-aiwashing-is-firmly-within-the-crossh.pdf
- Reading (OPINION, mine): regulators have so far policed two things: get-rich-quick "AI business" schemes sold to consumers, and companies claiming AI they did not have. Nate and Presto are cases of humans hidden behind an "AI" label, which is the inverse of the thesis's failure mode. No data found on the size of the AI course and advice market.

Viral statistics traced

- **"79% of Americans don't trust businesses to use AI responsibly."** Widely repeated (e.g. tech.co). It comes from an earlier Bentley-Gallup wave (2023, when trust was 21%), not the 2025 wave (31% trust). The 2026 figure is 27% trust. Anyone citing "79%" in 2026 is also at risk of confusing it with Gallup's separate 2026 figure of 79% expecting fewer jobs. Sources: https://www.bentley.edu/news/do-americans-trust-business-and-ai ; https://news.gallup.com/poll/712751/americans-cool-toward.aspx
- **"Nearly half of Americans say AI does more harm than good."** Decrypt's coverage of the 2026 Gallup release; Gallup's own figures as summarised are 39% (2026) and 31% (2025), with about half saying harm and good are equal. Use Gallup's number. https://news.gallup.com/poll/712751/americans-cool-toward.aspx
- **"Half of Americans are more concerned than excited."** Solid: Pew, probability panel, repeated since 2021, now 52%. Note it is about AI "in daily life", not about AI at work.
- **"EU fined companies EUR 18.4m under the AI Act in February 2026."** Appears on a low-quality web page; no official source found; timeline makes it implausible. Do not cite.
- **"Colorado's AI Act takes effect June 30, 2026."** Still repeated in compliance content. It was repealed and replaced by SB 26-189 in May 2026; the replacement takes effect 1 Jan 2027. https://www.ropesgray.com/en/insights/alerts/2026/05/colorado-scales-back-ai-law-with-targeted-implications-for-health-care (search extract)
- **"Anthropic paid $1.5bn for training on books."** Misleading. The trial court held training on lawfully acquired books was fair use; the liability exposure came from downloading and

storing pirated copies. https://www.jurist.org/news/2026/07/judge-approves-record-1-5-billion-settlement-involving-anthropic/ (search extract)

- **"Data centres use 4.4% of US electricity."** Widely attributed to LBNL's 2024 report; search extracts confirmed only the report's 6.7-12% projection for 2028 and the "more than doubled 2017-2023" finding. Check the report before citing 4.4%. https://escholarship.org/uc/item/32d6m0d1
- **"UK to introduce AI bill in 2026."** Two low-quality sites claim a frontier AI bill was introduced on 14 May 2026; law-firm sources say the May 2026 King's Speech contained none. Do not cite without bills.parliament.uk. https://www.lewissilkin.com/insights/2026/09/14/ai-judgment-day-on-the-horizon-while-uk-lawmakers-play-catch-up (search extract)

Strongest counter-evidence

Against the thesis ("AI fails because the business isn't ready; people resist a new way of working"):

- **Resistance is mostly about distribution and job security, not about learning a new way of working.** Rising concern tracks expected job loss (Pew 71%, Gallup 79%) and distrust of employers' motives (Gallup 27% trust; NNU 60% distrust employer's priorities). People are using AI at high rates while distrusting it (KPMG: 66% use, 46% trust). The barrier is "whose interest does this serve", which no amount of shared context fixes.
- **The CBA case is a technology and measurement failure.** The bank claimed a 2,000-calls-a-week reduction that its own staffing (overtime, team leaders on phones) contradicted. The business was "ready" enough to deploy; the bot did not do what was claimed. Unions caught it, not change management.
- **Workers in the Kaiser and NNU cases object to AI decisions they cannot override.** That is an argument about authority and accountability in the workflow, not about undocumented processes.
- **Distribution evidence makes resistance rational.** Labour share fell in 2025 (BLS, 54.4% in Q4) and the Danish register study finds chatbot use has not raised earnings or hours, with little of the time saved reaching pay. If the gains go to capital (hyperscaler capex of about $700bn+ in 2026 needs them to), workers' reluctance is a reasonable reading of incentives, not a readiness problem.
- **Some AI failures are vendor over-claiming, not buyer unreadiness.** The FTC's Air AI complaint (Aug 2025) alleges an agent sold to small businesses as able to replace staff "does not function as advertised". The SEC's Presto and Nate cases involved "AI" that was largely humans. When the product does not do what was sold, no amount of process documentation fixes it.

Against the mainstream reading of this track:

- **Europeans are more positive about AI at work than headline concern suggests** (Eurobarometer: >60% positive, >70% expect productivity gains). Concern about AI "in daily life" and acceptance of AI "at work" are different questions, and much of the "public hates AI" narrative rests on the former.
- **Asia is positive.** South Korea, India, South-East Asia and China report low concern and high trust. "Public backlash" is mostly an Anglosphere and southern European phenomenon in this data.
- **Unions are bargaining, not blocking.** Every major entertainment dispute ended in a ratified contract that permits AI with consent, disclosure and pay. The 2026 SAG-AFTRA deal reportedly gives up AI strike rights until 2030.

- **Regulation is not the binding constraint in 2026.** The EU delayed its heaviest obligations to 2027-2028 and has no known fines; the US executive branch is actively trying to weaken state rules. Colorado repealed its own flagship law before it took effect, Texas's TRAIGA was cut down before passage, the UK has no AI bill, and the one long-running US AI employment law (NYC LL144) was found poorly enforced. Litigation (copyright) and consumer-protection enforcement (FTC) have bitten harder than AI-specific statutes.
- **The energy critique is real but often overstated globally.** Data centres were about 1.5% of world electricity in 2024 (IEA). The force of the critique is local (US grids: nearly half of US demand growth to 2030), not global.

Open questions

Items still to source after the gap-fill pass:

1. UK: confirm on bills.parliament.uk that no AI bill has been introduced; confirm the Action Plan's recommendation count; ONS labour-share data.
2. US federal: whether the Commerce "onerous laws" list was published; FTC pre-emption policy statement; other Task Force suits; fate of xAI/DOJ v. Colorado after the repeal; any 2026 pre-emption bill.
3. US states: any enforcement action under SB 53, TRAIGA or Illinois HB 3773 (none found); whether DCWP's NYC LL144 enforcement changed after the audit; the audit PDF's own numbers.
4. Surplus capture: Q3 2026 capex from filings; whether 54.4% is a series low; Acemoglu and Johnson "Power and Progress" as OPINION.
5. Anti-AI case: NYT v. OpenAI and Getty v. Stability status; water-use estimates; Bender, Gebru et al. 2021; Marcus; Zitron. Label argument as OPINION.
6. Grifter layer: final outcomes for Ascend Ecom, FBA Machine, Air AI; any data on the size of the AI course and consultancy market.
7. Whether the 2026 WGA contract actually changed the AI terms (contract text vs IndieWire's "largely unchanged").
8. Outcome of the Kaiser NUHW dispute after April 2026.

What's genuinely uncertain

- Whether the 2026 souring in US opinion is durable or a cycle tied to layoff news, data-centre fights and the election year. Gallup's harm figure went 40% (2023), 31% (2025), 39% (2026): it has reversed before.
- Whether concern changes behaviour. Use keeps rising while trust falls. There is no good evidence linking public concern to slower enterprise adoption.
- Whether union contract language bites. None of the entertainment AI clauses has, to my knowledge, been tested in arbitration (not verified this session).
- Whether the US pre-emption push survives the courts and Congress, and whether the EU's delay becomes a pattern of further delay.
- Exact numbers in this file: none was checked against the primary page text in this session.
- Whether copyright liability for AI training will turn on acquisition (Bartz) or on market substitution (Ross). The two leading decisions point different ways and neither concerns a general-purpose model on appeal.

- Whether the 2025 fall in US labour share is cyclical or the start of an AI-driven shift. One year of data cannot tell.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
52%	US adults more concerned than excited about AI in daily life (37% in 2021)	Pew, n=3,488, Jun 22-28 2026	https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/	18 Aug 2026	Probability panel; seen via search summary only
71%	US adults expecting AI to mean fewer jobs over 20 years (64% in 2024)	Pew	same as above	18 Aug 2026	As above
55%	US adults under 30 more concerned than excited	Pew	same as above	18 Aug 2026	As above
+17 pts	Rise in Democrats expecting fewer jobs from AI	Pew	https://www.pewresearch.org/short-reads/2026/09/16/democrats-are-now-more-worried-than-republicans-about-ai-and-its-impact-on-jobs/	16 Sep 2026	Search summary only; base years not confirmed
34% / 42% / 16%	25-country median: more concerned / equally / more excited	Pew Global Attitudes, spring 2025	https://www.pewresearch.org/global/2025/10/15/how-people-around-the-world-view-ai/	15 Oct 2025	Strong method; search summary only
50% vs 16%	More concerned than excited: US (and Italy) vs South Korea	Pew Global Attitudes	same as above	15 Oct 2025	As above
39%	US adults saying AI does more harm than good (31% 2025, 40% 2023)	Gallup/ Bentley	https://news.gallup.com/poll/712751/americans-cool-toward.aspx	Jul 2026	Gallup web panel; Decrypt reports a conflicting higher figure
79%	US adults expecting AI to reduce US jobs over 10 years (73% 2025)	Gallup/ Bentley	same as above	Jul 2026	As above
27%	US adults trusting business to use AI responsibly (31% 2025, 21% 2023)	Gallup/ Bentley	same as above	Jul 2026	As above
66% / 46%	Regular AI use / willing to trust AI, 47 countries	KPMG & Univ. of Melbourne, n>48,000	https://figshare.unimelb.edu.au/articles/report/Trust_attitudes_and_use_of_artificial_intelligence_A_global_study_2025/28822919	2025 (fieldwork Nov 2024 to Jan 2025)	Large multi-country panel; consultancy co-author
52% / 53%	Excited / nervous about AI products, 30-country average	Ipsos AI Monitor 2025	https://www.ipsos.com/en-dk/ipsos-ai-monitor-2025	Jun 2025	Online panel; less representative in emerging markets
72% vs 32%	Trust in AI, China vs US	Edelman 2025 tech report	https://www.edelman.com/sites/g/files/aatuss191/files/fr/2025-08/Top%20Findings_2025%20EdelmanTrust%20Barometer_Insights%20for%20the%20Technology%20Sector%20%281%29.pdf	Aug 2025	PR-firm survey
>60% / 84%	Europeans positive on AI at work / want careful management	Eurobarometer via Commission	https://digital-strategy.ec.europa.eu/en/news/commission-survey-shows-most-europeans-support-use-artificial-intelligence-workplace	13 Feb 2025	Commission press framing
95.04%	SAG-AFTRA video game agreement ratification	PC Gamer	https://www.pcgamer.com/gaming-industry/the-videogame-voice-actors-strike-officially-ends-as-	Jul 2025	Union-reported tally

Figure	What it measures	Source	URL	Date	Reliability note
	vote, ending strike begun Jul 2024		performers-ratify-agreement-between-sag-aftra-and-publishers		
91.42% (19.25% turnout)	SAG-AFTRA 2026 TV/ Theatrical ratification	Fortune	https://fortune.com/2026/06/05/sag-aftra-approves-four-year-contract-studios-streamers/	5 Jun 2026	Union-reported tally
45 roles	CBA customer-service jobs cut for voice bot, then reinstated	Bloomberg; FSU	https://www.bloomberg.com/news/articles/2025-08-21/commonwealth-bank-reverses-job-cuts-decision-over-ai-chatbots	21 Aug 2025	Volume claims disputed between bank and union
~2,400	Kaiser NUHW mental-health workers in one-day AI-related strike	KQED; NUHW	https://www.kqed.org/news/12076753/northern-california-kaiser-therapists-hold-1-day-strike-over-ai-patient-care-concerns	18 Mar 2026	Some outlets say "more than 2,000"
60% / 69%	Nurses distrusting employer AI priorities / acuity scores mismatching own judgment	NNU survey, n>2,300	https://www.nationalnursesunited.org/press/national-nurses-united-survey-finds-ai-technology-undermines-patient-safety	May 2024	Union member survey, self-selected
2 Dec 2027 / 2 Aug 2028	New EU AI Act high-risk dates (standalone / embedded) after Digital Omnibus	Sidley; National Law Review	https://datamatters.sidley.com/2026/06/22/eu-lawmakers-reach-provisional-agreement-to-delay-key-eu-ai-act-obligations/	Jun-Jul 2026	Law-firm summaries; confirm on EUR-Lex
423-57-174	European Parliament vote on AI omnibus deal	Industry summary	https://www.europeanoutdoorgroup.com/members-knowledge-hub/european-parliament-adopts-provisional-agreement-on-ai-omnibus	16 Jun 2026	Secondary source; confirm on EP site
0	Known AI Act fines as of Nov 2026	Negative search result	n/a	Nov 2026	Absence of evidence; recheck
99-1	Senate vote to strip state-AI-law moratorium from budget bill	Nextgov/ Route Fifty	https://www.nextgov.com/artificial-intelligence/2025/07/senate-overwhelmingly-passes-amendment-removing-state-ai-moratorium/406448	1 Jul 2025	Search extract; roll call not checked on senate.gov
1 Jan 2027	Effective date of Colorado's replacement AI law (SB 26-189); original act never took effect	Ropes & Gray	https://www.ropesgray.com/en/insights/alerts/2026/05/colorado-scales-back-ai-law-with-targeted-implications-for-health-care	May 2026	Search extract; signing date 14 May (one source 18 May)
$1,500	Max NYC LL144 penalty per violation per day; enforcement found ineffective	NY State Comptroller; DLA Piper	https://www.osc.ny.gov/state-agencies/audits/2025/12/02/enforcement-local-law-144-automated-employment-decision-tools	2 Dec 2025	Search extract; audit PDF not read
$1m	Max SB 53 civil penalty per violation (large frontier developers, >$500m revenue)	FPF	https://fpf.org/blog/californias-sb-53-the-first-frontier-ai-law-explained/	Effective 1 Jan 2026	Search extract
54.4%	US nonfarm business labour share, Q4 2025 (down 1.7% y/y); 2025 productivity +2.2%	BLS Productivity and Costs	https://www.bls.gov/news.release/archives/prod2_03242026.htm	24 Mar 2026	Official statistic, search extract; subject to revision
~$700-745bn	Combined 2026 capex guidance, Microsoft, Alphabet, Amazon, Meta	Yahoo Finance; datacentres.com tracker	https://finance.yahoo.com/markets/article/magnificent-7-earnings-rush-reveals-ai-spending-surge-with-hyperscaler-capex-set-to-reach-725-billion-in-2026-224901707.html	Apr-Oct 2026	Company guidance via aggregators; definitions differ; not from filings
0 (CI rules out large effects)	Effect of AI chatbot adoption on earnings and hours, Denmark, 25,000 workers	Humlum & Vestergaard, NBER WP 33777	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	2025	Strong design; early period (2023-24); search extract
<=0.66 %			https://www.nber.org/papers/w32487	2024/2025	Academic forecast

Figure	What it measures	Source	URL	Date	Reliability note
	Upper bound on 10-year TFP gain from AI (0.53% adjusted)	Acemoglu, Economic Policy			
$1.5bn (~$3,000/work)	Bartz v. Anthropic class settlement, final approval 20 Jul 2026	JURIST	https://www.jurist.org/news/2026/07/judge-approves-record-1-5-billion-settlement-involving-anthropic/	Jul 2026	Search extract; docket not checked
2,243	Westlaw headnotes copied by Ross; Third Circuit: not fair use	Haynes Boone; LawNext	https://www.haynesboone.com/news/alerts/third-circuit-rejects-fair-use-for-rosss-copying-of-westlaw-headnotes	29 Sep 2026	Search extract; opinion not read
415 TWh (1.5%) to ~945 TWh	Global data-centre electricity 2024 to 2030 (base case)	IEA, Energy and AI	https://www.iea.org/reports/energy-and-ai/executive-summary	Apr 2025	Agency projection; search extract
6.7-12%	Projected US data-centre share of electricity by 2028	LBNL (Shehabi et al.)	https://energyanalysis.lbl.gov/publications/2024-lbnl-data-center-energy-usage-report	19 Dec 2024	Range reflects scenario uncertainty
$193,000	DoNotPay FTC final order (5-0)	FTC via Slashdot	https://tech.slashdot.org/story/25/02/11/1932223/	11 Feb 2025	Search extract
$25m / $15.9m / $14m	Alleged consumer losses: Ascend Ecom / FBA Machine / Click Profit	FTC complaints via press	https://siliconangle.com/2024/09/25/ftc-cracks-misleading-ai-powered-business-schemes-new-initiative/	2024-2025	Allegations, not findings
$225k + $175k	SEC penalties, Delphia and Global Predictions (first AI-washing cases)	Baker McKenzie; Ropes & Gray	https://www.ropesgray.com/en/insights/alerts/2025/01/ai-washing-enforcement-in-the-final-days-of-the-gensler-administration	Mar 2024	Search extract

Verbatim quotes

None. No source page could be opened in this session, so no wording could be confirmed. Quotes seen only in search snippets (for example Florida Governor DeSantis saying, in substance, that an executive order cannot pre-empt state legislation) are paraphrased in the findings above.

Sources worth reading in full

1. Pew, "How the U.S. Public and AI Experts View Artificial Intelligence" (Apr 2025): the cleanest data on the gap between builders and everyone else. https://www.pewresearch.org/internet/2025/04/03/how-the-us-public-and-ai-experts-view-artificial-intelligence/
2. Pew, "How People Around the World View AI" (Oct 2025): 25-country comparison with demographics and trust-in-regulator data. https://www.pewresearch.org/global/2025/10/15/how-people-around-the-world-view-ai/
3. KPMG and University of Melbourne, "Trust, attitudes and use of AI: A global study 2025": the use-vs-trust gap, including workplace questions, 47 countries. https://figshare.unimelb.edu.au/articles/report/Trust_attitudes_and_use_of_artificial_intelligence_A_global_study_2025/28822919
4. Gallup/Bentley, "Americans Cool Toward AI" (Jul 2026): the trend break in 2026 with trust-in-business series since 2023. https://news.gallup.com/poll/712751/americans-cool-toward.aspx
5. Pew, young adults and AI (Aug 2026) and partisan views (Sep 2026): where the shift is coming from. https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/
6. WGA 2023 MBA summary: the template most other AI bargaining copies. https://wgacontract2023.org/the-campaign/summary-of-the-2023-wga-mba

7. Finance Sector Union on the CBA reversal: a rare documented case of an AI redundancy reversed on the facts. https://www.fsunion.org.au/win-cba-backflips-on-customer-service-job-cuts-admits-they-got-it-wrong/
8. Gibson Dunn on the AI Act omnibus: what was and was not delayed. https://www.gibsondunn.com/eu-ai-act-omnibus-agreement-postponed-high-risk-deadlines-and-other-key-changes/
9. Ropes & Gray on the December 2025 pre-emption executive order: mechanics and legal weak points. https://www.ropesgray.com/en/insights/alerts/2025/12/trump-attempts-to-preempt-state-ai-regulation-through-executive-order
10. TUC draft AI bill: the most concrete worker-side legislative text in the UK. https://www.tuc.org.uk/research-analysis/reports/artificial-intelligence-regulation-and-employment-rights-bill
11. NY State Comptroller, audit of DCWP enforcement of Local Law 144 (Dec 2025): the only official audit of a US AI employment law in operation. https://www.osc.ny.gov/state-agencies/audits/2025/12/02/enforcement-local-law-144-automated-employment-decision-tools
12. Humlum and Vestergaard, "Large Language Models, Small Labor Market Effects" (NBER 2025) and the 2026 update: the strongest register-data evidence on wages and hours. https://econpapers.repec.org/RePEc:nbr:nberwo:33777
13. IEA, "Energy and AI" (Apr 2025): the reference numbers for the energy debate. https://www.iea.org/reports/energy-and-ai/executive-summary
14. Third Circuit opinion, Thomson Reuters v. Ross (Sep 2026), via LawNext: first appellate ruling on AI training and fair use. https://www.lawnext.com/2026/10/no-more-than-an-ordinary-copyright-case-3rd-circuit-unseals-its-opinion-in-thomson-reuters-case-against-ross.html
15. FTC complaint against Air AI (Aug 2025), via DLA Piper: what regulators treat as deceptive in agent marketing to small businesses. https://www.dlapiper.com/en/insights/publications/2025/08/ftcs-latest-ai-washing-case
16. Ropes & Gray on Colorado's repeal and replacement (May 2026): the clearest case of a state AI law announced and then withdrawn. https://www.ropesgray.com/en/insights/alerts/2026/05/colorado-scales-back-ai-law-with-targeted-implications-for-health-care

INCOMPLETE (narrower after the 2026-11-06 gap-fill pass). Sections (a) and (b) are sourced. Sections (c) "task change not job loss" and (d) "freedom to be more human backfires" now have sourced findings (radiology, translators, bank tellers, Autor; deskilling, intensification, algorithmic management, social penalty), all from search-engine extracts with no page fetched. Still missing: primary-source checks of BLS teller and radiologist counts and NRMP residency numbers; the Hui, Reshef and Zhou freelancer earnings paper; Autor and Thompson's 2025 "Expertise" paper; newer agent benchmark scores; Census BTOS firm AI-use rates; any verbatim quotes.

Track 10: Contrarian

Last updated: 2026-11-06. Labels: EVIDENCE = measured/primary data; FORECAST = projection; OPINION = argument, vendor or consultancy view.

Verification note (added 2026-11-06): page fetches were blocked for most hosts this session, so findings marked search summary, search extract or snippet rest on search-engine extracts of the cited URL. Quotation marks inside findings mark wording as reported by that source; only quotes in this file's "Verbatim quotes" section (and in 13-quotes.md) were checked against a fetched page. Corrections from the verification pass are in verification.md.

Access note. In this session every direct page fetch failed (curl and WebFetch both refused by the network proxy: nber.org, arxiv.org, metr.org, sequoiacap.com, budgetlab.yale.edu, ideas.repec.org, bls.gov, census.gov and others). Every finding below comes from web-search results that showed the source URL with an excerpt or summary. No page was read in full. Figures are therefore "seen in search excerpt, not re-read at source" and should be re-verified at the URL before publication. No verbatim quotes are given, because none could be checked against a fetched page.

Summary

- **The macro numbers do not yet show an AI boom.** US private nonfarm total factor productivity grew 0.8% in 2025, down from 1.5% in 2024, and labour productivity slowed from 3.0% to 2.2%, in the first full year of mass generative-AI use (BLS, 19 March 2026). UK output per hour fell 0.8% in Q2 2026. The fair counter: US labour productivity has run at about 2.1% a year since late 2019, which is respectable, and the productivity J-curve predicts measured gains lag heavy intangible investment.
- **The best micro evidence on ordinary workers shows small gains.** Self-reported time savings are about 1.4% of all US work hours (Bick, Blandin, Deming). In Denmark, chatbot adoption across about 25,000 workers in 11 exposed occupations produced precise null effects on earnings and recorded hours, with average time savings around 3% and only 3-7% of the benefit passed through to pay (Humlum and Vestergaard, NBER, May 2025).
- **The serious top-down macro forecast is modest.** Acemoglu's task-based estimate is at most 0.66% higher TFP over ten years, and below 0.53% once harder tasks are accounted for (Economic Policy, 2025). This is a model, not a measurement, and it excludes new tasks.
- **The spending-to-revenue gap is the strongest "less profitable" argument.** Hyperscaler capex for 2026 is reported at roughly $800bn; Bain (Sept 2025) says $2tn of annual revenue is needed by 2030 to fund the compute trend and projects an $800bn shortfall even after counting AI savings. Markets have had several sharp AI sell-offs in 2026 (February, late April, mid-October), but none has yet reduced planned capex: executives still guide higher for 2027.

- **Agents fail most where business work actually happens: multi-turn, policy-bound, repeated.** Salesforce's own CRMArena-Pro: about 58% success single-turn falling to about 35% multi-turn, with near-zero confidentiality awareness. On tau-bench, GPT-4o's retail success above 60% on one try fell to about 25% when it had to succeed 8 times in a row. TheAgentCompany: the best agent fully completed about 30% of simulated office tasks.
- **The headline "developers are 19% slower with AI" is real but narrow and now dated.** n=16 experienced open-source developers, early-2025 tools. METR's own 2026 redesign says its newer data cannot give a reliable estimate, partly because developers now refuse to work without AI. It has not been reversed; it has become unmeasurable by that design.
- **Gartner's "over 40% of agentic AI projects cancelled by end-2027" and "only about 130 genuine agentic vendors" are a forecast and an opinion,** but they cut against vendor hype from inside the analyst mainstream.
- **"Task change, not job loss" holds for radiology and salaried work so far, but not for freelance translation.** Practising US radiologists rose about 12% from 2010 to 2022 and pay rose, nine years after Hinton's "stop training radiologists". But on a large freelance platform, writing and coding posts fell about 21% within eight months of ChatGPT (Demirci, Hannane, Zhu, Management Science 2025), and most freelance translators reported price falls for the first time in 2025 (ELIS survey). The teller story shows task change can last two decades and then turn into job loss.
- **The "freedom to be more human" evidence is early and mixed, with real warning signs.** Endoscopists' unaided adenoma detection fell from 28.4% to 22.4% after routine AI use (Lancet Gastroenterology & Hepatology, Oct 2025; observational, 19 doctors, confounded). AI users are judged lazier and less competent (PNAS 2025, N=4,439). Workload surveys point to intensification, but the best register data show no rise in recorded hours.
- **For the thesis:** the contrarian evidence partly supports "the tech is the easy part" (adoption friction, unmeasured reorganisation) but partly contradicts it. Agent reliability on multi-turn, permission-sensitive tasks is a technology limit, not only a readiness problem. A business with perfectly documented workflows would still see a 35% multi-turn success rate on CRMArena-Pro.

Key findings

(a) Transformation slower, narrower or less profitable than claimed

Macro forecasts

- **[FORECAST]** Acemoglu estimates AI raises US TFP by no more than 0.66% over ten years, using task-exposure shares and early experimental productivity estimates aggregated via a version of Hulten's theorem. He then argues this is likely too high because early evidence comes from easy-to-learn tasks, cutting the estimate to below 0.53%. He also predicts AI widens the gap between capital and labour income. Acemoglu, "The Simple Macroeconomics of AI", NBER w32487 (May 2024); Economic Policy 40(121): 13-58 (2025). https://www.nber.org/papers/w32487 https://econpapers.repec.org/RePEc:oup:ecpoli:v:40:y:2025:i:121:p:13-58. https://dspace.mit.edu/handle/1721.1/166327
- **[OPINION]** Counter: critics argue the model excludes new tasks, new products and complementarity channels, so it is a floor for "cost savings on existing tasks" rather than a forecast of total effect. Re-expressed by one critic as about 0.064 percentage points of extra TFP growth a year. "Contra Acemoglu on AI", Effective Altruism Forum (2024). https://forum.effectivealtruism.org/posts/Qb8NdS88TKsXApLsa/contra-acemoglu-on-ai

Official productivity statistics

- **[EVIDENCE]** US private nonfarm business TFP rose 0.8% in 2025 (output +2.6%, combined inputs +1.7%), down from 1.5% in 2024. Labour productivity rose 2.2% in 2025 vs 3.0% in 2024. BLS Total Factor Productivity release, 19 March 2026. https://www.bls.gov/news.release/archives/prod3_03192026.pdf https://www.bls.gov/news.release/prod3.nr0.htm
- **[EVIDENCE]** US nonfarm business labour productivity rose 1.4% (annualised) in Q2 2026, against a 0.6% consensus; 2.2% year on year; 2.1% a year from Q4 2019 to Q2 2026. BLS Productivity and Costs, 6 Aug 2026 (revised 3 Sept 2026); BLS Economics Daily. https://www.bls.gov/news.release/archives/prod2_08062026.htm https://www.bls.gov/opub/ted/2026/productivity-up-2-2-percent-from-second-quarter-2025-to-second-quarter-2026.htm
- **[EVIDENCE]** BLS states AI enters its productivity measures only implicitly, via the capital measure of software used in production. So there is no official "AI contribution" series to point to. Seen in search excerpt of BLS productivity pages, 2026. https://www.bls.gov/lpc/lpc.htm
- **[EVIDENCE]** UK output per hour fell 0.8% quarter on quarter in Q2 2026 (expectation -0.5%), after +1.1% in Q1; +0.7% year on year. Two official measures disagree sharply: the tax-record (RTI) measure showed +2.1% year on year in Q1 2026 vs +0.4% on the Labour Force Survey measure. ONS flash estimate, 18 Aug 2026 (next release 17 Nov 2026), via Trading Economics and The Productivity Institute. https://www.ons.gov.uk/releases/productivityflashestimateandoverviewukapriltojune2026andjanuarytomarch2026 https://tradingeconomics.com/united-kingdom/productivity/news/576075 https://lab.productivity.ac.uk/insights/q2-2026-uk-productivity-growth/
- **[EVIDENCE, model-based]** Counter (the J-curve): Brynjolfsson, Rock and Syverson show that general-purpose technologies require large unmeasured intangible investment (process redesign, training, data), so measured TFP first dips and later overshoots. Adjusting for intangibles tied to computer hardware and software gives a TFP level 15.9% above official measures by end-2017. On AI specifically they found effects "small but growing" (search-summary wording). AEJ: Macroeconomics 13(1): 333-72 (2021); NBER w25148 (2018). https://www.aeaweb.org/doi/10.1257/mac.20180386 https://digitaleconomy.stanford.edu/publication/the-productivity-j-curve-how-intangibles-complement-general-purpose-technologies
- How to weigh it: the J-curve is the strongest argument that weak 2025 TFP is not evidence against AI. But it also cuts against fast-transformation claims: it says the payoff waits on organisational reinvention, which is slow. Both readings agree that 2026 statistics do not show the transformation yet.

Worker-level time savings and earnings

- **[EVIDENCE]** Generative AI saves time equal to about 1.4% of total US work hours (averaged across all workers, users and non-users). Between 1% and 5% of all work hours were assisted by generative AI in the February 2025 NBER version; 1-7% in the October 2025 St. Louis Fed revision. Method: Real-Time Population Survey, nationally representative, August and November 2024 waves, more than 10,000 respondents; self-reported. Bick, Blandin and Deming, "The Rapid Adoption of Generative AI", NBER w32966 (rev. Feb 2025); St. Louis Fed WP 2024-027 (rev. Oct 2025). A secondary report says it is published in Management Science (2026); not verified. https://www.nber.org/system/files/working_papers/w32966/w32966.pdf https://fedinprint.org/item/fedlwp/98805/102025 https://insurancenewsnet.com/oarticle/studies-from-federal-reserve-bank-update-current-data-on-management-science-the-rapid-adoption-of-generative-ai-business-management-science
- **[EVIDENCE, self-reported]** Counter, same team: a third of daily users reported saving at least four hours a week. Bick himself cautions that using self-reported time savings to

compute productivity could overestimate gains. St. Louis Fed Open Vault interview, Oct 2025. https://www.stlouisfed.org/open-vault/2025/oct/generative-ai-productivity-future-work

- **[EVIDENCE]** Danish administrative data linked to two surveys (late 2023 and 2024) of about 25,000 workers in 7,000 workplaces across 11 AI-exposed occupations: precise null effects of chatbot adoption on earnings and recorded hours in every occupation. Average reported time savings about 3% (2.8% of work hours in press coverage). Only 3-7% of productivity gains passed through to earnings. Workers report new tasks created, such as reviewing AI output and writing prompts. Humlum and Vestergaard, "Large Language Models, Small Labor Market Effects", NBER w33777 (May 2025); BFI WP 2025-56. https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf https://econpapers.repec.org/RePEc:nbr:nberwo:33777 https://flowingdata.com/2025/05/09/time-saved-with-ai-creates-more-work-elsewhere/
- **[EVIDENCE]** A later paper by the same authors (search results give the title "Still Waters, Rapid Currents", RFBerlin/CReAM working paper 26078, March 2026) reports precise nulls on earnings and hours that rule out effects larger than 2% two years after ChatGPT's launch. It says employers absorb AI through task reorganisation, and that adopters move into higher-paying occupations where chatbots are used more. Title and venue come from a search summary and could not be checked at source. https://ideas.repec.org/p/crm/wpaper/26078.html
- Fair counter: the Danish data cover the 2023-2024 chatbot era, before agents. Occupational mobility toward better-paid jobs is a benefit that a "no earnings effect within occupation" headline hides. Denmark has strong labour institutions, so pass-through may differ elsewhere.

Spending vs revenue

- **[OPINION]** Goldman Sachs Global Macro Research, "Gen AI: too much spend, too little benefit?" (Top of Mind, Issue 129, 25 June 2024): more than $1tn of AI capex expected in coming years with little to show so far. Jim Covello (head of global equity research) argues the technology is too expensive to justify the cost and asks what trillion-dollar problem it will solve; Acemoglu sees limited US upside over a decade. The same report includes a bullish counterview (software analyst Kash Rangan: today's AI firms are building foundations of a profitable industry). Exact wording not verified. https://www.goldmansachs.com/insights/goldman-sachs-exchanges/a-skeptical-look-at-ai-investment https://www.bespacific.com/?p=106220 https://research.contrary.com/report/too-much-spend-too-little-benefit
- **[OPINION]** Sequoia's David Cahn, "AI's $600B Question" (20 June 2024), following "AI's $200B Question" (Sept 2023): take Nvidia's data-centre run-rate revenue, double it for full data-centre cost, double again for end-user margin; the industry needs about $600bn of annual revenue to justify capex, leaving a gap of about $500bn. This is a back-of-envelope heuristic. One secondary source reports a later Cahn update at about $840bn; single source, unverified. https://www.sequoiacap.com/article/ais-600b-question/ https://www.tomshardware.com/tech-industry/artificial-intelligence/ai-industry-needs-to-earn-dollar600-billion-per-year-to-pay-for-massive-hardware-spend-fears-of-an-ai-bubble-intensify-in-wake-of-sequoia-report
- Counter to Cahn: the method counts only product revenue, not internal cost savings or productivity inside the hyperscalers themselves (critique quoted in the Tom's Hardware coverage above).
- **[FORECAST]** Bain Global Technology Report, 23 Sept 2025: meeting projected AI compute demand by 2030 needs about $2tn of annual revenue and about $500bn of annual capex; even if firms reinvested all AI-driven savings from sales, marketing, support and R&D, revenue would fall about $800bn short. Projects 200 GW of incremental global AI compute by 2030, half in the US. https://www.bain.com/ja/about-bain/media-center/press-releases/

20252/$2-trillion-in-new-revenue-needed-to-fund-ais-scaling-trend---bain--companys-6th-annual-global-technology-report/ https://www.theregister.com/2025/09/24/bain_ai_costs https://www.pymnts.com/news/artificial-intelligence/2025/ai-companies-face-800-billion-dollar-funding-shortfall-bain

- Counter: Bain's own tech lead publicly argued "the AI bubble won't pop" (City A.M.); the shortfall is a scenario conditional on the compute trend continuing, which may itself bend. https://www.cityam.com/?p=2370241
- **[EVIDENCE, secondary, not checked against filings]** After the late-October 2026 earnings round, combined 2026 calendar-year capex for the largest spenders is reported at $803-828bn, with executives guiding higher for 2027. Company figures in coverage: Amazon about $220bn cash capex (maintained); Microsoft about $170bn calendar-year, cut from a prior outlook because third-party data-centre capacity was delayed; Meta $135-145bn (shares fell 9.4% after); Alphabet $195-205bn (July guidance). Hargreaves Lansdown, Oct 2026; NextMSC, Oct 2026. The HL list oddly includes SpaceX, so its aggregate definition is unclear. https://www.hl.co.uk/news/ai-hyperscalers-are-expected-to-finish-the-year-spending-a-whopping-828-billion-on-capex https://www.nextmsc.com/industry-deep-dive/hyperscaler-capex-comparison-2026
- **[OPINION, weak]** An independent wiki estimate puts 2026 hyperscaler AI infrastructure spending around $700bn against direct AI service revenue of roughly $25-50bn, a 6-14x mismatch. Weak source; method not seen. https://www.longtermwiki.com/wiki/E513

Any 2026 investment correction?

- **[EVIDENCE, news]** Broad sell-off in AI-linked stocks in early February 2026 (Amazon, Oracle, chipmakers, data-centre REITs) over spending and disruption fears. NPR via WBAA, 6 Feb 2026. https://www.wbaa.org/2026-02-06/tech-stocks-plunge-as-investors-fret-over-ai-spending-and-disruption
- **[EVIDENCE, news, weak]** A leaked report on OpenAI internal revenue reportedly triggered a drop in AI infrastructure stocks on 28 April 2026. Single low-quality source. https://www.heygotrade.com/en/blog/ai-capex-risk-openai-revenue-report/
- **[EVIDENCE]** In mid-October 2026 investors turned on how AI data centres are financed, much of it with debt; the 20 most valuable AI and chip stocks lost more than $1tn in one week. Dealroom analysis, 30 Oct 2026. https://dealroom.co/charts/ai-selloff
- Counter: earlier sell-offs reversed quickly once hyperscalers reaffirmed spending (MetaTrader analysis), and planned 2027 capex is still rising. So far 2026 shows price corrections, not an investment correction. Microsoft's cut was attributed to capacity delays, not weaker demand. https://www.metatrader.com/fr/news/1732585-deja-vu-ai-overspending-fears-renew

(b) Agents further from production than hyped

METR developer RCT

- **[EVIDENCE]** Randomised trial, 16 experienced open-source developers, 246 real issues in their own mature repositories, each task randomly allowed or disallowed AI (mainly Cursor Pro with Claude 3.5/3.7 Sonnet, early 2025). With AI, tasks took 19% longer. Afterwards developers believed AI had sped them up by about 20%. METR, arXiv 2507.09089, July 2025. https://arxiv.org/abs/2507.09089
- **[EVIDENCE, secondary reports of primary]** METR's 24 Feb 2026 post ("We are changing our developer productivity experiment design"): a larger follow-up started August 2025 with newer tools gave an unreliable signal. Reported point estimate for returning developers is roughly an 18% speedup with a wide interval spanning slowdown and speedup. Main

problems: selection effects, developers increasingly unwilling to work without AI (so the hardest-to-replace AI tasks drop out), and difficulty measuring parallel agent use. metr.org was blocked; figures come from secondary write-ups and should be checked. https://birchtree.me/blog/an-update-from-the-study-that-said-devs-were-actually-slower-with-coding-agents/ https://blog.robbowley.net/2026/04/04/metrs-developer-productivity-research-2026-update/ https://particula.tech/blog/metr-reversed-19-percent-slower-ai-coding-study

- What survives: the perception gap (people feel faster than they are) is the more durable lesson than the 19% figure. It means self-reported time savings, the basis of Bick et al. and most vendor surveys, deserve a discount.

METR time horizons

- **[EVIDENCE, secondary]** METR's task-length measure: the length of task (in human expert time) a model completes with 50% success has doubled about every 7 months over 2019-2025; later estimates after 2023 are faster (Wikipedia reports about 131 days). Time Horizon 1.1 (Jan 2026) expanded the suite from 170 to 228 tasks and long tasks from 14 to 31, with few measured human baselines for long tasks. https://en.wikipedia.org/wiki/METR
- **[EVIDENCE, secondary]** The contrarian detail is the reliability gap: Claude Opus 4.6 is listed with a 50% horizon of about 14.5 hours but an 80% horizon of about 1 hour. A business process needs well above 80% reliability. Same source. https://en.wikipedia.org/wiki/METR
- Counter: the trend is steep and has not broken; both the 50% and 80% horizons are rising.

Enterprise-style agent benchmarks

- **[EVIDENCE]** TheAgentCompany (CMU et al.): simulated software-company environment with web, code, chat and colleague interaction. v1 (Dec 2024): best agent (Claude 3.5 Sonnet) fully completed 24% of tasks (34.4% partial credit). v2: best (Gemini 2.5 Pro) 30.3% full, 39.3% partial. A third-party aggregator lists 52.7% for a later model; unverified against the official leaderboard. https://www.alphaxiv.org/abs/2412.14161v2 https://benchmarklist.com/benchmarks/the_agent_company/
- **[EVIDENCE]** CRMArena-Pro (Salesforce AI Research, arXiv 2505.18878, May 2025): leading agents about 58% success single-turn, about 35% multi-turn with simulated users; near-zero inherent confidentiality awareness (they did not recognise information they should withhold). Best category, workflow execution, exceeded 83% single-turn for Gemini 2.5 Pro. Success in multi-turn correlated with the agent asking clarifying questions. https://arxiv.org/html/2505.18878v1 https://www.salesforce.com/blog/crmarena-pro/
- Note on incentive: Salesforce sells its own agent platform and publishes "why generic AI agents fall short"; the benchmark tests generic LLM agents. The result is still credible as a lower bound, but the framing serves the vendor. https://www.salesforce.com/blog/why-generic-ai-agents-fall-short/
- **[EVIDENCE]** tau-bench (Sierra, arXiv 2406.12045, June 2024): introduces pass^k, the chance an agent succeeds on all k repeats of the same task. GPT-4o: above 60% pass^1 in the retail domain, about 25% at pass^8; 35.2% pass^1 in airline. Older models; newer scores not retrieved this session. https://arxiv.org/pdf/2406.12045 https://sierra.ai/kr/blog/benchmarking-ai-agents

Analyst view

- **[FORECAST]** Gartner (25 June 2025): over 40% of agentic AI projects will be cancelled by end-2027 because of rising costs, unclear value and weak risk controls. "Agent washing": vendors rebranding assistants, RPA and chatbots as agents; Gartner estimates only about 130

of thousands of "agentic" vendors are genuine. Same release forecasts at least 15% of day-to-day work decisions made autonomously by agentic AI by 2028 (from 0% in 2024). https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027

- **[EVIDENCE, weak]** Gartner webinar poll, January 2025, n=3,412 attendees (self-selected): 19% reported significant agentic AI investment, 42% cautious investment. Same release.

Reasoning limits: both sides

- **[EVIDENCE]** Apple's "The Illusion of Thinking" (Shojaee et al., June 2025) used controllable puzzles (Tower of Hanoi, River Crossing and others) and reported that reasoning models' accuracy collapses beyond a complexity threshold and that their reasoning effort falls as problems get harder. (Paper not fetched; summarised from coverage.) https://arize.com/blog/the-illusion-of-thinking-what-the-apple-ai-paper-says-about-llm-reasoning/
- **[OPINION / EVIDENCE]** Rebuttal, "The Illusion of the Illusion of Thinking" (Alex Lawsen with Claude Opus as listed co-author, arXiv 2506.09250, June 2025): many Hanoi failures were models hitting output-token limits (15 disks need over 32,000 moves), the scoring marked truncated but correct strategies as failures, some River Crossing instances were unsolvable yet scored, and models asked for a recursive Lua function solved 15-disk Hanoi. The rebuttal does not claim models plan flawlessly. https://arxiv.org/html/2506.09250v1 https://9to5mac.com/2025/06/13/new-paper-pushes-back-on-apples-llm-reasoning-collapse-study/ https://venturebeat.com/ai/do-reasoning-models-really-think-or-not-apple-research-sparks-lively-debate-response
- How to weigh it: the rebuttal largely defuses the "models cannot reason" headline, but output-length limits and long-sequence execution errors are exactly the constraints agents face in long workflows. Either way, the paper is weak support for "agents won't work" and decent support for "long, exact multi-step execution is fragile".

(c) Displacement as task change, not job loss

Gap-fill pass (2026-11-06). All items below rest on search-engine extracts, not fetched pages. The Humlum and Vestergaard finding above (task reorganisation absorbs chatbots; no earnings or hours effect) remains the best-identified evidence for "task change, not job loss".

Radiology: the best-known failed displacement forecast

- **[OPINION, a forecast that has not come true]** Geoffrey Hinton said in 2016 that people should stop training radiologists because deep learning would outperform them within about five years (paraphrase; a widely circulated recording, described by Joshua Gans, who says he was in the room). Hinton later reframed the remark as about image analysis and AI working alongside radiologists. https://joshuagans.substack.com/p/radiologists-are-the-canaries-in (search extract); Fortune, 4 May 2026 https://fortune.com/2026/05/04/godfather-of-ai-geoffrey-hinton-radiologists-future-of-work-tech-ai-job-anxiety/ (search extract)
- **[EVIDENCE, secondary]** The number of practising US radiologists rose about 12% from 2010 to 2022, to 38,306, which Radiology Business reports as barely keeping pace with population growth. https://radiologybusiness.com/topics/healthcare-management/healthcare-staffing/radiologist-pipeline-barely-keeping-pace-population-growth (search extract)
- **[EVIDENCE, secondary]** Pay has risen, not fallen: Medscape's average radiologist pay was about $571,000 in 2025, up 9% year on year (as cited by Fortune, May 2026, link above). A single Substack post reports a record 1,208 diagnostic radiology residency positions in 2025 (+4% on 2024); not confirmed against NRMP or ACR data. A vendor analysis (Marit Health)

says radiology residency positions grew 33% from 2010 to 2025 against 69% for all residencies, i.e. training supply, not AI, is the binding constraint. https://www.marithealth.com/posts/radiology-shortage-ai-was-supposed-to-prevent (search extract)

- **[EVIDENCE]** Radiology is the most AI-penetrated specialty on paper: about 75-76% of FDA AI-enabled device authorisations, over 1,000 imaging tools by end-2025 (Applied Radiology). But one review found randomised-trial data for only 6 of 691 devices. https://www.appliedradiology.com/news/radiology-ai-tools-dominate-latest-fda-device-approvals ; https://www.intuitionlabs.ai/articles/ai-radiology-trends-2025 (search extract; the second is a secondary aggregator)
- Mainstream counterargument: radiology may not be a fair test. Cleared tools are narrow (one finding, one modality), imaging volumes have grown fast, and regulation and liability keep a radiologist in the loop. BLS still projects radiologist employment growth below the all-occupation average for 2024-2034 (figures seen only on aggregators; check bls.gov). "No job loss so far" is consistent with "slower hiring later".

Translators and interpreters: the strongest counterexample to "task change only"

- **[EVIDENCE]** BLS: about 75,300 interpreter and translator jobs in 2024, projected +2% to 2034 (slower than average), about 6,900 openings a year mostly from replacement, median pay $59,440 (May 2024). Slator reports BLS cut its projected growth from 3% to 2% between editions. The BLS category merges in-person interpreters (medical, legal) with written translators, so it masks any split. https://www.bls.gov/ooh/media-and-communication/interpreters-and-translators.htm (search extract); https://tech.slator.com/?p=105543 (search extract)
- **[EVIDENCE, peer-reviewed]** On a large online freelancing platform, posts for automation-prone jobs (writing, coding) fell about 21% within eight months of ChatGPT's release relative to manual-intensive jobs; 1,388,711 posts, July 2021 to July 2023. Writing was the most affected category. The remaining automation-prone jobs were more complex and better paid. Demirci, Hannane and Zhu, "Who Is AI Replacing?", Management Science, 2025. https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb ; https://www.diw.de/de/diw_01.c.906828.de/publikationen/externe_referierte_aufsaetze/2025_0010/who_is_ai_replacing__the_impact_of_generative_ai_on_online_freelancing_platforms.html (search extract)
- **[EVIDENCE, survey]** The ELIS 2025 survey (654 freelancers, 42 countries) found that, for the first time in the survey's history, most freelance translators reported price drops, despite inflation. Girletti and Lefer, 2025. https://research.dial.uclouvain.be/bitstreams/149e7a5a-6761-4c67-82e2-215b5fef51da/download (search extract)
- **[EVIDENCE, weak: self-selected member survey]** Society of Authors (UK), 2024: 787 respondents; 36% of translators said they had already lost work to generative AI and 43% said their income had fallen because of it. Seen only in secondary coverage. https://www.bookbusinessmag.com/aggregatedcontent/survey-finds-generative-ai-proving-major-threat-to-the-work-of-translators/ (search extract). A UK study of freelance translators (Penet et al., 2026) reports strong dissatisfaction with earnings and prospects. https://eprints.whiterose.ac.uk/id/eprint/237630/1/Penet%20et%20al.%202026.pdf (search extract)
- **[EVIDENCE, secondary review]** A 2025 literature review attributes to Qiao et al. (2024) a finding that translation freelancers completed fewer tasks and earned less after ChatGPT, while web developers on the same platform completed more and earned more. https://arxiv.org/pdf/2509.15265 (search extract)
- Mainstream counterargument: employee headcount in the BLS series is roughly flat, not collapsing, and Upwork's own analysis claims a positive effect on demand and earnings per

contract (vendor; it sells AI-era freelancing). https://www.upwork.com/research/ai-translation-future (search extract). The fair reading: for freelance written translation, the evidence points to lower prices and fewer jobs, which is income loss, not merely task change. For salaried and in-person interpreting, task change is the better description so far.

Bank tellers and ATMs: a story with a second half

- **[EVIDENCE, secondary]** James Bessen (Finance & Development, March 2015, "Toil and Technology") found that ATMs cut the tellers needed per urban branch from about 20 to about 13, but cheaper branches led banks to open about 43% more urban branches, so total teller employment rose while ATMs grew roughly fivefold from 1990 to 2010. Tellers moved towards sales and relationship tasks. https://conversableeconomist.com/2015/03/03/atms-and-a-rising-number-of-bank-tellers/ ; https://itif.org/publications/2016/01/07/why-computers-don%E2%80%99t-replace-workers/ (search extract; IMF primary not surfaced)
- **[EVIDENCE, secondary report of BLS]** The second half: BLS now projects teller employment to fall (reported as -13% over its latest ten-year window, from about 339,200), and attributes it to online and mobile banking reducing branch visits. Baseline counts for 2010 conflict across secondary sources (560,000 vs 332,000 depending on series). https://www.bls.gov/ooh/office-and-administrative-support/tellers.htm (search extract; via a bls.gov mirror and job sites); https://www.digitaltoday.co.kr/en/view/78231/mobile-banking-not-atms-behind-drop-in-us-bank-teller-staff (search extract)
- How to weigh it: both camps use the teller story. It shows task change can hold for two decades and then turn into job loss when a second technology removes the underlying activity (the branch visit). "Task change, not job loss" is a statement about a time window.

Autor's task framework

- **[OPINION, conceptual]** Autor (NBER w32140, 2024) argues AI could let workers with foundational training perform expert decision tasks now reserved for elites (medical care, document production, coding, teaching), rebuilding middle-skill jobs. He frames it explicitly as an argument about what is possible, not a forecast, and it depends on policy and on how firms deploy the tools. https://www.nber.org/papers/w32140 (search extract)
- **[EVIDENCE]** The empirical counterpart, already in Track 8 (09-what-will-people-do.md): Autor, Chin, Salomons and Seegmiller (QJE 2024) find augmentation innovations create new work titles while automation innovations reduce employment without creating new work. https://economics.mit.edu/sites/default/files/2026-04/New-Frontiers-QJE-v139%283%29-p1399-1465-Aug-2024.pdf
- Mainstream counterargument to the contrarian reading: the task framework does not say jobs are safe. It says the net effect depends on whether new tasks are created as fast as old ones are automated (the "reinstatement" side in Acemoglu and Restrepo's version). "Task change" is the mechanism by which job loss also happens. Not retrieved: Autor and Thompson's 2025 "Expertise" paper.

(d) "Freedom to be more human" fails or backfires

Gap-fill pass (2026-11-06); search extracts only. The earlier item stands: in the Danish data, workers report AI created new tasks (checking AI output, writing prompts) that absorb part of the time saved (Humlum and Vestergaard, 2025). https://flowingdata.com/2025/05/09/time-saved-with-ai-creates-more-work-elsewhere/

Deskilling and over-reliance

- **[EVIDENCE, observational]** Endoscopists at four Polish centres (part of the ACCEPT trial) who began using AI-assisted polyp detection then performed worse without it: the adenoma detection rate in standard, non-AI colonoscopies fell from 28.4% before AI was introduced to 22.4% after, an absolute fall of 6.0 points (95% CI -10.5 to -1.6). 19 experienced endoscopists. Budzyń et al., "Endoscopist deskilling after exposure to artificial intelligence in colonoscopy: a multicentre observational study", Lancet Gastroenterology & Hepatology 10(10): 896-903, Oct 2025, doi 10.1016/S2468-1253(25)00133-5. https://wrap.warwick.ac.uk/id/eprint/191005 (search extract)
- Counterargument: the Science Media Centre expert reaction notes that the study is observational, that major organisational changes happened in the department mid-study, and that the before/after window was short. A secondary post reports colonoscopy volume rose about 74% in the "after" period, an obvious confounder (not checked against the paper). The authors themselves call for randomised crossover trials, and a published letter (van de Sande et al.) treats it as a governance problem, not a reason to drop AI. https://www.sciencemediacentre.org/expert-reaction-to-observational-study-looking-at-detection-rate-of-precancerous-growths-in-colonoscopies-by-health-professionals-who-perform-them-before-and-after-the-routine-introduction-of-ai/ (search extract). This is the first clinical signal of deskilling, not proof of it.
- **[EVIDENCE, self-reported survey]** Survey of 319 knowledge workers (Microsoft Research and Carnegie Mellon): higher confidence in GenAI was associated with less critical-thinking effort, and higher self-confidence with more. Critical thinking shifted from producing work to verifying it, integrating responses and "task stewardship". Some participants could not verify outputs because they lacked domain knowledge. Lee, Sarkar et al., "The Impact of Generative AI on Critical Thinking", CHI 2025. https://www.microsoft.com/en-us/research/?p=1135061 ; https://advait.org/files/lee_2025_ai_critical_thinking_survey.pdf (search extract)
- Counterargument: cross-sectional and self-reported, so it shows association, not decline. The authors' own framing is that work shifts towards oversight, which can be read as task change rather than deskilling. Headlines such as "AI makes people dumber" overstate it.
- **[EVIDENCE, weak: preprint, small n]** MIT Media Lab, "Your Brain on ChatGPT" (Kosmyna et al., arXiv 2506.08872, June 2025): 54 participants wrote essays over three sessions across four months with ChatGPT, a search engine, or no tools, under EEG. The no-tool group showed the strongest neural connectivity and the ChatGPT group the weakest; in a fourth session, former ChatGPT users showed weaker recall and engagement. The authors call this "cognitive debt". https://arxiv.org/abs/2506.08872 (search extract, via LIRNEasia journal-club review https://lirneasia.net/2025/12/lirneasia-journal-club-your-brain-on-chatgpt-exploring-the-cognitive-debt-of-ai-assisted-writing)
- Counterargument: not peer reviewed; small, student-heavy sample; a lab essay task; EEG connectivity is not a measure of ability. One reported result points the other way for sequencing: people who first wrote unaided did well when later given AI. The fair lesson is "build the skill before outsourcing it", not "AI damages brains".
- Related, already in Track 8: four online experiments (N=3,562, Scientific Reports, April 2025) found GenAI collaboration raised performance but lowered intrinsic motivation and raised boredom on later solo work.

Work intensification

- **[EVIDENCE, weak: vendor survey]** Upwork Research Institute (July 2024): 77% of employees using AI said it had added to their workload; 39% spent more time reviewing or moderating AI content; 47% did not know how to achieve the productivity gains their

employer expected; 71% of full-time employees reported burnout. Sample: 2,500 respondents in the US, UK, Australia and Canada, of whom only 625 were full-time employees (plus 1,250 C-suite and 625 freelancers); fieldwork 16 April to 5 May 2024 (Walr / Workplace Intelligence). Upwork sells freelance labour and has an interest in the message that the full-time model is strained. https://investors.upwork.com/news-releases/news-release-details/upwork-study-finds-employee-workloads-rising-despite-increased-c (search extract)

- **[EVIDENCE, single-site qualitative]** Berkeley Haas eight-month ethnography (Ranganathan and Ye, HBR, Feb 2026): time saved became task expansion, blurred boundaries and longer hours. Details in Track 8. https://newsroom.haas.berkeley.edu/ai-promised-to-free-up-workers-time-uc-berkeley-haas-researchers-found-the-opposite/
- Counterargument: the Danish register data show no change in recorded hours, so intensification, if real, is not showing up as longer paid hours at scale. Self-reported time savings (Bick et al.) are positive for users. The Upwork number is a single question to a small employee subsample.

Algorithmic management and surveillance

- **[EVIDENCE, employer survey]** OECD employer survey (2025), over 6,000 firms in France, Germany, Italy, Japan, Spain and the US: managers mostly say algorithmic management tools improve their decisions and their own job satisfaction, but report concerns about unclear accountability, not being able to follow the tools' logic, and inadequate protection of workers' health. A union body (TUAC) summarises the cited worker surveys as showing more workload, stress and insecurity and less trust; it is an interested party. OECD AI Papers no. 31. https://ideas.repec.org/p/oec/comaaa/31-en.html ; https://oecd.ai/en/ai-publications/algorithmic-management-in-the-workplace ; https://tuac.org/news/new-oecd-report-on-algorithmic-management-reveals-urgent-need-for-worker-protections/ (search extract)
- **[EVIDENCE, partisan investigation]** US Senate HELP Committee majority staff (Sanders), December 2024, 18-month investigation: Amazon warehouses recorded about 30% more injuries than the warehouse industry average in 2023; internal safety staff recommended relaxing enforcement of productivity quotas and senior executives rejected this; the report alleges Amazon presented injury data selectively. Amazon called it wrong on the facts and based on outdated information. https://www.help.senate.gov/imo/media/doc/amazon_investigation.pdf (search extract)
- How to weigh it: algorithmic management predates generative AI, so this is not evidence about LLMs. It is evidence about what happens to freed human time when software allocates and paces work: it tends to be filled at a pace set by the system.

Social penalty for using AI

- **[EVIDENCE, experiments]** Four preregistered experiments (N=4,439): people expected to be, and were, judged lazier, less competent and less diligent when described as using generative AI rather than other tools; the penalty held across the described employee's age, gender and occupation. In a hiring simulation, managers who did not use AI themselves were less likely to pick candidates who used it regularly. Reif, Larrick and Soll, "Evidence of a social evaluation penalty for using AI", PNAS 122(19) e2426766122, May 2025. https://doi.org/10.1073/pnas.2426766122 ; https://www.psypost.org/people-who-use-ai-may-pay-a-social-price-according-to-new-psychology-research (search extract)
- Counterargument: vignette experiments, not field data. Coverage reports the penalty shrank when AI was clearly useful for the task and depended on evaluators' own AI use, which suggests it will fade as use becomes normal (as with earlier tools). It still matters for the

thesis: people hide AI use when it is penalised, which undermines "shared context" between humans and agents.

Viral statistics traced

- **"AI makes developers 19% slower."** Origin: METR RCT, July 2025, n=16 developers, 246 tasks, early-2025 tools, experienced maintainers on their own large repos. Solid as an RCT, narrow as a population. Often cited as a general claim about all developers or all AI tools, which METR itself does not claim. The 2026 follow-up did not "reverse" it (as some headlines said); METR said its newer data were too biased to estimate the effect.
- **"AI's $600B question."** Origin: one Sequoia partner's blog post, June 2024, using a 2x2 multiplier on Nvidia revenue. A useful heuristic, not a measurement. Repeated as if it were an accounting gap.
- **"Over 40% of agentic AI projects will be cancelled."** Origin: Gartner press release, June 2025. A prediction, not an observed cancellation rate. Often cited in the past tense.
- **"0.66% TFP over 10 years."** Origin: Acemoglu model. An upper bound under stated assumptions that exclude new tasks. Often repeated as "AI will add only 0.66% to GDP", which conflates TFP and GDP (the paper's GDP figure is a separate number not retrieved here).
- **"AI saves 1.4% of work hours."** Origin: Bick, Blandin, Deming. Averaged over all workers including non-users; self-reported. Users alone report much more. Easy to misuse in both directions.
- **"Agents fail 65% of multi-turn tasks."** Origin: CRMArena-Pro (Salesforce), generic LLM agents in a CRM sandbox, mid-2025 models. Real, but vendor-produced and dated.
- **"Hinton said stop training radiologists" (2016).** Real remark, widely recorded. Used now as proof AI forecasts fail. The radiologist count and pay rose, but the residency "record 1,208 positions in 2025" figure comes from a single Substack post; check NRMP before repeating it.
- **"ATMs increased bank teller jobs."** True for roughly 1990-2010 (Bessen). Usually cited without the second half: BLS now projects falling teller employment as mobile banking removes branch visits.
- **"77% of employees say AI added to their workload."** Upwork, July 2024. Based on 625 full-time employees inside a 2,500-person sample dominated by executives; vendor survey.
- **"AI deskills doctors" (endoscopy study).** One observational before/after study, 19 endoscopists, with a large concurrent rise in volume and organisational change. A signal, not a finding.
- **"ChatGPT rots your brain" (MIT).** A 54-person EEG preprint on essay writing. The authors' "cognitive debt" framing travelled much further than the data.

Strongest counter-evidence

Against this contrarian file (i.e. for faster transformation):

- US labour productivity at about 2.1% a year since 2019, and a strong Q2 2026 print, are better than the 2010s. If AI is starting to show up, this is where it would show first. (BLS, 2026; links above.)
- METR time horizons keep doubling every 4-7 months; benchmark failure rates from 2024-2025 models will not hold for 2026 models. TheAgentCompany and tau-bench results cited here are one to two model generations old.
- Hyperscalers are not cutting capex despite repeated sell-offs; they see demand they cannot meet (Microsoft's cut was supply-side).
- Users of AI report much larger time savings than the population average; the averages are dragged down by non-users, which is an adoption story, not a capability story.
- Radiology, the specialty with the most FDA-cleared AI tools, has more radiologists and higher pay than when displacement was predicted (Radiology Business; Fortune 2026). Humlum and Vestergaard find employers absorb AI by reorganising tasks.

Against the user's thesis ("tech is easy, transformation is hard"):

- The "freedom to be more human" half of the purpose statement has no supportive field evidence yet. The early signals are deskilling (Budzyń 2025), intensification (Upwork 2024; Berkeley Haas 2026) and a social penalty that pushes AI use into hiding (Reif et al. 2025). These are not readiness failures; they occur where adoption succeeds.
- The agent benchmarks show a real technical ceiling on multi-turn, permission-aware, repeated work. That is a tech problem. Shared context does not fix an agent that leaks confidential data or succeeds only one time in four across eight repeats.
- The Danish evidence suggests even where adoption happens and is encouraged by employers, earnings and hours barely move. That is not "business not ready"; it may be that the gains are small.

Open questions

Leads not verified this session (each needs its primary source fetched before use; no claim is made here):

- (c) Radiology: verify radiologist headcount (BLS OEWS; aggregators give 26,770 to 31,960 depending on year and definition) and NRMP diagnostic radiology positions 2016-2026 at source. Is AI slowing growth in reading volume per radiologist, or not used at scale?
- (c) Translators: find the platform paper behind a circulating "translator earnings fell 29.7% after ChatGPT" figure (seen in search results, not attributable to a URL, so not cited); get Hui, Reshef and Zhou's freelancer earnings estimates; read the Society of Authors 2024 survey report itself.
- (c) Bank tellers: pull BLS OEWS teller employment 2000-2025 directly; the secondary counts conflict.
- (c) Autor and Thompson, "Expertise" (2025); Yale Budget Lab occupational-mix tracker 2026 updates.

- (d) Endoscopy: a 2026 Norwegian prospective study (Pedersen et al., Endoscopy) reportedly found no significant fall in detection after AI was removed; seen only in a LinkedIn summary, so not cited. If confirmed it is the strongest counter to Budzyń.
- (d) Any field (not lab) study of AI and deskilling outside medicine; any peer-reviewed study of AI work intensification with hours data.
- (a) Census BTOS firm-level AI use rates for 2025-2026 (the most direct US firm adoption statistic); not retrieved.
- (b) Newer tau-bench / tau2-bench and TheAgentCompany scores for 2026 models; METR's own Feb 2026 post read in full.
- Whether the 2026 Humlum and Vestergaard paper is the published successor to NBER w33777, and its exact estimates.

What's genuinely uncertain

- Whether weak 2025 TFP is the dip of a J-curve or simply the absence of a large effect. The two are observationally similar for several years.
- Whether 2026 market sell-offs presage a capex cut. So far each one reversed and guidance rose.
- How much of the agent reliability gap closes with each model generation. The benchmark numbers here are 12-18 months old, which is a long time in this field.
- How far self-reported time savings overstate real savings. The METR perception gap suggests a lot for experts on familiar code; that may not generalise.
- The size of the revenue gap: Cahn and Bain are top-down heuristics and actual AI revenue disclosures are patchy.

Numbers for the facts file

Figure	What it measures	Source	URL	Date	Reliability note
≤0.66% (and <0.53%)	Upper-bound rise in US TFP over 10 years from AI	Acemoglu, Economic Policy	https://www.nber.org/papers/w32487	May 2024 (journal 2025)	Model-based; excludes new tasks. Not re-read at source
0.8% (vs 1.5% in 2024)	US private nonfarm TFP growth, 2025	BLS	https://www.bls.gov/news.release/archives/prod3_03192026.pdf	19 Mar 2026	Official; seen in search excerpt
2.2% (vs 3.0% in 2024)	US private nonfarm labour productivity growth, 2025	BLS	https://www.bls.gov/news.release/archives/prod3_03192026.pdf	19 Mar 2026	Official; seen in search excerpt
1.4% annualised; 2.2% y/y	US nonfarm business labour productivity, Q2 2026	BLS	https://www.bls.gov/opub/ted/2026/productivity-up-2-2-percent-from-second-quarter-2025-to-second-quarter-2026.htm	Aug/ Sep 2026	Official
2.1% a year	US labour productivity growth Q4 2019 to Q2 2026	BLS	https://www.bls.gov/opub/ted/2026/productivity-up-2-2-percent-from-second-quarter-2025-to-second-quarter-2026.htm	2026	Official
-0.8% q/q; +0.7% y/y	UK output per hour, Q2 2026	ONS via Trading Economics	https://tradingeconomics.com/united-kingdom/productivity/news/576075	18 Aug 2026	Official, but RTI vs LFS measures disagree
15.9%	TFP level above official, end-2017, after intangibles adjustment	Brynjolfsson, Rock, Syverson	https://www.aeaweb.org/doi/10.1257/mac.20180386	2021	Peer-reviewed, model-based
1.4%	Time saved by genAI as share of all US work hours	Bick, Blandin, Deming	https://www.nber.org/system/files/working_papers/w32966/w32966.pdf	Feb 2025	n>10,000, self-reported

Figure	What it measures	Source	URL	Date	Reliability note
1-5% / 1-7%	Share of all US work hours assisted by genAI	Bick, Blandin, Deming (two versions)	https://fedinprint.org/item/fedlwp/98805/102025	Feb / Oct 2025	Range changed between versions; cite version
~3% (2.8%)	Average time savings among Danish chatbot users	Humlum & Vestergaard	https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf	Apr/ May 2025	Admin + survey data, ~25,000 workers
3-7%	Share of productivity gains passed through to earnings	Humlum & Vestergaard	https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf	2025	Seen in coverage
>$1tn	Expected AI capex "in coming years"	Goldman Sachs Top of Mind 129	https://www.goldmansachs.com/insights/goldman-sachs-exchanges/a-skeptical-look-at-ai-investment	25 Jun 2024	Opinion/forecast
$600bn / $500bn gap	Annual AI revenue needed to justify capex / gap	Sequoia (Cahn)	https://www.sequoiacap.com/article/ais-600b-question/	20 Jun 2024	Heuristic (Nvidia rev x2 x2)
$2tn; $800bn	Annual revenue needed by 2030; shortfall after savings	Bain Global Technology Report	https://www.theregister.com/2025/09/24/bain_ai_costs	23 Sep 2025	Consultancy forecast
$803-828bn	2026 capex of largest AI spenders	Hargreaves Lansdown	https://www.hl.co.uk/news/ai-hyperscalers-are-expected-to-finish-the-year-spending-a-whopping-828-billion-on-capex	Oct 2026	Secondary; includes SpaceX; check filings
>$1tn in one week	Market value lost by top 20 AI/chip stocks, mid-Oct 2026	Dealroom	https://dealroom.co/charts/ai-selloff	30 Oct 2026	Secondary analysis
19% slower	Task time with AI, experienced OSS developers	METR RCT	https://arxiv.org/abs/2507.09089	Jul 2025	RCT, n=16, 246 tasks, early-2025 tools
~20% faster (believed)	Developers' own post-hoc estimate of AI speedup	METR RCT	https://arxiv.org/abs/2507.09089	Jul 2025	Same study
50%: ~14.5h vs 80%: ~1h	Claude Opus 4.6 METR time horizons	Wikipedia (METR)	https://en.wikipedia.org/wiki/METR	2026	Secondary; check metr.org
24% / 30.3%	Best agent full task completion, TheAgentCompany v1/v2	Xu et al.	https://www.alphaxiv.org/abs/2412.14161v2	Dec 2024 / 2025	Benchmark; older models
~58% / ~35%	Best agents single-turn / multi-turn success, CRMArena-Pro	Salesforce AI Research	https://arxiv.org/html/2505.18878v1	May 2025	Vendor benchmark; generic agents
>60% to ~25%	GPT-4o retail pass^1 to pass^8	tau-bench (Sierra)	https://arxiv.org/pdf/2406.12045	Jun 2024	Old model
28.4% to 22.4%	Adenoma detection rate in non-AI colonoscopies before vs after routine AI use	Budzyń et al., Lancet Gastroenterol Hepatol	https://wrap.warwick.ac.uk/id/eprint/191005	Oct 2025	Observational, 19 endoscopists, 4 centres; confounded; search extract
319	Knowledge workers surveyed on GenAI and critical thinking	Lee, Sarkar et al., CHI 2025	https://www.microsoft.com/en-us/research/?p=1135061	Apr 2025	Self-reported, cross-sectional
54	Participants in MIT "Your Brain on ChatGPT" EEG study	Kosmyna et al., arXiv 2506.08872	https://arxiv.org/abs/2506.08872	Jun 2025	Preprint; small n
N=4,439	Participants in four experiments on the social penalty for AI use	Reif, Larrick, Soll, PNAS	https://doi.org/10.1073/pnas.2426766122	May 2025	Preregistered vignette experiments
77% (n=625	Employees saying AI added to their workload		https://investors.upwork.com/news-releases/news-release-details/upwork-	Jul 2024	

Figure	What it measures	Source	URL	Date	Reliability note
employees)		Upwork Research Institute	study-finds-employee-workloads-rising-despite-increased-c		Vendor survey; 2,500 total incl. 1,250 C-suite
>6,000 firms	OECD employer survey on algorithmic management, 6 countries	OECD AI Papers no. 31	https://ideas.repec.org/p/oec/comaaa/31-en.html	2025	Employer (manager) perceptions
~30%	Amazon warehouse injuries above warehouse industry average, 2023	US Senate HELP majority staff	https://www.help.senate.gov/imo/media/doc/amazon_investigation.pdf	Dec 2024	Partisan report; Amazon disputes
+12% to 38,306	Practising US radiologists, 2010 to 2022	Radiology Business (reporting a study)	https://radiologybusiness.com/topics/healthcare-management/healthcare-staffing/radiologist-pipeline-barely-keeping-pace-population-growth	2025	Secondary; search extract
~$571,000 (+9%)	Average US radiologist pay, 2025	Medscape via Fortune	https://fortune.com/2026/05/04/godfather-of-ai-geoffrey-hinton-radiologists-future-of-work-tech-ai-job-anxiety/	May 2026	Secondary
75,300; +2%	US interpreter/translator jobs 2024; projected growth 2024-34	BLS OOH	https://www.bls.gov/ooh/media-and-communication/interpreters-and-translators.htm	2025	Official; search extract
-21%	Freelance platform posts for automation-prone (writing, coding) jobs, 8 months after ChatGPT	Demirci, Hannane, Zhu, Management Science	https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb	2025	Peer-reviewed; 1.39m posts; one platform
36%; 43%	UK translators who lost work / income to GenAI	Society of Authors	https://www.bookbusinessmag.com/aggregatedcontent/survey-finds-generative-ai-proving-major-threat-to-the-work-of-translators/	2024	Self-selected, 787 respondents; secondary
~20 to ~13 tellers; +43% branches	ATM effect per urban branch; urban branch growth	Bessen via Conversable Economist	https://conversableeconomist.com/2015/03/03/atms-and-a-rising-number-of-bank-tellers/	Mar 2015	Secondary report of IMF F&D article
>40%	Agentic AI projects Gartner predicts cancelled by end-2027	Gartner	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027	25 Jun 2025	Forecast, not measurement
~130	Genuine agentic vendors out of "thousands"	Gartner	same	25 Jun 2025	Opinion; method not disclosed

Verbatim quotes

None. No page could be fetched in this session, so no wording could be confirmed. Candidate quotes to fetch and verify later: Covello's "$1tn problem" line (Goldman Sachs Top of Mind 129, June 2024); METR's Feb 2026 statement that its new data give an unreliable signal (metr.org); Gartner's "agent washing" definition (Gartner press release, 25 June 2025).

Sources worth reading in full (ranked)

1. Humlum & Vestergaard, NBER w33777 and the 2026 follow-up: the best-identified evidence that real adoption has produced small labour-market effects. https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf
2. METR RCT (arXiv 2507.09089) plus METR's Feb 2026 redesign post: the cleanest causal study of AI on expert work, and an honest account of why measuring it got harder. https://arxiv.org/abs/2507.09089

3. CRMArena-Pro (arXiv 2505.18878): the most business-relevant agent benchmark; multi-turn and confidentiality results matter for any agentic-platform claim. https://arxiv.org/html/2505.18878v1
4. Acemoglu, "The Simple Macroeconomics of AI": the serious low estimate and its assumptions. https://www.nber.org/papers/w32487
5. Brynjolfsson, Rock, Syverson, "The Productivity J-Curve": the best framework for why statistics lag. https://www.aeaweb.org/doi/10.1257/mac.20180386
6. BLS TFP release, March 2026: the official 2025 slowdown. https://www.bls.gov/news.release/archives/prod3_03192026.pdf
7. Bick, Blandin, Deming (St. Louis Fed revision, Oct 2025): national adoption and time-savings baseline. https://fedinprint.org/item/fedlwp/98805/102025
8. tau-bench (arXiv 2406.12045): introduces pass^k, the right reliability metric for production agents. https://arxiv.org/pdf/2406.12045
9. Bain Global Technology Report 2025: the revenue-gap arithmetic. https://www.theregister.com/2025/09/24/bain_ai_costs
10. Apple "Illusion of Thinking" and the Lawsen rebuttal, read together. https://arxiv.org/html/2506.09250v1
11. Budzyń et al., Lancet Gastroenterology & Hepatology, Oct 2025: the first clinical deskilling signal; read with the Science Media Centre critique. https://wrap.warwick.ac.uk/id/eprint/191005
12. Reif, Larrick and Soll, PNAS 2025: why people hide AI use at work. https://doi.org/10.1073/pnas.2426766122
13. Demirci, Hannane and Zhu, Management Science 2025: the cleanest evidence of actual demand loss for writing work. https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb
14. Lee, Sarkar et al., CHI 2025: how GenAI shifts critical thinking towards verification. https://www.microsoft.com/en-us/research/?p=1135061

12. Facts and figures

Compiled 2026-11-06 from the "Numbers for the facts file" tables in track files 02 to 11. Every row has a source, URL and date. Read the reliability note before using any number.

How to read this file

- **Basis.** This session's network policy blocked page fetches for most hosts. Rows whose reliability note says "search summary", "search extract" or "snippet" were taken from search-engine extracts that cite the URL, not from reading the page. Rows from anthropic.com, microsoft.com and raw.githubusercontent.com were read directly.
- **Verification.** 29 of the most striking figures were re-checked in `verification.md`: 9 confirmed by re-fetching the page, 17 corroborated by an independent search, 2 corrected, 1 not confirmed. The corrected and unconfirmed ones are listed in DO NOT CITE at the end.
- **Labels.** Vendor and consultancy survey numbers are weak evidence at best. Exposure measures and analyst predictions are forecasts. The track files carry the EVIDENCE / FORECAST / OPINION labels for each finding.
- **Duplicates.** Some figures appear in more than one track table (e.g. Humlum and Vestergaard, MIT NANDA). They are left in place because each track notes a different use.

Headline figures at a glance (all verified or corroborated in the verification pass)

Figure	What it measures	Source	URL	Date	Status
5%	Custom enterprise GenAI tools reaching production with P&L impact (self-selected sample)	MIT NANDA, GenAI Divide	https://www.theregister.com/2025/08/18/generative_ai_zero_return_95_percent	Jul-Aug 2025	Corroborated; weak source
46%	AI POCs progressed to production, n=3,120	IDC for Lenovo CIO Playbook 2026	https://www.businesswire.com/news/home/20260127255538/en/	Jan 2026	Corroborated; vendor-sponsored
37%	Respondents attributing any EBIT impact to AI, n=1,719	McKinsey State of AI 2026	https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai	Aug 2026	Corroborated; self-report
27% to 40%	$1B+ firms scaling AI agents, 2025 to 2026	McKinsey State of AI 2026	https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai	Aug 2026	Corroborated; self-report
46% / 7%	Small employer firms using AI / users "fully integrated"	Fed Small Business Credit Survey	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	2026	Corroborated (7% via secondary)
~10% to ~17%	US firms using AI before/after BTOS wording change, Nov 2025	St Louis Fed on Census BTOS	https://www.stlouisfed.org/on-the-economy/2026/jun/measuring-ai-adoption-firms-how-you-ask-matters	Jun 2026	Corroborated; corrected from 9-18%
2.8% (about 3%)	Self-reported time savings; earnings and hours effects precisely zero	Humlum and Vestergaard, NBER w33777	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	May 2025	Corroborated
0.8% vs 1.5%	US private nonfarm TFP growth 2025 vs 2024	BLS	https://www.bls.gov/news.release/archives/prod3_03192026.pdf	19 Mar 2026	Corroborated
~9% / ~80%		Hosseini and Lichtinger	https://raw.githubusercontent.com/s-mahdihosseini/Website/	May 2026	Confirmed (re-fetched)

Figure	What it measures	Source	URL	Date	Status
	Junior employment fall after six quarters / fewer junior hires per quarter at GenAI adopters		b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf		
~19%	22-25 year olds in AI-exposed jobs below counterfactual (ADP)	Stanford Digital Economy Lab	https://digitaleconomy.stanford.edu/news/canariesaug26/	12 Aug 2026	Corroborated; descriptive
86%	Analysts following an agent's coin-flipped lean, n=50	Anthropic RCT	https://www.anthropic.com/research/agents-judgment-calls	27 Oct 2026	Confirmed; vendor research
0.692 to 0.462	Agent success, one try vs all four tries (tau-bench retail, 2024 model)	Sierra tau-bench	https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md	2024	Confirmed
~58% to ~35%	LLM agent success single-turn vs multi-turn CRM tasks	Salesforce CRMArena-Pro	https://arxiv.org/abs/2505.18878	May 2025	Corroborated; vendor benchmark
17.0% vs 91.2%	Enterprise text-to-SQL vs academic benchmark success	Spider 2.0	https://arxiv.org/abs/2411.07763	Nov 2024	Corroborated
78% / 52% / 53%	AI users bringing own AI / reluctant to admit use / worry about looking replaceable, n=31,000	Microsoft and LinkedIn WTI 2024	https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part	May 2024	Confirmed; vendor survey
10% vs 5%	Skilled AI users blocked by org vs ready orgs with unready people, n=20,000	Microsoft WTI 2026	https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization	May 2026	Confirmed; vendor survey
59% / 0-20%	Share of work using Claude / share engineers can fully delegate	Anthropic internal study	https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic	Dec 2025	Confirmed; vendor self-study
57%	Employees hiding AI use, n=48,000+, 47 countries	KPMG and University of Melbourne	https://mbs.edu/news/global-study-reveals-trust-of-ai-remains-a-critical-challenge	Apr 2025	Corroborated
52%	US adults more concerned than excited about AI, n=3,488	Pew Research Center	https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/	Aug 2026	Corroborated
30M+	Paid Microsoft 365 Copilot seats	Microsoft FY26 Q4 call	https://www.uctoday.com/unified-communications/microsoft-365-copilot-passes-30-million-paid-seats-as-cloud-and-ai-growth-power-record-quarter	29 Jul 2026	Corroborated; secondary
5,527 to 3,422	Klarna full-time employees, end-2022 to end-2024	Klarna F-1	https://www.bankingdive.com/news/klarna-buy-now-pay-later-bnpl-ai-workforce-ipo/742703/	Mar 2025	Corroborated
~1,600	Court decisions involving AI-hallucinated content, 35 countries	Charlotin database via AFP	https://www.malaymail.com/news/tech-gadgets/2026/06/16/do-your-job-global-courts-face-rising-tide-of-ai-generated-fake-cases-and-quotes/223865	Jun 2026	Corroborated; crowd-sourced

Track 1. AI-native definitions (from `02-ai-native-definitions.md`)

Figure	What it measures	Source	URL	Date	Reliability note
77% vs 12%	Share of Anthropic first-party API transcripts showing automation vs augmentation patterns	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	2025-09-15	Vendor data, 1M transcripts (Aug 2025), classifier-based. Fetched + grep-verified
0.38	Elasticity of output to input token length (basis for the	Anthropic Economic Index	same	2025-09-15	Inferential; author interpretation. Fetched

Figure	What it measures	Source	URL	Date	Reliability note
	"context bottleneck" inference)				
59% (from 28%)	Share of their work Anthropic engineers use Claude for, self-reported	Anthropic	https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic	2025-12-02	Self-report, n=132. Fetched + grep-verified
0-20%	Share of work most Anthropic engineers say they can "fully delegate"	Anthropic	same	2025-12-02	Self-report. Fetched + grep-verified
27%	Claude-assisted work that "wouldn't have been done otherwise"	Anthropic	same	2025-12-02	Self-estimate. Fetched + grep-verified
844 of 31,000	Leaders classed as Frontier Firm, of total respondents	Microsoft WTI 2025	https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born	2025-04-23	Vendor survey (Edelman); self-classification. Fetched
24% / 12%	Leaders reporting org-wide AI deployment / still in pilot	Microsoft WTI 2025	same	2025-04-23	Vendor survey. Fetched + grep-verified
19%	AI users in the "Frontier" zone (high individual and organisational readiness)	Microsoft WTI 2026	https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization	2026-05-05	n=20,000 AI users, 10 markets, self-report. Fetched
67% vs 32%	Share of modelled "AI impact" from organisational vs individual factors	Microsoft WTI 2026	same	2026-05-05	Random-forest importance on self-report; associational. Fetched
13%	AI users rewarded for reinventing work with AI even if results aren't met	Microsoft WTI 2026	same	2026-05-05	Self-report. Fetched
15x	Year-on-year growth in active agents in Microsoft 365	Microsoft WTI 2026	same	2026-05-05	Vendor telemetry, no absolute counts. Fetched
25% vs 14%	Frontier vs other professionals saying agent workflows/handoffs are documented at org level	Microsoft WTI 2026	same	2026-05-05	Self-report. Fetched
700	FTE customer service agents Klarna said its AI assistant replaced	Klarna via eMarketer	https://www.emarketer.com/content/klarna-backtracks-ai-customer-service-plans	2024-02 (claim)	Company claim; search summary only
~9,000 to ~5,000	Salesforce support headcount, per Benioff	Fortune	https://fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales/	2025-09-02	CEO podcast remark; no timeframe; search summary only
"a couple of hundred"	IBM HR roles replaced by AI	Entrepreneur citing WSJ	https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341	2025-05	CEO statement; search summary only
$400,000 ($225k + $175k)	SEC civil penalties, Delphia and Global Predictions AI-washing settlements	SEC press release 2024-36	https://www.sec.gov/newsroom/press-releases/2024-36	2024-03-18	Regulatory action; search extract
>70%	Share of Presto Voice orders needing human intervention, per SEC case	SEC settlement via law-firm summaries	https://www.dandodiary.com/2025/01/articles/artificial-intelligence/sec-files-ai-washing-enforcement-action-against-restaurant-technology-company/	2025-01-14	Regulatory finding as reported; SEC primary not seen
$40M+ / ~0%	Raised by Nate on AI claims / alleged real automation rate	DOJ SDNY press release	https://www.justice.gov/usao-sdny/pr/tech-ceo-charged-artificial-intelligence-investment-fraud-scheme	2025-04-09	Allegations in an indictment, not findings
5	Companies in FTC Operation AI Comply	FTC press release		2024-09-25	Regulatory action; search extract

Figure	What it measures	Source	URL	Date	Reliability note
			https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes		
~130	Vendors with genuine agentic capability, of thousands claiming it	Gartner press release	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027	2025-06-25	Analyst estimate, no method
5,527 to 3,422	Klarna full-time employees, end-2022 to end-2024	Klarna F-1, via press	https://www.sec.gov/Archives/edgar/data/2003292/000162828025012824/klarnagroupplcf-1.htm	2025-03	Filing figure seen only in secondary reports
2.3M / two-thirds	Chats handled by Klarna's assistant in first month / share of all service chats	Klarna via Euronews	https://www.euronews.com/2024/02/28/klarnas-ai-bot-is-doing-the-work-of-700-employees-what-will-happen-to-their-jobs	2024-02-28	Company claim
~10%	Decline in Duolingo US daily active users linked to the AI-first memo	Morgan Stanley via iTiger	https://www.itiger.com/news/2549526150	2025	Sell-side analyst attribution
721	Firms in MIT CISR AI maturity survey (2022)	MIT CISR	https://cisr.mit.edu/sites/default/files/2024-12/EnterpriseAIMaturityModel_CISRpressrelease_121924_0.pdf	2024-12-19	Pre-genAI; search summary only

Track 2. SMB (from `03-smb.md`)

Figure	What it measures	Source	URL	Date	Reliability note
~10% to ~17% (corrected in verification)	US firms using AI, before vs after BTOS question change (17 Nov 2025)	Ramp AI Index; Census wording note	https://ramp.com/data/ai-index-january-2026	14 Jan 2026	Level shift is a wording artefact. Search-summary only
~18%	US firms using AI, year-end 2025 (new wording)	Fed FEDS Note	https://www.federalreserve.gov/econres/notes/feds-notes/monitoring-ai-adoption-in-the-u-s-economy-20260403.html	3 Apr 2026	Official staff analysis of BTOS. Search-summary only
~19.8%	US firms using AI, period ending 3 May 2026	Census Bureau story	https://census.gov/library/stories/2026/05/ai-use-businesses.html	May 2026	Official. Search-summary only
<20%	US firms with 1-4 employees using AI (ceiling, not point)	Census Bureau story	https://census.gov/library/stories/2026/05/ai-use-businesses.html	May 2026	Census publishes as a band
46%	US small employer firms whose business or employees use AI	Fed SBCS 2026 Report on Employer Firms	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	Mar 2026 (fieldwork Sep-Nov 2025)	Convenience sample, weighted. Includes employees' use
7%	Share of SBCS AI-using firms "fully integrated" (~3% of all employer firms, my arithmetic)	Fed SBCS 2026	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	Mar 2026	Self-assessed. Search-summary only
71%	SBCS AI users reporting higher productivity	Fed SBCS 2026; Cook speech	https://www.federalreserve.gov/newsevents/speech/cook20260624a.htm	24 Jun 2026	Self-report
>50% / 30%	Of SBCS firms with no AI plans: "not applicable" / "prefer not to use"	Fed SBCS 2026	https://www.fedsmallbusiness.org/2026-report-on-employer-firms	Mar 2026	Search-summary only
24%		NFIB Small Business and	https://www.nfib.com/?p=38888	25 Jun 2025	Random draw of 20,000 members

Figure	What it measures	Source	URL	Date	Reliability note
	US small employers using AI (21% of <10 employees; 48% of >50)	Technology Survey			by email. Size split from secondary summary
98%	NFIB AI users reporting no change in employee numbers	NFIB	https://www.nfib.com/?p=38888	25 Jun 2025	Search-summary only
58%	US small businesses saying they use generative AI	US Chamber / Teneo 2025	https://www.uschamber.com/technology/the-impact-of-technology-on-u-s-small-business	2025	n and method not found
98%	US small businesses using any AI-enabled tool	US Chamber 2024 edition	https://www.uschamber.com/technology/the-impact-of-technology-on-u-s-small-business	Sep 2024	Measures software features, not adoption. Do not cite as adoption
55% (from 39%)	Small businesses "currently using" AI	Thryv	https://investor.thryv.com/news/news-details/2025/AI-Adoption-Among-Small-Businesses-Surges-41-in-2025-According-to-New-Survey-from-Thryv/	17 Jul 2025	Vendor, n=540, revenue floor $100k
76%	Goldman 10KSB participants using AI	Goldman Sachs 10KSB Voices (Babson / David Binder)	https://www.businesswire.com/news/home/20260317141697/en/Survey-Small-Businesses-Embrace-AI-But-Need-Training-and-Support-to-Fully-Harness-It	17 Mar 2026	n=1,256 programme alumni. Not representative
46.6%	US businesses with paid AI spend on Ramp	Ramp AI Index	https://ramp.com/data/ai-index-january-2026	14 Jan 2026	Ramp customers only, tech-forward
69%	Firms using AI, executive panels in US/UK/DE/AU	NBER w34836	https://www.nber.org/papers/w34836	Feb/Mar 2026	n about 6,000. Working paper
>80%	Firms in that panel reporting no AI impact on employment or productivity	NBER w34836	https://www.nber.org/papers/w34836	Feb/Mar 2026	Search-summary only
~12% to ~35%	UK firms with 10+ employees using AI, late 2023 to mid-2026	ONS AI in UK businesses 2023-2026	https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/articles/artificialintelligenceinukbusinesses/2023to2026	20 Jul 2026	Official. Search-summary only
29% / 28% / 49%	UK firms using AI, June 2026: all / 0-9 employees / 250+	ONS (BICS)	same as above	Jul 2026	Official. Search-summary only
1.6	AI tools per UK adopter	TNW on ONS	https://thenextweb.com/news/uk-ai-adoption-widening-not-deepening-ons-2026	Jul 2026	Not confirmed in ONS tables
4%	UK AI-using firms reporting headcount fell due to AI	ONS BICS	https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/bulletins/businessinsightsandimpactontheukeconomy/8january2026	8 Jan 2026	Attribution to bulletin inferred
39% / 21% / 16%	UK barriers 2023: use cases / cost / skills	ONS management practices article	https://www.ons.gov.uk/economy/economicoutputandproductivity/productivitymeasures/articles/managementpracticesandtheadoptionoftechnologyandartificialintelligenceinukfirms2023/2025-03-24	24 Mar 2025	Official. Search-summary only
16% overall; 14% micro	UK businesses using AI, 2025	DSIT AI Adoption Research (IFF / Technopolis, n=3,500 CATI)	https://assets.publishing.service.gov.uk/media/6960f3924343a0da370869ba/AI_Adoption_Research_Report.pdf	2025 (pub. ~Jan 2026)	Official-commissioned. Search-summary only

Figure	What it measures	Source	URL	Date	Reliability note
32% (vs 16%)	UK businesses using at least one AI technology, 2026	DSIT AI Adoption Survey 2026 (n=3,503 CATI)	https://www.gov.uk/government/publications/ai-adoption-survey-2026/ai-adoption-survey-report-2026-executive-summary	29 Oct 2026	Micro figure not obtained
20.0% (13.5% in 2024)	EU enterprises with 10+ employees using AI, 2025	Eurostat	https://ec.europa.eu/eurostat/statistics-explained/index.php/Use_of_artificial_intelligence_in_enterprises	Dec 2025	Excludes <10 employees
17.0% / 30.4% / 55.0%	EU AI use 2025: small / medium / large	Eurostat	same	Dec 2025	Search-summary only
26.7%	Scottish businesses using AI, Sep 2025 (14.0% Sep 2023)	Scottish Government BICS	https://www.gov.scot/publications/scottish-economic-bulletin-december-2025/pages/6/	Dec 2025	Scotland only
~11% / 21% / 41%	EU AI use 2024: small / medium / large (so +6 vs +14 points to 2025)	Eurostat via secondary	https://ioplus.nl/en/posts/ai-usage-soars-in-european-companies	Dec 2025	Secondary; check isoc_eb_ai table
31% (24% JP to 39% DE)	SMEs using generative AI, 7 OECD countries, late 2024	OECD, Generative AI and the SME Workforce, n>5,000	https://www.oecd.org/en/publications/generative-ai-and-the-sme-workforce_2d08b99d-en.html	Dec 2025	Official IO survey; frame not seen. Search extract
57%	Non-adopting SMEs citing poor fit with their work	OECD, same	same	Dec 2025	Search extract
83% / 9% / 6%	Adopting SMEs: no staffing effect / decrease / increase	OECD, same	same	Dec 2025	Search extract
$10,600/yr	Mean SMB spend on AI tools, services and initiatives (n=241)	Clarify Capital	https://clarifycapital.com/small-business-ai-spending	Apr 2026	Lender survey; mean skewed by big spenders
71% (from 41%)	SMBs paying for AI (n=363)	Clarify Capital	https://clarifycapital.com/small-business-ai-spending	Jul 2026	Lender survey. Search extract
£93 to £288	Mean AI spend per UK small business, Q2 2024 to Q2 2026	Capital on Tap via PPC Land	https://ppc.land/one-in-three-small-business-owners-raided-personal-savings-to-pay-for-ai/	2026	Card provider; method not seen; secondary
~€160/month	Average AI tool bill per UK/EU SMB customer, 2024	Spendesk	https://ffnews.com/blog/2024/12/13/spending-on-ai-tools-by-uk-smbs-grows-fivefold-in-past-two-years-according-to-spendesk-data	Dec 2024	Vendor customer data
$1,500-$25,000 + $50-$2,000/mo	Agency-quoted SMB AI automation build + running cost	Agency blogs	https://www.solulab.com/ai-automation-cost-for-small-businesses/	2025-26	OPINION; sellers' price lists
48% / 13%	MSPs ranking AI as clients' top need / earning meaningful AI revenue	Kaseya State of the MSP (n>1,000)	https://www.itpro.com/technology/artificial-intelligence/ai-demand-soars-as-traditional-deal-sizes-shrink-kaseya-reports	2025-26	Vendor survey of MSPs. Search extract
~50%	MSPs feeling ready to guide SMBs in AI deployment	OpenText (n>1,000 MSPs)	https://s23.q4cdn.com/197378439/files/doc_news/OpenText-Cybersecurity-Finds-92-of-Managed-Service-Providers-See-AI-Driven-Growth-But-Readiness-Gap-Widens-2025.pdf	2025	Vendor survey. Search extract
42% to 28%	SMBs using AI, 2024 to 2025	NEXT Insurance (n=1,500)	https://stacker.com/stories/small-business/small-business-ai-adoption-declines-just-28	Apr 2025	Insurer survey; conflicts with all official series
14%		Goldman 10KSB via Fortune		Mar 2026	Programme panel

Figure	What it measures	Source	URL	Date	Reliability note
	Goldman 10KSB owners with AI embedded across core operations		https://fortune.com/2026/03/18/small-business-ai-slow-integration-across-operations/		
74% / 22%	SMBs using or testing AI / fully confident AI can do routine tasks alone	Bluevine (n=942)	https://www.bluevine.com/newsroom/bluevine-study-shows-ai-is-saving-nearly-half-of-smbs-4-hours-a-week	Apr 2026	Lender survey
+15% / -8%	Performance effect of GPT-4 mentor on high / low baseline performers; average effect not distinguishable from zero	Otis et al., Management Science (RCT, n=640 Kenyan entrepreneurs)	https://pubsonline.informs.org/doi/10.1287/mnsc.2024.06909	2025	Peer-reviewed RCT; figures vary by version

Track 3. Enterprise (from `04-enterprise.md`)

Figure	What it measures	Source	URL	Date	Reliability note
5% (95% "fail")	Orgs whose custom or embedded GenAI tools reached production with measurable P&L impact (funnel 60% evaluated, 20% piloted, 5% production)	MIT NANDA, GenAI Divide	https://getdx.com/blog/the-ai-divide/ (summary); PDF at mlq.ai	Jul/ Aug 2025	Weak: preliminary, 52 interviews + 153 surveys + ~300 public cases, ~6-month P&L window
30%	GenAI projects to be abandoned after PoC by end-2025	Gartner press release	https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025	29 Jul 2024	FORECAST
>40%	Agentic AI projects cancelled by end-2027	Gartner press release	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027	25 Jun 2025	FORECAST
~130	Vendors with genuine agentic capability, out of "thousands"	Gartner, same release	same	25 Jun 2025	Analyst estimate, no method
19% (n=3,412)	Orgs with significant agentic investment	Gartner webinar poll, Jan 2025	same	25 Jun 2025	Self-selected webinar audience
88% / 4 of 33	Observed AI POCs not reaching widescale deployment	IDC CIO Playbook 2025 (Lenovo-sponsored)	https://www.cio.com/article/3850763/88-of-ai-pilots-fail-to-reach-production-but-thats-not-all-on-it.html	Feb/ Mar 2025	Vendor-sponsored, n unverified
46%	AI POCs progressed into production	Lenovo CIO Playbook 2026 (IDC, n=3,120, fieldwork 16 Sep-17 Oct 2025)	https://www.businesswire.com/news/home/20260127255538/en/	27 Jan 2026	Vendor-sponsored
42% (from 17%)	Companies abandoning most of their AI initiatives	S&P Global Market Intelligence (n≈1,006, NA + Europe)	https://www.ciodive.com/news/AI-project-fail-data-SPGlobal/742590/	2025	Primary report not seen, fieldwork dates conflict
46%	Average share of AI POCs scrapped before production	S&P Global MI	same	2025	As above
39%	Respondents attributing any EBIT impact to AI (most <5%)	McKinsey State of AI 2025 (n=1,993)	https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai	Nov 2025	Self-reported, consultancy

Figure	What it measures	Source	URL	Date	Reliability note
≤10%	Max share scaling agents in any single function	McKinsey 2025	same	Nov 2025	Self-reported
40% (from 27%)	$1B+ revenue orgs scaling AI agents	McKinsey State of AI 2026 (n=1,719)	https://www.mckinsey.com/~/media/mckinsey/business%20functions/quantumblack/our%20insights/the%20state%20of%20ai/the-state-of-ai-in-2026-on-the-road-to-roi-aug-2026.pdf	Aug 2026	Via secondary summaries, verify in PDF
25%	Respondents who moved 40% + of AI pilots into production	Deloitte State of AI in the Enterprise 2026 (n>3,000)	https://www.deloitte.com/us/en/about/press-room/state-of-ai-report-2026.html	early 2026	Consultancy survey
5%	"Future-built" companies achieving AI value at scale	BCG Widening AI Value Gap (n>1,250)	https://web-assets.bcg.com/35/20/636ea73746f8bc91f106ca56ee7c/jpr-2501023-widening-ai-value-gap-vf.pdf	Sep 2025	Consultancy maturity typology
74%	Enterprise leaders reporting positive GenAI ROI	Wharton/GBK (n>800 US)	https://knowledge.wharton.upenn.edu/special-report/2025-ai-adoption-report	28 Oct 2025	Self-reported ROI
80%	"AI projects fail", cited by RAND	RAND RRA2680-1 citing "some estimates"	https://www.rand.org/pubs/research_reports/RRA2680-1.html	13 Aug 2024	Second-hand estimate, not RAND's measurement
30M+	Paid M365 Copilot seats (~6.6-6.7% of paid commercial seats, analyst estimate)	Microsoft Q4 FY26 earnings call	https://www.uctoday.com/unified-communications/microsoft-365-copilot-passes-30-million-paid-seats-as-cloud-and-ai-growth-power-record-quarter	Jul 2026	Company metric, secondary report
15M	Paid M365 Copilot seats, of ~450M commercial	Microsoft Q2 FY26	https://www.directionsonmicrosoft.com/blog/microsoft-claims-15-million-paid-m365-copilot-seats	28-29 Jan 2026	Company metric
9,500+ / >$500M	Paid Agentforce deals / Agentforce ARR	Salesforce Q3 FY26 release	https://www.salesforce.com/news/press-releases/2025/12/03/fy26-q3-earnings/	3 Dec 2025	Company metric
$1.7B	ServiceNow 2026 AI net new ACV target (raised from $1B)	ServiceNow Q3 2026 8-K	https://www.sec.gov/Archives/edgar/data/0001373715/000137371526000092/erq3fy26.htm	Oct 2026	Company target, not revenue
~17% to 21.3% (vs 91.2% Spider 1.0)	Best model success on 632 real enterprise text-to-SQL tasks	Spider 2.0, arXiv 2411.07763	https://arxiv.org/abs/2411.07763	Nov 2024 / v2 Mar 2025	Peer-reviewed benchmark; leaderboard later ~41%; search extract
~30%	Accuracy drop of commercial assistants and long-context LLMs on long-term memory	LongMemEval, arXiv 2410.10813 (ICLR 2025)	https://arxiv.org/abs/2410.10813	2024-25	Relative to oracle retrieval; search extract
7.2% / 5.5%	Open-source MCP servers with general vulnerabilities / tool poisoning (n=1,899)	Hasan et al., arXiv 2506.13538	https://arxiv.org/pdf/2506.13538v2	2025	Academic scan; search extract
9.3	CVSS score of EchoLeak (CVE-2025-32711), zero-click M365 Copilot prompt injection	Microsoft / Aim Security	https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html	Jun 2025	Vendor CVE; fixed, no known exploitation; search extract
~1,000	Asana customers potentially exposed by MCP server bug	Asana spokesperson to BleepingComputer	https://bleepingcomputer.com/news/security/asana-warns-mcp-ai-feature-exposed-customer-data-to-other-orgs/	Jun 2025	Company statement via press; Asana declined to confirm to other outlets

Figure	What it measures	Source	URL	Date	Reliability note
>10,000 / ~97M	Active public MCP servers (Dec 2025) / monthly SDK downloads (Mar 2026)	Anthropic, via secondary blogs	https://nordicapis.com/10-interesting-mcp-statistics/	2025-26	Primary not seen; downloads are not production use
~200,000	JPMorgan employees onboarded to LLM Suite	Bloomberg via American Banker	https://www.americanbanker.com/news/how-jpmorganchase-democratized-employee-access-to-gen-ai	2025	Company figure via press; ~half active daily
3B+ / ~50M	Erica cumulative client interactions / users since 2018	Bank of America press release	https://newsroom.bankofamerica.com/content/newsroom/press-releases/2025/08/a-decade-of-ai-innovation--bofa-s-virtual-assistant-erica-surpas.html	20 Aug 2025	Company-reported, cumulative
~46,000	Goldman staff given GS AI Assistant (after 10,000 pilot)	Goldman via press	https://www.americanbanker.com/news/goldman-sachs-staff-now-write-a-million-gen-ai-prompts-a-month	24 Jun 2025	Company figure via press
+70% / +40%	Walmart Sparky users YoY / spend per order of Sparky users	Walmart Q2 FY27 call via CIO Dive	https://www.ciodive.com/news/walmart-AI-wins-sparky-john-furner/812606/	Aug 2026	Company-reported; correlation, not causal; Feb 2026 figure was 35%
82.9%	Organisations with AI projects delayed by security/ compliance review	Protegrity / EMA	https://securitybrief.co.uk/story/enterprise-ai-deployments-slowed-by-security-reviews	Jul 2026	Vendor-sponsored; n not seen
~30%	Of delayed projects, share stalled 4+ months	Protegrity / EMA	same	Jul 2026	As above
~25% (~3 hrs/ week)	Reduction in email time from Copilot; no change in meeting time	Dillon et al., NBER WP 33795, RCT, 66 firms, 7,137 workers	https://www.nber.org/system/files/working_papers/w33795/w33795.pdf	2025	Randomised; Microsoft-affiliated authors; search extract
21% to 75%	Weekly Copilot use among treated workers, range across firms	Dillon et al.	same	2025	Randomised access, so variation is organisational
14-15% (34% novices)	Resolutions per hour from AI assistant, customer support	Brynjolfsson, Li, Raymond, QJE 2025 (n=5,172)	https://www.nber.org/system/files/working_papers/w31161/w31161.pdf	2025	Peer-reviewed quasi-experiment
26 min/ day	Self-reported time saved, UK cross-government Copilot experiment (20,000 staff)	GDS findings report	https://www.gov.uk/government/publications/microsoft-365-copilot-experiment-cross-government-findings-report/microsoft-365-copilot-experiment-cross-government-findings-report-html	Jun 2025	Self-report, not randomised
No evidence of productivity gain	DBT Copilot pilot (~1,000 licences)	UK DBT via The Register	https://www.theregister.com/2025/09/04/m365_copilot_uk_government	Sep 2025	Government evaluation; primary not seen
7,600+ / 1 in 3 daily	Australian public servants in Copilot trial / share using daily	DTA, Nous Group evaluation	https://www.digital.gov.au/initiatives/copilot-trial/summary-evaluation-findings/cts-executive-summary	Oct 2024	Non-random sample
122%-408%	Projected ROI range, Teams with M365 Copilot	Forrester TEI commissioned by Microsoft	https://tei.forrester.com/go/Microsoft/TeamsandCopilot/	2025	OPINION: vendor-commissioned projection on a composite org

Track 4. Inertia and education (from `05-inertia-and-education.md`)

Figure	What it measures	Source	URL	Date	Reliability note
0 valid studies	Empirical support for "70% of change initiatives fail" found by Hughes across 5 cited sources	Hughes, J. Change Mgmt 11(4)	https://research.brighton.ac.uk/en/publications/do-70-per-cent-of-all-organizational-change-initiatives-really-fa/	2011	Peer-reviewed; contested by Burnes & Jackson 2011. [search-summary]
78 %	AI users bringing their own AI tools to work (80% at SMBs)	Microsoft/ LinkedIn WTI 2024, n=31,000	https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part	8 May 2024	Vendor survey, online knowledge workers. Verified curl+grep.
52 %	AI users reluctant to admit using it for their most important tasks	WTI 2024	same	8 May 2024	Vendor survey. Verified.
53 %	AI users worried that using it on important tasks makes them look replaceable	WTI 2024	same	8 May 2024	Vendor survey. Verified.
39 %	AI users who have received AI training from their company	WTI 2024	same	8 May 2024	Vendor survey. Verified.
67 % vs 40 %	Leaders vs employees familiar with AI agents	WTI 2025, n=31,000	https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born	23 Apr 2025	Vendor survey. Verified.
43 %	Workers globally who fear AI will take their jobs	WTI 2025	same	23 Apr 2025	Vendor survey. Verified.
77 %	AI-using employees saying AI added to their workload	Upwork Research Institute	https://investors.upwork.com/news-releases/news-release-details/upwork-study-finds-employee-workloads-rising-despite-increased-c	23 Jul 2024	Vendor; base is about 1,250 non-executives in 4 countries. [search-summary]
52 % / 36 %	US workers worried / hopeful about AI at work	Pew, n=5,273, Oct 2024	https://www.pewresearch.org/internet/2025/02/25/u-s-workers-are-more-worried-than-hopeful-about-future-ai-use-in-the-workplace/	25 Feb 2025	Probability panel; strong. [search-summary]
50 % / 28 %	US employees using AI at work at all / frequently	Gallup, n=23,717, Feb 2026	https://www.tomshardware.com/tech-industry/artificial-intelligence/half-of-all-us-employees-now-use-artificial-intelligence-at-work-crossing-landmark-threshold-for-first-time-gallup-data-shows-daily-and-weekly-usage-hitting-all-time-high-of-28-percent-in-q1-2026-with-65-percent-feeling-positive-about-its-impact-on-productivity	Q1 2026	Strong panel; secondary report. [search-summary]
54.6%	US adults 18–64 who have used generative AI	Bick, Blandin & Deming, St Louis Fed	https://www.stlouisfed.org/on-the-economy/2025/nov/state-generative-ai-adoption-2025	13 Nov 2025	Nationally representative. [search-summary]
4.1 % to	Share of all US work hours spent using generative	same	same	13 Nov 2025	[search-summary]

Figure	What it measures	Source	URL	Date	Reliability note
5.7%	AI, Nov 2024 to Aug 2025				
about 3%	Average time saved by chatbot users in Denmark; earnings and hours effects about 0	Humlum & Vestergaard, NBER w33777	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	2025	Survey linked to registers; strong. [search-summary]
about 20%	How much less likely women are than men to use generative AI	Otis, Cranney, Delecourt & Koning	https://www.hks.harvard.edu/centers/cid/publications/global-evidence-gender-gaps-and-generative-ai	Oct 2024 (rev.)	Meta-analysis of 16–18 studies, more than 140,000 people. [search-summary]
51%	Frontline (white-collar) employees who are regular AI users; stalled	BCG AI at Work 2025	https://www.bcg.com/press/26june2025-beyond-ai-adoption-full-potential	26 Jun 2025	Consultancy survey. [search-summary]
15% to 55%	Positive sentiment toward generative AI, without vs with strong leadership support	BCG AI at Work 2025	same	26 Jun 2025	Cross-sectional. [search-summary]
27% to 22%	Global manager engagement, 2024 to 2025	Gallup State of the Global Workplace 2026	https://www.unleash.ai/strategy-and-leadership/analysis/gallups-state-of-the-global-workplace-2026-report-three-key-decisions-for-hr-leaders	Apr 2026	Secondary. [search-summary]
21%	Generative AI users who have fundamentally redesigned at least some workflows	McKinsey State of AI	https://artificialintellitools.blogspot.com/2025/04/briefing-document-state-of-ai-how.html	Mar 2025	Consultancy; secondary blog. [search-summary]
15.9%	US TFP level above official figures by end-2017 after adjusting for IT intangibles	Brynjolfsson, Rock & Syverson, AEJ Macro	https://www.aeaweb.org/doi/10.1257/mac.20180386	2021	Peer-reviewed. [search-summary]
69%	Professionals interviewed who mentioned social stigma of AI use at work	Anthropic Interviewer, n=1,250	https://www.anthropic.com/research/anthropic-interviewer	4 Dec 2025	Vendor; crowdworker-recruited sample. Verified.
57%	Employees who hide AI use and present AI output as their own	KPMG / Univ. of Melbourne 2025, n≈48,000, 47 countries	https://mbs.edu/news/global-study-reveals-trust-of-ai-remains-a-critical-challenge	Apr 2025 (fieldwork Nov 2024 to Jan 2025)	Academic-led, Big Four sponsor. Search extract
~50%	Employees using AI in ways that contravene company policy	same	same	2025	"Almost half". Search extract
47% / 40%	Employees with AI training / with a workplace genAI policy	same	same	2025	Search extract
		Gallup			

Figure	What it measures	Source	URL	Date	Reliability note
80 % vs 44 %	Frequent AI use where managers actively support vs low support (private sector, AI-adopting orgs)		https://www.gallup.com/workplace/699689/ai-use-at-work-rises.aspx	late 2025	Cross-sectional. Search extract
1.7 x / 36 %	Frequent-use multiplier with strong manager support / share with strong support	Gallup AI indicator	https://www.gallup.com/699797/indicator-artificial-intelligence.aspx	May 2026	Search extract
21 % to 75 %	Weekly Copilot use among treated workers, range across 66 firms	Dillon et al., NBER w33795, n=7,137	https://www.nber.org/system/files/working_papers/w33795/w33795.pdf	2025	Randomised access; firm effects. Microsoft co-authors
68 %	Gap RCT participants at maximum quality score (all arms), so no ITT difference was detectable	arXiv 2604.08678, n=388	https://arxiv.org/abs/2604.08678	Apr 2026	Preprint, Microsoft authors, no IRB
10.9 to 12.1	Average direct reports per manager, 2024 to 2025	Gallup via Fast Company	https://www.fastcompany.com/91601567/the-great-flattening-rolls-on-as-uber-lays-off-middle-managers	2025-26	Secondary
-42 %	Middle-manager job postings, Apr 2022 peak to Oct 2025	Revelio Labs via secondary	https://pipedot.org/6SNX5	late 2025	Postings, not employment. Secondary
40 %	US employees feeling lack of direction due to fewer middle managers	Korn Ferry 2025	https://www.cfodive.com/news/cuts-middle-management-stressing-72-percent-senior-executives-korn-ferry-HR-WFH/745822/	2025	Consultancy survey
7 / 0.27	Claude usage per working-age capita relative to expected (Israel / India)	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	15 Sep 2025	One vendor's consumer users. Verified.

Track 5. What goes wrong (from `06-what-goes-wrong.md`)

Figure	What it measures	Source	URL	Date	Reliability note
65	Data scientists/engineers interviewed for RAND root-cause study	RAND RR-A2680-1	https://www.rand.org/pubs/research_reports/RRA2680-1.html	2024-08-13	Qualitative interviews; seen via search summary
>80 %	"Some estimates" of AI project failure, cited by RAND	RAND RR-A2680-1	https://www.rand.org/pubs/research_reports/RRA2680-1.html	2024-08-13	Secondhand within RAND; not a RAND measurement
30 to 50%	"As many as" this share of initial RPA projects fail	EY, Get ready for robots	https://www.eyfinancialservicesthoughtgallery.ie/wp-content/uploads/2016/11/ey-get-ready-for-robots.pdf	2016	Consultancy opinion, no method; origin of viral stat
~1,598		Charlotin database via	https://haqq.ai/blog/ai-legal-hallucination-audit.html	2026-06-09	

Figure	What it measures	Source	URL	Date	Reliability note
	Court decisions noting AI-hallucinated content (cumulative, worldwide)	HAQQ write-up			Third-party count; check live database
~8/day	New hallucination cases added (implied rate)	HAQQ write-up	https://haqq.ai/blog/ai-legal-hallucination-audit.html	2026-06	Third-party estimate
A$440,000	Deloitte DEWR report contract value (final instalment refunded)	The Register	https://www.theregister.com/2025/10/06/deloitte_ai_report_australia/	2025-10-06	Seen via search summary; refund amount unconfirmed
45	CBA roles cut then reinstated after FSU challenge	Bloomberg	https://www.bloomberg.com/news/articles/2025-08-21/commonwealth-bank-reverses-job-cuts-decision-over-ai-chatbots	2025-08-21	Seen via search summary
~2,000/week	Call reduction CBA claimed for voice bot (disputed by FSU)	ACS Information Age	https://ia.acs.org.au/article/2025/cba-reverses-ai-driven-job-cuts--admits--error-.html	2025-08	Company claim, disputed
~100	McDonald's restaurants in IBM voice-ordering test, ended by 26 Jul 2024	Nation's Restaurant News	https://www.nrn.com/quick-service/mcdonald-s-is-ending-its-ai-drive-thru-test-with-ibm	2024-06	Seen via search summary
500+ / 2m+	Taco Bell drive-thrus with voice AI / orders handled	Jalopnik (citing WSJ)	https://www.jalopnik.com/1956939/taco-bell-drive-through-18000-waters/	2025-08	Secondary of WSJ
~$304m	Zillow Offers inventory write-down, Q3 2021	Inman	https://www.inman.com/2021/11/02/zillow-to-shut-down-ibuying-program-zillow-offers/	2021-11-02	Some reports say $340m
~25%	Zillow workforce cut with Offers wind-down	CNN	https://edition.cnn.com/2021/11/09/tech/zillow-ibuying-home-zestimate	2021-11-09	Seen via search summary
$62.1m	MD Anderson spend on Watson Oncology Expert Advisor (IBM ~$39.2m, PwC ~$23m)	The Register on UT System audit	https://www.theregister.com/2017/02/20/watson_cancerbusting_trial_on_hold_after_damning_audit_report/	2017-02-20	Audit figure; PwC split varies by outlet
~1.1 billion	Applications Workday said were rejected using its tools in the class period	Holland & Knight	https://www.hklaw.com/en/insights/publications/2025/05/federal-court-allows-collective-action-lawsuit-over-alleged	2025-05-27	Scale figure from filings, not a discrimination finding
9.3	CVSS score of EchoLeak (CVE-2025-32711)	The Hacker News	https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html	2025-06	Microsoft rating
0.692 → 0.462	τ-bench retail pass^1 → pass^4, best listed model (claude-3-5-sonnet-20241022)	tau-bench repo	https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md	2024 to 2025	Fetched; primary
0.604 → 0.383	τ-bench retail pass^1 → pass^4, gpt-4o	tau-bench repo	https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md	2024	Fetched; primary
75+	Task fixes Sierra made to its own benchmark tasks in τ³-bench	tau2-bench repo	https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md	2026	Fetched; primary
30.3%	Best full-task completion on TheAgentCompany (Gemini 2.5 Pro), v2	alphaXiv listing of arXiv 2412.14161v2	https://www.alphaxiv.org/abs/2412.14161v2	2025-05	Seen via search summary
~7 months	METR 50% time-horizon doubling time (2019 to 2025)	METR repo README	https://raw.githubusercontent.com/METR/eval-analysis-public/main/README.md	2025	Fetched; primary
4x to 10x	Ratio of 50% to 80% time horizon across 20 models	My refit of METR TH1.1 runs	https://raw.githubusercontent.com/METR/eval-analysis-public/main/reports/time-horizon-1-1/data/raw/runs.jsonl	computed 2026-11-06	Own approximate analysis; not METR's published figures

Figure	What it measures	Source	URL	Date	Reliability note
$5,000	Rule 11 sanction in Mata v. Avianca (total vs per lawyer disputed in reports)	S.D.N.Y. Opinion and Order on Sanctions	https://law.justia.com/cases/federal/district-courts/new-york/nysdce/1:2022cv01461/575368/54	2023-06-22	Primary court document; search extract, not fetched
~50-60%	Average conformance between designed and actual purchase-order process	Mavim vendor case study	https://www.mavim.com/hubfs/Assets/Miscellaneous/Case%20study%20-%20Conformance%20checking%20-%20EN.pdf?hsLang=en	n.d.	Vendor case; search extract
600+	Variants of one process in one business group	Mavim, same	same	n.d.	Vendor case; depends on log granularity
18% / ~38 FTE	P2P transactions with manual changes / staff tied up in top-5 rework activities	Celonis / AkzoNobel case (IEEE TF PM site)	https://www.win.tue.nl/ieeetfpm/casestudies/celonis_akzonobel_-_customer_success_story.pdf	n.d. (c. 2010s)	Vendor case; search extract
75+ / 50+	Contributors / organisations behind the Process Mining Manifesto	IEEE Task Force on Process Mining	https://research.wu.ac.at/en/publications/process-mining-manifesto-3/	2011	Bibliographic; search extract
84% → 51%	GPT-4 prime-identification accuracy, March vs June 2023	Chen, Zaharia, Zou	https://arxiv.org/abs/2307.09009	2023-07	Seen via search summary; interpretation contested

Track 6. Working with agents (from `07-working-with-agents.md`)

Figure	What it measures	Source	URL	Date	Reliability note
86%	Share of analytical choices where analysts followed a coin-flipped agent lean (50% = no effect)	Anthropic, Herdağdelen, Joren, Tamkin	https://www.anthropic.com/research/agents-judgment-calls	2026-10-27	RCT, n=50 analysts, vendor-run, lab task
60% vs 6%	Switch rate away from own stated preference when agent disagreed vs agreed	Same	https://www.anthropic.com/research/agents-judgment-calls	2026-10-27	Only 38 cases; CI on gap 24 to 82 pp
77% / 15%	Analysts reusing agent arguments / crediting the agent	Same	https://www.anthropic.com/research/agents-judgment-calls	2026-10-27	Self-reported explanations
93%	Share of manual permission prompts users accept in Claude Code	Anthropic engineering	https://www.anthropic.com/research/claude-code-auto-mode	2026-03-25	Vendor product data, method not detailed
17%	False-negative rate of auto-mode classifier on real overeager actions	Same	https://www.anthropic.com/research/claude-code-auto-mode	2026-03-25	n=52 real cases, small
~20% to >40%	Sessions with full auto-approve, new (<50 sessions) vs experienced (750) Claude Code users	Anthropic	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	Vendor telemetry; default is manual approval
5% to ~9%	Per-turn interrupt rate, new vs experienced users	Same	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	Vendor telemetry

Figure	What it measures	Source	URL	Date	Reliability note
73% / 80% / 0.8%	API tool calls with a human in the loop / with at least one safeguard / irreversible	Same	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	Inferred by Claude from single tool calls; cannot see downstream review
~50%	Software engineering share of agentic tool calls on Anthropic public API	Same	https://www.anthropic.com/research/measuring-agent-autonomy	2026-02-18	n=998,481 tool calls
70% / 20%	User share of planning / execution decisions in Claude Code sessions	Anthropic	https://www.anthropic.com/research/claude-code-expertise	2026-06-16	~400k sessions, classifier-based
88% vs 89%	At least partial success, non-software vs software occupations, code-producing sessions	Same	https://www.anthropic.com/research/claude-code-expertise	2026-06-16	Success partly user-confirmed
77% vs 12%	Automation vs augmentation share of 1P API transcripts	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	2025-09-15	Share of transcripts, not of businesses; later reports show decline
27% to 39%	Directive (full delegation) share of Claude.ai conversations, late 2024 to Aug 2025	Same	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report	2025-09-15	Reversed by Nov 2025 data
52% / 45%	Augmented / automated share of Claude.ai conversations, Nov 2025	Anthropic Economic Index	https://www.anthropic.com/research/anthropic-economic-index-january-2026-report	2026-01-15	Vendor telemetry
0.24 SD	Team advantage over individuals without AI; closed by individual with AI	Dell'Acqua et al., NBER 33641	https://www.nber.org/papers/w33641	2025-03	n=776, one-day task; snippet only this session
12-16 %	Time saved by AI groups, P&G experiment	Same	https://www.nber.org/papers/w33641	2025-03	Snippet only
14% (WP) / 15% (QJE)	Productivity gain, customer support agents	Brynjolfsson, Li, Raymond	https://econpapers.repec.org/RePEc:nbr:nberwo:31161	2023 / 2025	One firm, ~5,170 agents; snippet only
3 hours, 25%	Weekly email time reduction for Copilot users (ITT 1.4 h); no change in meetings	Dillon et al., Microsoft	https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/	2025-04	RCT, 6,000 workers, 56 firms; Microsoft authors
46.1%	Tasks where workers were positive about agent automation	Shao et al., Stanford WORKBank	https://arxiv.org/abs/2506.06576	2025-06	Survey, 1,500 workers; snippet only
>20%	Inference cost increase from repository context files, with lower or no better success	Gloaguen et al.	https://arxiv.org/abs/2602.11988	2026-02-12	Snippet only; check full paper for effect sizes
-28.64%	Median agent runtime with AGENTS.md (124 PRs, 10 repos)	Lulla et al.	https://arxiv.org/abs/2601.20404	2026-01-28	Small, efficiency only; snippet only
				2025-12-09	

Figure	What it measures	Source	URL	Date	Reliability note
60,000+	Projects reported to use AGENTS.md	AAIF launch coverage	https://sdtimes.com/ai/linux-foundation-forms-agentic-ai-foundation-to-be-new-home-for-mcp-goose-and-agents-md/		Self-reported by backers; another outlet said 40,000
1M / ~280M	Agent Skills published / installed	Anthropic	https://www.anthropic.com/news/donating-agent-skills-to-the-agentic-ai-foundation	2026-10-22	Vendor claim, no method
10,000+ / 97M+	Active public MCP servers / monthly SDK downloads	Anthropic	https://www.anthropic.com/news/donating-the-model-context-protocol-and-establishing-of-the-agentic-ai-foundation	2025-12-09	Vendor claim; downloads include CI
~9,000 to ~5,000	Salesforce support headcount per CEO	Benioff on Logan Bartlett Show via Fortune	https://dc.fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales	2025-09-02	Spoken claim, no timeframe; snippet only
"hundreds"	IBM HR roles replaced by AI agents	Krishna to WSJ via Entrepreneur	https://www.entrepreneur.com/business-news/ibm-ceo-ai-replaced-hundreds-of-human-resources-staff/491341	2025-05	CEO statement; snippet only
33% / 78%	Leaders considering headcount cuts / AI-specific hires	Microsoft WTI 2025	https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born	2025-04-23	Vendor survey, n=31,000, intentions only
CAD ~812	Total award in Moffatt v Air Canada	2024 BCCRT 149 via McCarthy Tétrault	https://www.mccarthy.ca/en/insights/blogs/techlex/moffatt-v-air-canada-misrepresentation-ai-chatbot	2024-02	Snippet only; not binding precedent
>=6 months	Minimum log retention for deployers of high-risk AI (AI Act Art. 26)	Regulation (EU) 2024/1689 via Mishcon de Reya	https://mishcon.com/eu-uk-ai-navigator/article-26	2024 (applies Dec 2027 for Annex III after omnibus)	Search extract; check EUR-Lex
9 Dec 2026	Transposition deadline, Product Liability Directive 2024/2853 (software and AI as products)	Norton Rose Fulbright	https://connections.nortonrosefulbright.com/post/102jpjp/revised-product-liability-directive-introducing-rules-on-strict-liability-for-ai	OJ 18 Nov 2024	Search extract
6 Oct 2025	Reported OJ date of formal AI Liability Directive withdrawal (announced Feb 2025)	EAPIL	https://eapil.org/2025/10/09/european-commission-withdraws-two-proposals-assignments-of-claims-regulation-and-ai-liability-directive/	9 Oct 2025	Single source for the exact date
$25m	Per-insured limit, Armilla Lloyd's-backed AI liability cover (launched 30 Apr 2025)	Armilla	https://www.armilla.ai/resources/armilla-ai-raises-lloyds-backed-coverage-to-25m-as-traditional-insurers-retreat-from-ai-risk	2025-2026	Vendor announcement; no claims data
50 (11 in 2025)	ISO/IEC 42001 certificates issued by BSI	BSI	https://www.bsigroup.com/en-GB/insights-and-media/media-centre/press-releases/2026/november/bsi-awards-its-50th-iso-iec-42001-certificate/	Nov 2026	One certifier only; no global registry
~1.19M /	GitHub code-search hits for files named	GitHub search API,	https://api.github.com/search/code?q=filename:AGENTS.md	2026-11-06	DO NOT CITE: GitHub's total_count is approximate and includes forks

Figure	What it measures	Source	URL	Date	Reliability note
~1.03M	AGENTS.md / CLAUDE.md	queried by this agent			and vendored copies; same query gave ~7M for SKILL.md, which looks implausible

Track 7. Displacement (from `08-displacement.md`)

Figure	What it measures	Source	URL	Date	Reliability note
No significant rise; detectable effect about 1pp	Differential unemployment, top-quartile exposed vs zero-exposure US workers, post-ChatGPT	Massenkoff and McCrory, Anthropic	https://www.anthropic.com/research/labor-market-impacts	2026-03-05	Fetched. CPS DiD. Author is an AI vendor
-14%	Job-finding rate for 22-25 year olds into most exposed occupations vs 2022	Same	https://www.anthropic.com/research/labor-market-impacts	2026-03-05	Fetched. "Just barely statistically significant"
75% / 67%	Claude task coverage: computer programmers / customer service reps and data entry keyers	Same	https://www.anthropic.com/research/labor-market-impacts	2026-03-05	Fetched. Exposure, not effect
about 13% (Aug 2025)	Relative employment decline, 22-25 in most exposed occupations	Brynjolfsson, Chandar, Chen, as cited by Hosseini and Lichtinger	https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf	2025-08	Fetched secondary citation; ADP data; descriptive
19% (by June 2026)	Same gap, Aug 2026 revision	Stanford Digital Economy Lab	https://digitaleconomy.stanford.edu/news/canariesaug26/ ; https://digitaleconomy.stanford.edu/app/uploads/2026/08/Canaries_August2026.pdf	2026-08-12	Search extract; now corroborated by a second result (PDF + aiweekly). Revision gives 15% for the July 2025 vintage. 4.6m workers
about 0.1pp	Max contribution of young exposed workers' employment decline to US unemployment rise since Nov 2022	Dallas Fed, Atkinson and Yamco	https://www.dallasfed.org/research/economics/2026/0106	2026-01-06	Search extract. Upper-bound scale check
about 1% / 35%	NY-region service firms reporting AI layoffs / AI-using service firms retraining	NY Fed regional surveys	https://libertystreeteconomics.newyorkfed.org/2025/09/are-businesses-scaling-back-hiring-due-to-ai/	2025-09-04	Search extract. Regional firm self-report
0.47-0.57	Correlation of occupational AI exposure/adoption with unemployment increase 2022-25	St. Louis Fed, Ozkan and Sullivan	https://fedinprint.org/item/l00001/101478	2025-08-26	Search extract. Non-causal; summaries disagree on figure
5.7% (Q1 2026)	US recent-graduate (22-27) unemployment	NY Fed college labor market	https://www.newyorkfed.org/research/college-labor-market	2026	Search extract; Q2 reported as 5.6%
about 7% of hires	New-grad share of hires at largest tech firms, down >50% vs 2019, about 25% vs 2023	SignalFire State of Tech Talent 2025	https://www.signalfire.com/blog/signalfire-state-of-talent-report-2025	2025-05	Search extract. VC firm's own data
-15 to -16%	US entry-level postings on Handshake, Aug 2024 to Aug 2025;	Handshake via Fortune	https://fortune.com/2025/11/19/next-level-leadership-gen-z-college-graduates-tough-	2025-11	Search extract; secondary

Figure	What it measures	Source	URL	Date	Reliability note
	applications per post +26-30%		market-how-to-stand-out-according-to-jobs-experts/		
40% / 60%	Share of jobs exposed to AI, world / advanced economies	IMF SDN/2024/001 (Cazzaniga et al.)	https://wiki.ubc.ca/images/b/bd/IMF_GenAI_Future_of_Work.pdf	2024-01-14	FORECAST (exposure, not job loss). Search extract
6-7% (range 3-14%)	Share of US workers eventually displaced by GenAI over adoption period; peak unemployment +0.5pp	Goldman Sachs (Hatzius; Briggs and Dong)	https://www.staffingindustry.com/news/global-daily-news/ais-impact-on-jobs-likely-temporary-goldman-sachs	2025	FORECAST. Secondary; Goldman primary not seen
-15% / -10% / -6%	UK online adverts over about 3 years, high / mid / low AI-exposure occupations	Bank of England staff (Bank Underground, Schlicht)	https://bankunderground.co.uk/2026/08/06/canaries-in-the-column-ai-exposure-and-the-uks-hiring-slowdown/	2026-08-06	Search extract. Staff blog; postings; not causal; pre-ChatGPT trend disputed
-7% y/y	UK graduate postings, lowest for time of year since 2020	Indeed Hiring Lab UK	https://www.hiringlab.org/uk/?p=14268	2026-08	Search extract; postings
-31.9%	UK entry-level adverts since Nov 2022	Adzuna	https://hrreview.co.uk/hr-news/recruitment/entry-level-jobs-have-declined-since-chatgpt-launch	2025-06	Search extract; vendor postings; no control group
-8% (2024/25), +4% (2025/26)	Graduate hiring by ISE member employers	Institute of Student Employers	https://ise.org.uk/_userfiles/pages/files/reports/student_recruitment_survey_2025.pdf	2025 / 2026	Search extract. 155 large employers, about 31k hires
6%	UK AI-using firms (10+ staff) saying AI reduced headcount	ONS BICS via secondary site	https://www.ons.gov.uk/releases/aiinukbusinesses	2026-09	UNVERIFIED secondary figure; not seen in ONS bulletin
about 15%	Euro-area AI-using firms citing lower labour costs as a motive; AI-intensive firms more likely to hire	ECB blog, Lebastard and Sondermann (SAFE, about 5,000 firms)	https://www.ecb.europa.eu/press/blog/date/2026/html/ecb.blog20260304~d9e34fc95f.en.html	2026-03-04	Search extract. Firm self-report
+1.3pp	Expected 12-month employment growth, significant AI users vs non-users, euro area	ECB Economic Bulletin 2/2026 box	https://ecb.europa.eu/press/economic-bulletin/focus/2026/html/ecb.ebbox202602_06~c9f485c478.en.html	2026	Search extract. Expectations
about -9% after 6 quarters	Junior employment at GenAI-adopting firms vs non-adopters	Hosseini and Lichtinger	https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf	2026-05-06	Fetched. 66m workers, 281k firms; adoption proxied by postings
about 80%	Fall in junior hires per quarter at adopters vs pre-period mean (4.0 of 5.06)	Same	same	2026-05-06	Fetched
3.7% / 16%	Share of firms classed as GenAI adopters / their share of employment	Same	same	2026-05-06	Fetched. Adoption by Mar 2025
-5.5%	22-25 employment in high-exposure occupations within Swedish employers	Lodefalk et al. 2026, cited by Hosseini and Lichtinger	same	2026	Fetched secondary citation; original not seen

Figure	What it measures	Source	URL	Date	Reliability note
>35%	Fall in US entry-level postings since Jan 2023	Revelio Labs (Simon 2025), cited by Hosseini and Lichtinger	https://www.reveliolabs.com/news/macro/is-ai-responsible-for-the-rise-in-entry-level-unemployment/	2025	Vendor data; secondary citation
-21%	Freelance writing and coding posts vs manual posts, 8 months after ChatGPT	Demirci, Hannane, Zhu, Management Science	https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb	2025	Search summary; 1.39m posts
-5.2%	Monthly earnings of Upwork writing freelancers after ChatGPT	Hui, Reshef, Zhou	https://www.cesifo.org/DocDL/cesifo1_wp10601.pdf	2023 WP	Search summary
<1%	Upper bound on earnings/hours effects, Danish exposed occupations	Humlum and Vestergaard, NBER 33777	https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf	2025	Search summary; 25k workers, 7k workplaces
<2%	Same bound, two years after ChatGPT, 2026 version ("Still Waters, Rapid Currents")	Humlum and Vestergaard, RFBerlin DP 26078	https://ideas.repec.org/p/crm/wpaper/26078.html	2026-03	Search extract (abstract)
121,401 (about 20%) [DO NOT CITE: unconfirmed, see verification.md]	US announced job cuts citing AI, Jan-Oct 2026	Challenger, Gray & Christmas	https://www.challengergray.com/blog/october-layoffs-at-three-year-low-seasonal-hiring-plans-up-26/	2026-11-05	Search summary; announcements, company-stated reasons
54,836	US announced job cuts citing AI, full year 2025	Challenger (via secondary)	https://aiunderstanding.org/news/challenger-says-employers-cited-ai-in-112-713-u-s-job-cuts-through-july-2026	2026-08	Search summary; secondary
about 1.1%/month, 13.4% annualised	US layoffs and discharges rate, 12 months to Apr 2026 (BLS JOLTS)	Anthropic citing BLS	https://www.anthropic.com/research/economic-index-june-2026-report	2026-06-26	Fetched; check against BLS directly
net -8% (23% cut, 15% created)	Headcount change reported by UK firms using AI 1+ year, five sectors	Morgan Stanley via eWeek/ITPro	https://www.eweek.com/news/ai-job-losses-uk-morgan-stanley-report/	2026-01	Search summary; non-representative firm survey; NOT 8% of UK jobs
5.6%	Unemployment, US recent college grads 22-27, Q1 2026	NY Fed via Bloomberg Law	https://news.bloomberglaw.com/new-york-brief/unemployment-for-recent-college-grads-remains-high-ny-fed-says	2026	Search summary
12.2%	2030 unemployment forecast by US workers expecting AI to do 80% + of their work	Anthropic / Morning Consult, n=10,980	https://www.anthropic.com/research/ai-expectations-gap	2026-10-16	Fetched. Expectation, not outcome
10%	Claude users rating own job loss in next year likely/very likely	Anthropic Economic Index Survey, n≈9,700	https://www.anthropic.com/research/economic-index-june-2026-report	2026-06-26	Fetched. Not representative

Track 8. What will people do (from `09-what-will-people-do.md`)

Figure	What it measures	Source	URL	Date	Reliability note
~60% (some versions 63%)	Share of 2018 US employment in job titles that did not exist in 1940	Autor, Chin, Salomons, Seegmiller, QJE	https://economics.mit.edu/sites/default/files/2026-04/New-Frontiers-QJE-v139%283%29-p1399-1465-Aug-2024.pdf	Aug 2024	Peer-reviewed. Search summary only; not checked against PDF
41% to 1.9%	Agriculture share of US workforce, 1900 to 2000	USDA ERS (compiling HSUS and Census)	https://webharvest.gov/peth04/20041105150816/http://www.ers.usda.gov/publications/agoutlook/jun2000/ao272h.pdf	Jun 2000	Official, but seen via repost; other series give 2.4% for 2000
>25% fall; vs ~780 hrs	Fall in men's annual market hours 1931-2011, vs Keynes's implied target	Crafts, CEPR DP16140	https://repec.cepr.org/repec/cpr/ceprdp/DP16140.pdf	2021	Working paper; search summary only
47%	Share of US employment at "high risk" of computerisation	Frey and Osborne via ITIF	https://itif.org/publications/2022/09/30/oops-the-predicted-47-percent-of-job-loss-from-ai-didnt-happen/	2013 (forecast); 2022 (review)	Forecast; ITIF is an advocacy think tank
+16.4%	Employment growth 2013-2021 in insurance underwriters (highest F&O risk)	ITIF	same as above	Sep 2022	Advocacy source using BLS data
-0.2 pp; -0.42%	Employment-to-population and wage effect per extra robot per 1,000 workers	Acemoglu and Restrepo, JPE	https://econpapers.repec.org/RePEc:ucp:jpolec:doi:10.1086/705716	2020	Peer-reviewed; other versions give different magnitudes
>50%	Share of US phone network AT&T automated 1920-1940	Feigenbaum and Gross, QJE	https://econpapers.repec.org/RePEc:oup:qjecon:v:139:y:2024:i:3:p:1879-1939.	Aug 2024	Peer-reviewed; search summary only
~1.3 hrs/week; ~2 pp	Fall in hours and employment, $1,000/month recipients vs controls	OpenResearch / Vivalt et al., NBER WP 32719 (via Fortune)	https://www.nber.org/system/files/working_papers/w32719/w32719.pdf	Jul 2024	RCT; figures from secondary coverage; a conference version reports 3.9 pp participation fall; reconcile
28% to 40% (control 32% to 37%)	Full-time employment, Stockton SEED recipients, Feb 2019 to Feb 2020	Martin-West and Castro Baker	https://catalog.results4america.org/case-studies/guaranteed-income-stockton	Mar 2021	n=125 treated; small
€560/month, ~2,000 people	Finland basic income experiment design	Finnish Ministry of Social Affairs and Health / Kela	https://stm.fi/en/-/perustulokokeilun-tulokset-tyollisyysvaikutukset-vahaisia-toimeentulo-ja-psyykkinen-terveys-koettiin-paremmaksi	May 2020	RCT; employment effect "small"
0 employment effect; +1.8 pp (17%) part-time	Alaska PFD labour effects	Jones and Marinescu, AEJ: Policy	https://www.aeaweb.org/doi/10.1257/pol.20190299	2022	Peer-reviewed; synthetic control; part-time figure conflicts with one text fragment
61 firms, ~2,900 staff; 56 continue, 18 permanent	UK four-day-week pilot	Autonomy / 4 Day Week Global, via WEF	https://www.weforum.org/stories/2023/03/four-day-work-week-uk-trial/	Feb 2023	Self-selected firms; advocacy-run
>2,500 workers, ~1% of workforce, 35-36 hrs	Iceland public-sector hours-reduction trials	Autonomy and Alda	https://issuu.com/4dayweekglobal/docs/iceland_4dw	Jun 2021	Advocacy report; not a four-day week

Figure	What it measures	Source	URL	Date	Reliability note
86%	Icelandic workers on contracts with shortened hours or the right to shorten	Autonomy and Alda	https://issuu.com/4dayweekglobal/docs/iceland_4dw	Jun 2021	Advocacy report; size of cuts varies
10-15% per year; 1.0-1.5 yrs	Mortality hazard increase 20 years after displacement; implied life-expectancy loss	Sullivan and von Wachter, QJE	https://econpapers.repec.org/RePEc:oup:qjecon:v:124:y:2009:i:3:p:1265-1306.	2009	Peer-reviewed; 50-100% first-year figure is from a draft only
~30% vs ~1%	Share of lost work hours going to sleep plus TV vs job search, Great Recession	Aguiar, Hurst, Karabarbounis (NBER Digest)	https://www.nber.org/digest/nov11/time-use-during-recessions	Nov 2011	ATUS-based; secondary summary
N=3,562; ~-11% motivation, ~+20% boredom	GenAI collaboration followed by solo work	Scientific Reports (4 experiments)	https://fortune.com/2025/06/03/ai-boost-worker-productivity-hurts-motivation-boredom	Apr 2025 (paper)	Online experiments; effect sizes from secondary coverage
3-7%	Share of chatbot productivity gains passed through to earnings	Humlum and Vestergaard, NBER WP 33777	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	2025	Danish register data, 25,000 workers; working paper
~5,671,836	US business applications, 2025 (annual record)	Census BFS (annual sum via secondary)	https://www.census.gov/econ/bfs/pdf/historic/bfs_2025m12.pdf	Jan 2026	Official monthly data; annual total from secondary sum
37.6% to 30.1%	High-propensity share of US business applications, 2019 to 2025	Consultant calculation from Census BFS	https://founderreports.com/how-many-businesses-are-started-each-year/	2026	Secondary calculation
29.8m (2022); 30.4m (2023)	US nonemployer businesses	Census NES / NES-D	https://www.census.gov/library/stories/2025/07/nonemployer-business-growth.html	Jul / Nov 2025	Official; search extract
4.9%; 4.7%	Nonemployer growth 2021, 2022 (employer: 1.9%, 1.8%)	Census	https://www.census.gov/library/stories/2025/07/nonemployer-business-growth.html	Jul 2025	Official; pre-AI period
~4.38m vs ~5.0m	UK self-employed, Q4 2025 vs end-2019	ONS LFS via Statista	https://statista.com/statistics/318234/united-kingdom-self-employed	2026	LFS volatile; series disagree
~$401m; ~$65m	Medvi 2025 revenue; net profit	NYT via Dealroom	https://dealroom.co/news/127216-from-20k-to-1-8b-in-revenue-how-one-founder-built-a-telehealth-giant-wit/	Apr 2026	Journalist saw financials; contested; $1.8bn is a projection
~$80m + earn-out	Wix purchase of Base44 (6-8 staff, ~6 months old)	Globes	https://en.globes.co.il/en/article-wix-acquires-israeli-vibe-coding-co-base44-1001513267	Jun 2025	Press report
6.2%; ~96,000	England adult social care vacancy rate and vacancies, 2025/26	Skills for Care via LocalGov	https://www.localgov.co.uk/Social-care-vacancy-rate-falls-to-decade-low-but-challenges-ahead/64563	Jun 2026	Sector data body; secondary report
£12.60 (39p above NLW)	Median care worker hourly pay, England, Dec 2025	Skills for Care via Healthcare Management	https://www.healthcare-management.uk/care-workers-pay-minimum-wage-reveals	2026	Sector data; NLW rose to £12.71 in Apr 2026
$34,900; +17%; 765,800/yr	US home health and personal care aides: median pay 2024; growth 2024-34; annual openings	BLS OOH	https://www.bls.gov/ooh/healthcare/home-health-aides-and-personal-care-aides.htm	2025	Official; search extract
45 firms, ~900 staff;	Germany four-day-week pilot 2024	University of Münster /	https://www.thelocal.de/20241106/what-are-the-results-	Nov 2024	Self-selected; ~20% reverted

Figure	What it measures	Source	URL	Date	Reliability note
~73% continue		Intraprenör via The Local	of-germanys-four-day-work-week-trial		
9.5%	Traffic fall during Valencia's four Monday holidays, 2023	City-commissioned evaluation via The Mayor	https://www.themayor.eu/en/a/view/valencia-s-government-says-4-day-week-trial-was-good-for-the-people-and-the-city-12172	2023	Four weeks; holidays not work redesign

Track 9. Social and political (from `10-social-political.md`)

Figure	What it measures	Source	URL	Date	Reliability note
52%	US adults more concerned than excited about AI in daily life (37% in 2021)	Pew, n=3,488, Jun 22-28 2026	https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/	18 Aug 2026	Probability panel; seen via search summary only
71%	US adults expecting AI to mean fewer jobs over 20 years (64% in 2024)	Pew	same as above	18 Aug 2026	As above
55%	US adults under 30 more concerned than excited	Pew	same as above	18 Aug 2026	As above
+17 pts	Rise in Democrats expecting fewer jobs from AI	Pew	https://www.pewresearch.org/short-reads/2026/09/16/democrats-are-now-more-worried-than-republicans-about-ai-and-its-impact-on-jobs/	16 Sep 2026	Search summary only; base years not confirmed
34% / 42% / 16%	25-country median: more concerned / equally / more excited	Pew Global Attitudes, spring 2025	https://www.pewresearch.org/global/2025/10/15/how-people-around-the-world-view-ai/	15 Oct 2025	Strong method; search summary only
50% vs 16%	More concerned than excited: US (and Italy) vs South Korea	Pew Global Attitudes	same as above	15 Oct 2025	As above
39%	US adults saying AI does more harm than good (31% 2025, 40% 2023)	Gallup/ Bentley	https://news.gallup.com/poll/712751/americans-cool-toward.aspx	Jul 2026	Gallup web panel; Decrypt reports a conflicting higher figure
79%	US adults expecting AI to reduce US jobs over 10 years (73% 2025)	Gallup/ Bentley	same as above	Jul 2026	As above
27%	US adults trusting business to use AI responsibly (31% 2025, 21% 2023)	Gallup/ Bentley	same as above	Jul 2026	As above
66% / 46%	Regular AI use / willing to trust AI, 47 countries	KPMG & Univ. of Melbourne, n>48,000	https://figshare.unimelb.edu.au/articles/report/Trust_attitudes_and_use_of_artificial_intelligence_A_global_study_2025/28822919	2025 (fieldwork Nov 2024 to Jan 2025)	Large multi-country panel; consultancy co-author
52% / 53%	Excited / nervous about AI products, 30-country average	Ipsos AI Monitor 2025	https://www.ipsos.com/en-dk/ipsos-ai-monitor-2025	Jun 2025	Online panel; less representative in emerging markets
72% vs 32%	Trust in AI, China vs US	Edelman 2025 tech report	https://www.edelman.com/sites/g/files/aatuss191/files/fr/2025-08/Top%20Findings_2025%20EdelmanTrust%20Barometer_Insights%20for%20the%20Technology%20Sector%20%281%29.pdf	Aug 2025	PR-firm survey

Figure	What it measures	Source	URL	Date	Reliability note
>60% / 84%	Europeans positive on AI at work / want careful management	Eurobarometer via Commission	https://digital-strategy.ec.europa.eu/en/news/commission-survey-shows-most-europeans-support-use-artificial-intelligence-workplace	13 Feb 2025	Commission press framing
95.04%	SAG-AFTRA video game agreement ratification vote, ending strike begun Jul 2024	PC Gamer	https://www.pcgamer.com/gaming-industry/the-videogame-voice-actors-strike-officially-ends-as-performers-ratify-agreement-between-sag-aftra-and-publishers	Jul 2025	Union-reported tally
91.42% (19.25% turnout)	SAG-AFTRA 2026 TV/ Theatrical ratification	Fortune	https://fortune.com/2026/06/05/sag-aftra-approves-four-year-contract-studios-streamers/	5 Jun 2026	Union-reported tally
45 roles	CBA customer-service jobs cut for voice bot, then reinstated	Bloomberg; FSU	https://www.bloomberg.com/news/articles/2025-08-21/commonwealth-bank-reverses-job-cuts-decision-over-ai-chatbots	21 Aug 2025	Volume claims disputed between bank and union
~2,400	Kaiser NUHW mental-health workers in one-day AI-related strike	KQED; NUHW	https://www.kqed.org/news/12076753/northern-california-kaiser-therapists-hold-1-day-strike-over-ai-patient-care-concerns	18 Mar 2026	Some outlets say "more than 2,000"
60% / 69%	Nurses distrusting employer AI priorities / acuity scores mismatching own judgment	NNU survey, n>2,300	https://www.nationalnursesunited.org/press/national-nurses-united-survey-finds-ai-technology-undermines-patient-safety	May 2024	Union member survey, self-selected
2 Dec 2027 / 2 Aug 2028	New EU AI Act high-risk dates (standalone / embedded) after Digital Omnibus	Sidley; National Law Review	https://datamatters.sidley.com/2026/06/22/eu-lawmakers-reach-provisional-agreement-to-delay-key-eu-ai-act-obligations/	Jun-Jul 2026	Law-firm summaries; confirm on EUR-Lex
423-57-174	European Parliament vote on AI omnibus deal	Industry summary	https://www.europeanoutdoorgroup.com/members-knowledge-hub/european-parliament-adopts-provisional-agreement-on-ai-omnibus	16 Jun 2026	Secondary source; confirm on EP site
0	Known AI Act fines as of Nov 2026	Negative search result	n/a	Nov 2026	Absence of evidence; recheck
99-1	Senate vote to strip state-AI-law moratorium from budget bill	Nextgov/ Route Fifty	https://www.nextgov.com/artificial-intelligence/2025/07/senate-overwhelmingly-passes-amendment-removing-state-ai-moratorium/406448	1 Jul 2025	Search extract; roll call not checked on senate.gov
1 Jan 2027	Effective date of Colorado's replacement AI law (SB 26-189); original act never took effect	Ropes & Gray	https://www.ropesgray.com/en/insights/alerts/2026/05/colorado-scales-back-ai-law-with-targeted-implications-for-health-care	May 2026	Search extract; signing date 14 May (one source 18 May)
$1,500	Max NYC LL144 penalty per violation per day; enforcement found ineffective	NY State Comptroller; DLA Piper	https://www.osc.ny.gov/state-agencies/audits/2025/12/02/enforcement-local-law-144-automated-employment-decision-tools	2 Dec 2025	Search extract; audit PDF not read
$1m	Max SB 53 civil penalty per violation (large frontier developers, >$500m revenue)	FPF	https://fpf.org/blog/californias-sb-53-the-first-frontier-ai-law-explained/	Effective 1 Jan 2026	Search extract
54.4%	US nonfarm business labour share, Q4 2025 (down 1.7% y/y); 2025 productivity +2.2%	BLS Productivity and Costs	https://www.bls.gov/news.release/archives/prod2_03242026.htm	24 Mar 2026	Official statistic, search extract; subject to revision
~$700-745bn	Combined 2026 capex guidance, Microsoft, Alphabet, Amazon, Meta	Yahoo Finance; datacentres.com tracker	https://finance.yahoo.com/markets/article/magnificent-7-earnings-rush-reveals-ai-spending-surge-with-hyperscaler-capex-set-to-reach-725-billion-in-2026-224901707.html	Apr-Oct 2026	Company guidance via aggregators; definitions differ; not from filings

Figure	What it measures	Source	URL	Date	Reliability note
0 (CI rules out large effects)	Effect of AI chatbot adoption on earnings and hours, Denmark, 25,000 workers	Humlum & Vestergaard, NBER WP 33777	https://econpapers.repec.org/RePEc:nbr:nberwo:33777	2025	Strong design; early period (2023-24); search extract
<=0.66 %	Upper bound on 10-year TFP gain from AI (0.53% adjusted)	Acemoglu, Economic Policy	https://www.nber.org/papers/w32487	2024/2025	Academic forecast
$1.5bn (~$3,000/work)	Bartz v. Anthropic class settlement, final approval 20 Jul 2026	JURIST	https://www.jurist.org/news/2026/07/judge-approves-record-1-5-billion-settlement-involving-anthropic/	Jul 2026	Search extract; docket not checked
2,243	Westlaw headnotes copied by Ross; Third Circuit: not fair use	Haynes Boone; LawNext	https://www.haynesboone.com/news/alerts/third-circuit-rejects-fair-use-for-rosss-copying-of-westlaw-headnotes	29 Sep 2026	Search extract; opinion not read
415 TWh (1.5%) to ~945 TWh	Global data-centre electricity 2024 to 2030 (base case)	IEA, Energy and AI	https://www.iea.org/reports/energy-and-ai/executive-summary	Apr 2025	Agency projection; search extract
6.7-12%	Projected US data-centre share of electricity by 2028	LBNL (Shehabi et al.)	https://energyanalysis.lbl.gov/publications/2024-lbnl-data-center-energy-usage-report	19 Dec 2024	Range reflects scenario uncertainty
$193,000	DoNotPay FTC final order (5-0)	FTC via Slashdot	https://tech.slashdot.org/story/25/02/11/1932223/	11 Feb 2025	Search extract
$25m / $15.9m / $14m	Alleged consumer losses: Ascend Ecom / FBA Machine / Click Profit	FTC complaints via press	https://siliconangle.com/2024/09/25/ftc-cracks-misleading-ai-powered-business-schemes-new-initiative/	2024-2025	Allegations, not findings
$225k + $175k	SEC penalties, Delphia and Global Predictions (first AI-washing cases)	Baker McKenzie; Ropes & Gray	https://www.ropesgray.com/en/insights/alerts/2025/01/ai-washing-enforcement-in-the-final-days-of-the-gensler-administration	Mar 2024	Search extract

Track 10. Contrarian (from `11-contrarian.md`)

Figure	What it measures	Source	URL	Date	Reliability note
≤0.66% (and <0.53%)	Upper-bound rise in US TFP over 10 years from AI	Acemoglu, Economic Policy	https://www.nber.org/papers/w32487	May 2024 (journal 2025)	Model-based; excludes new tasks. Not re-read at source
0.8% (vs 1.5% in 2024)	US private nonfarm TFP growth, 2025	BLS	https://www.bls.gov/news.release/archives/prod3_03192026.pdf	19 Mar 2026	Official; seen in search excerpt
2.2% (vs 3.0% in 2024)	US private nonfarm labour productivity growth, 2025	BLS	https://www.bls.gov/news.release/archives/prod3_03192026.pdf	19 Mar 2026	Official; seen in search excerpt
1.4% annualised; 2.2% y/y	US nonfarm business labour productivity, Q2 2026	BLS	https://www.bls.gov/opub/ted/2026/productivity-up-2-2-percent-from-second-quarter-2025-to-second-quarter-2026.htm	Aug/ Sep 2026	Official
2.1% a year	US labour productivity growth Q4 2019 to Q2 2026	BLS	https://www.bls.gov/opub/ted/2026/productivity-up-2-2-percent-from-second-quarter-2025-to-second-quarter-2026.htm	2026	Official
-0.8% q/q; +0.7% y/y	UK output per hour, Q2 2026	ONS via Trading Economics	https://tradingeconomics.com/united-kingdom/productivity/news/576075	18 Aug 2026	Official, but RTI vs LFS measures disagree
15.9%	TFP level above official, end-2017, after intangibles adjustment	Brynjolfsson, Rock, Syverson	https://www.aeaweb.org/doi/10.1257/mac.20180386	2021	Peer-reviewed, model-based

Figure	What it measures	Source	URL	Date	Reliability note
1.4%	Time saved by genAI as share of all US work hours	Bick, Blandin, Deming	https://www.nber.org/system/files/working_papers/w32966/w32966.pdf	Feb 2025	n>10,000, self-reported
1-5% / 1-7%	Share of all US work hours assisted by genAI	Bick, Blandin, Deming (two versions)	https://fedinprint.org/item/fedlwp/98805/102025	Feb / Oct 2025	Range changed between versions; cite version
~3% (2.8%)	Average time savings among Danish chatbot users	Humlum & Vestergaard	https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf	Apr/ May 2025	Admin + survey data, ~25,000 workers
3-7%	Share of productivity gains passed through to earnings	Humlum & Vestergaard	https://bfi.uchicago.edu/wp-content/uploads/2025/04/BFI_WP_2025-56-3.pdf	2025	Seen in coverage
>$1tn	Expected AI capex "in coming years"	Goldman Sachs Top of Mind 129	https://www.goldmansachs.com/insights/goldman-sachs-exchanges/a-skeptical-look-at-ai-investment	25 Jun 2024	Opinion/forecast
$600bn / $500bn gap	Annual AI revenue needed to justify capex / gap	Sequoia (Cahn)	https://www.sequoiacap.com/article/ais-600b-question/	20 Jun 2024	Heuristic (Nvidia rev x2 x2)
$2tn; $800bn	Annual revenue needed by 2030; shortfall after savings	Bain Global Technology Report	https://www.theregister.com/2025/09/24/bain_ai_costs	23 Sep 2025	Consultancy forecast
$803-828bn	2026 capex of largest AI spenders	Hargreaves Lansdown	https://www.hl.co.uk/news/ai-hyperscalers-are-expected-to-finish-the-year-spending-a-whopping-828-billion-on-capex	Oct 2026	Secondary; includes SpaceX; check filings
>$1tn in one week	Market value lost by top 20 AI/chip stocks, mid-Oct 2026	Dealroom	https://dealroom.co/charts/ai-selloff	30 Oct 2026	Secondary analysis
19% slower	Task time with AI, experienced OSS developers	METR RCT	https://arxiv.org/abs/2507.09089	Jul 2025	RCT, n=16, 246 tasks, early-2025 tools
~20% faster (believed)	Developers' own post-hoc estimate of AI speedup	METR RCT	https://arxiv.org/abs/2507.09089	Jul 2025	Same study
50%: ~14.5h vs 80%: ~1h	Claude Opus 4.6 METR time horizons	Wikipedia (METR)	https://en.wikipedia.org/wiki/METR	2026	Secondary; check metr.org
24% / 30.3%	Best agent full task completion, TheAgentCompany v1/v2	Xu et al.	https://www.alphaxiv.org/abs/2412.14161v2	Dec 2024 / 2025	Benchmark; older models
~58% / ~35%	Best agents single-turn / multi-turn success, CRMArena-Pro	Salesforce AI Research	https://arxiv.org/html/2505.18878v1	May 2025	Vendor benchmark; generic agents
>60% to ~25%	GPT-4o retail pass^1 to pass^8	tau-bench (Sierra)	https://arxiv.org/pdf/2406.12045	Jun 2024	Old model
28.4% to 22.4%	Adenoma detection rate in non-AI colonoscopies before vs after routine AI use	Budzyń et al., Lancet Gastroenterol Hepatol	https://wrap.warwick.ac.uk/id/eprint/191005	Oct 2025	Observational, 19 endoscopists, 4 centres; confounded; search extract
319	Knowledge workers surveyed on GenAI and critical thinking	Lee, Sarkar et al., CHI 2025	https://www.microsoft.com/en-us/research/?p=1135061	Apr 2025	Self-reported, cross-sectional
54	Participants in MIT "Your Brain on ChatGPT" EEG study	Kosmyna et al., arXiv 2506.08872	https://arxiv.org/abs/2506.08872	Jun 2025	Preprint; small n
N=4,439	Participants in four experiments on the social penalty for AI use	Reif, Larrick, Soll, PNAS	https://doi.org/10.1073/pnas.2426766122	May 2025	Preregistered vignette experiments

Figure	What it measures	Source	URL	Date	Reliability note
77% (n=625 employees)	Employees saying AI added to their workload	Upwork Research Institute	https://investors.upwork.com/news-releases/news-release-details/upwork-study-finds-employee-workloads-rising-despite-increased-c	Jul 2024	Vendor survey; 2,500 total incl. 1,250 C-suite
>6,000 firms	OECD employer survey on algorithmic management, 6 countries	OECD AI Papers no. 31	https://ideas.repec.org/p/oec/comaaa/31-en.html	2025	Employer (manager) perceptions
~30%	Amazon warehouse injuries above warehouse industry average, 2023	US Senate HELP majority staff	https://www.help.senate.gov/imo/media/doc/amazon_investigation.pdf	Dec 2024	Partisan report; Amazon disputes
+12% to 38,306	Practising US radiologists, 2010 to 2022	Radiology Business (reporting a study)	https://radiologybusiness.com/topics/healthcare-management/healthcare-staffing/radiologist-pipeline-barely-keeping-pace-population-growth	2025	Secondary; search extract
~$571,000 (+9%)	Average US radiologist pay, 2025	Medscape via Fortune	https://fortune.com/2026/05/04/godfather-of-ai-geoffrey-hinton-radiologists-future-of-work-tech-ai-job-anxiety/	May 2026	Secondary
75,300; +2%	US interpreter/translator jobs 2024; projected growth 2024-34	BLS OOH	https://www.bls.gov/ooh/media-and-communication/interpreters-and-translators.htm	2025	Official; search extract
-21%	Freelance platform posts for automation-prone (writing, coding) jobs, 8 months after ChatGPT	Demirci, Hannane, Zhu, Management Science	https://spiral.imperial.ac.uk/entities/publication/88113a91-b994-4717-8526-dadac657f1cb	2025	Peer-reviewed; 1.39m posts; one platform
36%; 43%	UK translators who lost work / income to GenAI	Society of Authors	https://www.bookbusinessmag.com/aggregatedcontent/survey-finds-generative-ai-proving-major-threat-to-the-work-of-translators/	2024	Self-selected, 787 respondents; secondary
~20 to ~13 tellers; +43% branches	ATM effect per urban branch; urban branch growth	Bessen via Conversable Economist	https://conversableeconomist.com/2015/03/03/atms-and-a-rising-number-of-bank-tellers/	Mar 2015	Secondary report of IMF F&D article
>40%	Agentic AI projects Gartner predicts cancelled by end-2027	Gartner	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027	25 Jun 2025	Forecast, not measurement
~130	Genuine agentic vendors out of "thousands"	Gartner	same	25 Jun 2025	Opinion; method not disclosed

DO NOT CITE

These figures are widely repeated but either trace to a weak origin, are routinely misread, or failed verification in this session. Each line says what to use instead, where there is something.

Failed or not confirmed in the verification pass

Figure	Problem	Use instead
121,401 AI-cited US announced job cuts, Jan to Oct 2026 (Challenger)	Could not be confirmed by independent search; October 2026 report not found	112,713 through July 2026 (Challenger July report), with the caveats in 08
Census BTOS "9% to 18%"	Corrected: about 10% to about 17% per St Louis Fed	"Roughly doubled because the question changed"
"Context files make agents worse" (Gloaguen et al.) without qualification	Applies to LLM-generated files; developer-written files slightly helped	State the split

Viral statistics with weak or misread origins

Figure	Where it really comes from	Why not to cite it
"95% of AI pilots fail"	MIT NANDA, July 2025, self-described preliminary; 52 interviews, 153 conference-collected survey responses, 300 public initiatives	It measures custom enterprise tools reaching production with P&L impact in about six months, in a self-selected sample. Not a failure rate for "AI pilots".
"70% of change programmes fail"	Hammer and Champy's self-described "unscientific estimate"; Hughes (2011) found no valid empirical basis for any of five common sources	No empirical basis. Contested by Burnes and Jackson (2011), but nobody has measured it.
"More than 80% of AI projects fail" (RAND 2024)	RAND attributes it to "some estimates"; it is not RAND's finding	Cite RAND for its root causes (65 interviews), not for the 80%.
"30 to 50% of RPA projects fail" (EY)	One 2016 EY thought-leadership PDF: "as many as" 30 to 50% of initial projects, no method published	Hedge dropped and misdated in later citations.
"88% of AI POCs don't reach production" (IDC/Lenovo 2025)	4 of 33 POCs in a vendor-sponsored study	Same sponsor reported 46% a year later. Neither is a stable estimate.
"30% of GenAI projects abandoned after PoC by end of 2025" and ">40% of agentic projects cancelled by 2027" (Gartner)	Gartner press releases, July 2024 and June 2025	FORECASTS, routinely cited as findings or in the past tense. Cite only as Gartner's prediction.
"Only about 130 real agentic AI vendors" (Gartner)	Press release; method not stated	An analyst estimate with no published method.
"The UK lost 8% of jobs to AI" (Morgan Stanley, Jan 2026)	Survey of firms already using AI for a year, in five sectors, reporting own headcount changes	Not a labour-market statistic.
"40% of jobs will be affected by AI" (IMF 2024)	Exposure measure (task overlap), SDN/2024/001	Exposure, not displacement; about half of exposed jobs judged likely to benefit. Cite only as exposure.
"47% of jobs at risk" (Frey and Osborne 2013)	Model-based automation probability by occupation	Not a forecast of job losses; task-based re-estimates are far lower.
"AI makes developers 19% slower"	METR RCT, July 2025, n=16 experienced maintainers, early-2025 tools	Solid but narrow; METR's 2026 redesign says newer data cannot give a reliable estimate. Never cite as a general claim.
"77% of employees say AI added to their workload" (Upwork)	Vendor survey; the 77% covers employees using AI among 625 employees in a 2,500 sample with 1,250 executives	Usable only with its scope. Do not drop "using AI".
"Big Tech cut new-grad hiring by 50%" (SignalFire)	VC firm's own talent data; change in new-grad share of hires vs 2019	Comparison year chosen for effect; versus 2023 the fall is much smaller.
"UK entry-level jobs down a third since ChatGPT" (Adzuna)	Job-board adverts Nov 2022 to mid-2025	Overlaps monetary tightening and NI and minimum-wage rises. Cannot be attributed to AI.
"98% of small businesses use AI" (US Chamber/Teneo 2024)	Counts any software with AI features	Not deliberate AI use.
"AI adoption among small businesses surged 41%" (Thryv 2025)	Relative growth 39% to 55% in a 540-person vendor panel	Misread as a level.
"68% / 76% of small businesses use AI" (Goldman Sachs 10KSB)	Alumni of Goldman's own programme	Not a general SMB sample.
"Klarna's AI does the work of 700 agents"	Klarna's own Feb 2024 announcement, unaudited	Use only with the 2025 quality reversal and the F-1 headcount figures attached.
"IBM fired 8,000 people and replaced them with AI"	Blend of separate events and viral posts	CEO said HR roles replaced were "in the hundreds"; total headcount rose.
"Salesforce cut 4,000 jobs because of AI"	Benioff podcast remark (support about 9,000 to 5,000), no timeframe or method	CEO statement, not a filing. Cite as his claim only.
"77% of businesses use AI for automation" (Anthropic)	77% is the share of sampled API transcripts showing automation patterns	Not a share of businesses.

Figure	Where it really comes from	Why not to cite it
"Iceland's four-day week was an overwhelming success"	Advocacy report (Autonomy and Alda, 2021): 35-36 hour weeks for about 2,500 public-sector workers	Not a four-day week; about 1% of the workforce.
"92% of UK four-day-week firms kept it"	56 of 61 self-selected firms continuing in some form; 18 permanent	Self-selected sample run by an advocacy organisation.
"The one-person billion-dollar company is here" (Medvi)	NYT profile, April 2026; $1.8bn is the founder's projection	About $401m 2025 revenue; FDA warning letter and class action.
"AI's $600B question" (Sequoia)	One partner's blog heuristic, June 2024	Heuristic, not an accounting gap.
"AI will add only 0.66% to GDP" (Acemoglu)	0.66% is a TFP upper bound over ten years, model-based, excluding new tasks	Do not conflate TFP with GDP.
"AI deskills doctors"	One observational study, 19 endoscopists, confounded	A signal, not a finding.
"ChatGPT rots your brain" (MIT Media Lab)	54-person EEG preprint on essay writing	Not peer-reviewed at release; framing outran data.
"79% of Americans don't trust businesses to use AI responsibly"	2023 Bentley-Gallup wave	Out of date; use the 2026 figure (27% trust).
"EU fined companies EUR 18.4m under the AI Act"	Low-quality web page; no official source	No AI Act fine found to date.
"UK introduced a frontier AI bill in May 2026"	Two low-quality sites	Law-firm sources say the King's Speech contained none.
"Data centres use 4.4% of US electricity" (LBNL)	Attributed to LBNL 2024; not confirmed in extracts	Check the report first.
"Record 1,208 radiology residency positions in 2025"	Single Substack post	Check NRMP data.
"Hundreds of millions of hours saved" style vendor ROI claims, including Forrester TEI studies	Vendor-commissioned models	OPINION, not evidence.

Unverified leads (named in track files but not sourced this session)

- Sculley et al.'s "only a small fraction of a real ML system is ML code" (Figure 1): check the NeurIPS PDF.
- Any "process debt" origin claim: not traced.
- Keynes's "fifteen-hour week" wording, Covello's "$1tn problem" line, Gartner's "agent washing" definition, the Moffatt tribunal's "remarkable submission": fetch originals before quoting.

13. Quotes

Every quote below was copied from a page fetched in this session and re-fetched and matched word for word in the verification pass (see verification.md). Format: quote (speaker, source, date) URL. Context: one line.

Read this before using these quotes. Only three hosts could be fetched in this session: anthropic.com, microsoft.com and raw.githubusercontent.com. So every verified quote comes from Anthropic, Microsoft or a benchmark repository. Anthropic and Microsoft sell AI and have an interest in adoption. This is a sampling artefact of the session's network limits, not a judgement that these are the best or most balanced quotes. Before publishing, balance them with the candidate quotes at the end once their sources can be opened.

AI-native, shared context and the Frontier Firm (from track file 02)

- "The upshot is that deploying AI for complex tasks might be constrained more by access to information than on underlying model capabilities." (Anthropic Economic Index report, Anthropic, 15 Sep 2025) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report. Context: the authors' reading of input-length data from 1M API transcripts. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "This could represent a barrier to broader enterprise deployment for some important tasks that rely on dispersed context that is not already centralized or digitized. Correcting for this bottleneck may require firms to restructure their organization, invest in new data infrastructure, and centralize information for effective model deployment." (Anthropic Economic Index, 15 Sep 2025) same URL. Context: the closest primary statement to the user's shared-context definition. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "For software development with centralized code repositories, access to this information is in principle straightforward." (Anthropic Economic Index, 15 Sep 2025) same URL. Context: why coding leads agent adoption. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "I work way more with Claude than with any of my colleagues." (Anonymous Anthropic engineer, "How AI is transforming work at Anthropic", 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: reduced collaboration. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "More junior people don't come to me with questions as often." (Anonymous senior Anthropic engineer, same source, 2 Dec 2025) same URL. Context: mentorship loss. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Structured around on-demand intelligence and powered by "hybrid" teams of humans + agents, these companies scale rapidly, operate with agility, and generate value faster." (Microsoft, 2025 Work Trend Index, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: the Frontier Firm definition. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "In many cases, people are ready. The systems around them are not." (Microsoft, 2026 Work Trend Index, 5 May 2026) https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization. Context: framing for the

"Transformation Paradox". Verified: re-fetched and matched in the verification pass, 2026-11-06.

- "They capture these signals and encode them into shared routines, improving future work while preserving accountability and control." (Microsoft, 2026 Work Trend Index, 5 May 2026) same URL. Context: what Frontier Firms do with agent outputs; the nearest Microsoft gets to "shared context". Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Organizations that can answer these questions are building Owned Intelligence" (Microsoft, 2026 Work Trend Index, 5 May 2026) same URL. Context: the page then defines Owned Intelligence as institutional know-how unique to the firm. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Adoption, hidden use and stigma (from track file 05)

- "Without guidance or clearance from the top, employees are taking things into their own hands and keeping AI use under wraps: 78% of AI users are bringing their own AI tools to work (BYOAI)" (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: framing of shadow AI. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "52% of people who use AI at work are reluctant to admit to using it for their most important tasks." (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: hidden use. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Only 39% of people globally who use AI at work have gotten AI training from their company." (Microsoft and LinkedIn, Work Trend Index 2024, 8 May 2024) https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part. Context: training deficit. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Last year, employees led the AI wave. This year, it's flipped." (Microsoft, Work Trend Index 2025, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: leaders now ahead of employees on agent familiarity. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "69% of professionals mentioned the social stigma that can come with using AI tools at work" (Anthropic, "Introducing Anthropic Interviewer", 4 Dec 2025) https://www.anthropic.com/research/anthropic-interviewer. Context: 1,250 AI-conducted interviews; vendor source. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "A colleague recently said they hate AI and I just said nothing. I don't tell anyone my process because I know how a lot of people feel about AI." (Fact-checker interviewed by Anthropic Interviewer, 4 Dec 2025) https://www.anthropic.com/research/anthropic-interviewer. Context: hidden AI use because of peer stigma. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "When producing output is so easy and fast, it gets harder and harder to actually take the time to learn something." (Anthropic employee, quoted in "How AI is transforming work at Anthropic", 2 Dec 2025) https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic. Context: skill-atrophy concern in an AI-native firm. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Agent reliability and benchmarks (from track file 06)

- "**Key finding**: AI agent time horizons have been doubling approximately every 7 months." (METR, eval-analysis-public repository README, accessed 2026-11-06) https://raw.githubusercontent.com/METR/eval-analysis-public/main/README.md. Context: METR's own summary of its time-horizon method. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "τ³-bench added voice full-duplex evaluation, the `banking_knowledge` domain and 75+ task fixes" (Sierra Research, tau2-bench repository README, accessed 2026-11-06) https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md. Context: the benchmark's maintainers repaired dozens of their own task definitions. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "2.0 English runs are **not** input-identical to v1.0.1, so 2.0 scores are not comparable with the 1.x leaderboard." (Sierra Research, tau2-bench README, accessed 2026-11-06) https://raw.githubusercontent.com/sierra-research/tau2-bench/main/README.md. Context: agent scores are tied to benchmark versions, so cross-time comparisons need care. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "But how performant are AI agents at helping to accelerate or even autonomously perform work-related tasks?" (TheAgentCompany authors, repository README, accessed 2026-11-06) https://raw.githubusercontent.com/TheAgentCompany/TheAgentCompany/main/README.md. Context: framing question of the benchmark. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Working with agents: oversight, deference and division of labour (from track file 07)

- "Across the four analytical choices, analysts went with their agent’s lean 86% of the time." (Amaç Herdağdelen, Hailey Joren, Alex Tamkin, Anthropic, "How AI agents shape experts’ judgment calls", 27 Oct 2026) https://www.anthropic.com/research/agents-judgment-calls. Context: RCT where agent preferences were set by coin flip. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "a human reviewer is only an independent check if their judgment is their own" (same authors, same page, 27 Oct 2026) https://www.anthropic.com/research/agents-judgment-calls. Context: why deference matters for oversight. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Manual prompts sit in the middle, and in practice users accept 93% of them anyway." (Anthropic engineering, "How we built Claude Code auto mode", 25 Mar 2026) https://www.anthropic.com/research/claude-code-auto-mode. Context: approval fatigue in step-by-step oversight. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The 17% false-negative rate on real overeager actions is the honest number." (same page, 25 Mar 2026) https://www.anthropic.com/research/claude-code-auto-mode. Context: residual risk of classifier-based approvals. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "effective oversight doesn’t require approving every action but being in a position to intervene when it matters" (Anthropic, "Measuring AI agent autonomy in practice", 18 Feb 2026) https://www.anthropic.com/research/measuring-agent-autonomy. Context: Anthropic's interpretation of rising auto-approve and interrupt rates. Verified: re-fetched and matched in the verification pass, 2026-11-06.

- "existing models are capable of more autonomy than they exercise in practice" (same page, 18 Feb 2026) https://www.anthropic.com/research/measuring-agent-autonomy. Context: the "deployment overhang" argument. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "On average, people make about 70% of the planning decisions but only 20% of the execution decisions" (Anthropic, "Agentic coding and persistent returns to expertise", 16 Jun 2026) https://www.anthropic.com/research/claude-code-expertise. Context: division of labour in ~400k sessions. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "77% of API transcripts show automation patterns (especially full task delegation) versus just 12% for augmentation" (Anthropic Economic Index, 15 Sept 2025) https://www.anthropic.com/research/anthropic-economic-index-september-2025-report. Context: enterprise API use. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "primarily impacted behaviors that could be changed independently and not behaviors that required coordination to change" (Dillon, Jaffe, Immorlica, Stanton, Microsoft Research abstract, Apr 2025) https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/. Context: Copilot RCT, email vs meetings. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "someone who builds, delegates to, and manages agents to amplify their impact" (Microsoft Work Trend Index 2025, definition of "agent boss", 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: vendor framing (OPINION). Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Some may blend HR and IT or create new leadership roles" (same, 23 Apr 2025) https://www.microsoft.com/en-us/worklab/work-trend-index/2025-the-year-the-frontier-firm-is-born. Context: org-design forecast. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "So far, 1 million skills have been published and nearly 280 million installed." (Anthropic, "Donating Agent Skills to the Agentic AI Foundation", 22 Oct 2026) https://www.anthropic.com/news/donating-agent-skills-to-the-agentic-ai-foundation. Context: vendor adoption claim. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Displacement and expectations (from track file 08)

- "We find no systematic increase in unemployment for highly exposed workers since late 2022, though we find suggestive evidence that hiring of younger workers has slowed in exposed occupations" (Maxim Massenkoff and Peter McCrory, Anthropic economists, "Labor market impacts of AI: A new measure and early evidence", 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: key-findings summary of CPS analysis. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The averaged estimate in the post-ChatGPT era is a 14% drop in the job finding rate compared to that in 2022 in the exposed occupations, although this is just barely statistically significant. (There is no such decrease for workers older than 25.)" (same source, 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: young-worker section. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The young workers who are not hired may be remaining at their existing jobs, taking different jobs, or returning to school." (same source, 5 Mar 2026) https://www.anthropic.com/research/labor-market-impacts. Context: authors' caveat on the hiring result. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Following adoption, junior employment declines sharply in adopting firms relative to non-adopters, while senior employment trends remain largely unchanged." (Seyed M. Hosseini

and Guy Lichtinger, Harvard, "Generative AI as Seniority-Biased Technological Change", version 6 May 2026, abstract) https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf. Context: headline result. Verified: re-fetched and matched in the verification pass, 2026-11-06.

- "The decline is driven primarily by slower hiring rather than increased separations." (same, abstract) Same URL. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "One potential implication is that if firms overestimate the near-term automation capabilities of GenAI, the resulting decline in junior employment may ultimately prove temporary." (same, introduction) Same URL. Context: authors' forward-looking-hiring interpretation. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Notably, respondents were on average more worried about job loss for others than for themselves." (Anthropic Economic Index report "Cadences", 26 Jun 2026) https://www.anthropic.com/research/economic-index-june-2026-report. Context: survey of about 9,700 Claude users. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "The Economic Index Survey is not representative of the general population." (same) Same URL. Context: authors' own sample caveat. Verified: re-fetched and matched in the verification pass, 2026-11-06.
- "Across US workers, ten more percentage points of work affected is associated with 0.4 points more forecasted unemployment." (Sacher, Cotter, Jones, Korinek, McCrory, "The AI expectations gap", Anthropic, 16 Oct 2026) https://www.anthropic.com/research/ai-expectations-gap. Context: Morning Consult survey, n=10,980. Verified: re-fetched and matched in the verification pass, 2026-11-06.

Candidate quotes NOT verified (do not use in quotation marks yet)

These were seen only in search extracts or remembered from secondary coverage. They are paraphrased here. Fetch the source and copy the exact wording before quoting.

- **Moffatt v. Air Canada, 2024 BCCRT 149 (14 Feb 2024):** the tribunal described Air Canada's argument that its chatbot was a separate legal entity responsible for its own actions as a "remarkable submission". Source to fetch: https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html
- **Sebastian Siemiatkowski (Klarna CEO), 2025:** said the cost focus had produced lower-quality customer service. Exact wording not verified.
- **Commonwealth Bank of Australia, Aug 2025:** reported as conceding it had not adequately considered all relevant business factors when cutting 45 roles. Exact wording not verified.
- **Amjad Masad (Replit CEO), July 2025:** called the agent's deletion of a production database unacceptable and something that should never be possible. Exact wording not verified.
- **Satya Nadella, Microsoft FY26 Q4 call (29 Jul 2026):** said Microsoft has over 30 million paid Microsoft 365 Copilot seats with net seat adds more than doubling quarter over quarter. Wording appears in search extracts only; check the transcript.
- **Gartner (Anushree Verma), 25 Jun 2025:** described most agentic AI projects as early-stage experiments driven by hype. Check the press release.
- **Jim Covello (Goldman Sachs), Top of Mind 129, June 2024:** framed AI spending as a large investment problem that needs AI to solve complex problems to pay back. Check the PDF.

- **Hammer and Champy (1993):** described their 50 to 70% re-engineering failure figure as an unscientific estimate (as reported by Hughes 2011).
- **John Maynard Keynes (1930):** suggested three-hour shifts or a fifteen-hour week would satisfy most people's need to work. Check the essay text.
- **Will Stronge (Autonomy), July 2021:** called the Iceland trials an overwhelming success. Advocacy source.
- **Hosseini and Lichtinger cite Jerome Powell (Sept 2025 FOMC press conference)** on AI and young graduates. Check the Fed transcript.
- **Ron DeSantis (Florida Governor), 2025-26:** said in substance that an executive order cannot pre-empt state legislation. Check a primary transcript.

14. Reading list

Ranked by how much each source would change how you write about the AI-native transition, weighted towards primary sources and against vendor material. One line each on why. Each track file has a longer track-specific list under "Sources worth reading in full".

Status in brackets: [read] means the page was fetched this session; [not opened] means the network blocked it and it is known only from search extracts, so reading it in full matters more.

Top 25

1. **Humlum and Vestergaard, "Large Language Models, Small Labor Market Effects", NBER w33777 (May 2025), plus the 2026 update.** The best-identified test of your thesis: wide adoption with employer encouragement, tiny measured effects. If you read one thing, read this. [not opened] https://econpapers.repec.org/RePEc:nbr:nberwo:33777
2. **RAND, "The Root Causes of Failure for AI Projects and How They Can Succeed" (Aug 2024).** The only non-vendor study of failure causes, and the best support for "nobody can describe the problem". Read the method limits too. [not opened] https://www.rand.org/pubs/research_reports/RRA2680-1.html
3. **Microsoft Work Trend Index 2026, "Agents, human agency and the opportunity for every organization".** Vendor, but the data contradict "people resist": people ready, systems not. [read] https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization
4. **Hosseini and Lichtinger, "Generative AI as Seniority-Biased Technological Change" (May 2026 version).** Best-identified evidence of the junior-hiring effect, and the argument that firms cut in anticipation. [read] https://raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf
5. **Anthropic, "How AI agents shape experts' judgment calls" (Oct 2026).** RCT showing human oversight of agents is mostly deference. Essential if you sell "humans and agents together". [read] https://www.anthropic.com/research/agents-judgment-calls
6. **MIT NANDA, "The GenAI Divide: State of AI in Business 2025".** The most cited and most misread source in the field. Read the methodology page before ever using "95%". [not opened] https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
7. **Federal Reserve, Small Business Credit Survey 2026 Report on Employer Firms.** The only large US small-firm survey that measures depth of AI integration and reasons for non-use. [not opened] https://www.fedsmallbusiness.org/2026-report-on-employer-firms
8. **ONS, "Artificial intelligence in UK businesses: 2023 to 2026" (July 2026).** Official breadth-vs-depth data, including micro firms. [not opened] https://www.ons.gov.uk/businessindustryandtrade/business/businessservices/articles/artificialintelligenceinukbusinesses/2023to2026
9. **Dillon, Jaffe, Immorlica, Stanton, "Shifting Work Patterns with Generative AI" (Microsoft Research).** 6,000-worker RCT showing individual work changes and coordinated work does not. The cleanest evidence for "transformation is the hard part". [read] https://www.microsoft.com/en-us/research/publication/shifting-work-patterns-with-generative-ai/

10. **Anthropic Economic Index, September 2025.** The best large-sample support for the shared-context idea ("constrained more by access to information"). Vendor; inferred, not measured. [read] https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
11. **Gloaguen et al., "Evaluating AGENTS.md" (Feb 2026).** Controlled test of context files: LLM-generated ones hurt, human-written ones help a little. Directly relevant to Studio. [not opened] https://arxiv.org/abs/2602.11988
12. **CRMArena-Pro (Salesforce AI Research, 2025).** The most business-shaped agent benchmark: multi-turn success falls to about 35%. [not opened] https://arxiv.org/abs/2505.18878
13. **tau-bench paper and repository.** pass^k is the right way to talk about agent reliability in business processes. [read] https://arxiv.org/abs/2406.12045 and https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md
14. **Anthropic, "How AI is transforming work at Anthropic" (Dec 2025).** A rare inside view of an AI-native firm, including mentoring and collaboration costs. [read] https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic
15. **Hughes, "Do 70 Per Cent of All Organizational Change Initiatives Really Fail?", Journal of Change Management (2011).** Retires the most-used change statistic. [not opened] https://research.brighton.ac.uk/en/publications/do-70-per-cent-of-all-organizational-change-initiatives-really-fa/
16. **Brynjolfsson, Rock and Syverson, "The Productivity J-Curve" (AEJ: Macro 2021).** The theory that best explains why transformation looks unprofitable at first. [not opened] https://www.aeaweb.org/doi/10.1257/mac.20180386
17. **Massenkoff and McCrory (Anthropic), "Labor market impacts of AI" (Mar 2026).** Clear null on unemployment with stated power; young-worker hiring signal. [read] https://www.anthropic.com/research/labor-market-impacts
18. **Brynjolfsson, Chandar, Chen, "Canaries in the Coal Mine?" August 2026 revision.** The most cited displacement dataset; read the revision, not the 2025 headlines. [not opened] https://digitaleconomy.stanford.edu/news/canariesaug26/
19. **Autor, Chin, Salomons, Seegmiller, "New Frontiers: The Origins and Content of New Work, 1940-2018" (QJE 2024).** Where new work comes from, and why augmentation vs automation decides which history repeats. [not opened] https://economics.mit.edu/sites/default/files/2026-04/New-Frontiers-QJE-v139%283%29-p1399-1465-Aug-2024.pdf
20. **Kasy and Lehner, "Employing the unemployed of Marienthal" (IZA DP 16088).** Direct test of what work gives people beyond income. Essential for "freedom to be more human". [not opened] https://docs.iza.org/dp16088.pdf
21. **Vivalt et al., OpenResearch unconditional cash study, NBER w32719.** What people actually did with money and time. [not opened] https://www.nber.org/system/files/working_papers/w32719/w32719.pdf
22. **Moffatt v. Air Canada, 2024 BCCRT 149.** Short ruling; the clearest statement that a company owns what its agent says. [not opened] https://www.canlii.org/en/bc/bccrt/doc/2024/2024bccrt149/2024bccrt149.html
23. **Census BTOS "AI Question Wording Updates" note.** Required before citing any US adoption trend across late 2025. [not opened] https://www.census.gov/hfp/btos/downloads/AI%20Question%20Wording%20Updates.pdf
24. **Ericsson, "Defining AI native" white paper.** The only maturity model built around the word, and it admits the term has no clear definition. Useful for positioning your own definition. [not opened] https://www.ericsson.com/en/reports-and-papers/white-papers/ai-native

25. **Gartner press release, 25 June 2025 (agentic AI cancellations and "agent washing").** Read it to see how thin the basis of a much-quoted forecast is. [not opened] https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027

Next 10, by theme

- **Agents in production:** Anthropic, "Measuring AI agent autonomy in practice" (Feb 2026), the only large dataset on how people supervise agents. [read] https://www.anthropic.com/research/measuring-agent-autonomy
- **Teaming:** Dell'Acqua et al., "The Cybernetic Teammate" (NBER w33641, 2025), P&G field experiment. [not opened] https://www.nber.org/papers/w33641
- **Macro scepticism:** Acemoglu, "The Simple Macroeconomics of AI" (NBER w32487). [not opened] https://www.nber.org/papers/w32487
- **Developer productivity:** METR RCT (arXiv 2507.09089), with METR's 2026 redesign post. [not opened] https://arxiv.org/abs/2507.09089
- **Enterprise data reality:** Spider 2.0 (arXiv 2411.07763). [not opened] https://arxiv.org/abs/2411.07763
- **Public opinion:** Pew, "How the U.S. Public and AI Experts View Artificial Intelligence" (Apr 2025). [not opened] https://www.pewresearch.org/internet/2025/04/03/how-the-us-public-and-ai-experts-view-artificial-intelligence/
- **Trust vs use:** KPMG and University of Melbourne, "Trust, attitudes and use of AI: a global study 2025". [not opened] https://figshare.unimelb.edu.au/articles/report/Trust_attitudes_and_use_of_artificial_intelligence_A_global_study_2025/28822919
- **Clerical automation history:** Feigenbaum and Gross, "Answering the Call of Automation" (QJE 2024). [not opened] https://www.nber.org/system/files/working_papers/w28061/w28061.pdf
- **UK labour market:** Bank Underground, "Canaries in the column" (Aug 2026). [not opened] https://bankunderground.co.uk/2026/08/06/canaries-in-the-column-ai-exposure-and-the-uks-hiring-slowdown/
- **Counter-narrative on enterprise ROI:** Wharton/GBK "Accountable Acceleration" (2025). [not opened] https://knowledge.wharton.upenn.edu/special-report/2025-ai-adoption-report

15. Open questions

What to research next, and where this pack is thin. Blunt by request.

What was thin, and why

The whole pack rests mostly on search extracts, not on reading the sources. The cloud environment's network policy blocked page fetches for almost every host. Only anthropic.com, microsoft.com and raw.githubusercontent.com could be opened. A shared cap of about 200 web searches per turn also cut the first research wave short: every one of the ten track files was first written INCOMPLETE, and a second gap-filling pass narrowed but did not close the gaps. Concretely:

- No government statistics release, regulator text, court ruling, NBER/arXiv paper (other than those mirrored on GitHub), consultancy report or news article was read in full. Their numbers are as reported by search extracts that cite them. The verification pass corroborated 17 of the most important ones by independent search and corrected two. The rest of the roughly 330 rows in 12-facts-and-figures.md have not been independently checked.
- **The verified quotes all come from Anthropic and Microsoft (plus benchmark READMEs).** Both sell AI. That skews the quotable record towards vendor framing. Do not mistake it for balance.
- **Vendor research is over-represented in the evidence on agents.** The best data on how people actually supervise agents (deference, approval rates, planning vs execution decisions) are Anthropic's studies of its own products. There is almost no independent field evidence.
- Several tracks have sections that rest on two or three sources. Weakest: SMB implementation cost and abandonment (no credible figure exists), process-documentation quality as a cause of AI failure (no study found), the UK and EU displacement evidence, the anti-AI case beyond energy and copyright, and the AI-advice "grifter" market (enforcement cases only, no market size).

To fix the access problem: add the blocked hosts to the cloud environment's allowed network domains (environment settings, Network access) and rerun the verification pass on 12-facts-and-figures.md. Priority hosts: rand.org, nber.org, arxiv.org, bls.gov, census.gov, ons.gov.uk, fedsmallbusiness.org, federalreserve.gov, gartner.com, mckinsey.com, pewresearch.org, canlii.org, eur-lex.europa.eu, oecd.org, budgetlab.yale.edu, digitaleconomy.stanford.edu.

The questions that matter most for your thesis

1. **Does inability to describe the workflow actually predict AI project failure?** This is the thesis's central causal claim and nothing found tests it. RAND's interviews are the closest. The process-mining literature documents gaps between designed and actual processes but does not link them to AI outcomes. You may be best placed to answer this with your own client data: before-and-after process maps, time to production, and what broke.
2. **Does a shared, stored context improve agent outcomes outside coding?** The only controlled evidence (Gloaguen et al. 2026) is for coding agents: LLM-generated context files hurt, developer-written ones help slightly. Nobody has tested SOPs, skills or a shared business context store for finance, ops or customer service. A controlled comparison in Studio would be original evidence.

3. **Is the junior-hiring drop permanent, or an anticipation effect that reverses?** Hosseini and Lichtinger raise this explicitly. It decides whether "AI-native" firms are quietly removing the apprenticeship layer. Where the missing young hires go is unstudied.
4. **Is the binding constraint readiness, or the size of the gain?** Danish administrative data show about 3% time savings and no earnings or hours effect despite adoption and encouragement. If gains are small for typical tasks, "the business isn't ready" may be a comforting misdiagnosis. Watch for the 2026 Humlum and Vestergaard update and any replication outside Denmark.
5. **Does human oversight of agents work in the field?** The 86% deference result is from a lab-style RCT with a vendor. Is there any field study of reviewer accuracy on agent output in finance, legal or HR?
6. **What does a credible "AI-native" 10-person firm look like, with numbers?** Only vendor case studies were found. An independent set of 5 to 10 documented small firms (cost, time to build, what broke, what staff now do) would be valuable and publishable.
7. **Does freed time go to "more human" work anywhere?** Early evidence points to intensification (Berkeley Haas ethnography) and lower intrinsic motivation (Scientific Reports 2025). Is there any setting, perhaps four-day-week trials combined with AI, where freed time went to care, learning or community? Nothing found.

Specific items to check next (from the track files)

Definitions (02): first business use of "AI-native" (VC usage traced only to 2023); the wording of Klarna's F-1 risk factors on AI and service quality; any SEC or FTC AI-washing actions after April 2025; Gartner, BCG and Deloitte maturity models in their own words.

SMB (03): the full DSIT 2026 adoption report (published 29 Oct 2026, not fetched); the SF Fed's July 2026 qualitative SBCS brief; the size of the "AI automation agency" market; SMB AI spend split between subscriptions, integration and outside help.

Enterprise (04): does the MIT NANDA report itself say purchased tools succeed about twice as often as internal builds, and that the core barrier is a "learning gap"? Why did IDC/Lenovo's POC-to-production figure move from 12% to 46% in a year? Salesforce Agentforce renewal and churn data. Any controlled evidence that semantic or context layers work. Procurement time from pilot to signed contract.

Inertia (05): an independent RCT of AI training (not access) on business outcomes; Gallup's "8.7-fold" manager effect at source; any measurement of adoption or productivity after de-layering.

What goes wrong (06): a systematic dataset of customer-service AI reversals (a rate, not a list); the lawyer vs self-represented split in the hallucination database over time; operational agent incidents with post-mortems.

Working with agents (07): firm filings that quantify AI-attributed headcount change with a method; whether a general-purpose business agent is "high-risk" under EU AI Act Annex III; UK and US liability positions for agent errors.

Displacement (08): a UK replication of the Canaries design using ONS/HMRC PAYE data; any AI-specific Bank of England Decision Maker Panel result; the OECD Employment Outlook 2026 chapter on AI and entry-level work; the October 2026 Challenger report (to replace the unconfirmed figure).

What people do (09): reconcile the OpenResearch labour-supply figures (1.3 hours and 2 points vs 3.9 points); detailed time-use tables; long-run Kenya UBI results; primary sources on Medvi.

Social and political (10): whether any EU AI Act fine has been issued; enforcement under California SB 53, Texas TRAIGA and Illinois HB 3773; Q3 2026 hyperscaler capex from filings; whether 54.4% is a series low for the US labour share; the strongest primary statements of the anti-AI case (Bender and Gebru, Marcus, Zitron) read in full rather than characterised.

Contrarian (11): BLS OEWS series for radiologists, translators and bank tellers pulled directly; NRMP radiology positions; METR's Feb 2026 redesign post read in full; the source of the circulating "translator earnings fell 29.7%" figure.

A blunt note on the method

Ten agents searching in parallel used up the search cap faster than one sequential analyst would have, and the network block meant depth could not substitute for breadth. The pack is broad, honestly labelled and usable for orientation and for deciding what to read. It is not yet citable for publication without reading the primary sources. Treat 01 and 99 as a well-informed first pass, not a final verdict.