Verification pass

Run 2026-11-06, after all track files were written and before the final files. Method notes first, then results.

Method and its limits

The brief asks for each URL to be re-fetched. In this session the environment's network policy blocked page fetches for almost every host. The hosts that could be fetched were www.anthropic.com, www.microsoft.com and raw.githubusercontent.com. So the verification used two methods:

- **RE-FETCHED:** the page was downloaded again with curl, converted to text (pdftotext for PDFs), and the number or quote was located in the text with its surrounding context. This is full verification.
- **CORROBORATED BY INDEPENDENT SEARCH:** for blocked hosts, a fresh web search (not the one the track agent used) was run and the figure, sample and date were checked against the extract it returned, which cites the primary page. This is weaker than a re-fetch: it confirms that the figure is reported consistently, not that I read it in the original. Treat these as "cite with the caveat that the original was not opened".
- **NOT CONFIRMED:** neither method confirmed the figure. These go to DO NOT CITE in 12-facts-and-figures.md.

Numbers checked (25)

#	Figure	Track file	Method	Result	Notes
1	tau-bench retail pass^1 0.692 to pass^4 0.462	06, 11, 01	Re-fetched README	CONFIRMED	Model is Claude 3.5 Sonnet (Oct 2024); GPT-4o 0.604 to 0.383. These are 2024 models. Caveat added to 01. https://raw.githubusercontent.com/sierra-research/tau-bench/main/README.md
2	Anthropic engineers use Claude in 59% of work; most can "fully delegate" 0-20%	02, 01	Re-fetched	CONFIRMED	Self-reported; 12 months earlier the figure was 28%. Productivity self-report +50%. https://www.anthropic.com/research/how-ai-is-transforming-work-at-anthropic
3	WTI 2026: 10% blocked agency, 5% unclaimed capacity, 13% rewarded for reinvention, n=20,000	02, 01	Re-fetched	CONFIRMED	Fielded by Edelman DXI, 18 Feb to 20 Apr 2026, 10 markets, AI users only. 16% "stalled". https://www.microsoft.com/en-us/worklab/work-trend-index/agents-human-agency-and-the-opportunity-for-every-organization
4	WTI 2024: 78% BYOAI, 52% reluctant to admit, 53% worry about looking replaceable, n=31,000	05, 01	Re-fetched	CONFIRMED	Edelman DXI, 31 markets, 15 Feb to 28 Mar 2024. https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part
5	86% of analysts followed agent's coin-flipped lean, n=50	07, 01	Re-fetched	CONFIRMED	Per-choice range 81% to 90%; 50% = no influence. https://www.anthropic.com/research/agents-judgment-calls
6	Junior employment at GenAI adopters down about 9% after six	08, 01	Re-fetch	CONFIRMED	4.0 fewer junior hires per firm-quarter relative to non-adopters, about 80% of the pre-period mean of 5.06. Résumé data, 65M workers, 280k firms. https://

#	Figure	Track file	Method	Result	Notes
	quarters; about 80% fewer junior hires per quarter		ed PDF		raw.githubusercontent.com/s-mahdihosseini/Website/b99fa2962dcc4c165bbc3c41dd283afa99c5ccfb/GenAI_SenBTC_May_2026.pdf
7	Those expecting AI to do 80%+ of their work forecast 12.2% unemployment in 2030, n=10,980	08	Re-fetched	CONFIRMED	Morning Consult, 11 to 23 Aug 2026, weighted. https://www.anthropic.com/research/ai-expectations-gap
8	Software engineering nearly 50% of agentic activity, n=998,481 tool calls	07	Re-fetched	CONFIRMED	https://www.anthropic.com/research/measuring-agent-autonomy
9	77% of API transcripts automation vs 12% augmentation	02, 07	Re-fetched	CONFIRMED	https://www.anthropic.com/research/anthropic-economic-index-september-2025-report
10	MIT NANDA: 52 interviews, 153 survey responses, 300+ initiatives; "5%"	04, 01	Independent search	CORROBORATED, with a correction of framing	Survey collected at four industry conferences (self-selected). The report's own sentence, as quoted by The Register, is that only 5% of custom enterprise AI tools reach production; elsewhere it says 95% of organisations get no measurable return. The two framings are both in the report. Use the narrow one. https://www.theregister.com/2025/08/18/generative_ai_zero_return_95_percent
11	Fed SBCS: 46% of small employer firms use AI; 7% of users "fully integrated"	03, 01	Independent search	CORROBORATED (partly)	46% confirmed via a June 2026 Fed Governor speech; 7% via an America's SBDC summary that does not name the survey year in the extract. Cite with caveat. https://www.fedsmallbusiness.org/2026-report-on-employer-firms
12	Census BTOS question change 17 Nov 2025; use roughly doubled	03, 01	Independent search	CORRECTED	Change date confirmed (new wording from 4 Dec 2025 data). St Louis Fed puts the jump at about 10% to 17%, not 9% to 18%. Track file 03 corrected. https://www.stlouisfed.org/on-the-economy/2026/jun/measuring-ai-adoption-firms-how-you-ask-matters
13	IDC/Lenovo 2026: 46% of AI POCs progressed to production, n=3,120	04, 01	Independent search	CORROBORATED	Lenovo press release; IDC fieldwork 16 Sep to 17 Oct 2025. Vendor-sponsored. https://www.businesswire.com/news/home/20260127255538/en/
14	Humlum and Vestergaard: about 3% time savings, precise zero effects on earnings and hours	05, 08, 11, 01	Independent search	CORROBORATED	Survey time saving is 2.8% (abstract rounds to 3%). The 3-7% figure is pass-through of gains to earnings; TechRadar misstates it. https://econpapers.repec.org/RePEc:nbr:nberwo:33777
15	US TFP +0.8% in 2025 vs 1.5% in 2024; labour productivity 2.2% vs 3.0%	11, 01	Independent search	CORROBORATED	Also: TFP grew 1.0% a year over 2019-25 vs 0.6% over 2007-19. That is a fair counterpoint and is now noted in 01. https://www.bls.gov/news.release/archives/prod3_03192026.pdf
16	McKinsey 2026: 37% attribute some EBIT to AI; $1B+ firms scaling agents 27% to 40%, n=1,719	04, 01	Independent search	CORROBORATED	Fielded 4 May to 8 June 2026, 97 countries. High performers about 6%, flat.
17	Stanford "Canaries" Aug 2026: 22-25 year olds in exposed jobs about 19% below counterfactual	08	Independent search	CORROBORATED	Not comparable with the 2025 version's 13% (different method). Authors call it descriptive, not causal. https://digitaleconomy.stanford.edu/news/canariesaug26/
		10			

#	Figure	Track file	Method	Result	Notes
18	Pew June 2026: 52% more concerned than excited; 71% expect fewer jobs		Independent search	CORROBORATED	n=3,488, 22 to 28 June 2026. https://www.pewresearch.org/short-reads/2026/08/18/young-adults-in-the-us-are-increasingly-wary-of-ai/
19	Challenger: 121,401 AI-cited US announced cuts Jan to Oct 2026	08	Independent search	NOT CONFIRMED	The independent search could not find the October 2026 report. Through July 2026 the figure was 112,713 (corroborated). Moved to DO NOT CITE; 08 annotated. https://www.challengergray.com/wp-content/uploads/2026/08/Challenger-Report-July-2026.pdf
20	CRMArena-Pro: about 58% single-turn, about 35% multi-turn	11, 01	Independent search	CORROBORATED	"Near-zero confidentiality awareness" appears in secondary summaries; the paper frames confidentiality as a separate test. Cite the 58/35 figures; soften the confidentiality phrasing. https://arxiv.org/abs/2505.18878
21	Gloaguen et al.: context files reduce success, cost +20%	07, 01	Independent search	CORRECTED	The reduction applies to LLM-generated context files. Developer-written files tended to slightly improve success. 01 and 07 corrected. This changes the reading in your favour. https://arxiv.org/abs/2602.11988
22	Microsoft 365 Copilot over 30M paid seats (July 2026)	04	Independent search	CORROBORATED	Nadella on FY26 Q4 call, 29 July 2026; about 7% of 450M+ commercial seats is an analyst estimate, not Microsoft's.
23	Spider 2.0: about 17% vs 91% on Spider 1.0	04, 01	Independent search	CORROBORATED	17.0% (o1-preview agent framework) on 632 tasks, Nov 2024. The "17-21%" range in 04 includes later models; the 17.0% figure is the one the paper states.
24	Charlotin database: about 1,600 decisions with hallucinated content by June 2026	06, 01	Independent search	CORROBORATED	AFP, 16 June 2026: about 1,600 across 35 countries; 1,116 in the US. Crowd-sourced, not definitive.
25	Klarna FTEs 5,527 (2022) to 3,422 (2024)	02, 01	Independent search	CORROBORATED	From the F-1. CEO says attrition and a hiring freeze did much of it, not AI alone. Later figure about 2,907 (2025).
26	Gartner: >40% of agentic projects cancelled by end-2027; poll of 3,412 webinar attendees; about 130 real vendors	04, 02	Independent search	CORROBORATED	FORECAST. Drivers: costs, unclear value, risk controls.
27	KPMG/Melbourne 2025: 57% hide AI use; 66% use regularly; 46% willing to trust	05, 10	Independent search	CORROBORATED	n=48,000+, 47 countries, Nov 2024 to Jan 2025.
28	Upwork 2024: 77% say AI added to workload	05, 11	Independent search	CORROBORATED, with scope note	Applies to employees using AI within a 2,500 sample that included 1,250 C-suite, 625 employees, 625 freelancers in four countries. Vendor survey.
		03	Independ	CORROBORATED	

#	Figure	Track file	Method	Result	Notes
29	ONS: 28% of firms with 0 to 9 employees use AI (June 2026)		ent search	(secondary)	From a summary of BICS wave 159; ONS bulletin itself shows 17% of 0-9 firms reporting extensive use vs 9% of 250+. Cite with caveat.

Totals: 9 confirmed by re-fetch; 17 corroborated by independent search (2 of them with framing caveats); 2 corrected; 1 not confirmed and moved to DO NOT CITE.

Quotes checked

Every quote in every track file's "Verbatim quotes" section (42 quotes across 02, 05, 06, 07, 08) was re-fetched from its URL with curl, converted to text and matched after normalising whitespace and curly quotes. **All 42 matched.** The script and page cache are in the session scratchpad, not the repo.

All 42 come from three hosts (anthropic.com, microsoft.com, raw.githubusercontent.com) because those were the only hosts reachable. **This makes the quote set heavily skewed towards two vendors with an interest in AI adoption.** That is a sampling artefact of this session, not a judgement that these are the best quotes. 13-quotes.md lists the candidate quotes from other sources (Gartner, Hughes, Keynes, Covello, the Moffatt tribunal) that should be fetched and added once the network allows.

Inline quotation marks in track-file findings that are not in the Verbatim quotes sections were NOT verified. A note to that effect was added under the title of each track file.

Corrections made to track files

- 03-smb.md: BTOS jump changed from "about 9% to about 17-18%" to "about 10% to about 17%".
- 07-working-with-agents.md and 01-thesis-stress-test.md: the Gloaguen et al. result qualified (LLM-generated vs developer-written context files).
- 08-displacement.md: Challenger Jan to Oct 2026 figure flagged NOT CONFIRMED; Jan to July figure (112,713) used instead.
- 01-thesis-stress-test.md: tau-bench model-vintage caveat added.

What a proper verification pass still needs

Open the blocked primary documents and check the remaining numbers in 12-facts-and-figures.md, starting with: the RAND report PDF, the MIT NANDA PDF, the Fed SBCS 2026 report, the ONS bulletin tables, Gartner's press releases, the BLS release, NBER w33777, and the CanLII text of Moffatt v. Air Canada. This needs the hosts added to the environment's allowed network domains.